



Vlaanderen
is supercomputing

VSC HPC Introduction

ICTS KU Leuven

VSC staff:

Ehsan Moravveji

Mag Selwa

Geert Jan Bex (UHasselt)

Jan Ooghe

Questions and problems troubleshooting

- Blackboard Collaborate session:
<https://eu.bbcollab.com/guest/ce4229a2f9eb4cc5b90f490d660eb1c2>



Material

- Everything is on Github:
<https://hpcleuven.github.io/HPC-intro/>
- Video Recordings
 - Scan the QR code
 - Recommended videos: ~ 2 hrs
 - Optional videos: ~1 hr



What is High Performance Computing?

- using supercomputers to solve advanced computation problems
- Reduce the computation time from days, years, decades, or centuries to minutes, hours, days, or weeks
- The key is parallelism



In practice, it is more like ...



The concept is simple: **Parallelism** = employing multiple processors for a single problem

Outline

- What is the VSC?
- What is a cluster?
- Genius Cluster
- Storage
- Login nodes
- Connection Setup
- Software environment
- How to submit jobs?
- Dedicated hardware
- How to choose resources?
- Optional material
 - Linux in brief
 - Conda for Python and R
 - Worker Framework



The background of the slide features a large, colorful mural painting. The mural depicts a futuristic cityscape with various floating spheres of different sizes and colors (red, blue, green). In the foreground, there is a large, stylized green leaf-like shape. A central element is a tablet or screen displaying a grid of numbers and data points, suggesting a theme of technology and data analysis.

VSC

(Vlaams Supercomputer Centrum)

VSC PARTNERSHIP



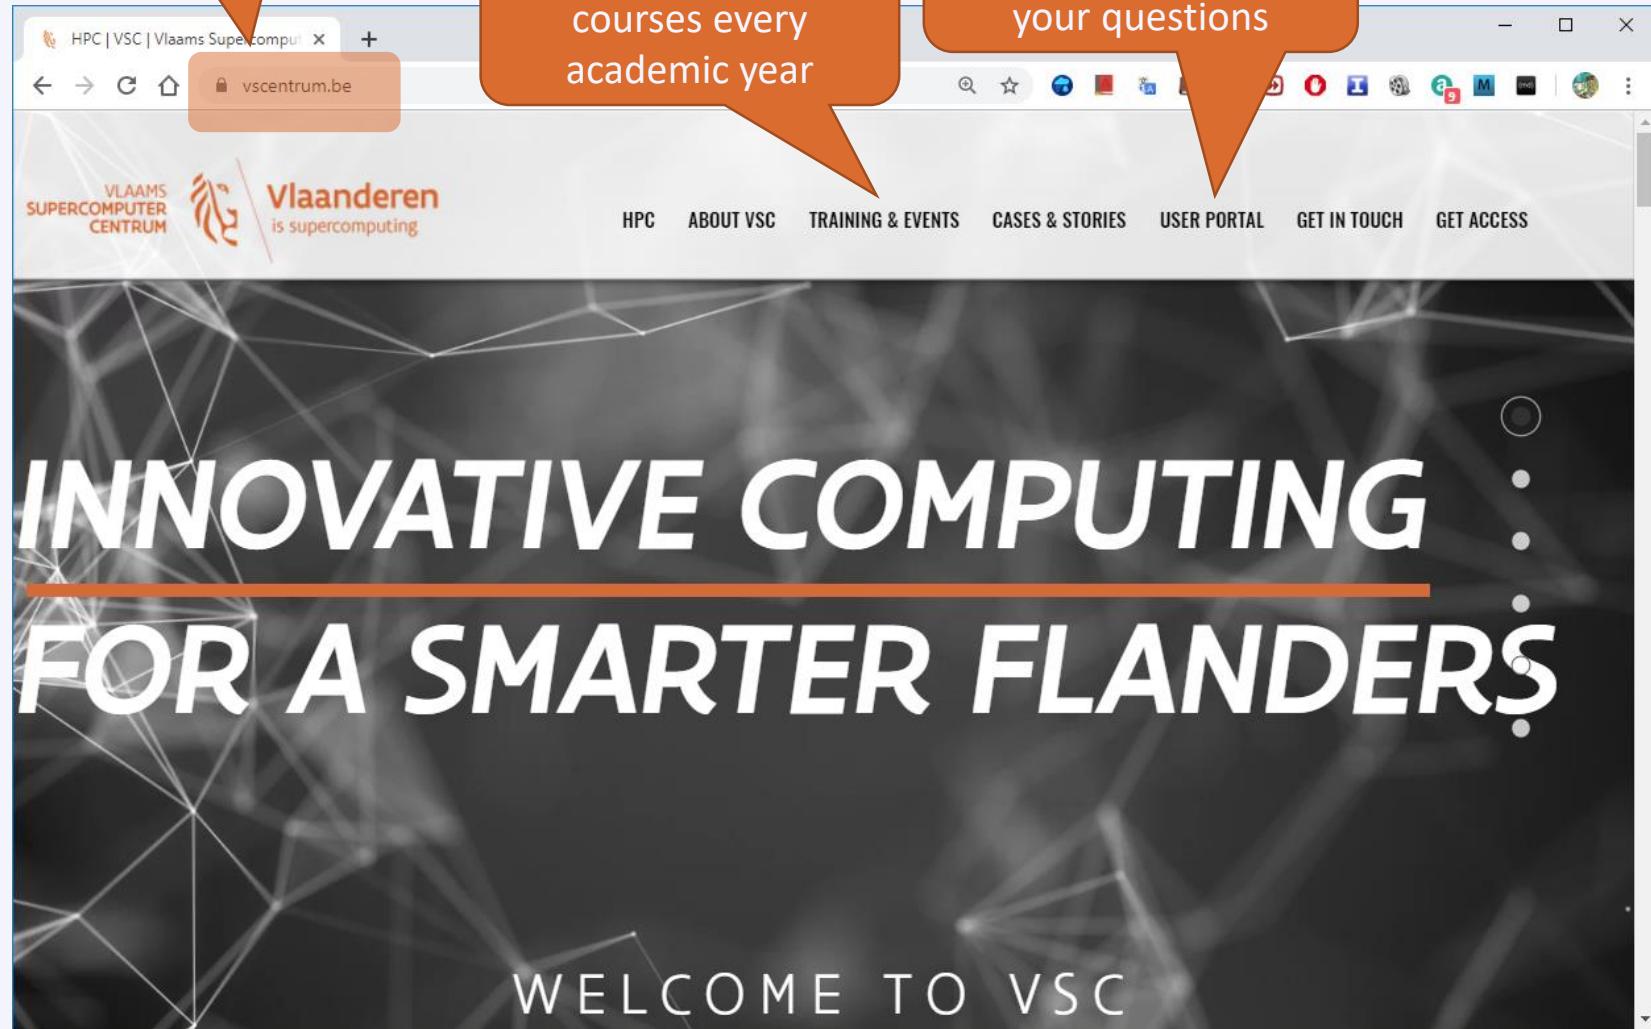
Supported by



VSC HPC Environments



www.vscentrum.be



Search your keywords here;
e.g. qsub

Welcome to VSC documentation

VSC documentation 1.0 documentation »

next | index

Next topic
Getting access

This Page
Show Source

Quick search

Go

- Getting access
 - VSC accounts
 - How to request an account?
 - Next steps
 - Additional information
- Access and data transfer
 - Logging in to a cluster
 - Data storage
 - Transferring data
 - GUI applications on the clusters
 - VPN
- Software stack
 - Using the module system
 - Specialized software stacks
- Running jobs
 - Job script
 - Submitting and monitoring a job
 - Job output
 - Troubleshooting
 - Advanced topics
- Software development

v: latest

Support and Services

Basic support

- Helpdesk (hpcinfo@kuleuven.be)
- Monitoring and reporting

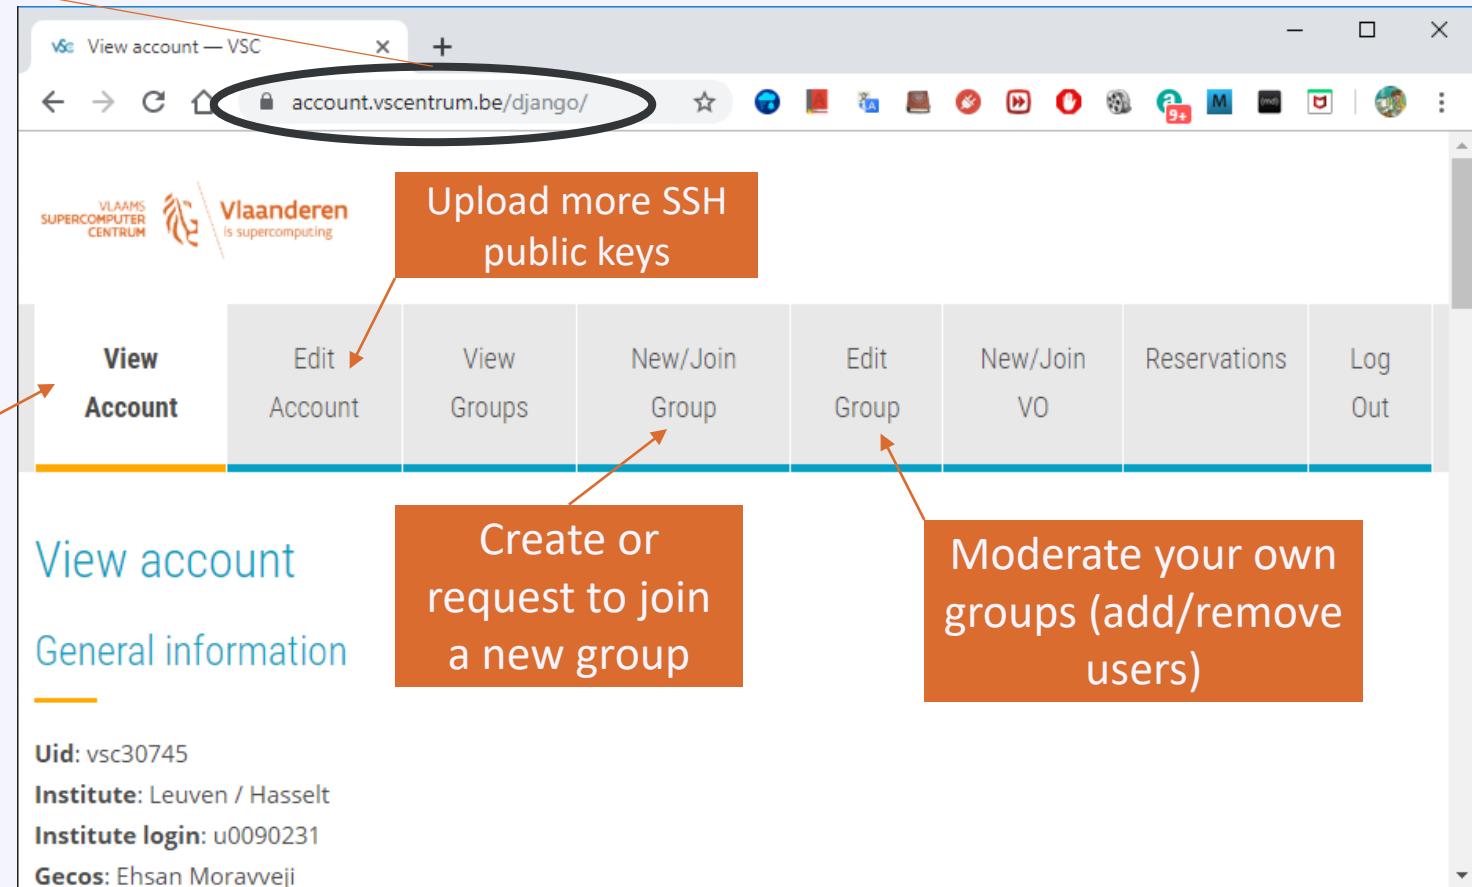
Application support

- Installation and porting
- Optimisation and debugging
- Benchmarking
- Workflows and best practices

Training

- Documentation and tutorials
- Scheduled trainings / workshops
- On request workshops
- One-to-one sessions

To manage your
VSC account:
account.vscentrum.be



Become a VSC user

- Create a secure (4096 bit) [SSH key pairs](#)
Upload it on the account page: www.account.vscentrum.be
- You need to [request a VSC account](#)
Normally processed swiftly
- Request [introductory credits](#) (2000 free credits for 6 months)
- Request [project credits](#) (for supervisors and project leaders)
 - You need to create a VSC group
 - Add users to the group to give them access to use credits
 - Fill out the request form
- Extra storage requests
 - Scratch extension: free of charge
 - Archive fileset: 70 € per TB per year
 - Staging fileset: 130 € per TB per year
- All service costs (compute and storage) are all explained
Go to ICTS service catalogus: <https://icts.kuleuven.be/sc>
Click on [High Performance Computing \(NL/EN\)](#)

The screenshot shows a web browser window titled 'View account — VSC'. The URL is 'account.vscentrum.be/djang...'. The page header includes the Vlaams Supercomputer Centrum logo and the text 'Vlaanderen is supercomputing'. A navigation menu at the top has tabs for 'View Account' (which is selected), 'Edit Account', 'View Groups', 'New/Join Group', 'Edit Group', 'New/Join VO', 'Reservations', and 'Log Out'. Below the menu, the main content area is titled 'View account' and 'General information'.

The screenshot shows a web browser window titled 'ICTS Servicecatalogus'. The URL is 'icts.kuleuven.be/sc'. The page header includes the KU Leuven logo and the text 'ICTS SERVICECATALOGUS'. It features a 'MENU' button, language options ('NL EN'), and a user profile ('Ehsan Moravveji'). The main content area includes sections for 'Home', 'ICTS SERVICECATALOGUS', 'WAT IS DE ICTS SERVICECATALOGUS?', 'OPGELET' (with a note about netto prijzen), and 'ZOEKEN' (Search). There is also a graphic of a computer mouse on a blue book.

VSC training 2020/2021

- Introductory



Linux

HPC intro

Linux for HPC

Make intro

Infosessions:

- Containers
- Notebooks

Linux scripting

Linux tools

Version control systems

- Intermediate

C++ for scientific computing

Fortran for programmers

C

- Python as a second language
- Python: System programming
- Scientific Python
- Python for Software engineering
- Python for data science
- Python for machine learning

worker/atools

- Advanced

High Performance Python

?

MPI

OpenMP

Debugging techniques

Code optimization

- Specialized track

PRACE MOOC Defensive programming and debugging: <https://www.futurelearn.com/courses/defensive-programming-and-debugging>

To Acknowledge VSC in publications

Why?

- a contractual obligation for the VSC
- helps VSC secure funding
- you will benefit from it in the long run

At KU Leuven

- add the relevant papers to the virtual collection "High Performance Computing" in Lirias

In het nederlands

De rekeninfrastructuur en dienstverlening gebruikt in dit werk, werd voorzien door het VSC (Vlaams Supercomputer Centrum), gefinancierd door het FWO en de Vlaamse regering – departement EWI.

In English

The computational resources and services used in this work were provided by the VSC (Flemish Supercomputer Center), funded by the Research Foundation - Flanders (FWO) and the Flemish Government – department EWI.



Tier-2 Clusters

Tier-2 Clusters @ KU Leuven

ThinKing (since 2014)

352 nodes: 7,616 cores



Genius (since 2018)

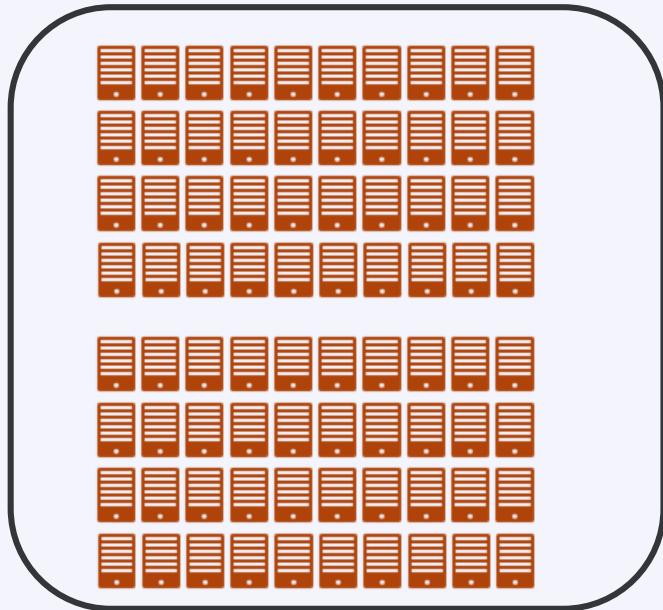
250 nodes: 8,936 cores



Tier-2 Overview



Compute nodes



Thinking:

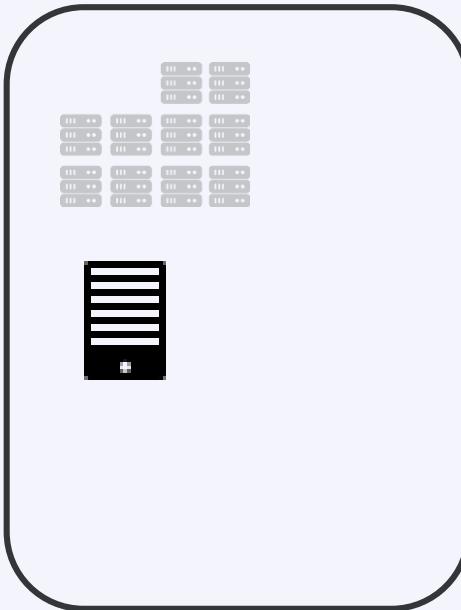
+ 48x + 86x Haswells 24c 64/128 GB

Genius:

+ 96x Skylake 36c 192 GB

+ 144x CascadeLake 36c 192 GB

Large memory nodes

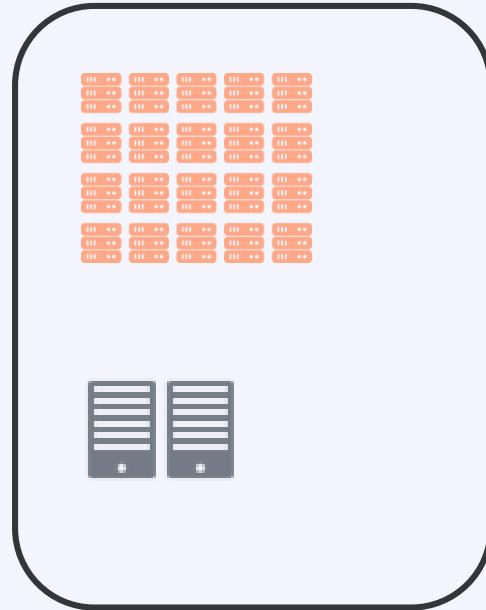


Genius:

+ 10x Skylake 36c 768GB

+ 1x Superdome 112c 6 TB

GPU nodes



Genius:

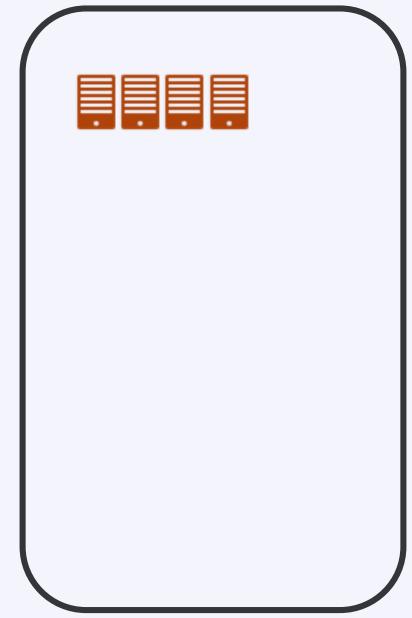
+ 20x Skylake 36c 192 GB

4x P100 16GB

+ 2x CascadeLake 36c 768GB

8x V100 32GB

Exp nodes



Genius:

4 AMD Naples 64c 256 GB

Technical Hardware Specifications

	Tier 2					Tier 1		
Cluster name	ThinKing		Genius			BrENIAC		
Processor type	Haswell		SkyLake			Broadwell		SkyLake
Cores per node	24		36			36		28
Base Clock Speed	2.5 GHz		2.3 GHz			2.6 GHz		2.6Ghz
Total nodes	48	96	86	10	144	580	408	
Node memory (GB)	64	96	192	768	192	128	256	192
Memory per core (GB)	2.5	3.8	5.2	21.2	5.3	4.4	9.0	6.7
Total cores	3,456		3,456			4,320		
Peak performance (Flops/cycle)	8 DP FLOPs/cycle: 4-wide FMA (fused multiply-add) instructions AVX2		16 DP FLOPs/cycle: 8-wide FMA (fused multiply-add) instructions AVX-512			16 DP FLOPs/cycle: 8-wide FMA (fused multiply-add) instructions AVX-512		8 DP FLOPs/cycle: 4-wide FMA (fused multiply-add) instructions AVX2
Network	Infiniband FDR		Infiniband EDR			Infiniband EDR		
Cache (L1 KB/L2 KB/L3 MB)	12x(32i+32d) / 12x256 / 30MB		18x(32i+32d) / 18x1024 / 25 MB			18x(32i+32d) / 18x1024 / 25 MB		14x(32i+32d) / 14x 256 / 35 MB

	Tier 2	Tier 1
Access	All researchers and students	Approved project experienced users
Start with	2,000 introduction credits	500 node days starting grant
Project Credits	Pay as you GO	Panel Review Project proposal Free
Computation Limit	None	500 - 5,000 node days
Application deadline	None	3 February 5 June 2 October
Max. Project Duration		8 months
Available walltime	1h, 24h, 72h, 7d	Max. 3 days



Genius

SkyLake (since 8/2018)
CascadeLake (since 1/2020)

Tier-2 Cluster: Genius

Type of node	CPU type	Inter-connect	# cores	installed mem	local discs	# nodes
SkyLake	Xeon 6140	IB-EDR	36	192 GB	800 GB	86
SkyLake large mem	Xeon 6140	IB-EDR	36	768 GB	800 GB	10
SkyLake GPU	Xeon 6140 4xP100 SXM2	IB-EDR	36	192 GB	800 GB	20
CascadeLake	Gold 6240	IB-EDR	36	192 GB	800 GB	144
CascadeLake GPU	Gold 6240 8xV100 SMX2	IB-EDR	36	768 GB	800 GB	2
SkyLake Superdome	Gold 6132	Flex Grid	14	6 TB	6 TB	8

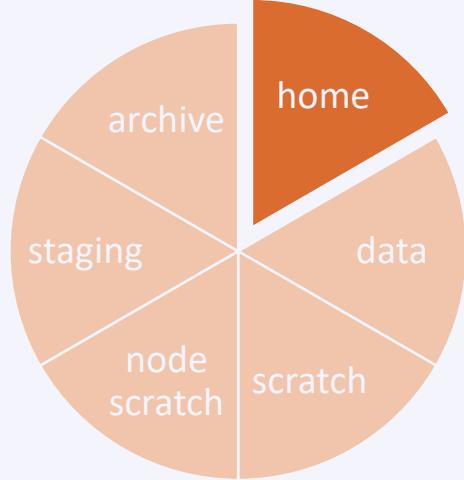


Storage

Overview of the storage infrastructure

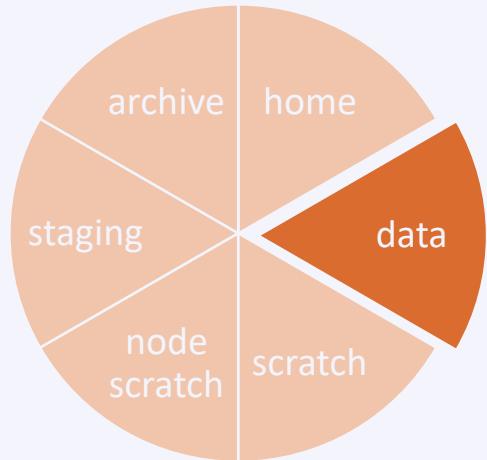
- Your files are owned only by you.
Other VSC users have no permission to read/write/execute your files (POSIX)
- A VSC account has 3 default storages (free of charge)
 - \$VSC_HOME
 - \$VSC_DATA
 - \$VSC_SCRATCH
- You can additionally request staging and archive storages
- Different storage volumes have different:
 - mount point
 - size and performance
 - use case
 - backup and maintenance policy
- More info on [ICTS Service Catalog](#) (EN/NL)

Storage



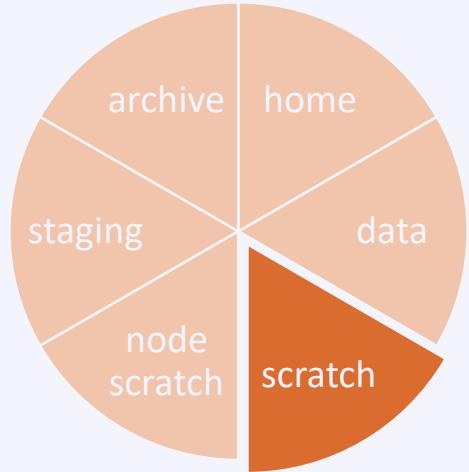
Storage	home folder
Env. Variable	\$VSC_HOME
Filesystem Type	NFS
Access	Global
Backup	Hourly, daily, weekly (until last month)
Default Quota	3 GB
Extension	Not possible
Usage	Only storing SSH keys, config files
Remarks	<ul style="list-style-type: none">- Stay away from using it- Can easily overflow:<ul style="list-style-type: none">+ Your jobs may crash+ Login issues

Storage



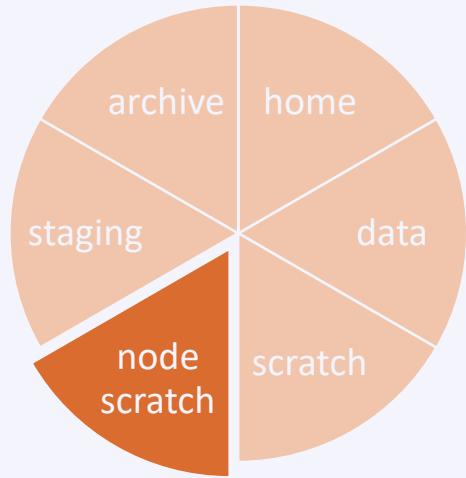
Storage	data folder
Env. Variable	\$VSC_DATA
Filesystem Type	NFS
Access	Global
Backup	Hourly, daily, weekly (until last month)
Default Quota	75 GB
Extension	On purchase
Usage	Your data, code, software, results
Remarks	<ul style="list-style-type: none">- Permanent storage for initial/final results- Not optimal for intensive or parallel I/O

Storage



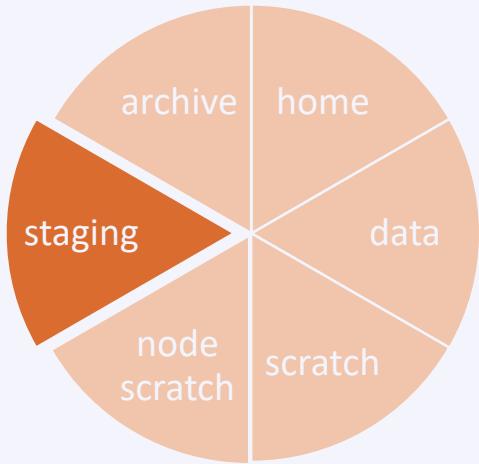
Storage	scratch folder
Env. Variable	\$VSC_SCRATCH
Filesystem Type	GPFS
Access	Local
Backup	delete after 28 days from last access
Default Quota	100 GB
Extension	Free
Usage	Intensive, parallel I/O, temporary files
Remarks	<ul style="list-style-type: none">- Recommended storage for all jobs- Copy scratch files to VSC_DATA or local storage after jobs are done- Deleted files cannot be recovered

Storage



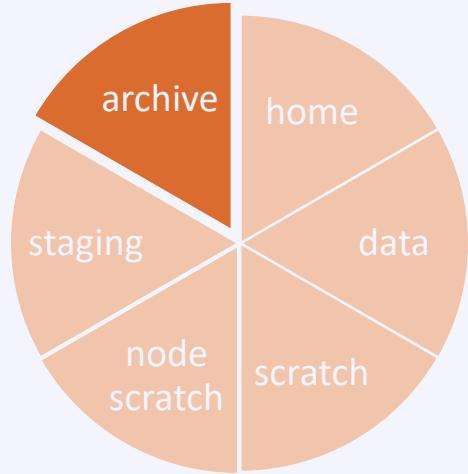
Storage	Node scratch folder
Env. Variable	\$VSC_SCRATCH_NODE
Filesystem Type	GPFS
Access	On compute node, only at runtime
Backup	None
Default Quota	200 GB
Extension	Read about beeOND
Usage	Temporary storage at runtime
Remarks	<ul style="list-style-type: none">- Fastest I/O, attached to the node- Is cleaned after job terminates- Copy the data to your home, scratch, or staging before job ends

Storage



Storage	Staging folder
Path	/staging/leuven/stg_000XX
Filesystem Type	GPFS
Access	On demand, only Tier-2@KUL
Backup	None
Default Quota	None
Extension	On purchase, from 1 TB
Usage	Permanent; share with a group
Remarks	<ul style="list-style-type: none">- Accessible from login/compute nodes- Fast, parallel I/O

Storage



Storage	Archive folder
Path	/archive/leuven/arc_000XX
Filesystem Type	NFS
Access	On demand, only Tier-2@KUL
Backup	None
Default Quota	None
Extension	On purchase, from 1 TB
Usage	Permanent; share with a group
Remarks	<ul style="list-style-type: none">- Accessible only from login nodes- Cannot use it from compute nodes- Slow, but durable filesystem

Storage

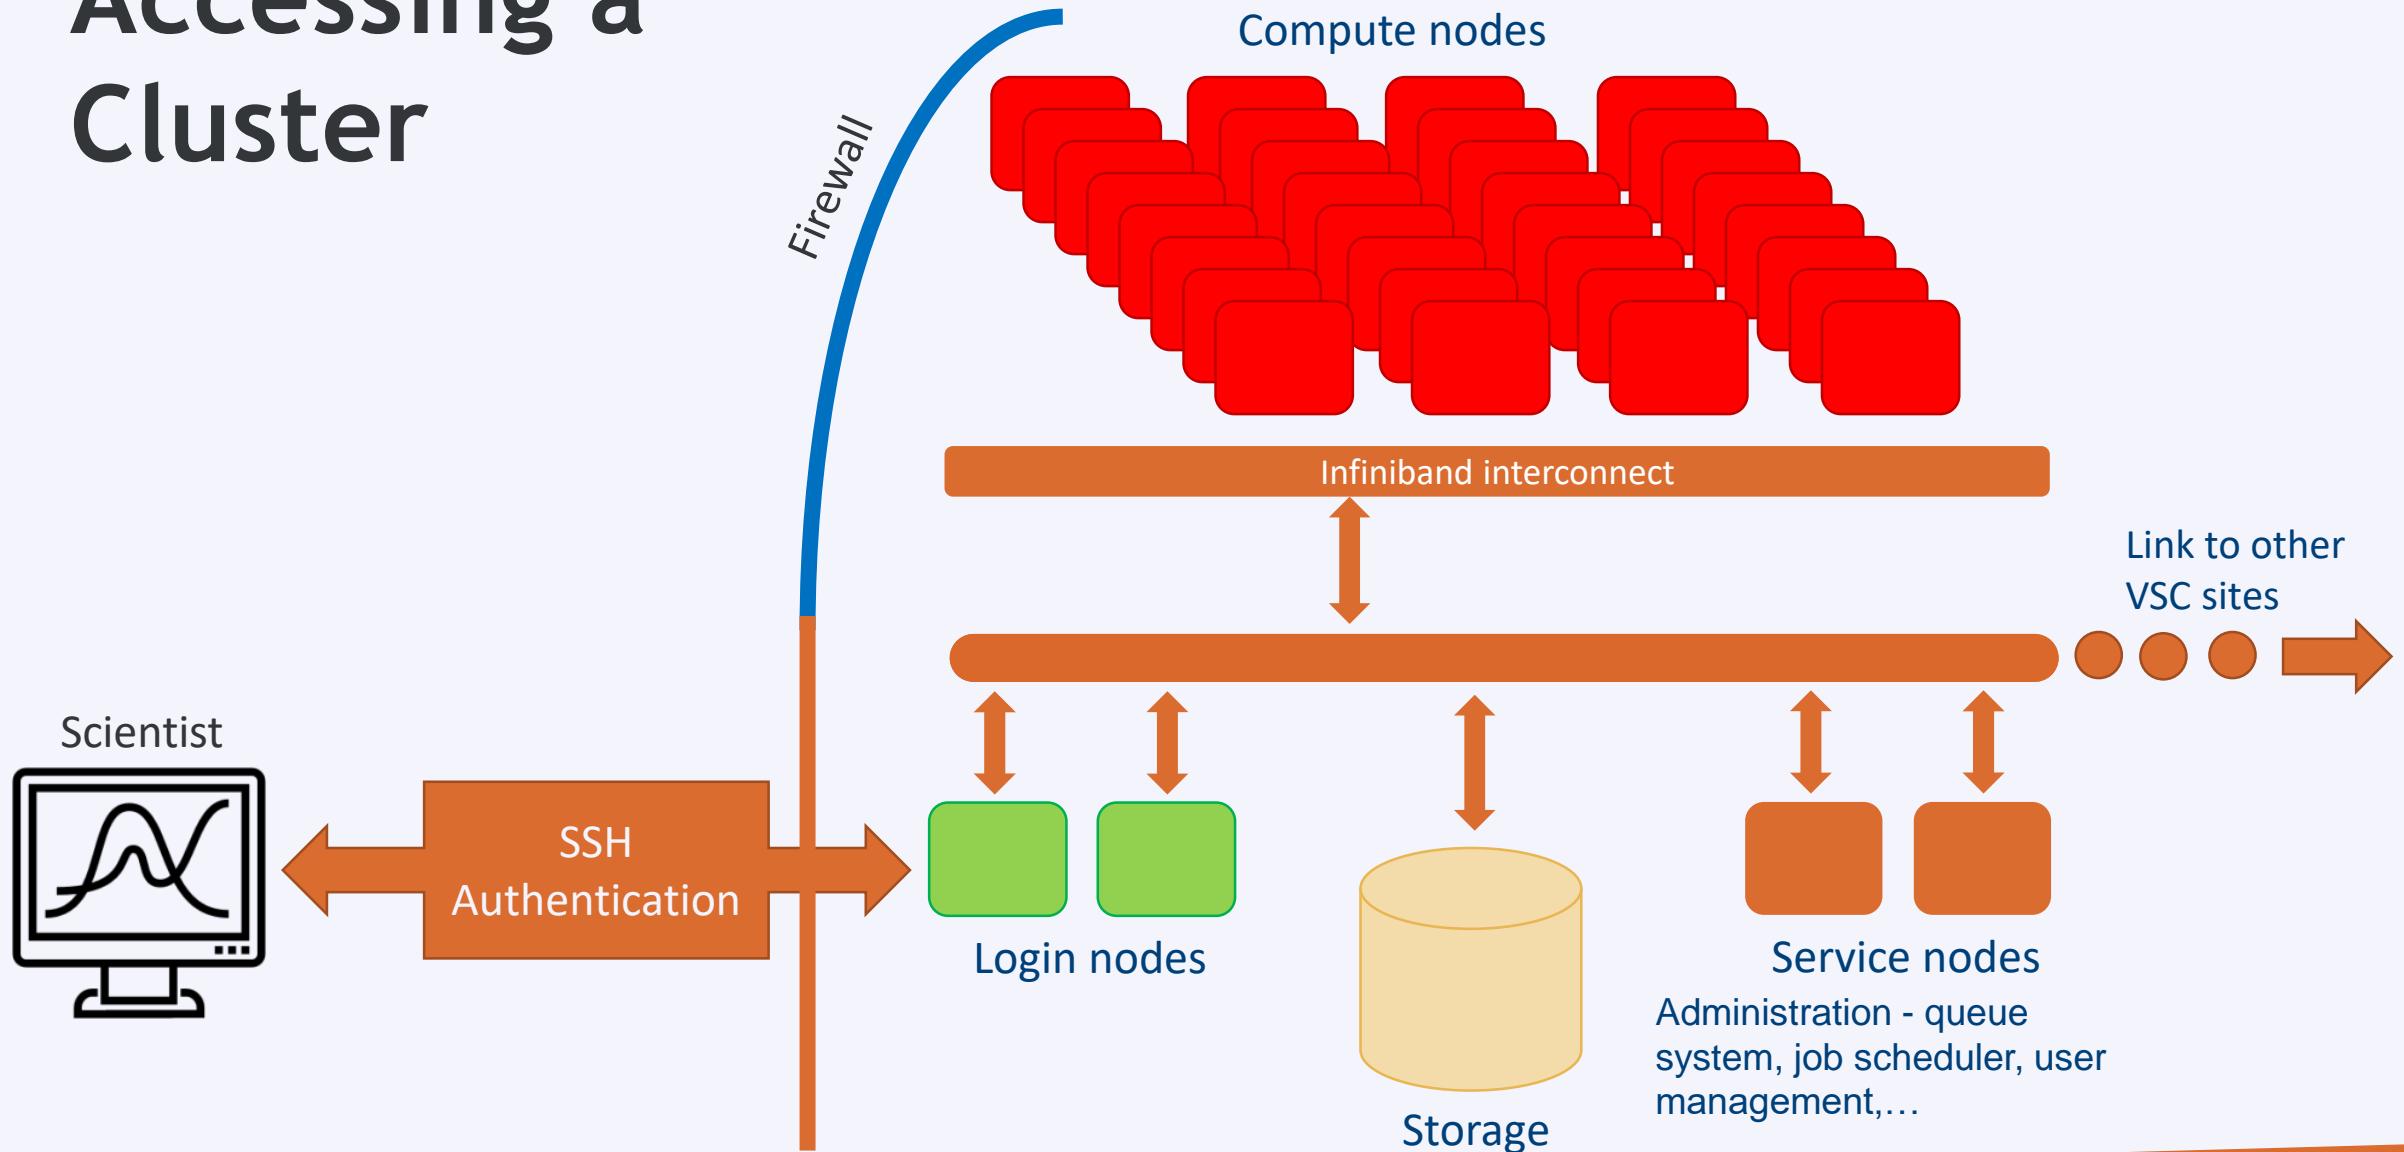
- [Request form for extra storage](#)
- [More information](#)
- Do not use /tmp**
It is only 10 GB and is reserved for the OS
and root processes
Your application can crash if using /tmp
- You are automatically logged into your home
folder upon login.
Make sure you immediately go to your other
storages, e.g.
`$> cd $VSC_DATA`
- Always check your storage balance using
myquota command

Example

```
$> myquota
file system $VSC_HOME
    Blocks: 1479M of 3072M
    Files: 12934 of 100k
file system $VSC_DATA
    Blocks: 102G of 225G
    Files: 1043k of 10000k
file system $VSC_SCRATCH
    Blocks: 15M of 1.5T
```

Login Nodes

Accessing a Cluster



Using Login Nodes

- To develop and/or compile code and/or software
- To check your storage and credit balance
- To manage jobs (submit, check status, debug, resubmit, ...)
- To move data around
 - within VSC: use data, scratch, staging, archive
 - outside VSC: copy/sync from/to your local storage
- To pre-process or post-process your data/jobs
- To visualize your data
- To share files/folders

Tips

- Login nodes are shared resources
- Do not** execute heavy-lifting tasks (core, memory)
- Instead, submit jobs

Warning

Login Hosts on Different Machines/partitions

- Windows: [PuTTY](#) or [MobaXterm](#) or NX
Linux/Mac: [terminal](#) or NX
- To login, you need an active VSC number and a hostname
\$> ssh -X vscXXXXX@<hostname>

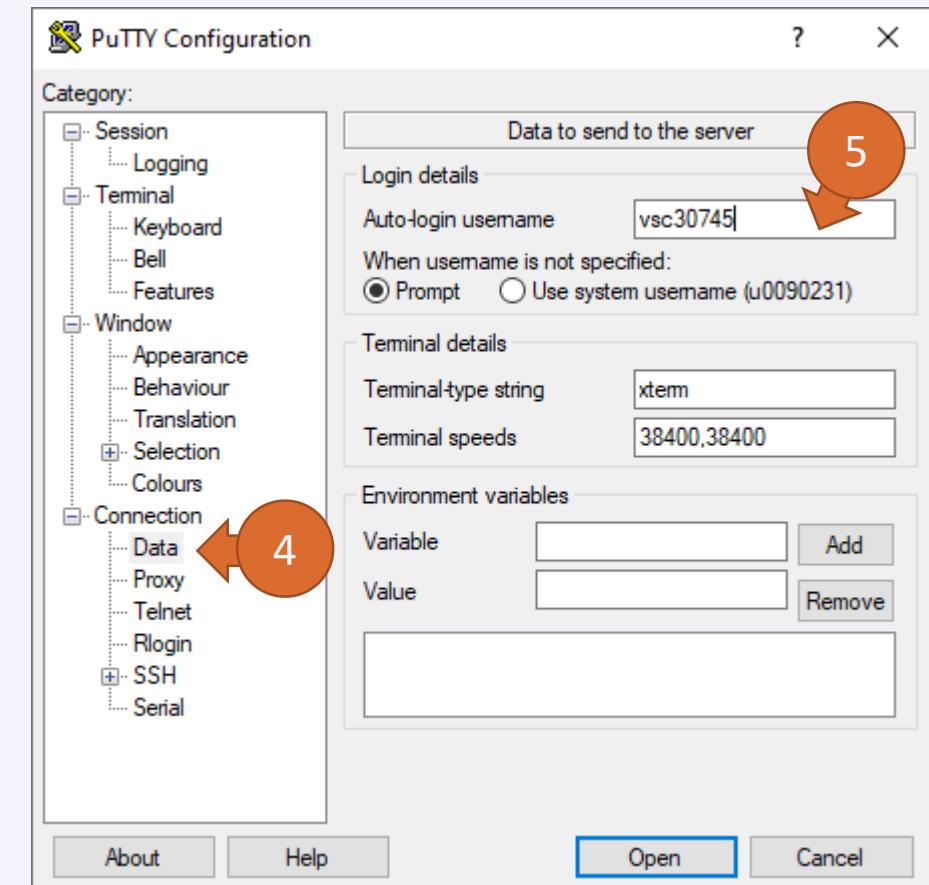
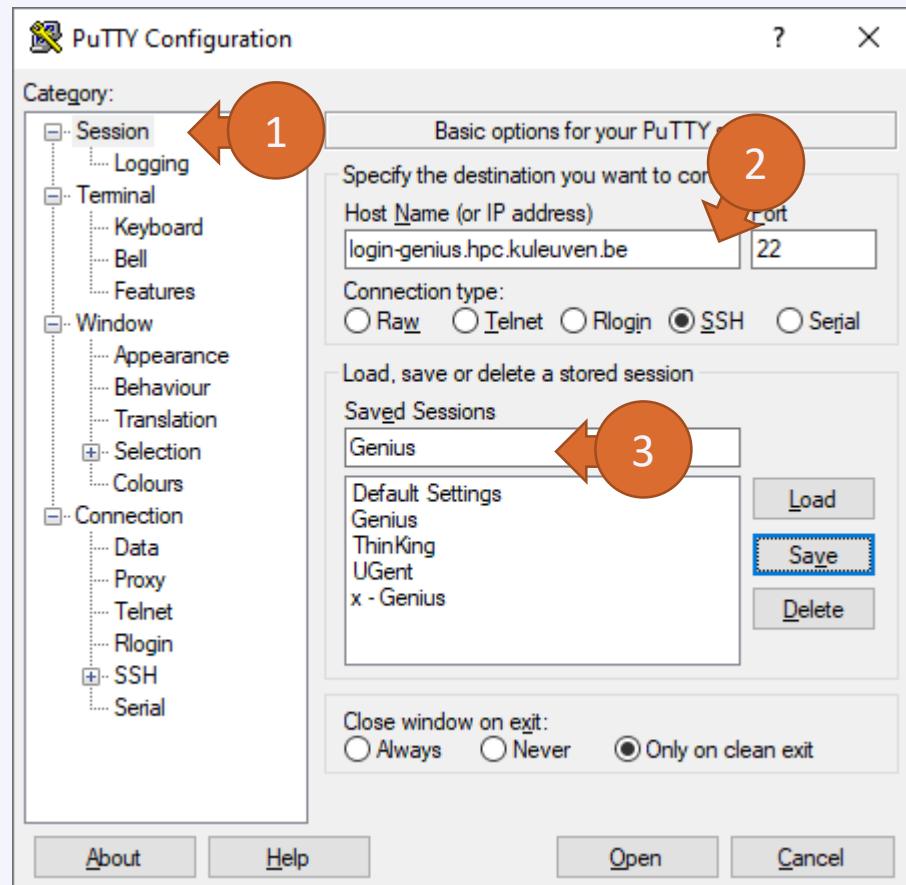
Cluster / Partition	<hostname>	Remark(s)
ThinKing	login-thinking.hpc.kuleuven.be	Recommended
	login7-tier2.hpc.kuleuven.be	Haswell partition
	login8-tier2.hpc.kuleuven.be	
Genius	login.hpc.kuleuven.be	Recommended
	login-genius.hpc.kuleuven.be	
	login{1,2}-tier2.hpc.kuleuven.be	No GPU
	login{3,4}-tier2.hpc.kuleuven.be	Nvidia Quadro P6000
Superdome	Any Genius login node (above)	-q qsuperdome -l partition=superdome
AMD	Any Genius login node (above)	-l partition=amd

Login Setup

- PuTTY
- Terminal

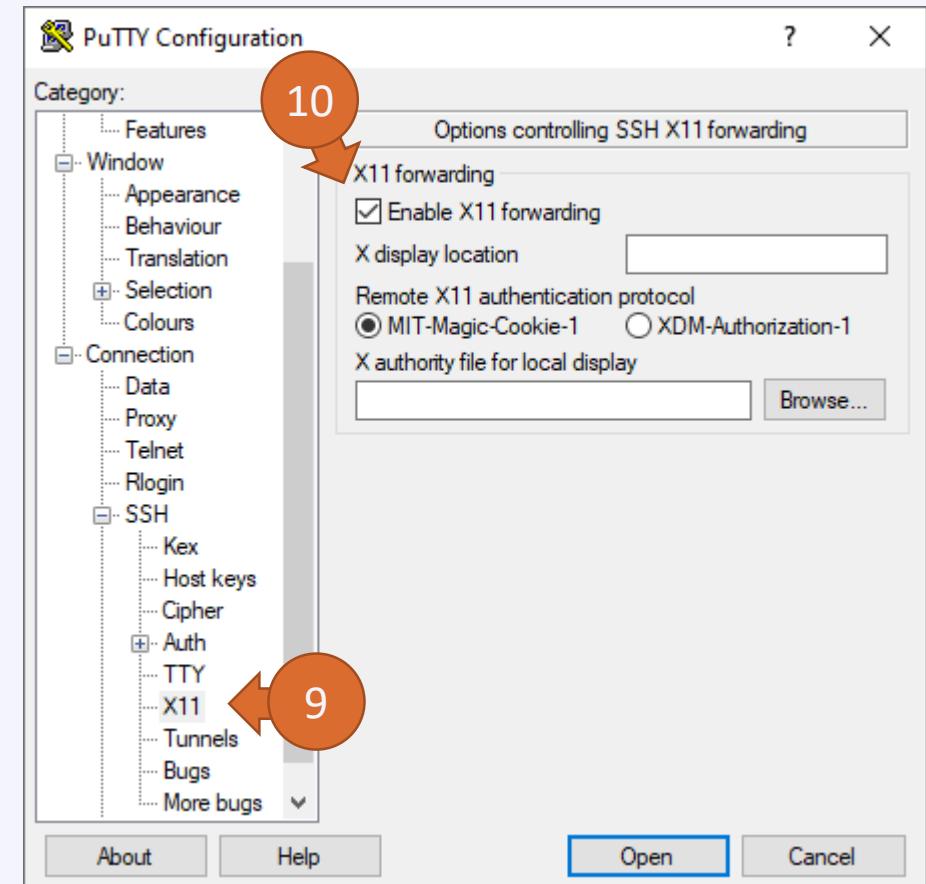
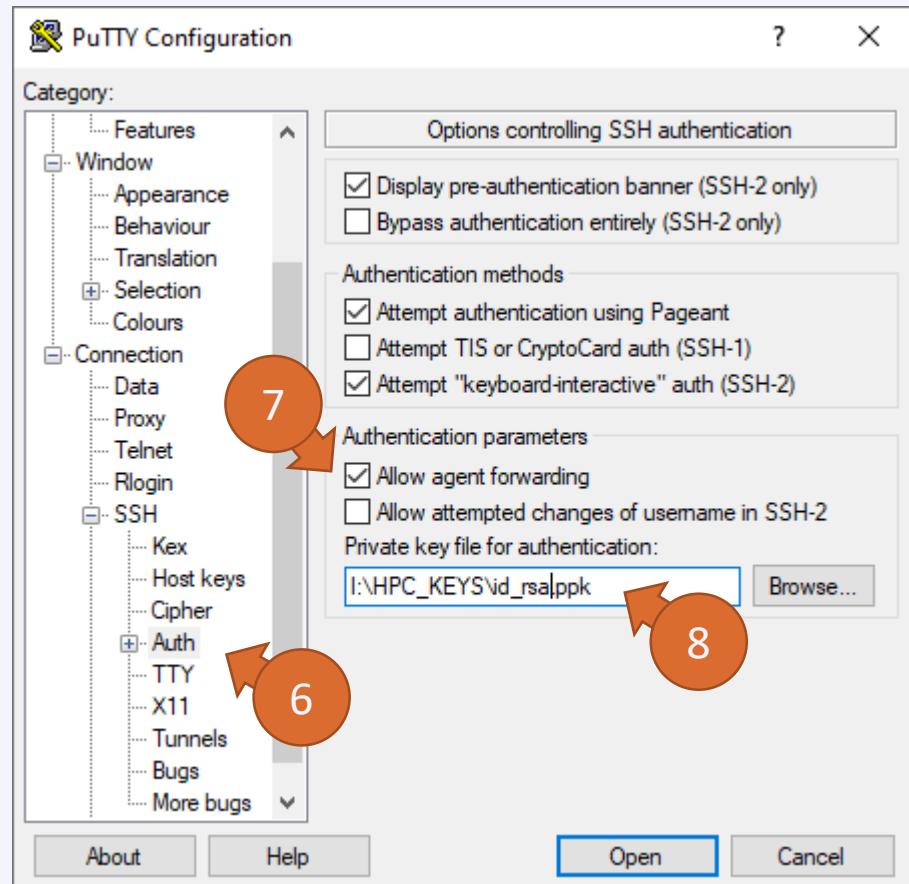
Setup PuTTY in 12 Clicks

(Windows only)



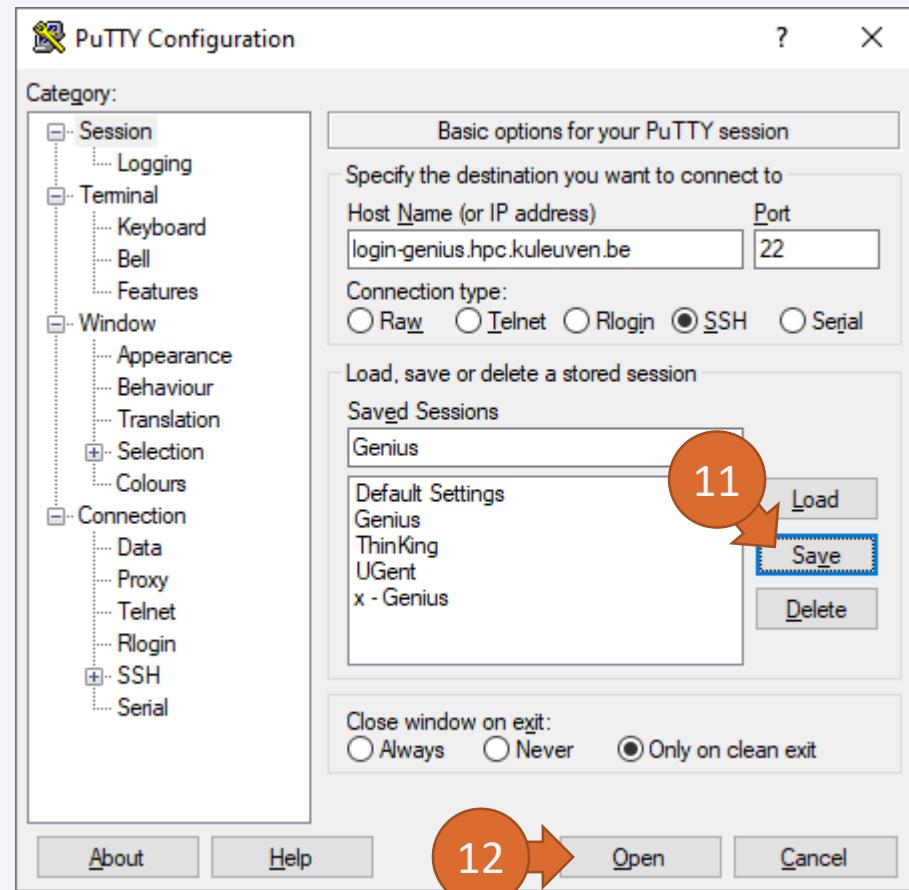
Setup PuTTY in 12 Clicks

(Windows only)



Setup PuTTY in 12 Clicks

(Windows only)



If PuTTY asks for **password**, exit immediately, and check the path to your private key. Else you will be blocked for 24h.

A screenshot of a PuTTY terminal window titled 'login4-tier2.hpc.kuleuven.be - PuTTY'. The window displays the following text:

```
Using username "vsc30745".
Authenticating with public key "rsa-key-20170705"
Passphrase for key "rsa-key-20170705":
```

A screenshot of a PuTTY terminal window titled 'login4-tier2.hpc.kuleuven.be - PuTTY'. The window displays the following text:

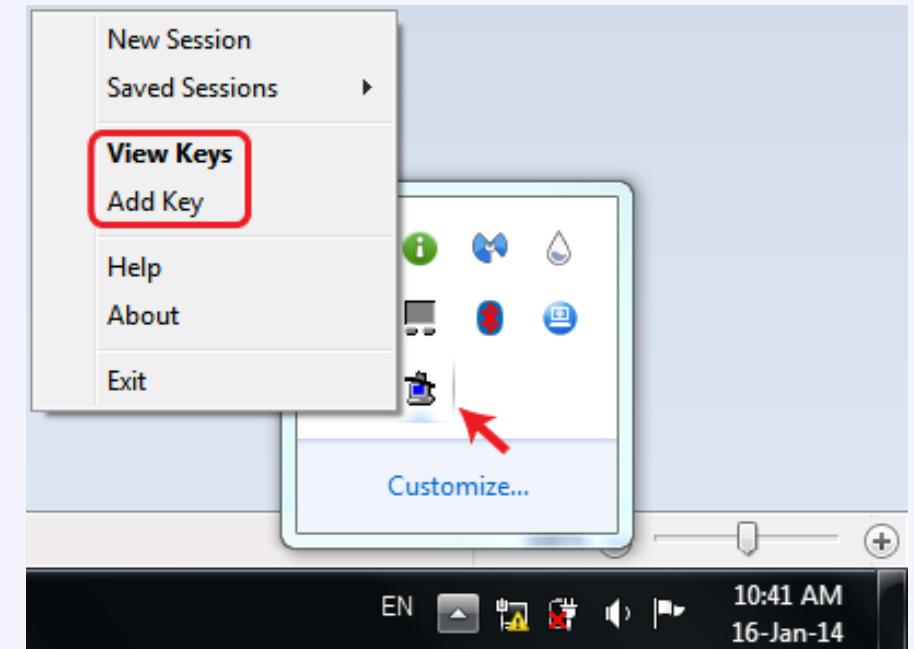
```
Using username "vsc30745".
Authenticating with public key "rsa-key-20170705"
Passphrase for key "rsa-key-20170705":
Last login: Wed Jan  9 15:23:52 2019
[REDACTED]
Deze server wordt (deels) met Puppet beheerd; alle wi
jzigingen die niet via
Puppet worden aangebracht, riskeren (automatisch) ong
edaan te worden gemaakt.

This server is managed by Puppet; all changes not imp
```

Activate Your SSH Agent

(Windows only)

- When you install PuTTY, the agent called **Pageant** is automatically installed
- Open Pageant; it hides inside your bottom-right tray
- Click on “Add Key” and brows to your private key(s) folder
- After choosing a key, you are asked for a passphrase
- If successful, agent always remembers your key



Connecting via Terminal

(Linux and Mac)

- Check your SSH Agent. Is your SSH key found?

```
$> ssh-add -l
```

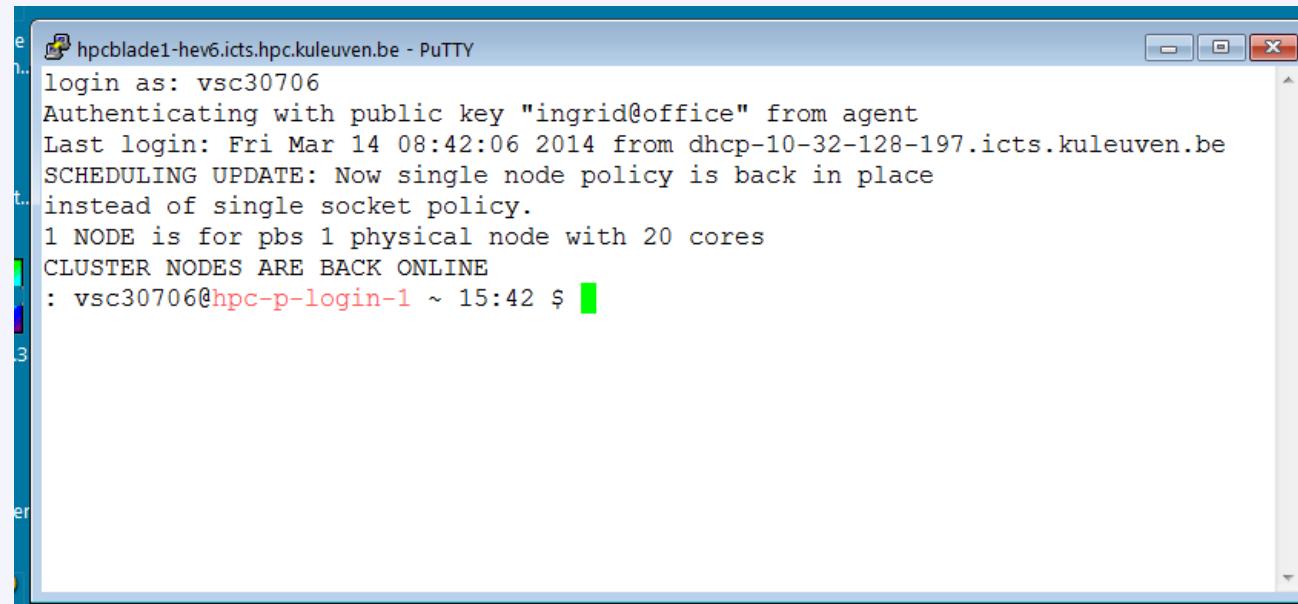
- If your key is not found, add it to the agent:

```
$> ssh-add </path/to/my/private/keyfile>
```

- Try to ssh now

```
$> ssh vscXXXXX@<hostname>
```

If asked for **password**,
please stop
connecting and
contact support,
otherwise after a few
attempts you will be
blocked for 24h.





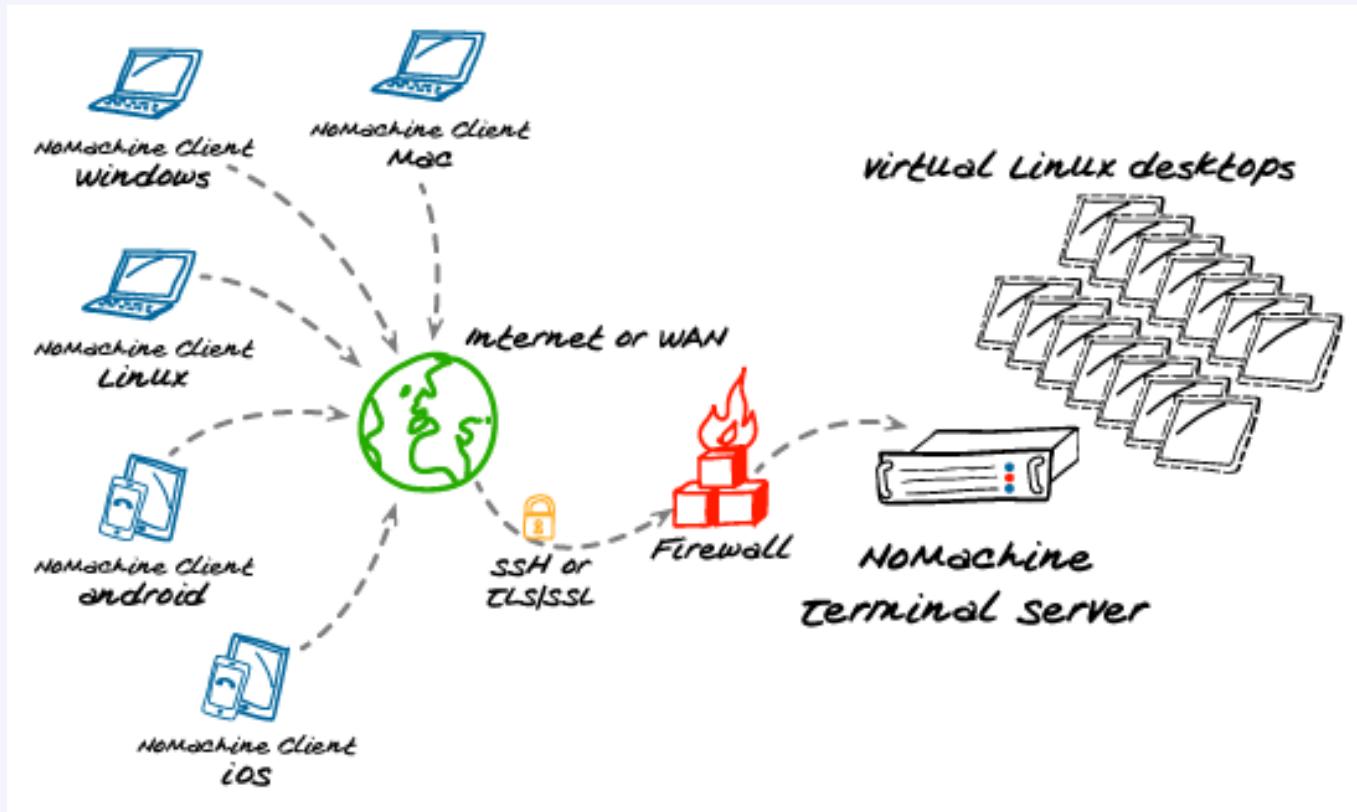
NX

- About
- Setup

NX – The Graphical Login

NX is:

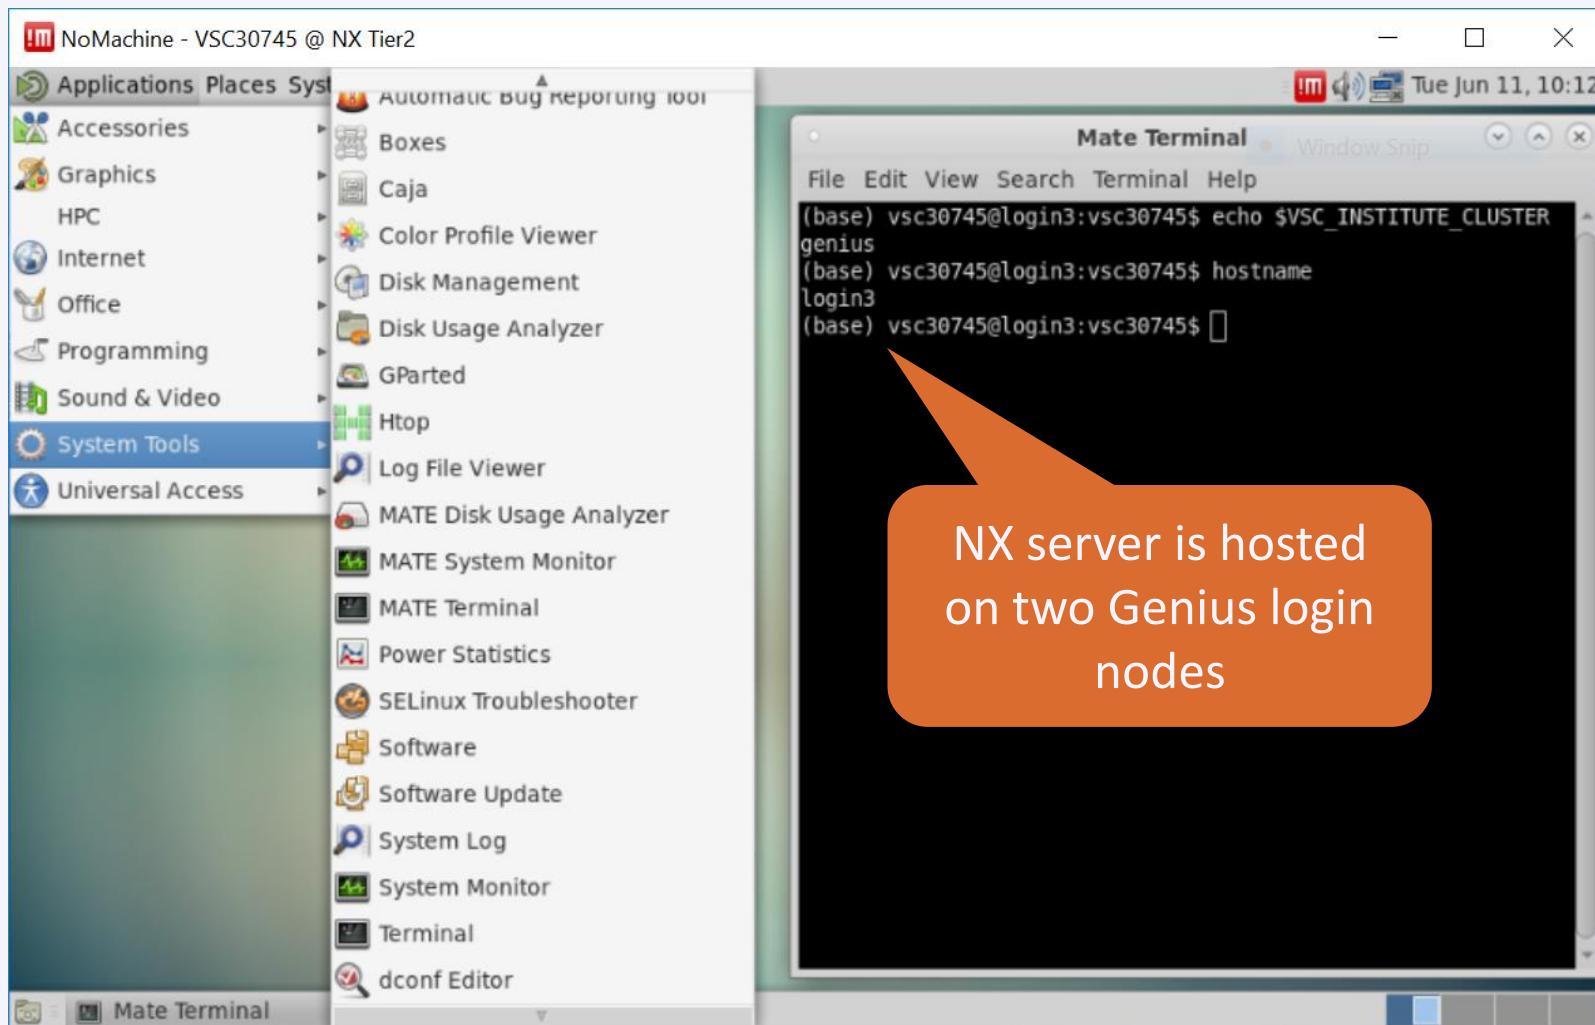
- a graphical remote desktop server
- login to VSC
- start a virtual Linux instance(s)
- Client: [NoMachine](#)
- [Configuration Guide](#)
- Using OpenSSH key
[Convert your key](#)



Advantages of NX

- ❑ Graphical login
- ❑ A session stays alive/active even if you close your client (e.g. laptop)
- ❑ Resume a session, and immediately keep working where you left off
- ❑ More interactive jobs
- ❑ Available apps:
 - Internet browser, terminal, ...
 - Text editor, PDF reader, Image viewer, ...
 - Matlab, RStudio, SAS
- ❑ **Remark:**
 - NX is shared among tens of users
 - Do not compute/test your code there**
 - Instead, submit jobs from NX to compute nodes

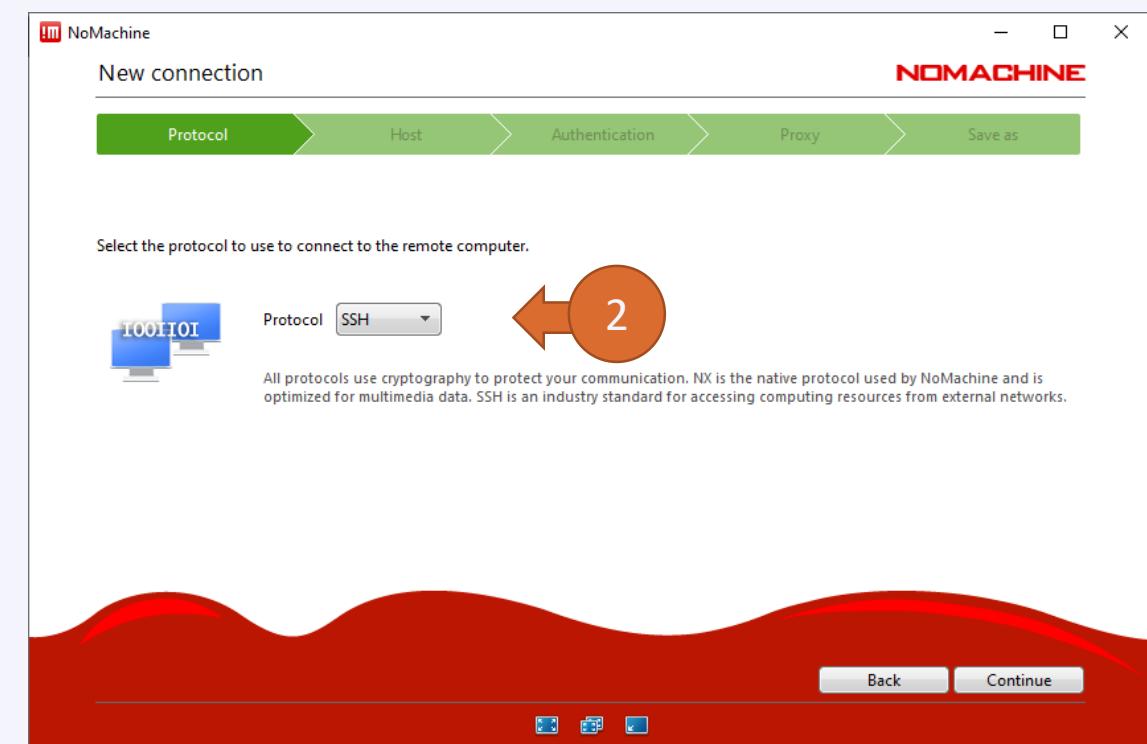
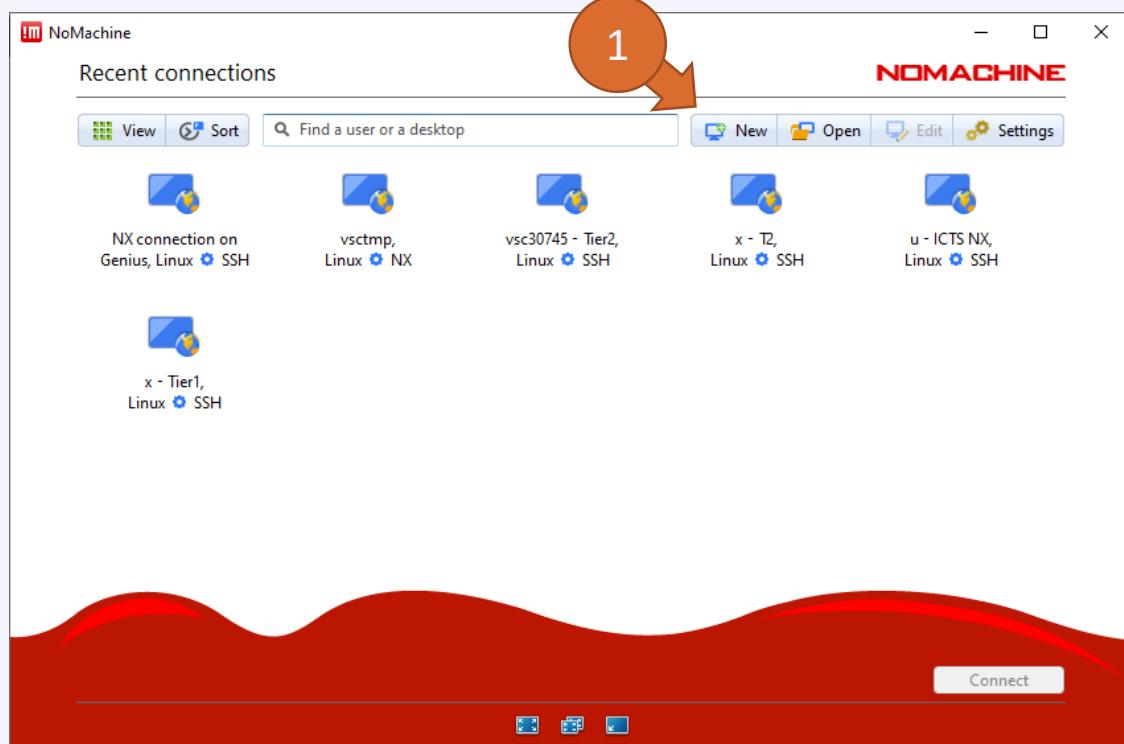
NX virtual desktop



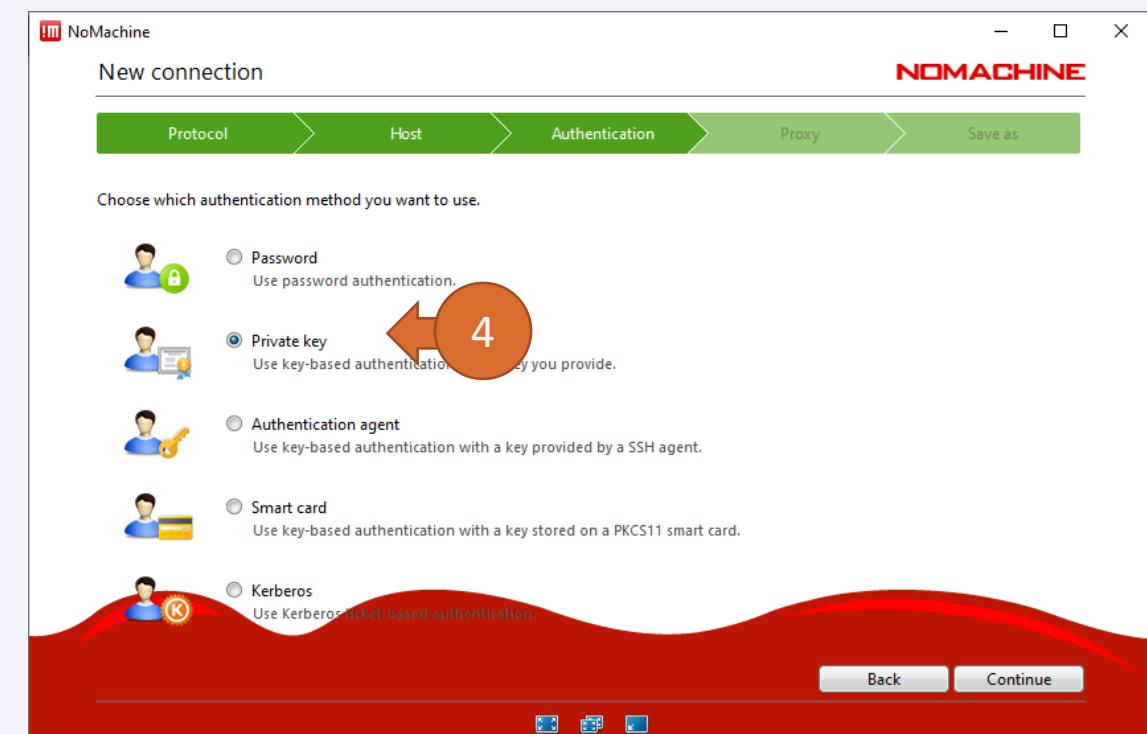
NX: available software

- Accessories:** Gedit, Vi IMproved, Emacs (dummy version), Calculator
- Graphics:** gThumb (picture viewer), Xpdf Viewer
- Internet:** Firefox
- HPC: Computation:** Matlab (2018a), RStudio, SAS
- Visualisation:** Paraview, VisIt, VMD
- Programming:** Meld Diff Viewer (visual diff and merge tool)
- System tools:** File Browser, Terminal
- Additionally:** Gnuplot (graphing utility), Filezilla (file transfer tool), Evince (PDF, PostScript, TIFF, XPS, DVI Viewer)
- Software launched through modules from Terminal.

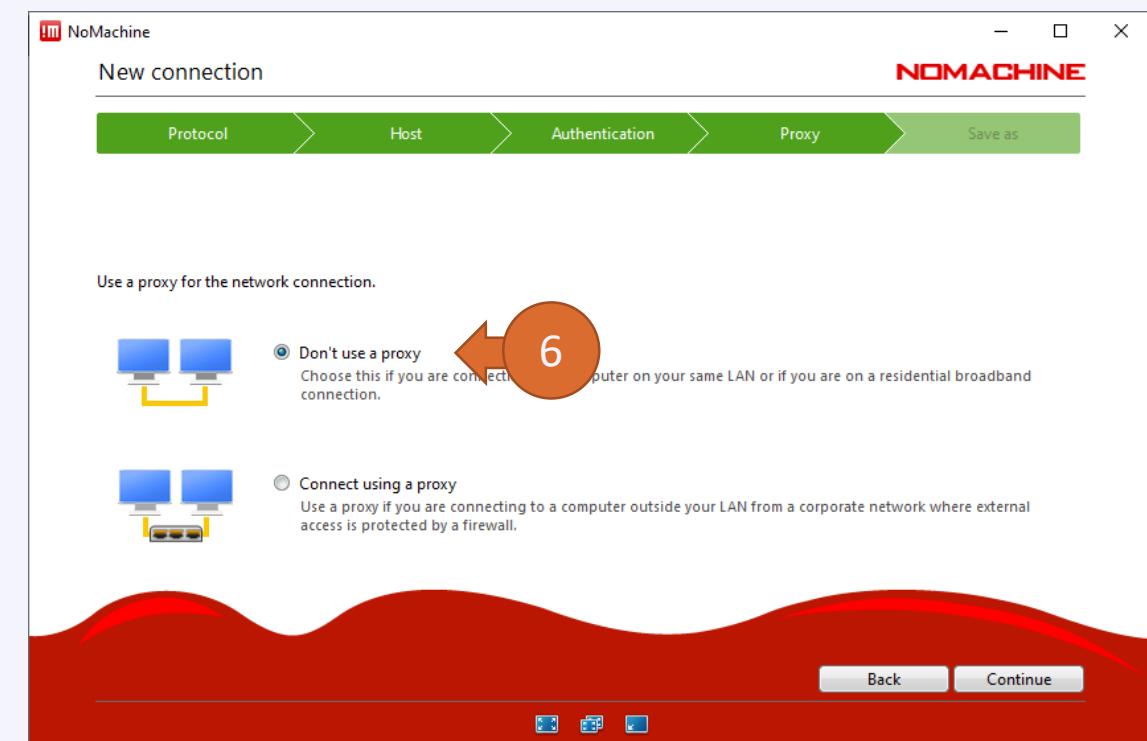
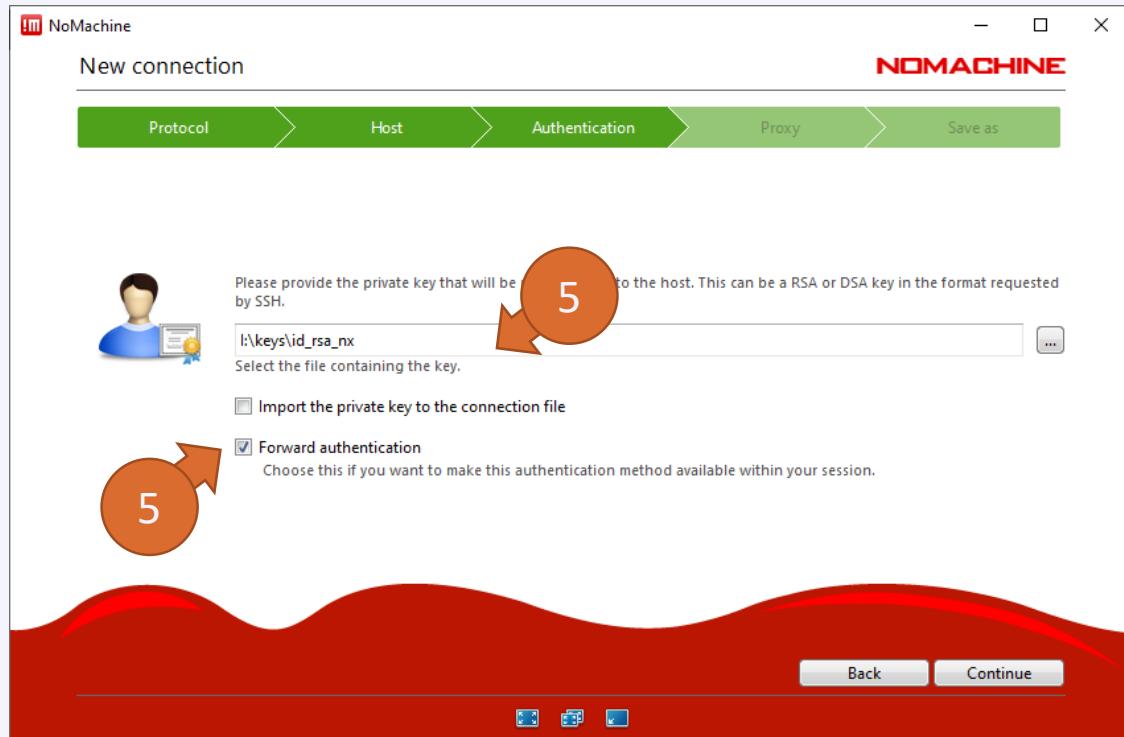
Setup NX in 10 Steps



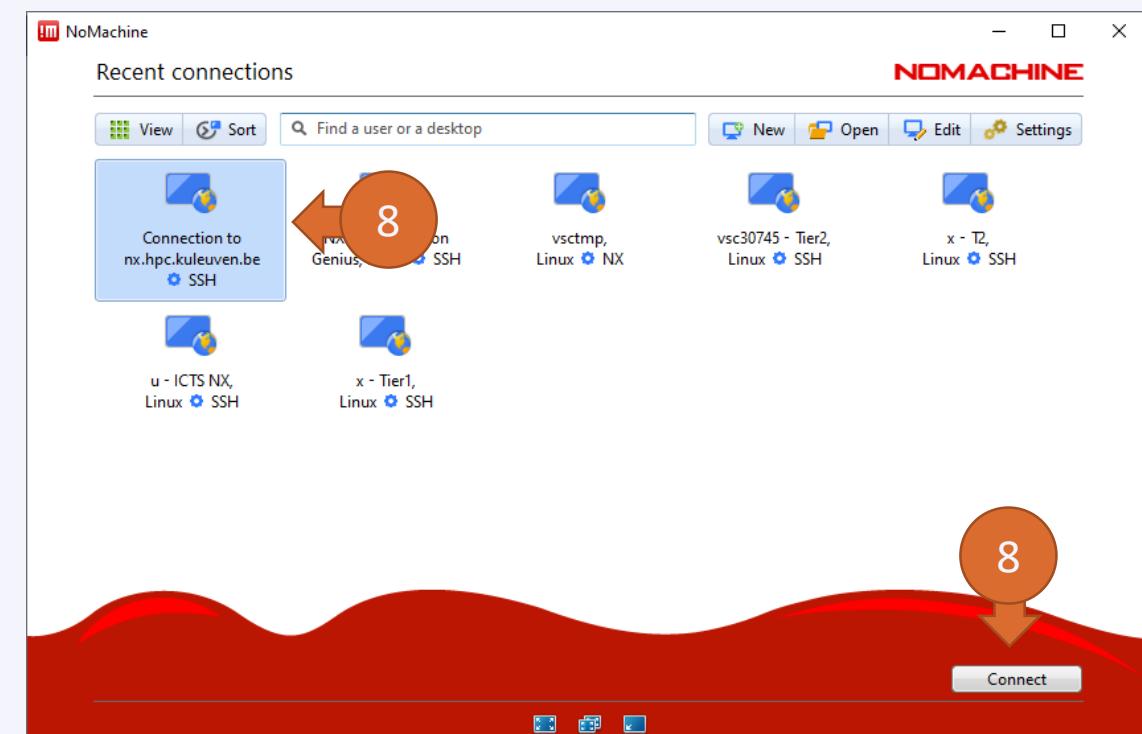
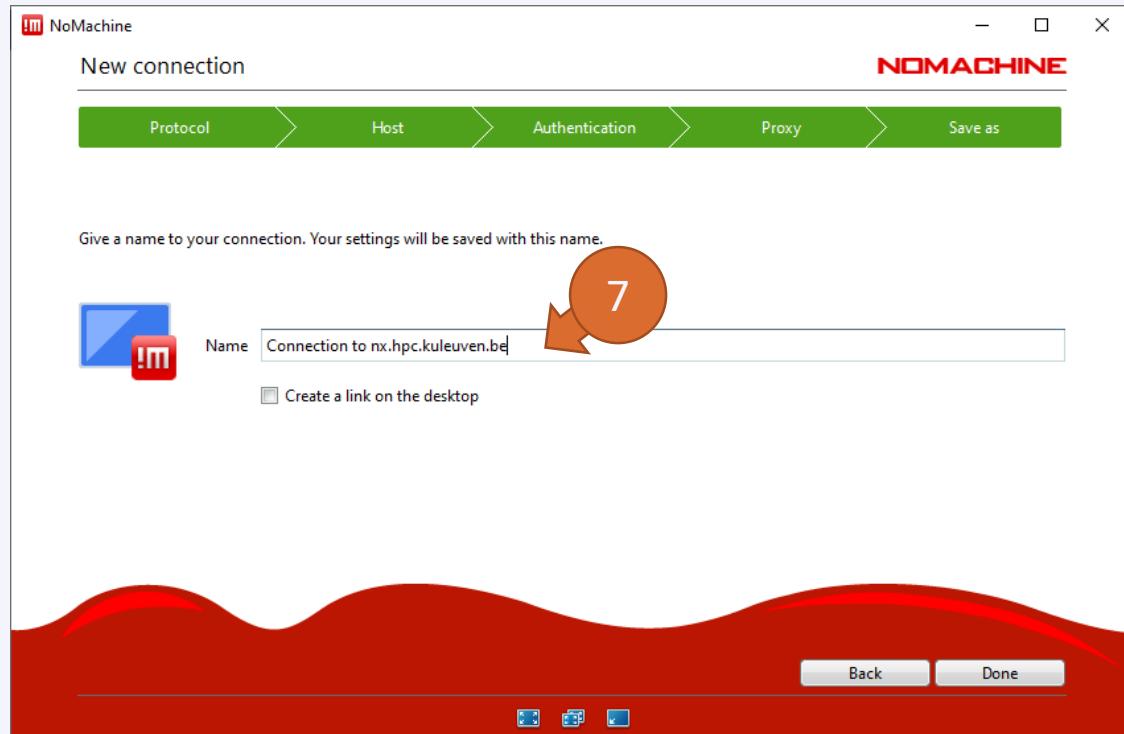
Setup NX in 10 Steps



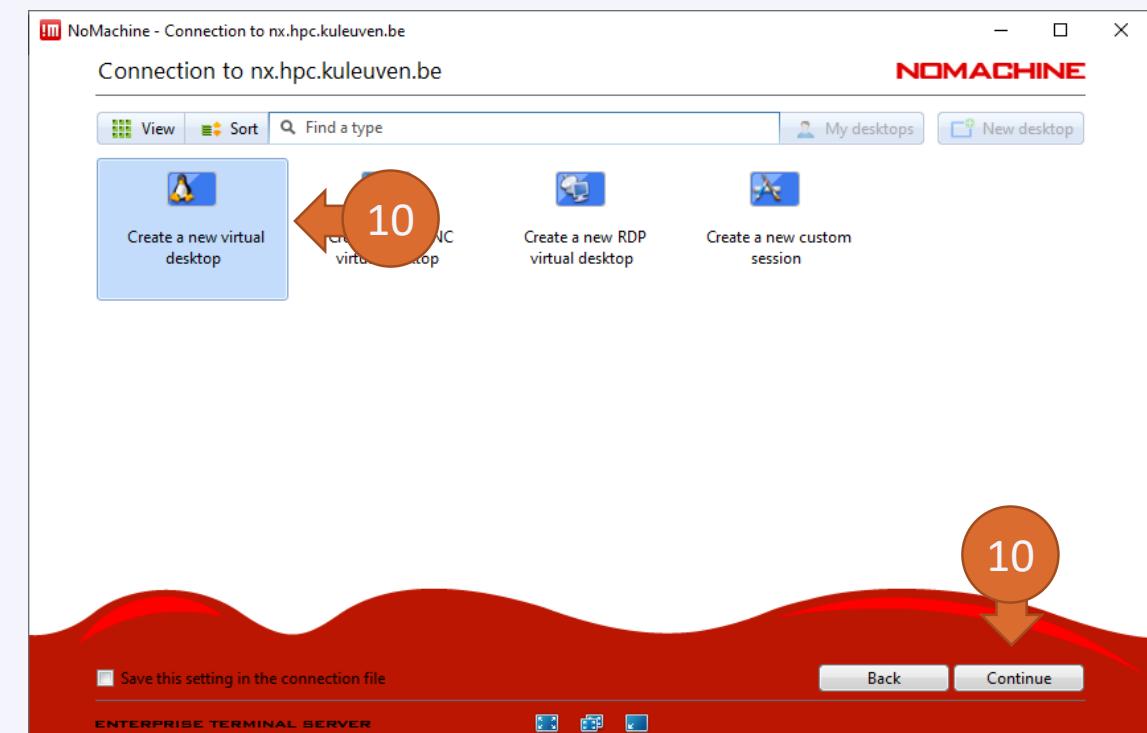
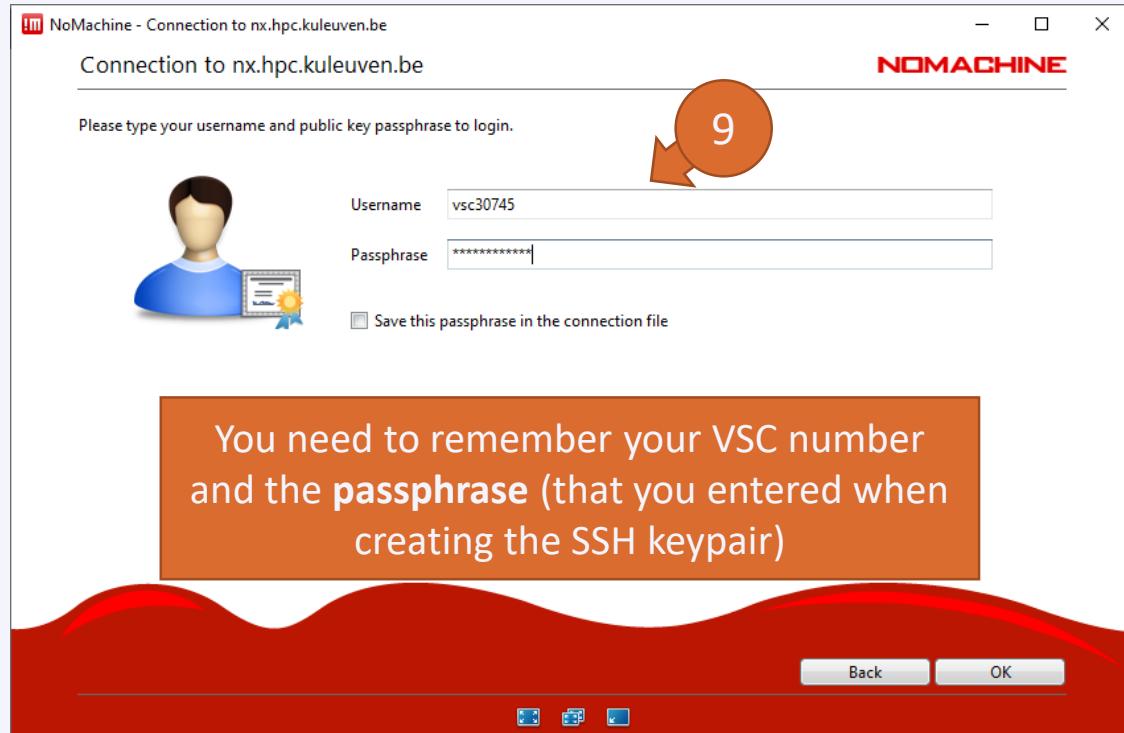
Setup NX in 10 Steps

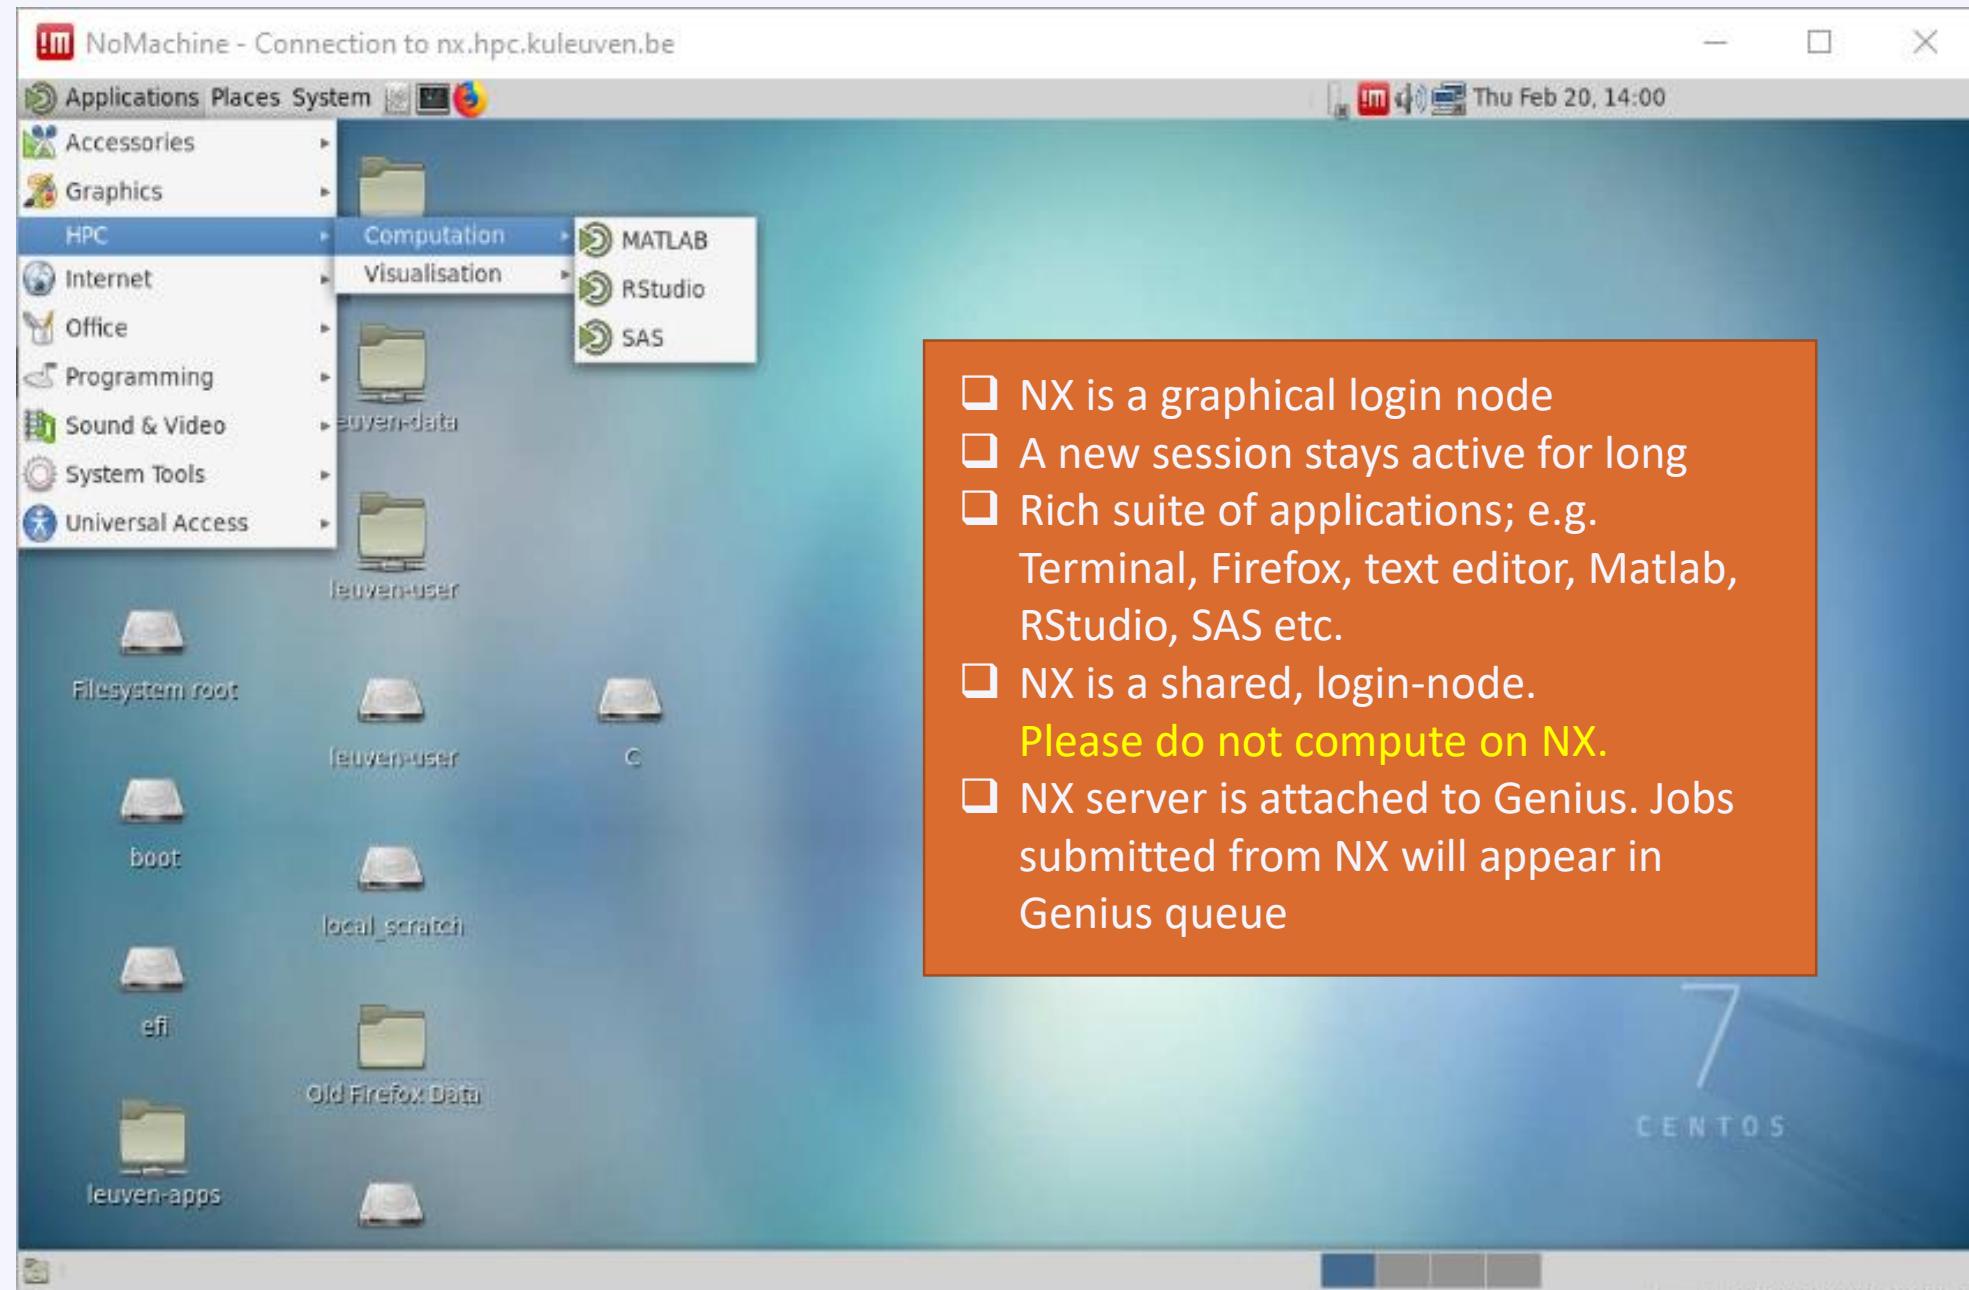


Setup NX in 10 Steps



Setup NX in 10 Steps

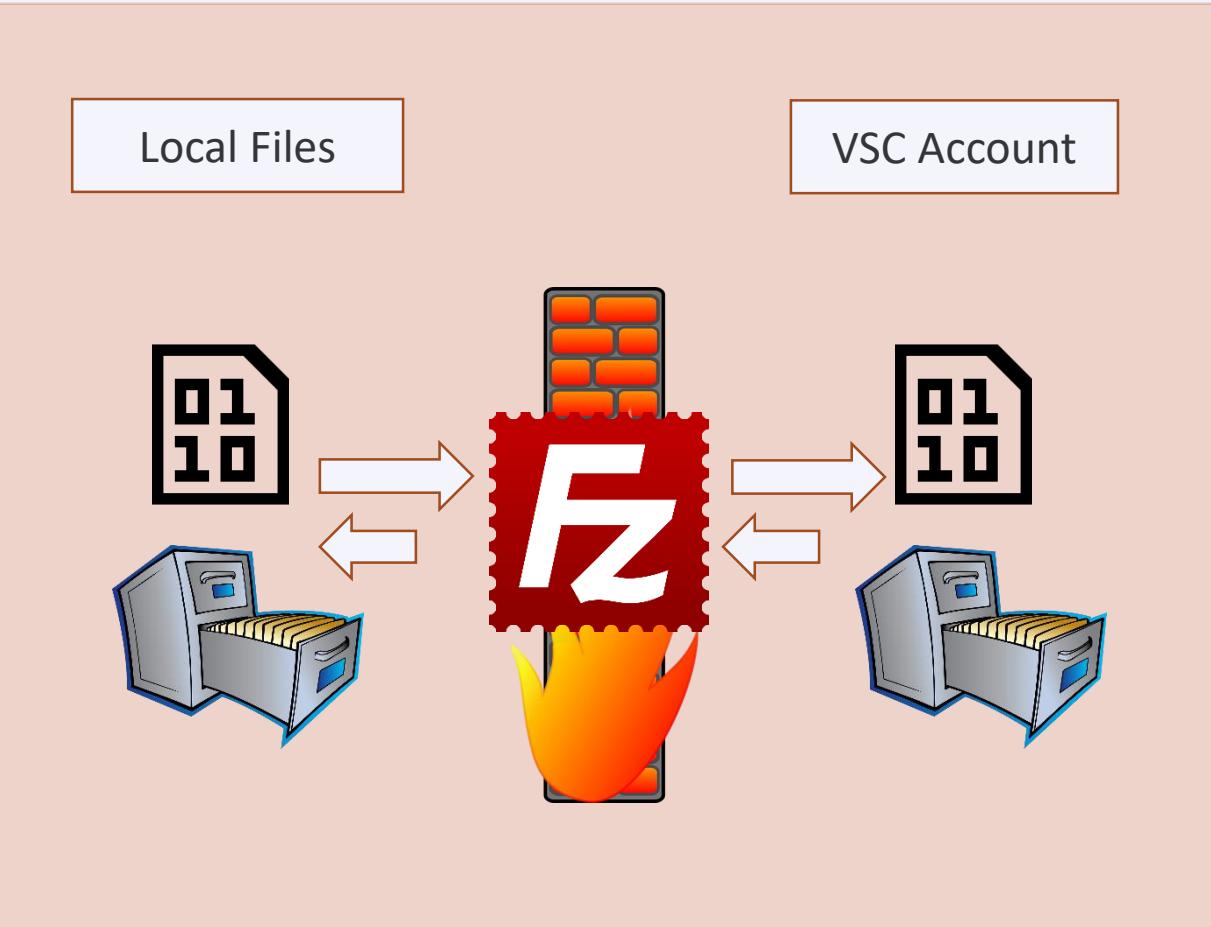




- NX is a graphical login node
- A new session stays active for long
- Rich suite of applications; e.g.
Terminal, Firefox, text editor, Matlab,
RStudio, SAS etc.
- NX is a shared, login-node.
Please do not compute on NX.
- NX server is attached to Genius. Jobs
submitted from NX will appear in
Genius queue

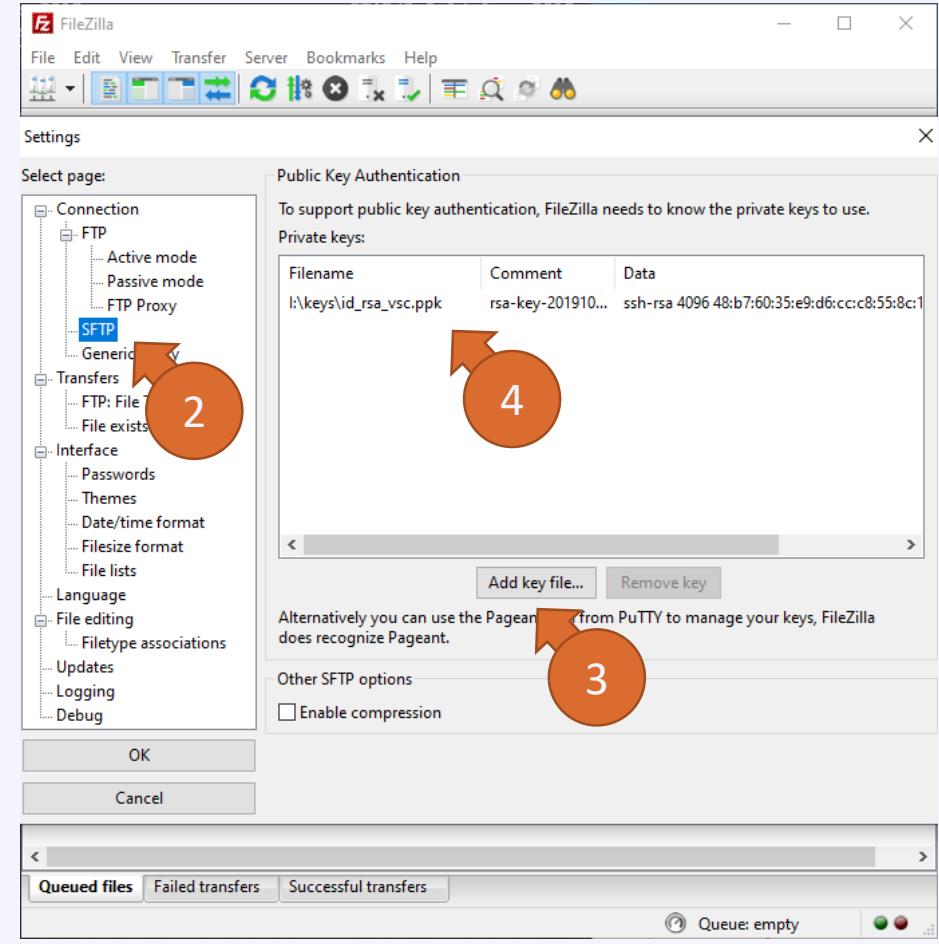
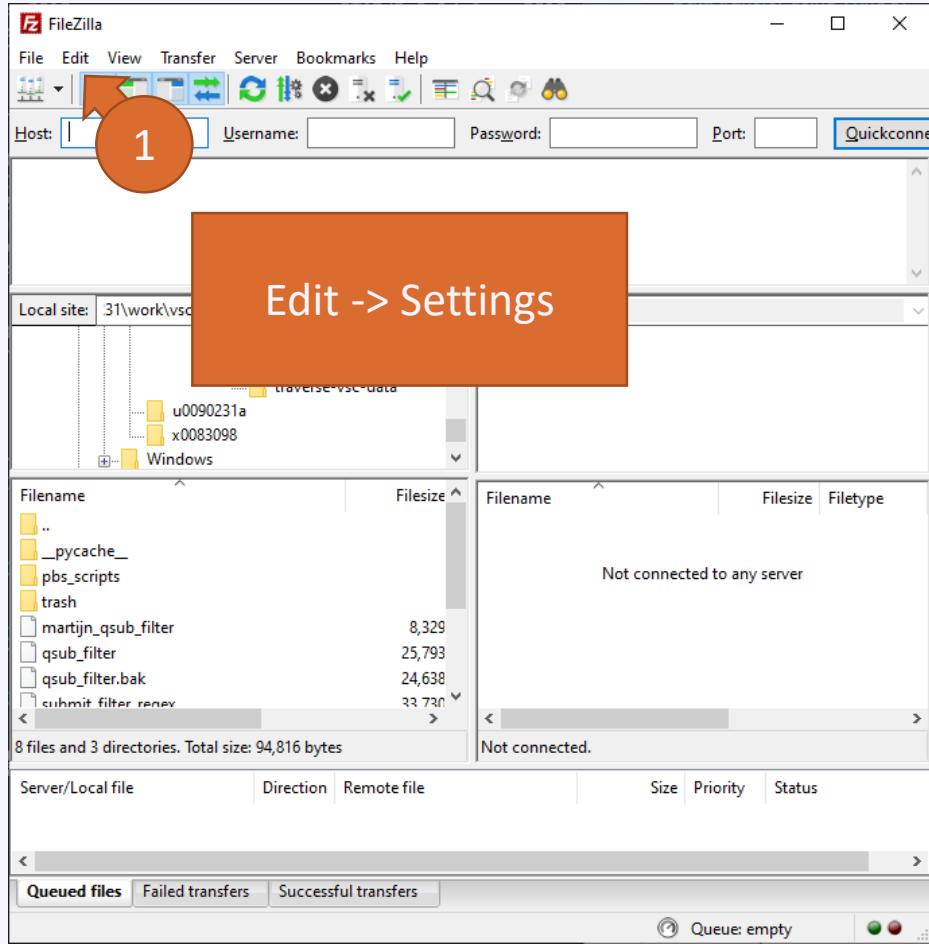
File Transfer with FileZilla

File transfer

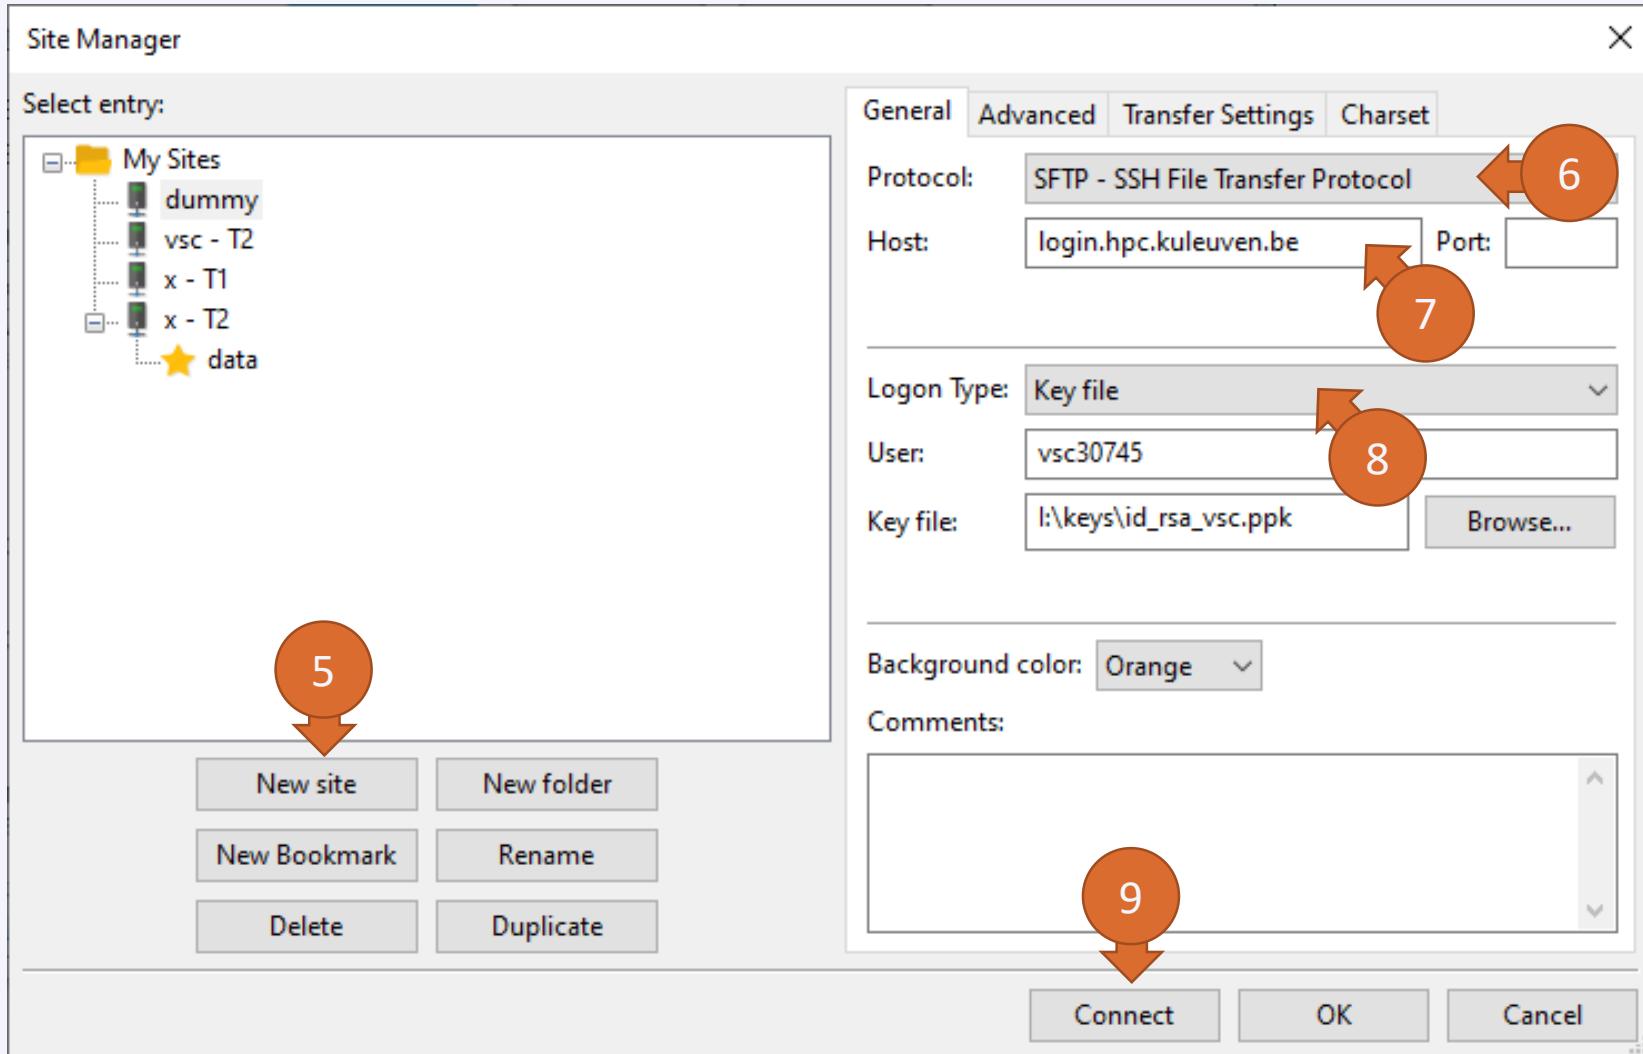


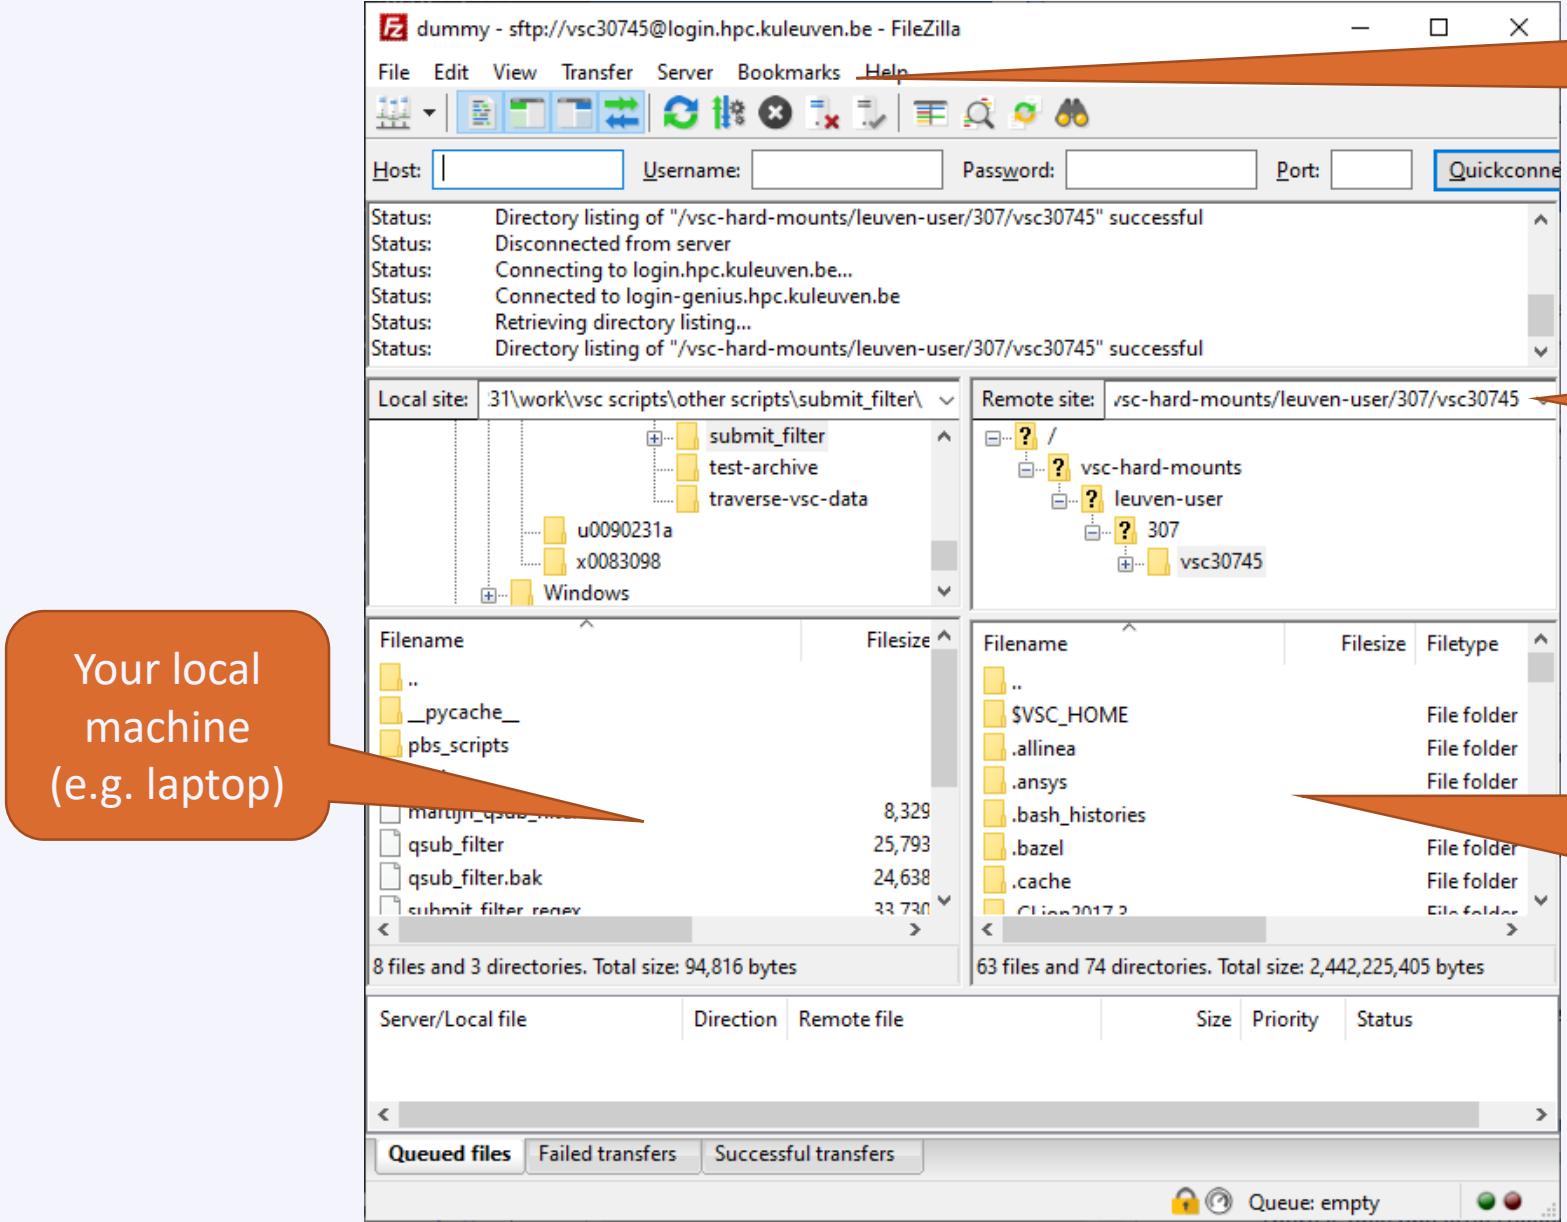
Application	OS
FileZilla	Windows, Linux, Mac
WinSCP	Windows
rsync, scp	Linux, Mac
Globus	Window, Linux, Mac

File transfer



File transfer





Your local
machine
(e.g. laptop)

For convenience, you'd
better bookmark your data
and scratch folders!

Instead, go to your
\$VSC_DATA folder, e.g.
/data/leuven/399/vsc39934

Your VSC storage (e.g.
\$VSC_DATA,
\$VSC_SCRATCH).
Note: by default you
arrive at your own
\$VSC_HOME. Copy
nothing there!

Software Stack

Software: Available Modules

- OS: Linux CentOS 7.7, Kernel 3.10.0-1127.18.2.el7.x86_64
- Compilers:
Intel 2018a (icc, icpc, ifort; Intel MPI; Intel MKL)
FOSS 2018a (gcc, g++, gfortran; OpenMPI; ScaLAPACK, OpenBLAS, FFTW)
- Note: **Never mix FOSS and Intel compilers (gives dependency conflict)**

Command	Remark
module av	List all installed modules
module av Python	List all Python-related modules
module load Python/3.6.4-intel-2018a	Load a specific module
module list	List all loaded modules and their dependencies
module unload Python/3.6.4-intel-2018a	Unload a module (but dependencies still stay)
module purge	Remove all modules from your work session

Software: Your Specific Needs

- You can always install your desired software in your \$VSC_DATA
Use Intel or FOSS toolchains
- If you cannot, ask us for help
- Python/R packages for AI and ML must be installed by the users themselves
- Install miniconda in your \$VSC_DATA

```
$> wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh
$> bash Miniconda3-latest-Linux-x86_64.sh -p $VSC_DATA/miniconda3
```
- Create conda environment per every project

```
$> conda create -n EnvTF tensorflow-gpu=2.0.0 pytorch
```
- Use your own conda/Python in your jobs/scripts

```
$> conda activate EnvTF
```



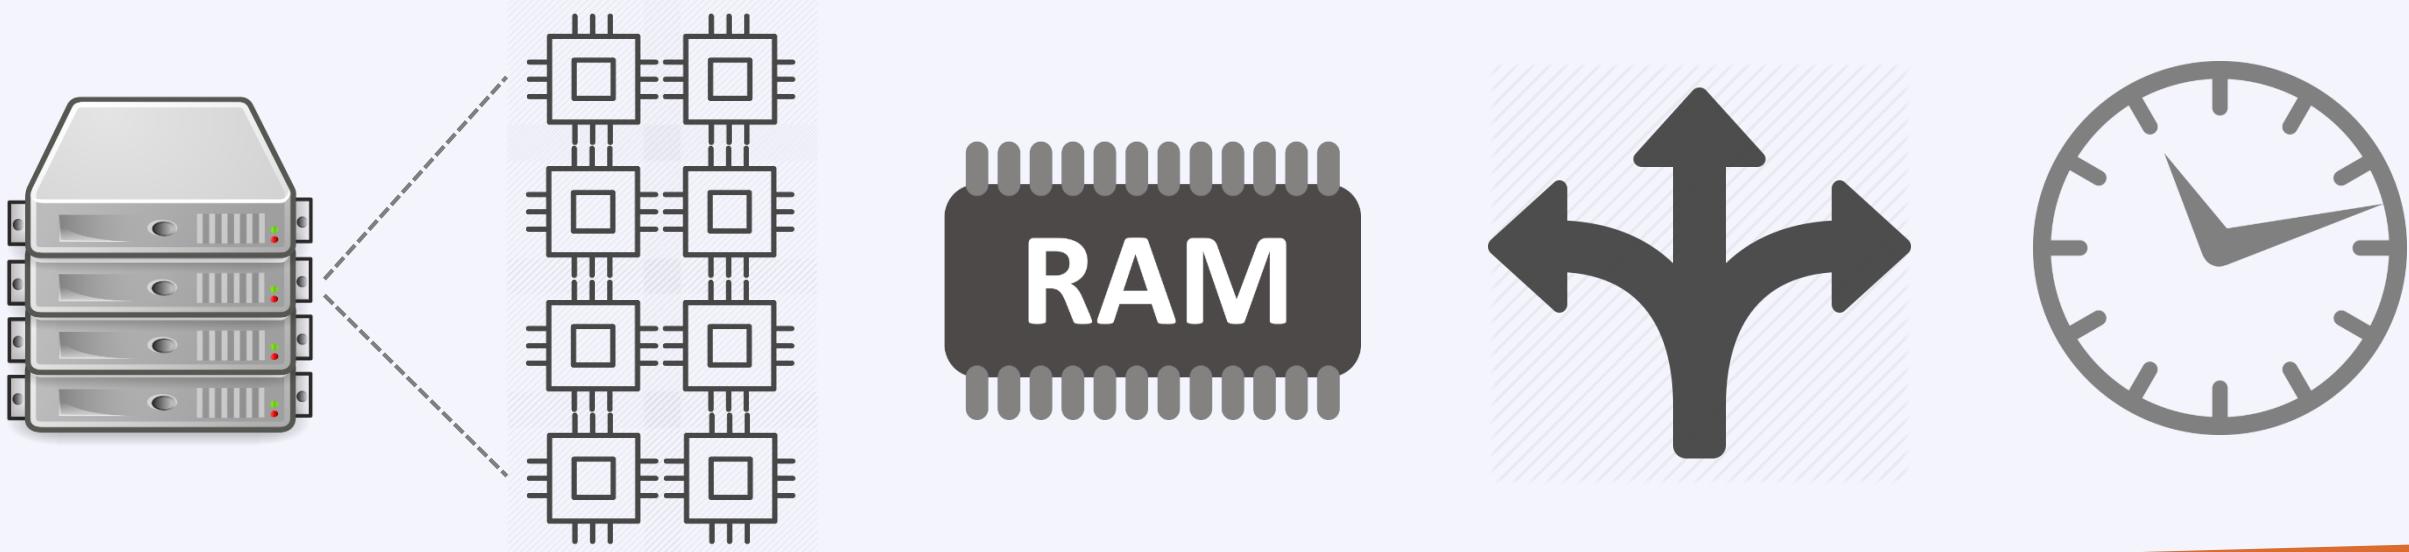
Starting to Compute

Extra services

- Request introduction credits
<https://admin.kuleuven.be/icts/onderzoek/hpc/request-introduction-credits>
- Request project credits (create your group first)
<https://admin.kuleuven.be/icts/onderzoek/hpc/request-project-credits>
https://icts.kuleuven.be/sc/forms/Aanvraagformulier_HPC_Credits
- Extra project credits (to add to existing project)
<https://admin.kuleuven.be/icts/onderzoek/hpc/extra-project-credits>
- Request extra storage
- <https://admin.kuleuven.be/icts/onderzoek/hpc/hpc-storage>

Resource Glossary

- Nodes:** how many compute servers to request?
- Cores:** how many cores per node to use?
- Memory requirement:** how much memory each core needs?
- Partition:** GPU, BigMem, Superdome, AMD
- Walltime:** how long to use resources?



Backfill

Nodes

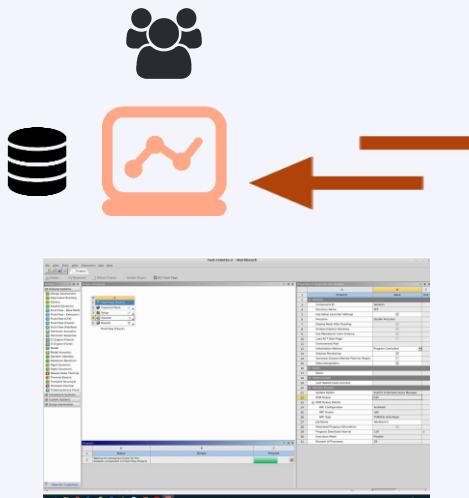


Job queue

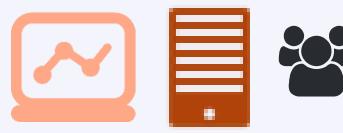


Ways to use the cluster

INTERACTIVE GUI WORKLOAD



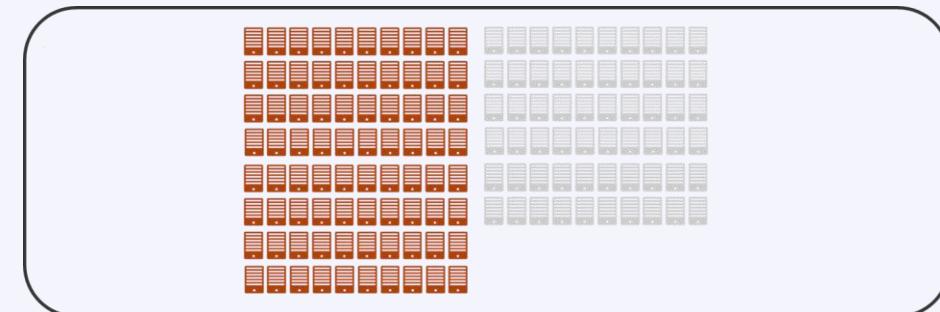
BATCH WORKLOAD



INTERACTIVE WORKLOAD



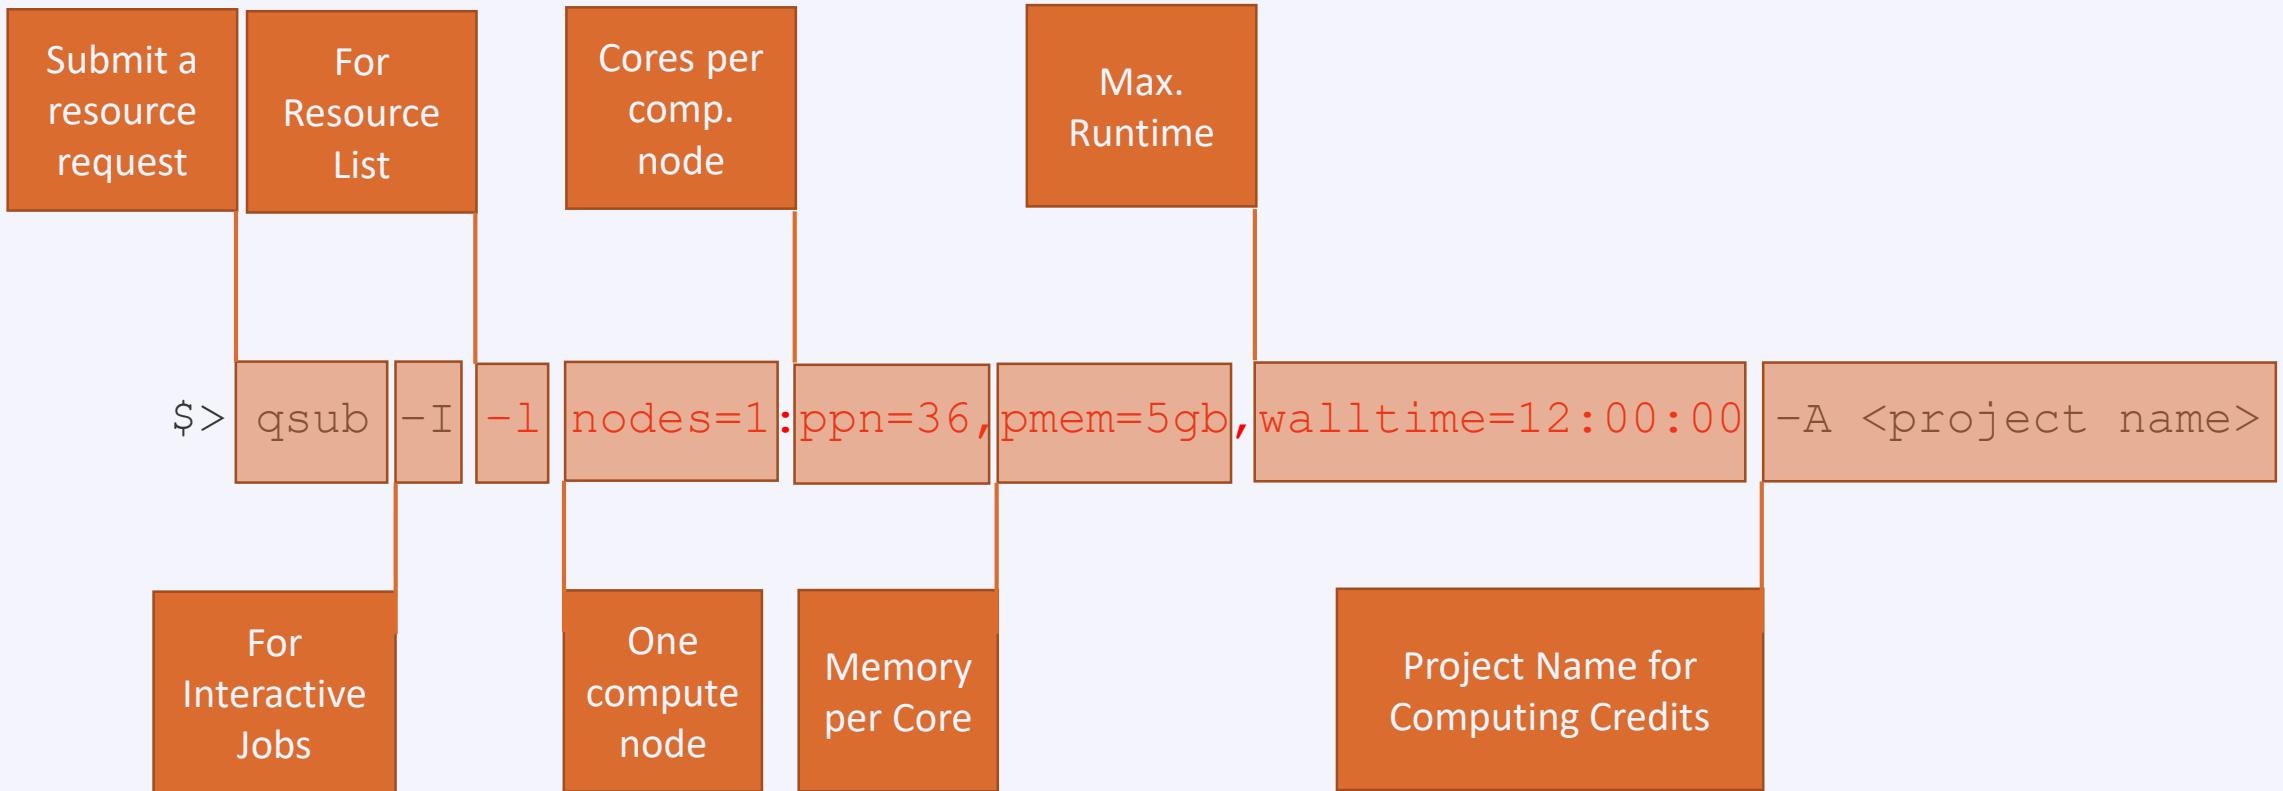
Storage



Types of Jobs

- 1) **Batch jobs** are by far the most common and allow for the most efficient use of the infrastructure. Essentially, a batch job is a bash shell script that is executed on a compute node, and that can spawn a parallel computation on many nodes. These jobs are placed in the queue, and the user can forget about it until it is finished.
- 2) **Interactive jobs** are intended to work on one or more compute nodes interactively. This can be useful in the context of software development for debugging and profiling applications, or for interactive calculations or visualizations. Basically, one gets a shell on one of the compute nodes.
- 3) To get GPU(s), you can use **JupyterHub** with your internet browser

Requesting Compute Resources Interactively



Remarks

- Specifying **<project name>** for credits is mandatory, e.g.: **-A lp_hpcinfo_training**
- If you request 2000 introductory credits, you use them like: **-A default_project**
- Implicit defaults are **nodes=1:ppn=1,pmem=5gb,walltime=1:00:00**

Interactive Jobs

- Interactive job: 1 core for 1 hour (default)

```
$> qsub -I -A lp_hpc
```

- Interactive job with X-forwarding

```
$> qsub -I -X -A lp_hpc
```

- Ask 2 SkyLake nodes, 36 cores each, 4 GB RAM per core for 12 hours

```
$> qsub -I -l nodes=2:ppn=36:skylake,pmem=4gb,walltime=12:00:00 -A lp_hpc
```

- Ask 2 CascadeLake nodes, 36 cores each, 4GB RAM per core for 30 min

```
$> qsub -I -l nodes=2:ppn=36:cascadelake,pmem=4gb,walltime=30:00 -A lp_hpc
```

- Request fraction of a node

```
$> qsub -I -l nodes=1:ppn=8:cascadelake -A lp_hpc
```

Batch Workload

Job Script

- Create a new text (ASCII) file
- Give it a good name!
E.g. simulation.pbs
- The first line is the shebang!
#!/bin/bash
- Put all commands line by line, e.g.
echo "Hello World!"
- Submit the job to the batch server
- Receive a unique JobID
- Error and output files

```
#!/bin/bash  
echo "Hello World!"
```

Command Line

```
$ qsub simulation.pbs -A default_project
```

JobID

50041238.tier2-p-
moab2.tier2.hpc.kuleuven.be

stderr, stdout

```
$ ls simulation*  
simulation.pbs  
simulation.pbs.e50041238  
simulation.pbs.o50041238
```

Standard Error File

- Always created
- <JobScript>.e<JobID>
- Contains all errors and warnings
- If empty: everything went well
- Always study it
- Address all warnings and errors (if you can)
- Typical error examples ...

stderr

```
$ ls *.e*
simulation.pbs.e50041238
```

Job Crash

```
forrtl: error (78): process killed (SIGTERM)
Stack trace terminated abnormally.
```

Short Walltime

```
PBS: job killed: walltime 432031 exceeded
limit 432000
```

Low Disk Space

```
IOError: [Errno 122] Disk quota exceeded
```

Standard Output File

- Always created
- <JobScript>.o<JobID>
- Contains all standard output (instead of screen)
- Always study it

```
$ ls *.o*
simulation.pbs.o50041238
```

stdout

time: 3600
nodes: 1
procs: 1
account string: lpt2_sysadmin
queue: q1h
=====
Hello World!
=====
Date: Fri Mar 20 14:32:01 CET 2020
Allocated nodes:
r26i27n11
Job ID: 50240161.tier2-p-moab-2.tier2.hpc. ...
Resource List:
pmem=5gb,walltime=01:00:00,neednodes=1:ppn=36
Resources Used:
cput=00:00:01,vmem=0kb,walltime=00:00:04,mem=0kb,energy_used=0
Queue Name: q1h

time: 4
nodes: 1
procs: 1
account: lpt2_sysadmin

Output File

Summary

stdout

Resources

summary

Example: PBS Job Script

```
#!/bin/bash  
#PBS -l walltime=00:10:00  
#PBS -l nodes=1:ppn=36  
#PBS -l pmem=5gb  
#PBS -N testjob  
#PBS -A lp_hpcinfo_training  
#PBS -m abe  
#PBS -M my.name@kuleuven.be
```

Shebang

```
module load intel/2018a  
which icc
```

Resource List

Module load(s)

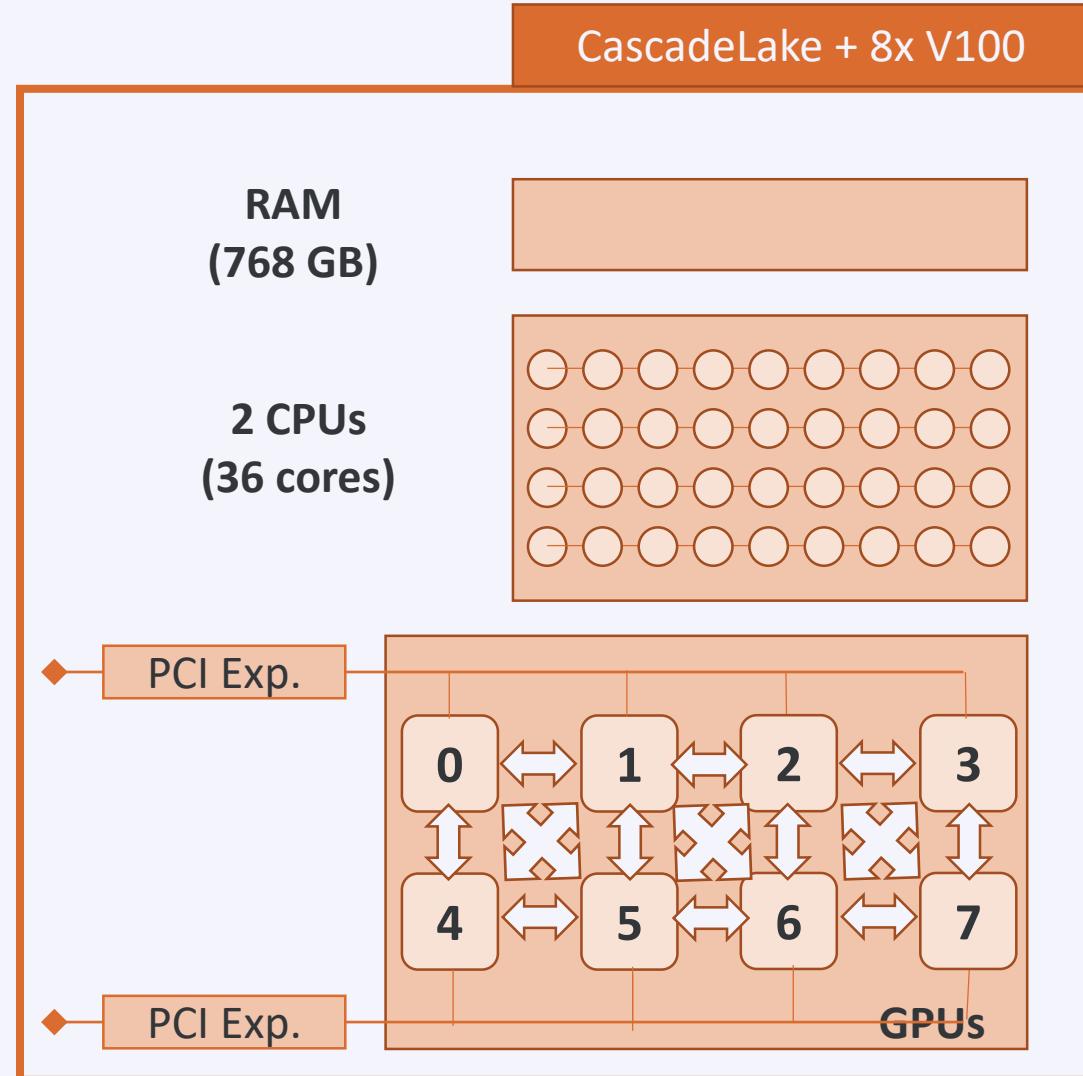
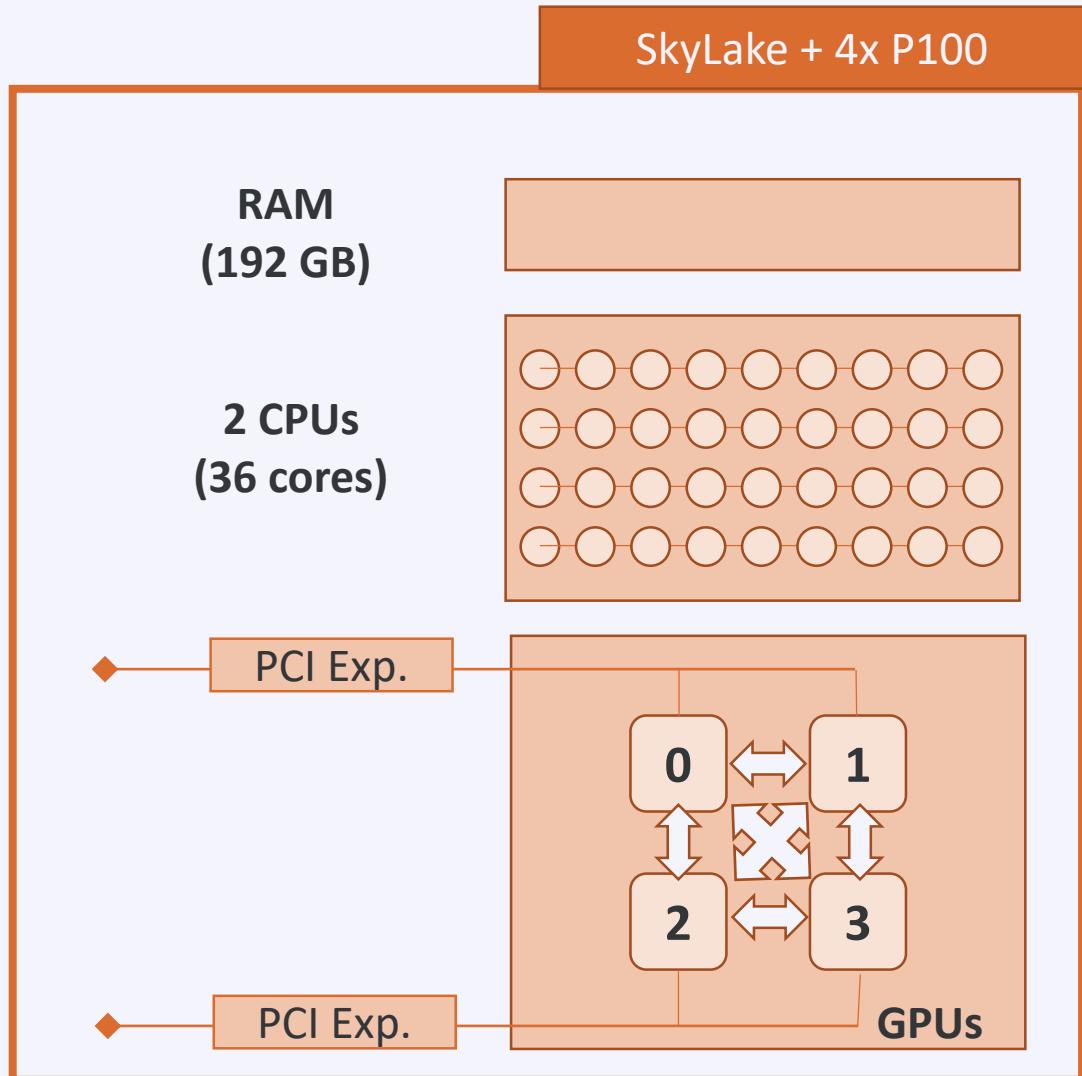
```
cd $PBS_O_WORKDIR  
cp /apps/leuven/training/HPC-intro/cpujob.pbs $VSC_SCRATCH  
touch output.log  
echo I am done
```

Execute commands

Move data

GPU Partition

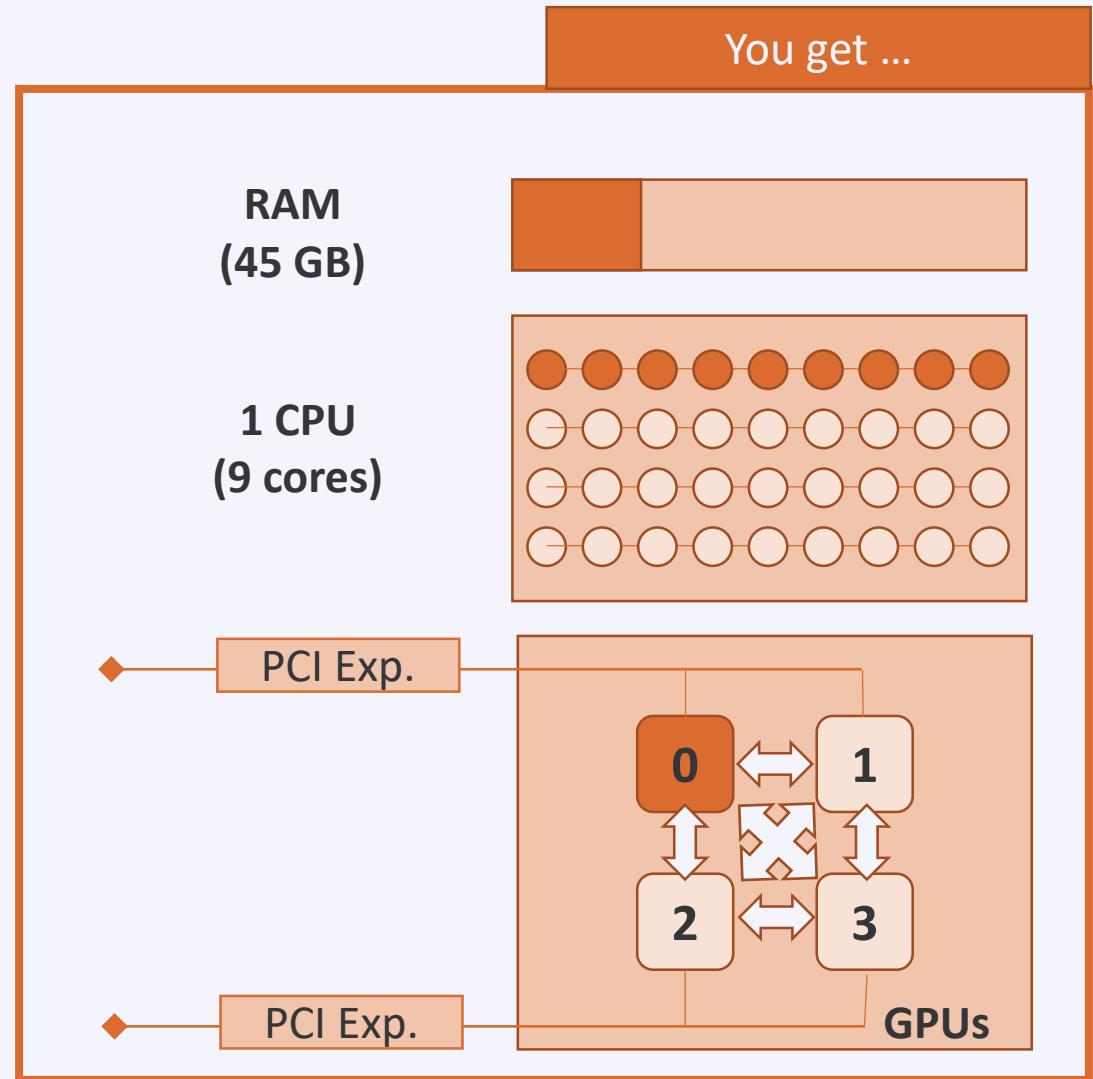
GPU Nodes



Using P100 GPU(s)

Using 1 GPU

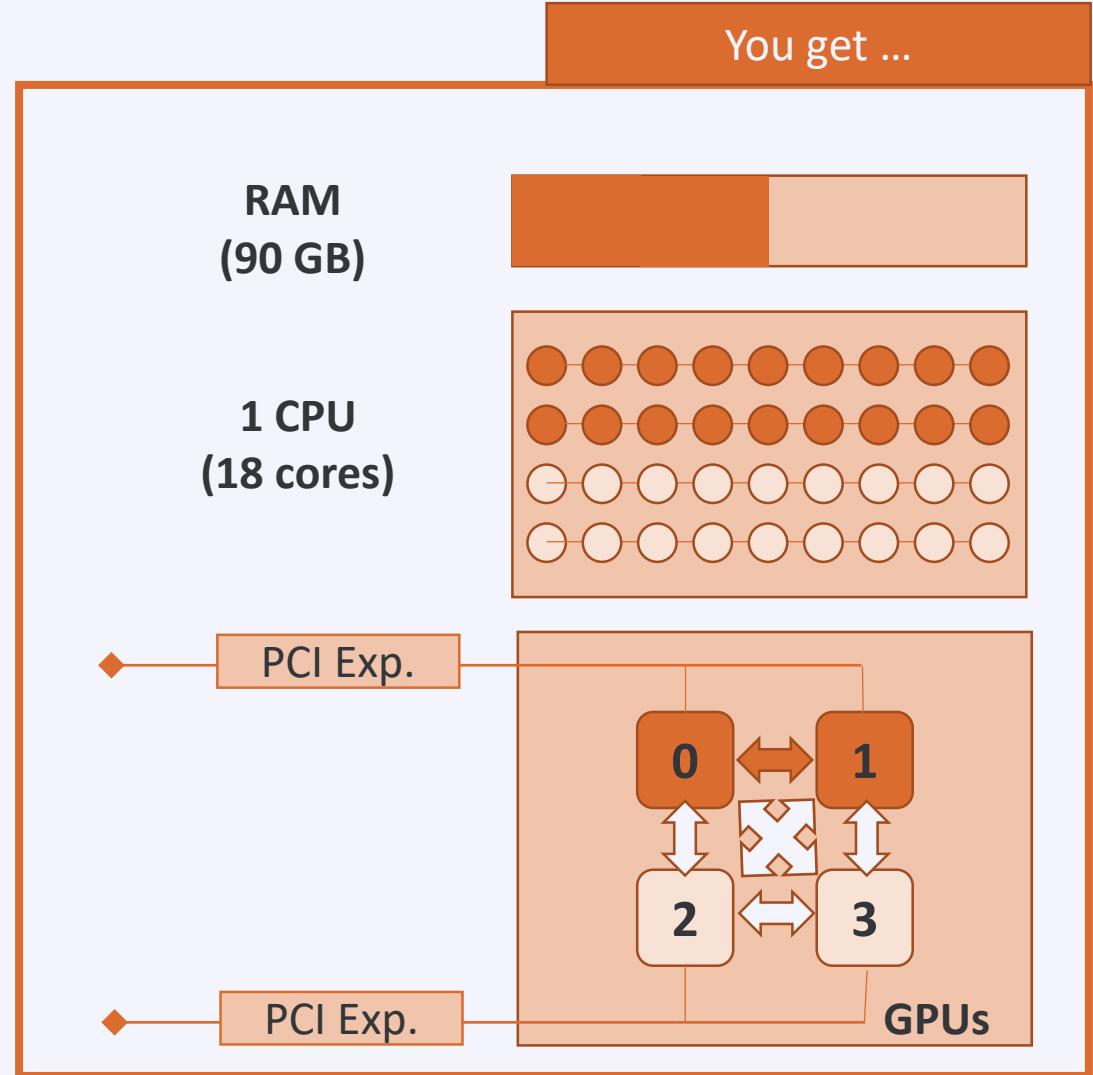
```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l nodes=1:ppn=9:gpus=1:skylake
#PBS -l pmem=5gb
#PBS -N testjob
#PBS -A default_project
```



Using P100 GPU(s)

```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l nodes=1:ppn=18:gpus=2:skylake
#PBS -l pmem=5gb
#PBS -N testjob
#PBS -A default project
```

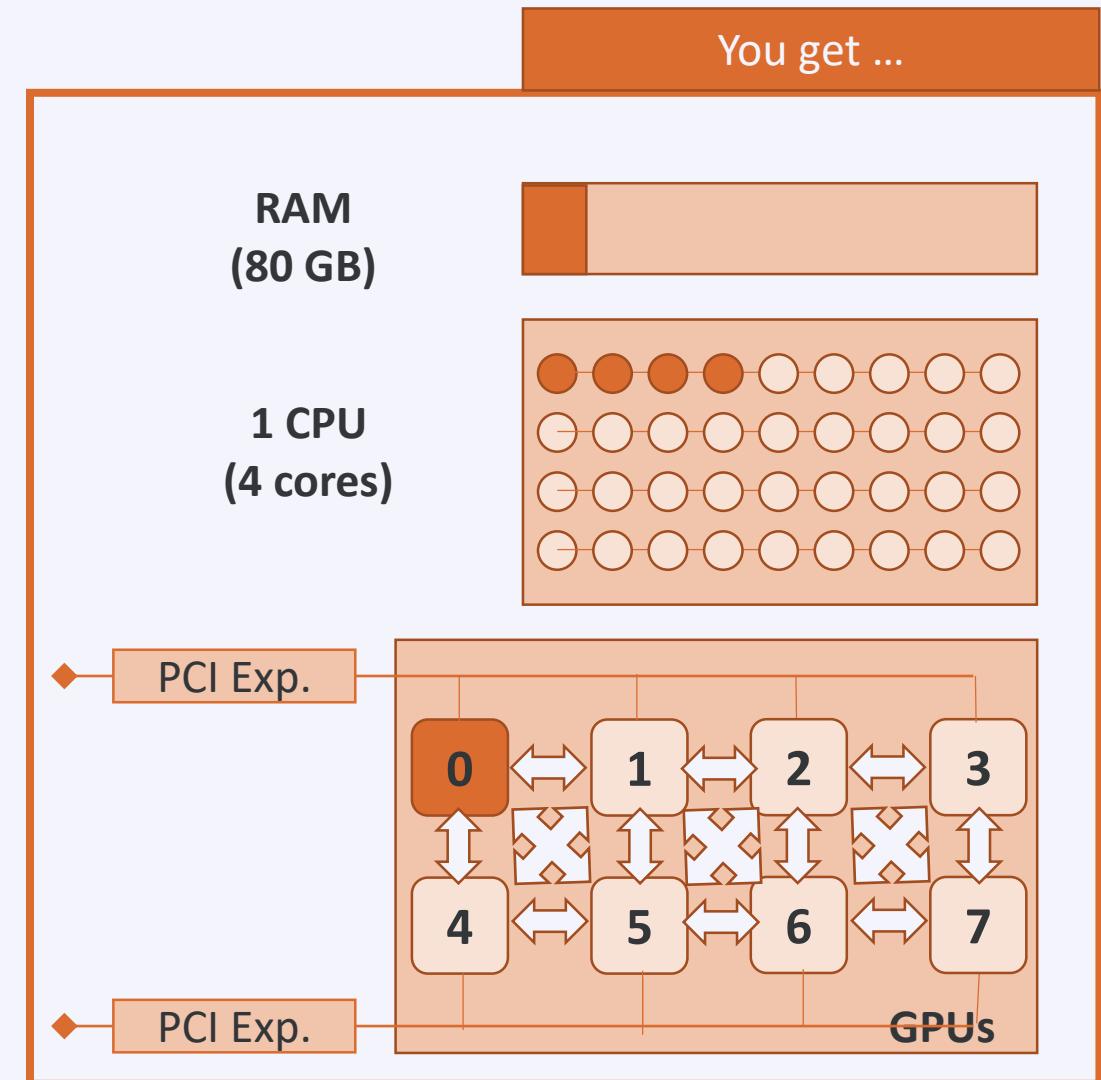
Using 2 GPUs



Using V100 GPU(s)

```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l
nodes=1:ppn=4:gpus=1:cascadelake
#PBS -l pmem=20gb
#PBS -N testjob
#PBS -A default_project
```

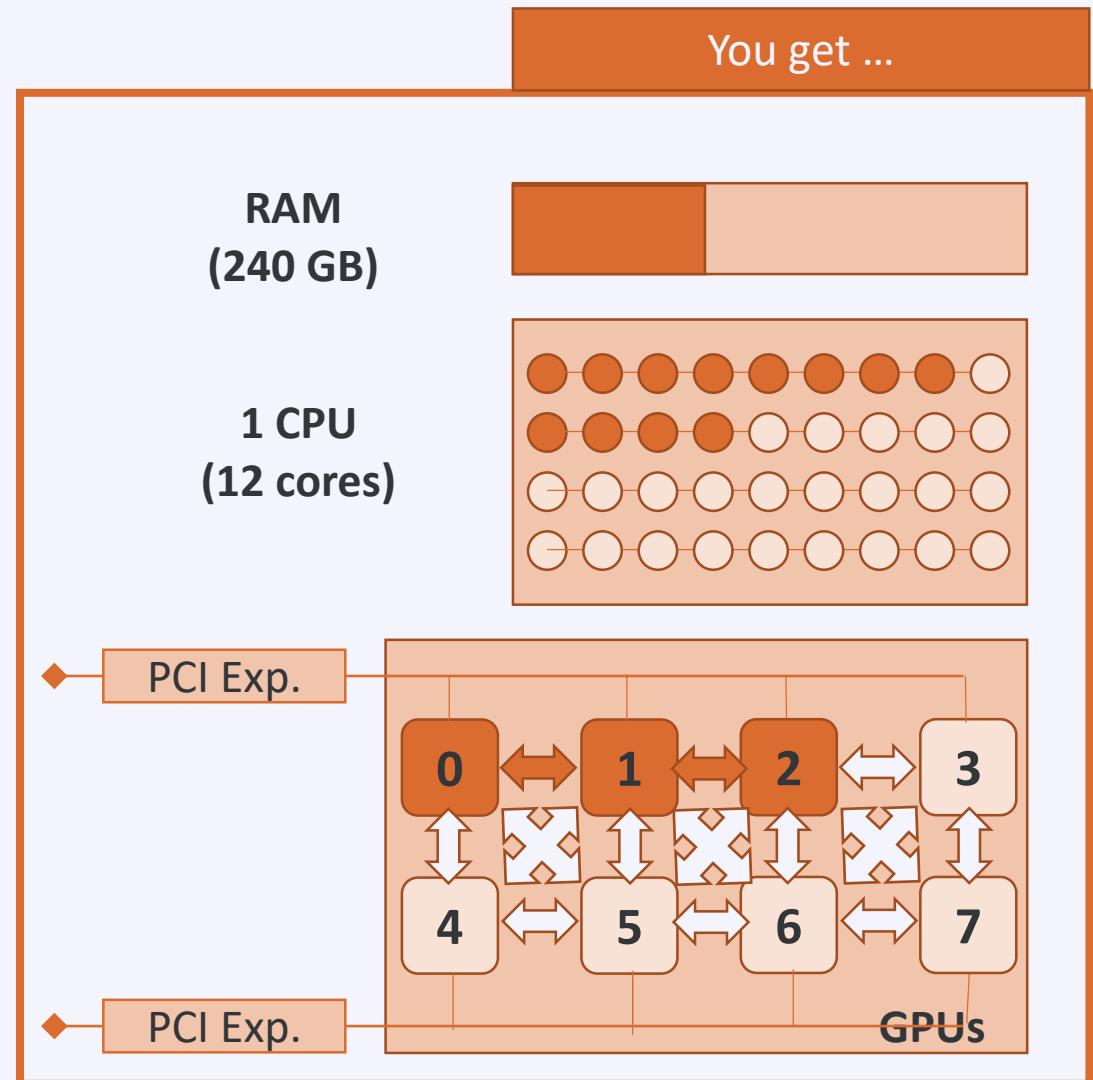
Using 1 GPU



Using V100 GPU(s)

```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l
nodes=1:ppn=12:gpus=3:cascadelake
#PBS -l pmem=20gb
#PBS -N testjob
#PBS -A default_project
```

Using 3 GPU





Other Partitions:

- Large Memory
- AMD
- Debugging Queue

Requesting Large Memory Nodes

- All Thin Nodes have 192 GB memory each
You need to leave 8 GB aside for the Operating System, and use the rest
- By default, every core can have
 $\text{pmem} = (192\text{GB} - 8\text{GB}) / (36 \text{ cores}) \approx 5\text{GB per-core}$
- There are 3 methods you can request larger memory than normal:
 - Use less cores (ppn), but larger memory-per-core (pmem)
 - Use bigmem partition
 - Use Superdome machine

First Method

- You sacrifice ppn, in favor of pmem
- E.g. if you need 20GB memory per core:
`$> qsub -I -l nodes=1:ppn=8,pmem=20gb ...`
- **Note:** the total memory requested must be less than $192\text{GB} - 8\text{GB} = 184\text{GB}$, i.e.
 $\text{ppn} \times \text{pmem} \leq 184\text{GB}$

Requesting Large Memory Nodes

Second Method

- We have 10 Large Memory nodes, with 768 GB memory each

- Specify -l partition=bigmem with qsub command

- Each core can have

$$\text{pmem} = (768 \text{ GB} - 18 \text{ GB}) / (36 \text{ cores}) \approx 20 \text{ GB per-core}$$

- E.g.

```
$> qsub -I -l partition=bigmem, nodes=1:ppn=36, pmem=20gb ...
```

- You can also compromise ppn for pmem

E.g.

```
$> qsub -I -l partition=bigmem, nodes=1:ppn=10, pmem=75gb ...
```

- Make sure $\text{ppn} \times \text{pmem} \leq 750\text{GB}$

Requesting Large Memory Nodes

Third Method

- Superdome: 8 nodes, 14 cores-per-node, **6 TB** total RAM
- Default pmem=50gb
- numanode is total CPUs to use (max. 8)
- lprocs is total number of cores to use

Request **1/8th** of the machine,
with 50 GB per core

```
$> module load superdome  
$> qsub -l partition=superdome -q qsuperdome \  
      -L tasks=1:lprocs=14:place=numanode=1
```

Request **3/8th** of the machine,
with 50 GB per core

```
$> qsub -l partition=superdome -q qsuperdome \  
      -L tasks=1:lprocs=42:place=numanode=3
```

If you want to use > 50GB/core
Eg 7 cores with 100GB/core

```
$> qsub -l partition=superdome -q qsuperdome \  
      -L tasks=1:lprocs=7:place=numanode=1:memory=700gb
```

(Experimental) AMD Nodes

- 4 nodes
- 2x EPYC 7501 CPUs (2.0 GHz, 32 cores)
- 8 memory channels / CPU
- Max Mem BW 158.95 GiB/s
- RAM: 256 GB
- Remark: pmem=3800mb
If not specified, the job will not start

Job Script

```
#!/bin/bash -l
#PBS -l walltime=01:30:00
#PBS -l partition=amd
#PBS -l nodes=1:ppn=64
#PBS -l pmem=3800mb
#PBS -N testjob
#PBS -A lp_my_project
```

Debugging / Testing Jobs

- To quickly test/debug your (parallel) application
- 2 dedicated nodes on ThinkKing and Genius
One GPU node and one CPU node
- Such jobs do **not** go to the normal queue, so they start faster
- Max. walltime is **30 minutes**
- You must specify Quality of Service (**qos**)

Request Debugging Nodes

```
$> qsub -l nodes=1:ppn=10 -l qos=debugging -l \
    walltime=30:00 -A default_project
```

Managing & Monitoring Jobs

Command	Purpose
\$> qsub	Submit a job (batch/interactive)
\$> qdel <JobID>	Delete a specific job
\$> checkjob -vvv <JobID>	Very detailed job info (very useful to diagnose issues)
\$> qstat -n	Status of all recent jobs
\$> qstat -Q -f	Info about available queues
\$> showstart <JobID>	Give a <i>rough</i> estimate of start time
\$> showq \$> showq -p gpu	Show minimal info about a queue or partition (-p)
\$> pbstop \$> clusterview \$> queueview	Overview of the cluster Overview of the nodes, cores and GPUs Overview of the job queues
\$> mam-balance	Overview of different credit projects that you can use (<code>qsub -A <Project></code>)
\$> mam-list-allocations	Detailed overview of your credit projects

Jupyter Notebook on NX

- ❑ Install Miniconda in your \$VSC_DATA

```
$> wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh  
$> bash Miniconda3-latest-Linux-x86_64.sh -p $VSC_DATA/miniconda3
```

- ❑ Create a conda env including Jupyter

```
$> conda create -n JupNtbk jupyter scikit-learn
```

- ❑ Start an interactive GPU job (qsub -I ...)

- ❑ Activate this env

```
$> conda activate JupNtbk
```

- ❑ In another terminal, launch the notebook to get URL

```
$> jupyter notebook --ip=`hostname -i` --port=${USER}:3 --no-browser  
--allow-root
```

- ❑ Copy-paste the URL into NX Firefox; done

Singularity Containers

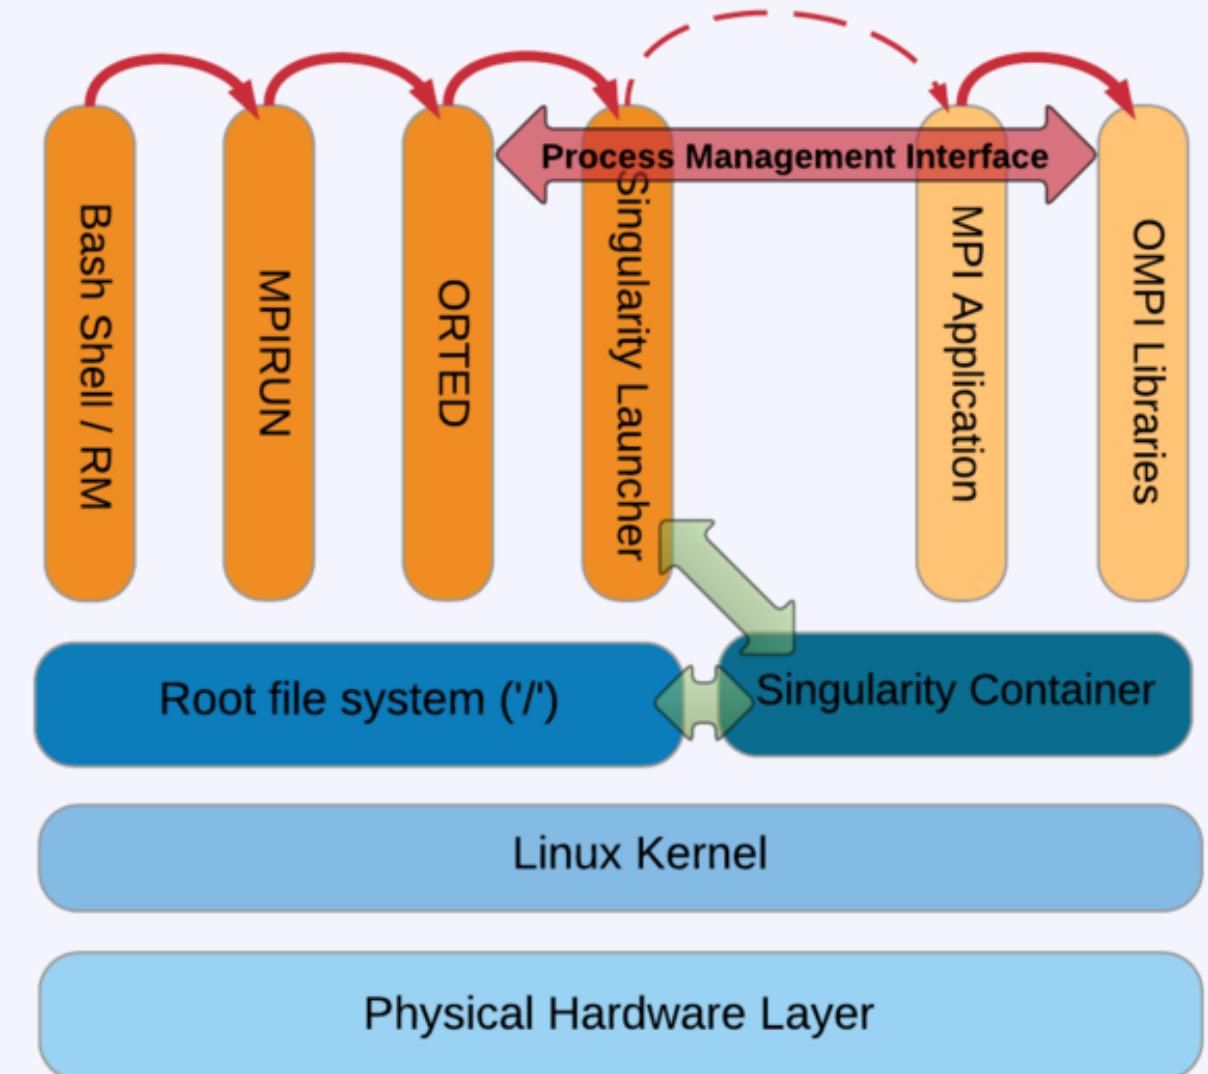
- What?
 - + Self-contained OS & software & data
- Why?
 - + fully resolved dependency chains
 - + portable workflow
- How?
 - + You create the image
 - + Run it on Genius
 - + MPI/OpenMP is supported

```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l nodes=4:ppn=36:pmem=5gb
#PBS -A lp_hpcinfo_training

cd $PBS_O_WORKDIR

singularity run Project.simg ./model.exe
```

PBS Script



The background of the slide is a hand-drawn illustration of an industrial setting. It features a network of brown pipes and valves in the foreground. Behind them is a large, circular metal drum with a green valve at the bottom. Further back, there's a wall with several blue circular valves and a control panel with a digital display showing various numbers and symbols. The style is artistic and somewhat abstract.

Where to run your jobs

Which system should I use?

Knowing yourself is
the beginning
of all wisdom



What kind of work I do?

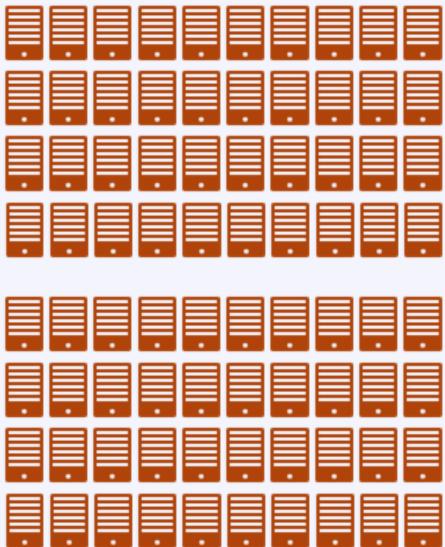
- CPU usage
 - Do you run single core jobs or parallel?
 - Single core
 - Is your code optimized to AVX ?
 - Parallel:
 - single node or multinode?
 - Single node: does my application scale? Up to how much cores ?
 - Multinode:
 - Does my application scale?
 - Memory usage
 - How much memory (per core) my jobs use?
 - Is your processing memory bound?

What kind of work I do?

- How to find out
 - Know your algorithm
 - Benchmark your typical workload on the different architectures
 - Use tools to track behavior of your code
 - Use [Monitor](#) tool
 - Use profiling tools eg. [ARM MAP](#) or [Scalasca](#)
 - Consult training agenda
 - Upcoming Allinea ARM Forge training

Which system for what?

Compute nodes (thin nodes)



Thinking:

- Haswell: 24 cores
64/128 GB

Single core jobs
Worker jobs
Single node parallel jobs

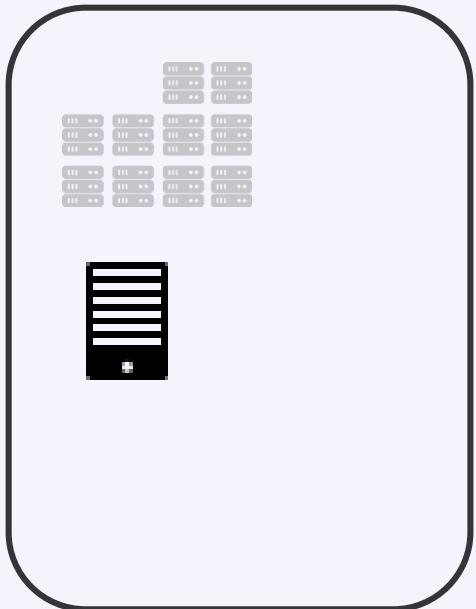
Genius:

- Skylake: 36 cores
- Cascade Lake: 36 cores
- 192 GB
- IB EDR

Multinode parallel jobs
AVX 512
Higher core/memory ratio
Higher memory bandwidth

Which system for what?

Compute nodes (Large memory)



Genius:

- Big memory 36 cores
- 768 GB

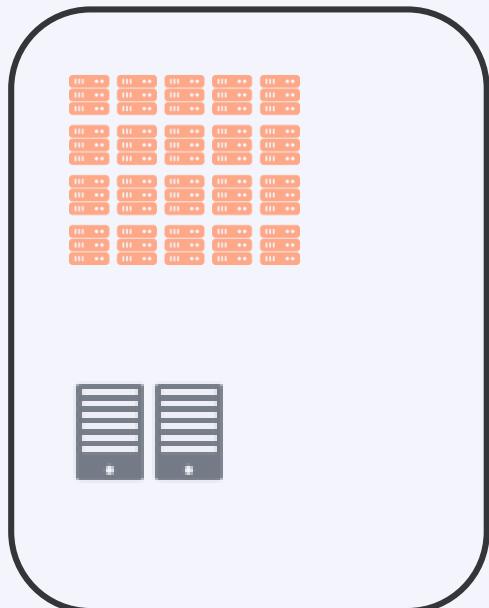
5GB<mem/core<20GB

- Superdome 112 cores
- 6TB

> 20GB mem/core
> 700 GB of total RAM
Single node scaling > 36 cores

Which system for what?

Compute nodes (GPU nodes)



Genius Skylake GPU: P100

Single GPU workload
Multi GPU workload with limited scaling

Genius Cascadelake GPU: V100

> 4 GPU workload
Specific ML workloads (convolutional networks)

Credits

Credits card concept:

- **Pre-authorization**: holding the balance as unavailable until the merchant clears the transaction
- Balance to be held as unavailable: based on requested resourced (walltime, nodes)
- **Actual charge** based on what was really used: used walltime (you pay only what you use, e.g. when job crashes)
- See output file

How to check available credits?

```
$ mam-balance
```

How to check the rates?

Check service catalogue

```
$ mam-list-chargerates
```

Credit Pricing

- You pay as you go!
- For academic projects:
1000 credits = 3.5 EUR
- Credits needed for a job:

#credits = Walltime (hr) × #nodes × Factor

- For shared nodes, you pay a fraction of the costs,
based on the ppn specified

Cluster / Partition	Credits/h
Genius Cascadelake	11.3
Genius Cascadelake 8 GPUs	40
Genius Skylake 4 GPUs	20
Genius Skylake BigMem	12
Genius Skylake	10
Genius / Superdome	10
Genius / AMD nodes	10
ThinKing / Haswell	6.68
ThinKing / IvyBridge	4.76

How to Manage Credits?

Command	Purpose
mam-balance	List active projects and available credits
mam-list-allocations	List the validity dates of different projects/allocations
mam-list-chargerates	List current charge rates for different nodes
mam-list-accounts -a <account-name>	List the members in an account

demo/test yourself

- ✓ Request membership to lp_hpcintro_training group (account.vscentrum.be)

- ✓ Login with putty
- ✓ Filetransfer with Filezilla
- ✓ Login with NX
- ✓ Check disk quota
- ✓ Check the credits
- ✓ Check/load/list/unload/purge module

demo/test yourself

- ✓ Copy intro training files (`/apps/leuven/training/HPC_intro/`) to your `$VSC_HOME`
- ✓ Submit cpujob to the cluster
- ✓ List all your jobs (`qstat`)
- ✓ Check the information about the cpujob (`checkjob`)
- ✓ Modify the mat.pbs script to request 1 node, 36 cores for 30 minutes and get the notification about job start/end by e-mail
- ✓ Check the status of all the jobs

demo - monitoring

- ✓ Submit an interactive job
 - Run your program on a compute node
 - Open a new terminal and ssh to a compute node
 - Check the resources usage (`top`)
- ✓ Submit an interactive job to GPU node
 - Run your program on a compute node
 - Open a new terminal and ssh to a compute node
 - Check the resources and GPU usage (`top`, `nvidia-smi` or `watch "nvidia-smi"`)
 - Submit a batch GPU job
- ✓ While the job is running get the information about the node (`checkjob`) and check usage of resources on the node (`ssh`, `top`, `nvidia-smi`)

demo – conda installation

✓ Install miniconda in your \$VSC_DATA directory

- \$ wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh
- \$ bash Miniconda3-latest-Linux-x86_64.sh -b -p \$VSC_DATA/miniconda3

✓ Add a PATH to conda:

- \$ export PATH="\$VSC_DATA/miniconda3/bin:\$PATH"

✓ Check if conda is added to your \$PATH (\$ which conda)

✓ Add it to \$PATH in your .bashrc

- \$ echo 'export PATH="\$VSC_DATA/miniconda3/bin:\$PATH"' >> .bashrc

demo – conda usage

- ✓ Create a conda environment including Jupyter

```
$ conda create -n science jupyter numpy scipy
```

Activate this environment

- ✓ \$ source activate science

Add matplotlib package to this environment

- ✓ \$ conda install matplotlib

Return to original environment

```
$ conda deactivate
```



demo – notebooks

- ✓ Start an interactive GPU job

```
$ qsub -I -l walltime=30:00 -l nodes=1:ppn=9:gpus=1 -l  
partition=gpu -A default_project
```

- ✓ Activate conda environment

```
$ source activate science
```

- ✓ Go to your working directory (you can use \$PBS_O_WORKDIR if you qsub from there)

- ✓ Start notebook

```
$ jupyter notebook --port ${USER:3} --ip $(hostname)
```

```
$ jupyter notebook --port 30468 --ip $(hostname)
```

- ✓ Open the link in the browser in NX and test your notebook

demo – worker

- ✓ Copy intro training files (/apps/leuven/training/worker/) to your \$VSC_HOME
- ✓ Go to exercise1 directory
- ✓ Submit worker job
- ✓ Check the output file

Questions

Helpdesk:

hpcinfo@kuleuven.be or https://admin.kuleuven.be/icts/HPCinfo_form/HPC-info-formulier

VSC web site:

<http://www.vscentrum.be/>

VSC documentation: <https://vlaams-supercomputing-centrum-vscdocumentation.readthedocs-hosted.com/en/latest/>

VSC agenda: training sessions, events



Systems status page:

<http://status.kuleuven.be/hpc>

