



Vlaanderen  
is supercomputing

# VSC HPC Introduction

ICTS KU Leuven

VSC staff:

Ehsan Moravveji

Mag Selwa

Geert Jan Bex (UHasselt)

Jan Ooghe

# What is High Performance Computing?

- ❑ using supercomputers to solve advanced computation problems
- ❑ Reduce the computation time from days, years, decades, or centuries to minutes, hours, days, or weeks
- ❑ The key is parallelism



# In practice, it is more like ...



The concept is simple: **Parallelism** = employing multiple processors for a single problem

# Outline

- What is the VSC?
- What is a cluster?
- Genius Cluster
- Storage
- Login nodes
- Connection Setup
- Software environment
- How to submit jobs?
- Dedicated hardware
- How to choose resources?
- Optional material
  - Linux in brief
  - Conda for Python and R
  - Worker Framework





# VSC

(Vlaams Supercomputer Centrum)



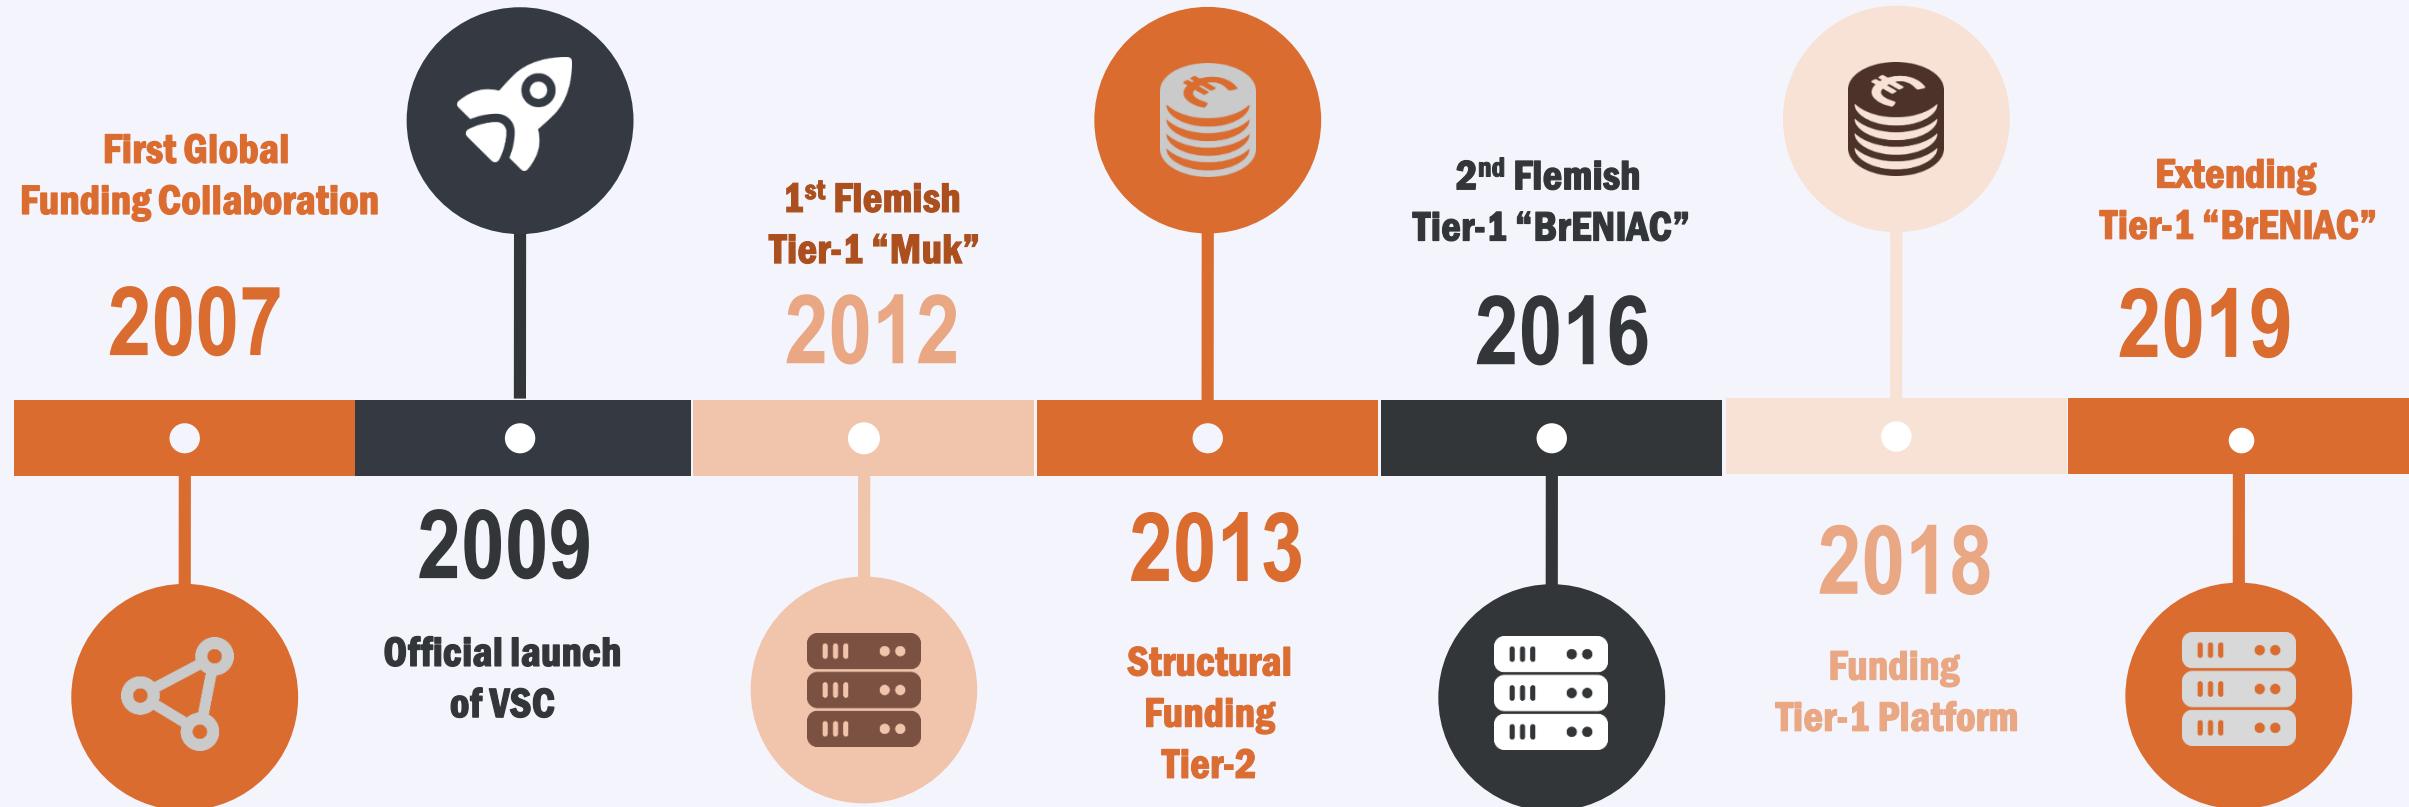
# VSC PARTNERSHIP



Supported by



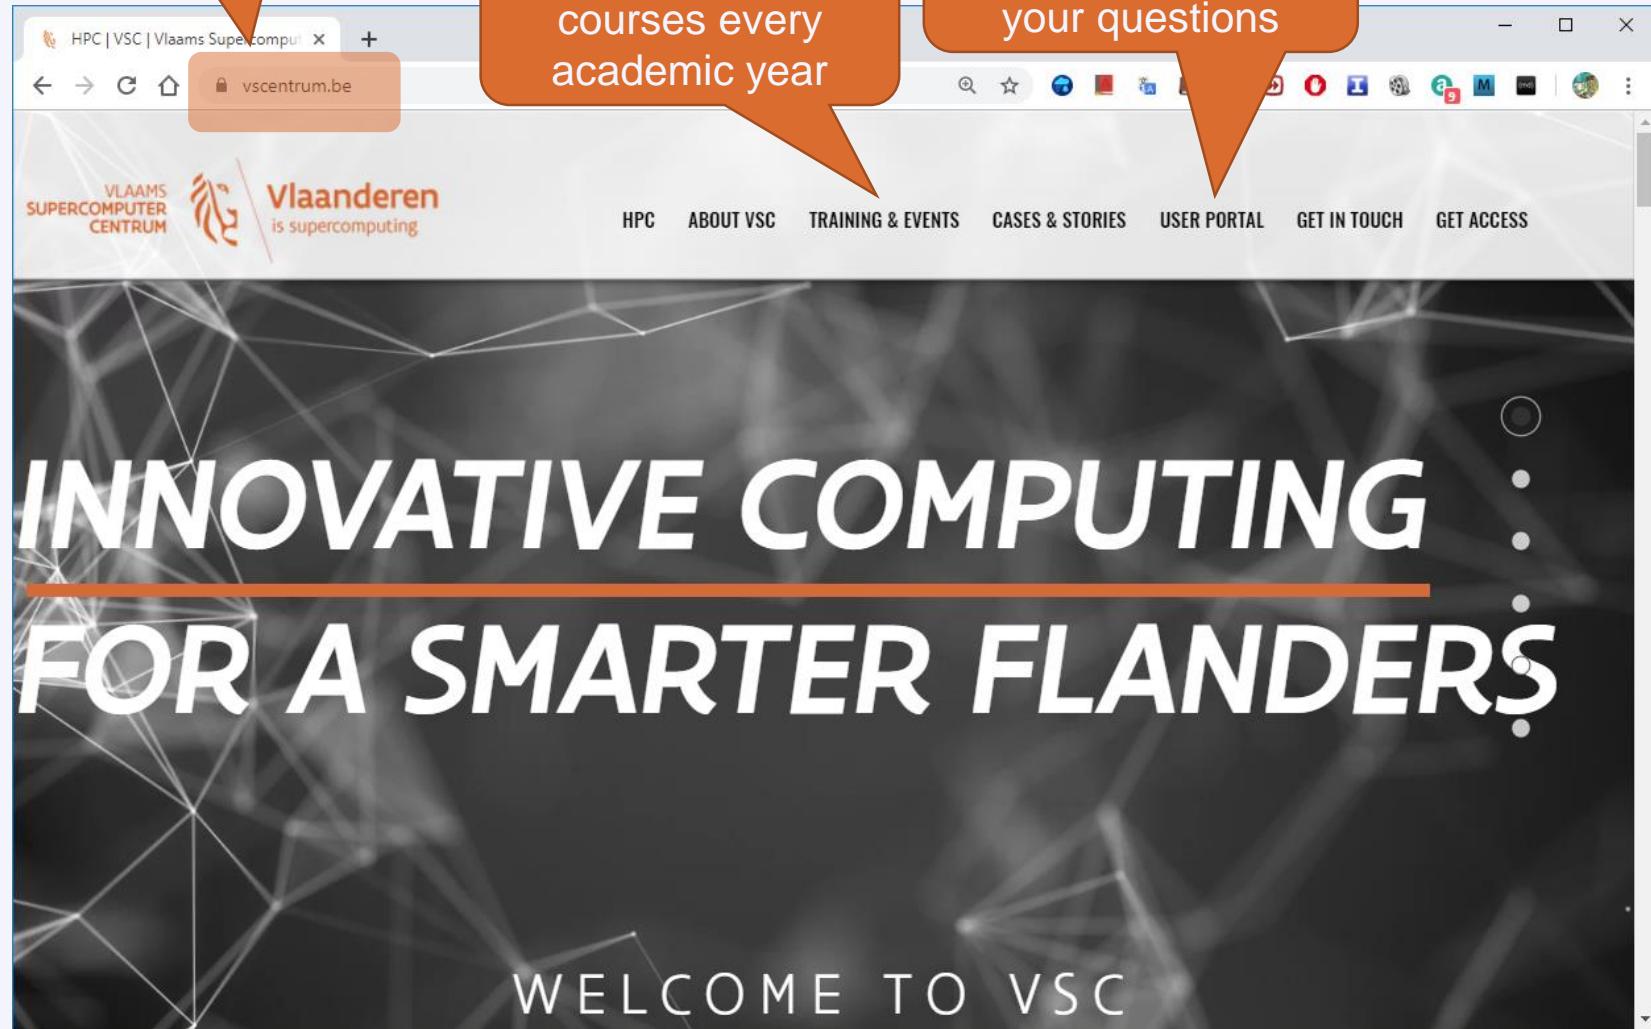
# VSC history



# VSC HPC Environments



[www.vscentrum.be](http://www.vscentrum.be)



Search your keywords here;  
e.g. qsub

Welcome to VSC documentation

next | index

Getting access

Access and data transfer

Software stack

Running jobs

Software development

V: latest

Getting access

- Getting access
  - VSC accounts
  - How to request an account?
  - Next steps
  - Additional information
- Access and data transfer
  - Logging in to a cluster
  - Data storage
  - Transferring data
  - GUI applications on the clusters
  - VPN
- Software stack
  - Using the module system
  - Specialized software stacks
- Running jobs
  - Job script
  - Submitting and monitoring a job
  - Job output
  - Troubleshooting
  - Advanced topics
- Software development

# Support and Services

## Basic support

- Helpdesk ([hpcinfo@kuleuven.be](mailto:hpcinfo@kuleuven.be))
- Monitoring and reporting

## Application support

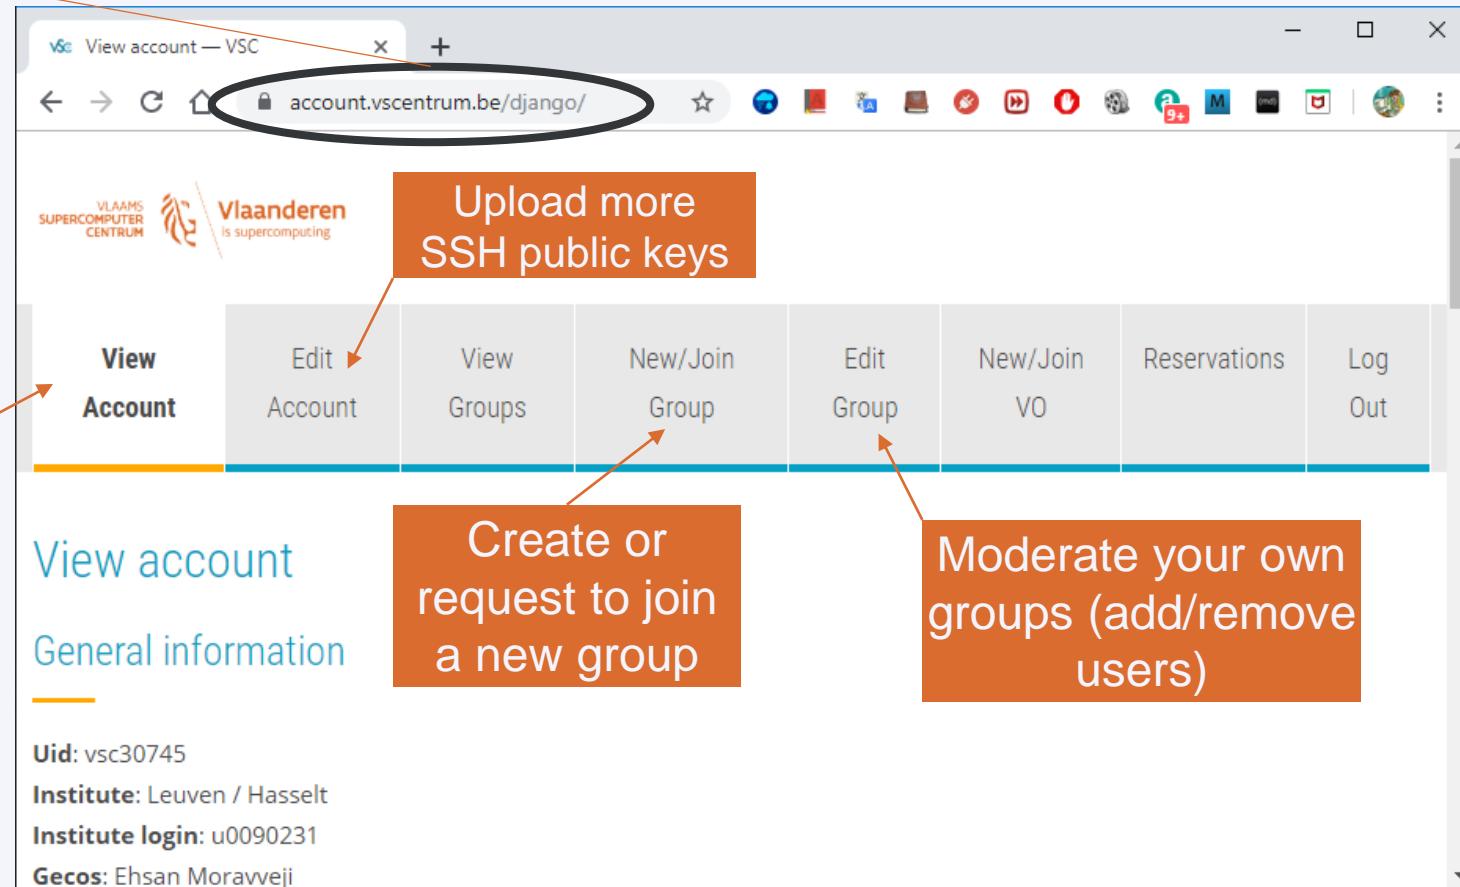
- Installation and porting
- Optimisation and debugging
- Benchmarking
- Workflows and best practices

## Training

- Documentation and tutorials
- Scheduled trainings / workshops
- On request workshops
- One-to-one sessions

To manage your  
VSC account:

[account.vscentrum.be](http://account.vscentrum.be)



Upload SSH Key  
Request account

# Become a VSC user

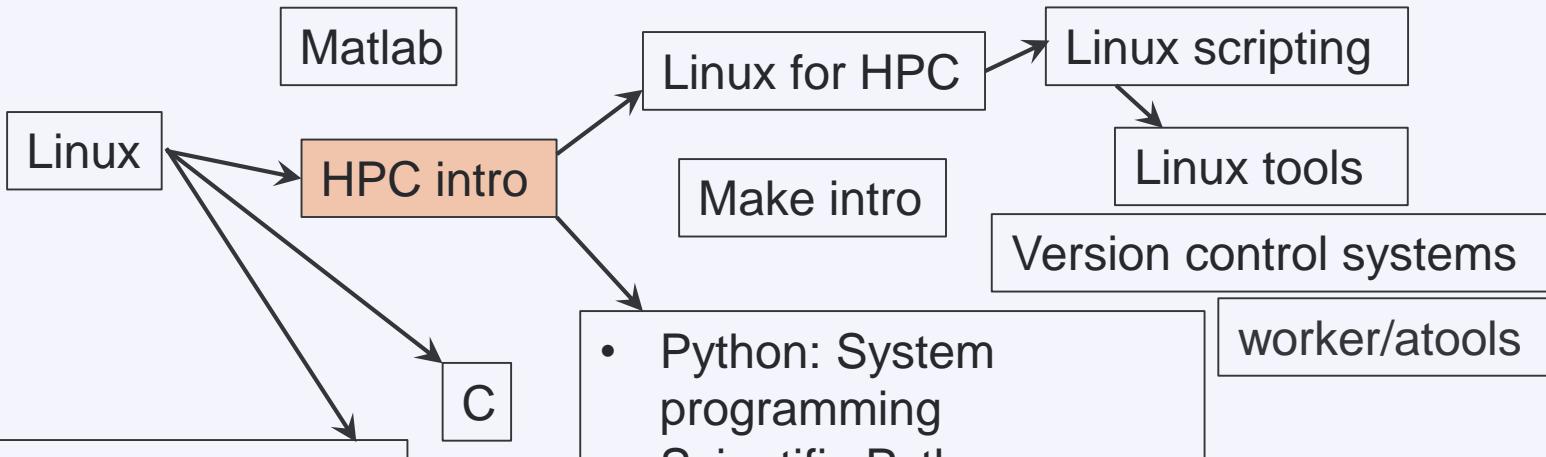
- Create a secure (4096 bit) [SSH key pairs](#)  
Upload it on the account page: [www.account.vscentrum.be](http://www.account.vscentrum.be)
- You need to [request a VSC account](#)  
Normally processed swiftly
- Request [introductory credits](#) (2000 free credits for 6 months)
- Request [project credits](#) (for supervisors and project leaders)
  - You need to create a VSC group
  - Add users to the group to give them access to use credits
  - Fill out the request form
- Extra storage requests
  - Scratch extension: free of charge
  - Archive fileset: 70 € per TB per year
  - Staging fileset: 130 € per TB per year
- All service costs (compute and storage) are all explained  
Go to ICTS service catalogus: <https://icts.kuleuven.be/sc>  
Click on [High Performance Computing](#) (NL/EN)

The screenshot shows a web browser window titled 'View account — VSC'. The URL is 'account.vscentrum.be/djang...'. The page header includes the Vlaams Supercomputer Centrum logo and the text 'Vlaanderen is supercomputing'. A navigation menu at the top has tabs for 'View Account' (which is selected), 'Edit Account', 'View Groups', 'New/Join Group', 'Edit Group', 'New/Join VO', 'Reservations', and 'Log Out'. Below the menu, the main content area displays 'View account' and 'General information'.

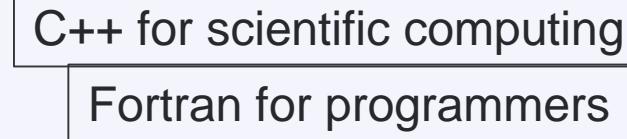
The screenshot shows a web browser window titled 'ICTS Servicecatalogus'. The URL is 'icts.kuleuven.be/sc'. The page header includes the KU Leuven logo and the text 'ICTS SERVICECATALOGUS'. A navigation menu has links for 'Home', 'ICTS SERVICECATALOGUS', 'WAT IS DE ICTS SERVICECATALOGUS?', 'OPGELET', and 'ZOEK IN DE SERVICECATALOGUS'. The 'OPGELET' section contains text about netto prijzen and overhead. On the right side, there is an image of a computer mouse on a blue book.

# VSC training 2018/2019

- Introductory



- Intermediate



- Advanced

High Performance Python

- Specialized track

?

Lunchbox sessions:

- Containers
- Notebooks

Stay up-to-date <https://www.vscentrum.be/en/education-and-trainings>

# To Acknowledge VSC in publications

## Why?

- a contractual obligation for the VSC
- helps VSC secure funding
- you will benefit from it in the long run

## At KU Leuven

- add the relevant papers to the virtual collection "High Performance Computing" in Lirias

## In het nederlands

*De rekeninfrastructuur en dienstverlening  
gebruikt in dit werk, werd voorzien door het  
VSC (Vlaams Supercomputer Centrum),  
gefincierd door het FWO en de Vlaamse  
regering – departement EWI.*

## In English

*The computational resources and services  
used in this work were provided by the VSC  
(Flemish Supercomputer Center), funded by  
the Research Foundation - Flanders (FWO)  
and the Flemish Government – department  
EWI.*

## Tier-2 Clusters

# Tier-2 Clusters @ KU Leuven

**ThinKing** (since 2014)  
352 nodes: 7,616 cores



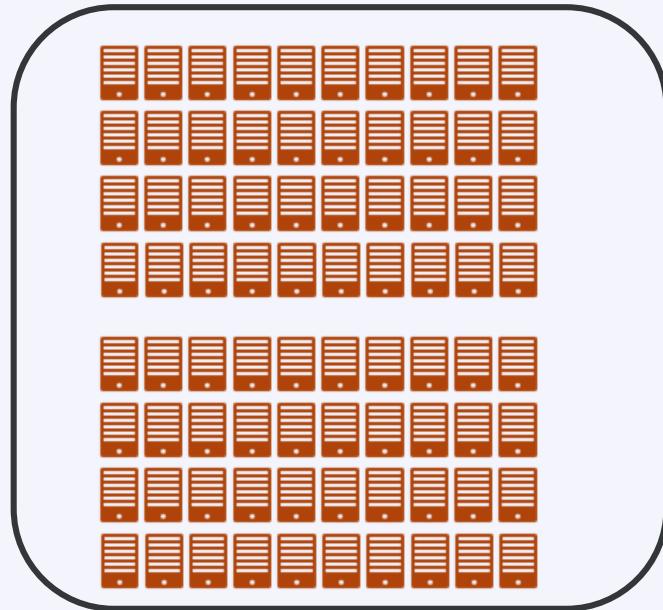
**Genius** (since 2018)  
250 nodes: 8,936 cores



# Tier-2 Overview



Compute nodes



**Thinking:**

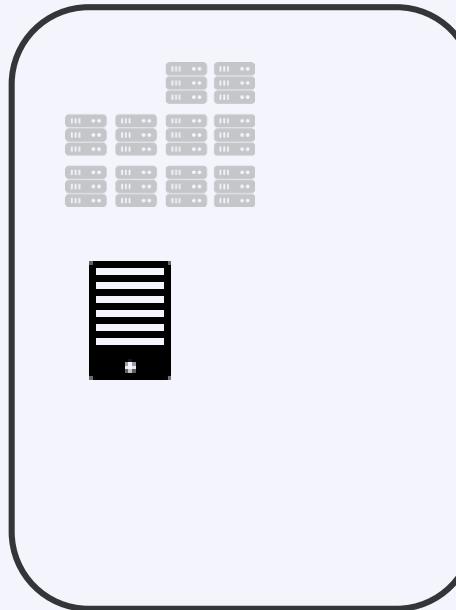
+ 176x + 32x Ivybridge 20c 64/128 GB  
**(EOL)**

+ 48x + 86x Haswells 24c 64/128 GB

**Genius:**

+ 96x Skylake 36c 192 GB  
+ 120x CascadeLake 36c 192 GB

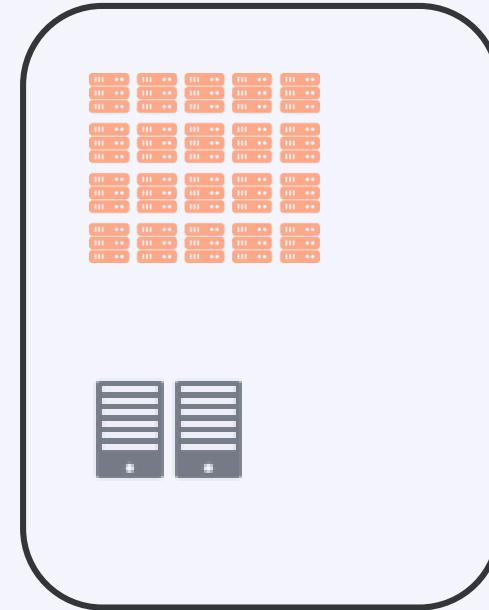
Large memory nodes



**Genius:**

+ 10x Skylake 36c 768GB  
+ 1x Superdome 112c 6 TB

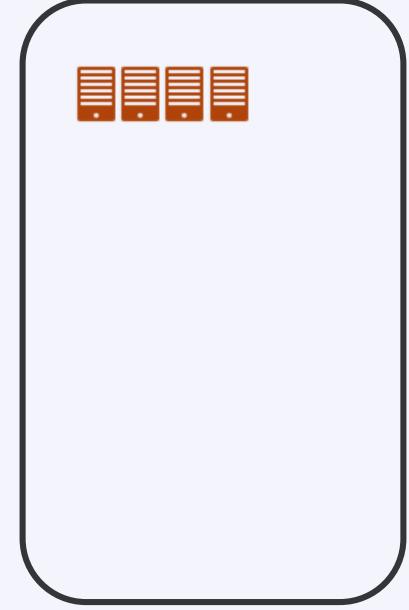
GPU nodes



**Genius:**

+ 20x Skylake 36c 192 GB  
4x P100 16GB  
+ 2x CascadeLake 36c 192GB  
8x V100 32GB

Exp nodes



**Genius:**

4 AMD Naples 64c 256 GB

# Technical Hardware Specifications

	Tier 2						Tier 1		
Cluster name	ThinKing			Genius			BrENIAC		
Processor type	Ivybridge		Haswell		SkyLake		Cascade Lake		Broadwell
Cores per node	20		24		36		36		SkyLake
Base Clock Speed	2.8 GHz		2.5 GHz		2.3 GHz		2.6 GHz		2.4 GHz
Total nodes	176	32	48	96	86	10	120	580	408
Node memory (GB)	64	128	64	96	192	768	192	128	256
Memory per core (GB)	3.0	6.2	2.5	3.8	5.2	21.2	5.3	4.4	9.0
Total cores	4,160		3,456		3,456		4,320		27,664
Peak performance (Flops/cycle)	4 DP FLOPs/cycle: 4-wide AVX addition OR 4-wide AVX multiplication		8 DP FLOPs/cycle: 4-wide FMA (fused multiply-add) instructions AVX2		16 DP FLOPs/cycle: 8-wide FMA (fused multiply-add) instructions AVX-512		16 DP FLOPs/cycle: 8-wide FMA (fused multiply-add) instructions AVX2		16 DP FLOPs/cycle: 8-wide FMA (fused multiply-add) instructions AVX-512
Network	Infiniband QDR 2:1		Infiniband FDR		Infiniband EDR		Infiniband EDR		Infiniband EDR
Cache (L1 KB/L2 KB/L3 MB)	10x(32i+32d) / 10x256 / 25 MB		12x(32i+32d) / 12x256 / 30MB		18x(32i+32d) / 18x1024 / 25 MB		18x(32i+32d) / 18x1024 / 25 MB		14x(32i+32d) / 14x 256 / 35 MB
									18x(32i+32d) / 18x1024 / 25 MB

	<b>Tier 2</b>	<b>Tier 1</b>
Access	All researchers and students	Approved project experienced users
Start with	2,000 introduction credits	500 node days starting grant
Project Credits	Pay as you GO	Panel Review Project proposal Free
Computation Limit	None	500 - 5,000 node days
Application deadline	None	3 February 5 June 2 October
Max. Project Duration		8 months
Available walltime	1h, 24h, 72h, 7d	Max. 3 days



# Genius

SkyLake (since 8/2018)

CascadeLake (since 1/2020)

# Tier-2 Cluster: Genius

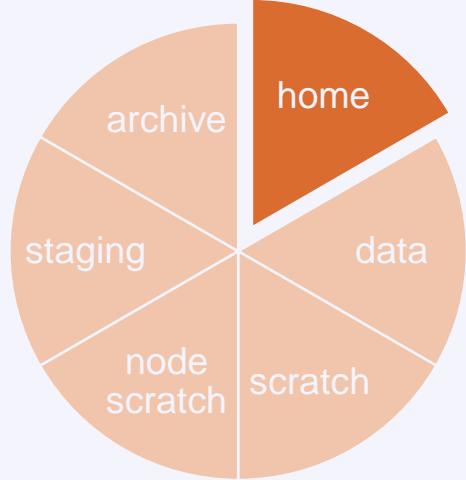
Type of node	CPU type	Inter-connect	# cores	installed mem	local discs	# nodes
SkyLake	Xeon 6140	IB-EDR	36	192 GB	800 GB	86
SkyLake large mem	Xeon 6140	IB-EDR	36	768 GB	800 GB	10
SkyLake GPU	Xeon 6140 4xP100 SXM2	IB-EDR	36	192 GB	800 GB	20
CascadeLake	Gold 6240	IB-EDR	36	192 GB	800 GB	120
CascadeLake GPU	Gold 6240 8xV100 SMX2	IB-EDR	36	192 GB	800 GB	2
SkyLake Superdome	Gold 6132	Flex Grid	14	6 TB	6 TB	8

# Storage

# Overview of the storage infrastructure

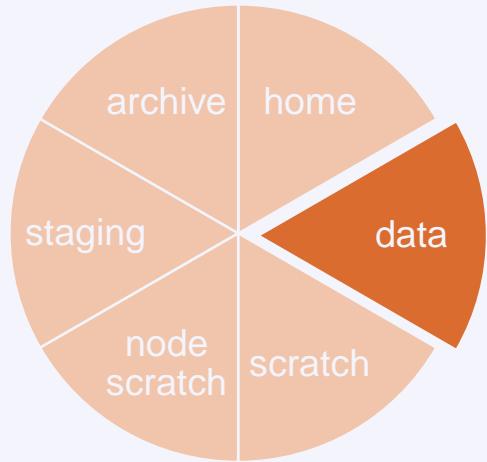
- Your files are owned only by you.  
Other VSC users have no permission to read/write/execute your files (POSIX)
- A VSC account has 3 default storages (free of charge)
  - \$VSC\_HOME
  - \$VSC\_DATA
  - \$VSC\_SCRATCH
- You can additionally request staging and archive storages
- Different storage volumes have different:
  - mount point
  - size and performance
  - use case
  - backup and maintenance policy
- More info on [ICTS Service Catalog](#) (EN/NL)

# Storage



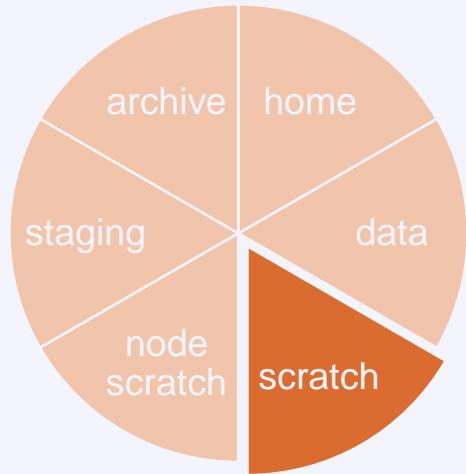
Storage	home folder
<b>Env. Variable</b>	\$VSC_HOME
<b>Filesystem Type</b>	NFS
<b>Access</b>	Global
<b>Backup</b>	Hourly, daily, weekly (until last month)
<b>Default Quota</b>	3 GB
<b>Extension</b>	Not possible
<b>Usage</b>	Only storing SSH keys, config files
<b>Remarks</b>	<ul style="list-style-type: none"><li>- Stay away from using it</li><li>- Can easily overflow:<ul style="list-style-type: none"><li>+ Your jobs may crash</li><li>+ Login issues</li></ul></li></ul>

# Storage



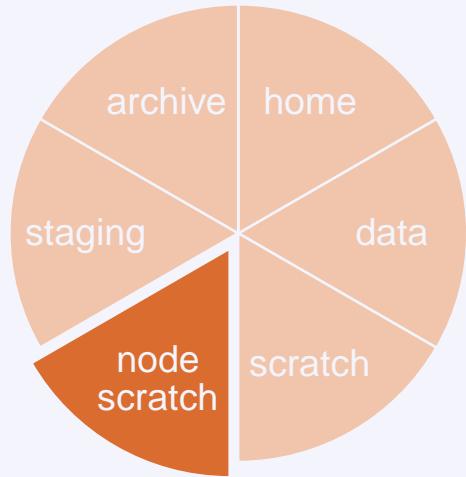
Storage	data folder
<b>Env. Variable</b>	\$VSC_DATA
<b>Filesystem Type</b>	NFS
<b>Access</b>	Global
<b>Backup</b>	Hourly, daily, weekly (until last month)
<b>Default Quota</b>	75 GB
<b>Extension</b>	On purchase
<b>Usage</b>	Your data, code, software, results
<b>Remarks</b>	<ul style="list-style-type: none"><li>- Permanent storage for initial/final results</li><li>- Not optimal for intensive or parallel I/O</li></ul>

# Storage



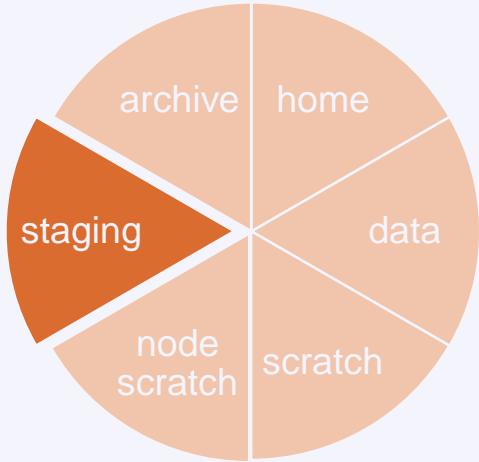
Storage	scratch folder
<b>Env. Variable</b>	\$VSC_SCRATCH
<b>Filesystem Type</b>	GPFS
<b>Access</b>	Global
<b>Backup</b>	delete after 28 days from last access
<b>Default Quota</b>	100 GB
<b>Extension</b>	Free
<b>Usage</b>	Intensive, parallel I/O, temporary files
<b>Remarks</b>	<ul style="list-style-type: none"><li>- Recommended storage for all jobs</li><li>- Copy scratch files to VSC_DATA or local storage after jobs are done</li><li>- Deleted files cannot be recovered</li></ul>

# Storage



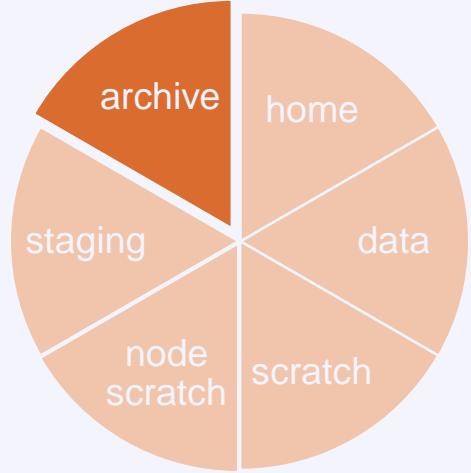
Storage	Node scratch folder
<b>Env. Variable</b>	\$VSC_SCRATCH_NODE
<b>Filesystem Type</b>	GPFS
<b>Access</b>	On compute node, only at runtime
<b>Backup</b>	None
<b>Default Quota</b>	200 GB
<b>Extension</b>	<a href="#">Read about beeOND</a>
<b>Usage</b>	Temporary storage at runtime
<b>Remarks</b>	<ul style="list-style-type: none"><li>- Fastest I/O, attached to the node</li><li>- Is cleaned after job terminates</li><li>- Copy the data to your home, scratch, or staging before job ends</li></ul>

# Storage



Storage	Staging folder
<b>Path</b>	/staging/leuven/stg_000XX
<b>Filesystem Type</b>	GPFS
<b>Access</b>	On demand, only Tier-2@KUL
<b>Backup</b>	None
<b>Default Quota</b>	None
<b>Extension</b>	On purchase, from 1 TB
<b>Usage</b>	Permanent; share with a group
<b>Remarks</b>	<ul style="list-style-type: none"><li>- Accessible from login/compute nodes</li><li>- Fast, parallel I/O</li></ul>

# Storage



Storage	Archive folder
<b>Path</b>	/archive/leuven/arc_000XX
<b>Filesystem Type</b>	NFS
<b>Access</b>	On demand, only Tier-2@KUL
<b>Backup</b>	snapshots
<b>Default Quota</b>	None
<b>Extension</b>	On purchase, from 1 TB
<b>Usage</b>	Permanent; share with a group
<b>Remarks</b>	<ul style="list-style-type: none"><li>- Accessible only from login nodes</li><li>- Cannot use it from compute nodes</li><li>- Slow, but durable filesystem</li></ul>

# Storage

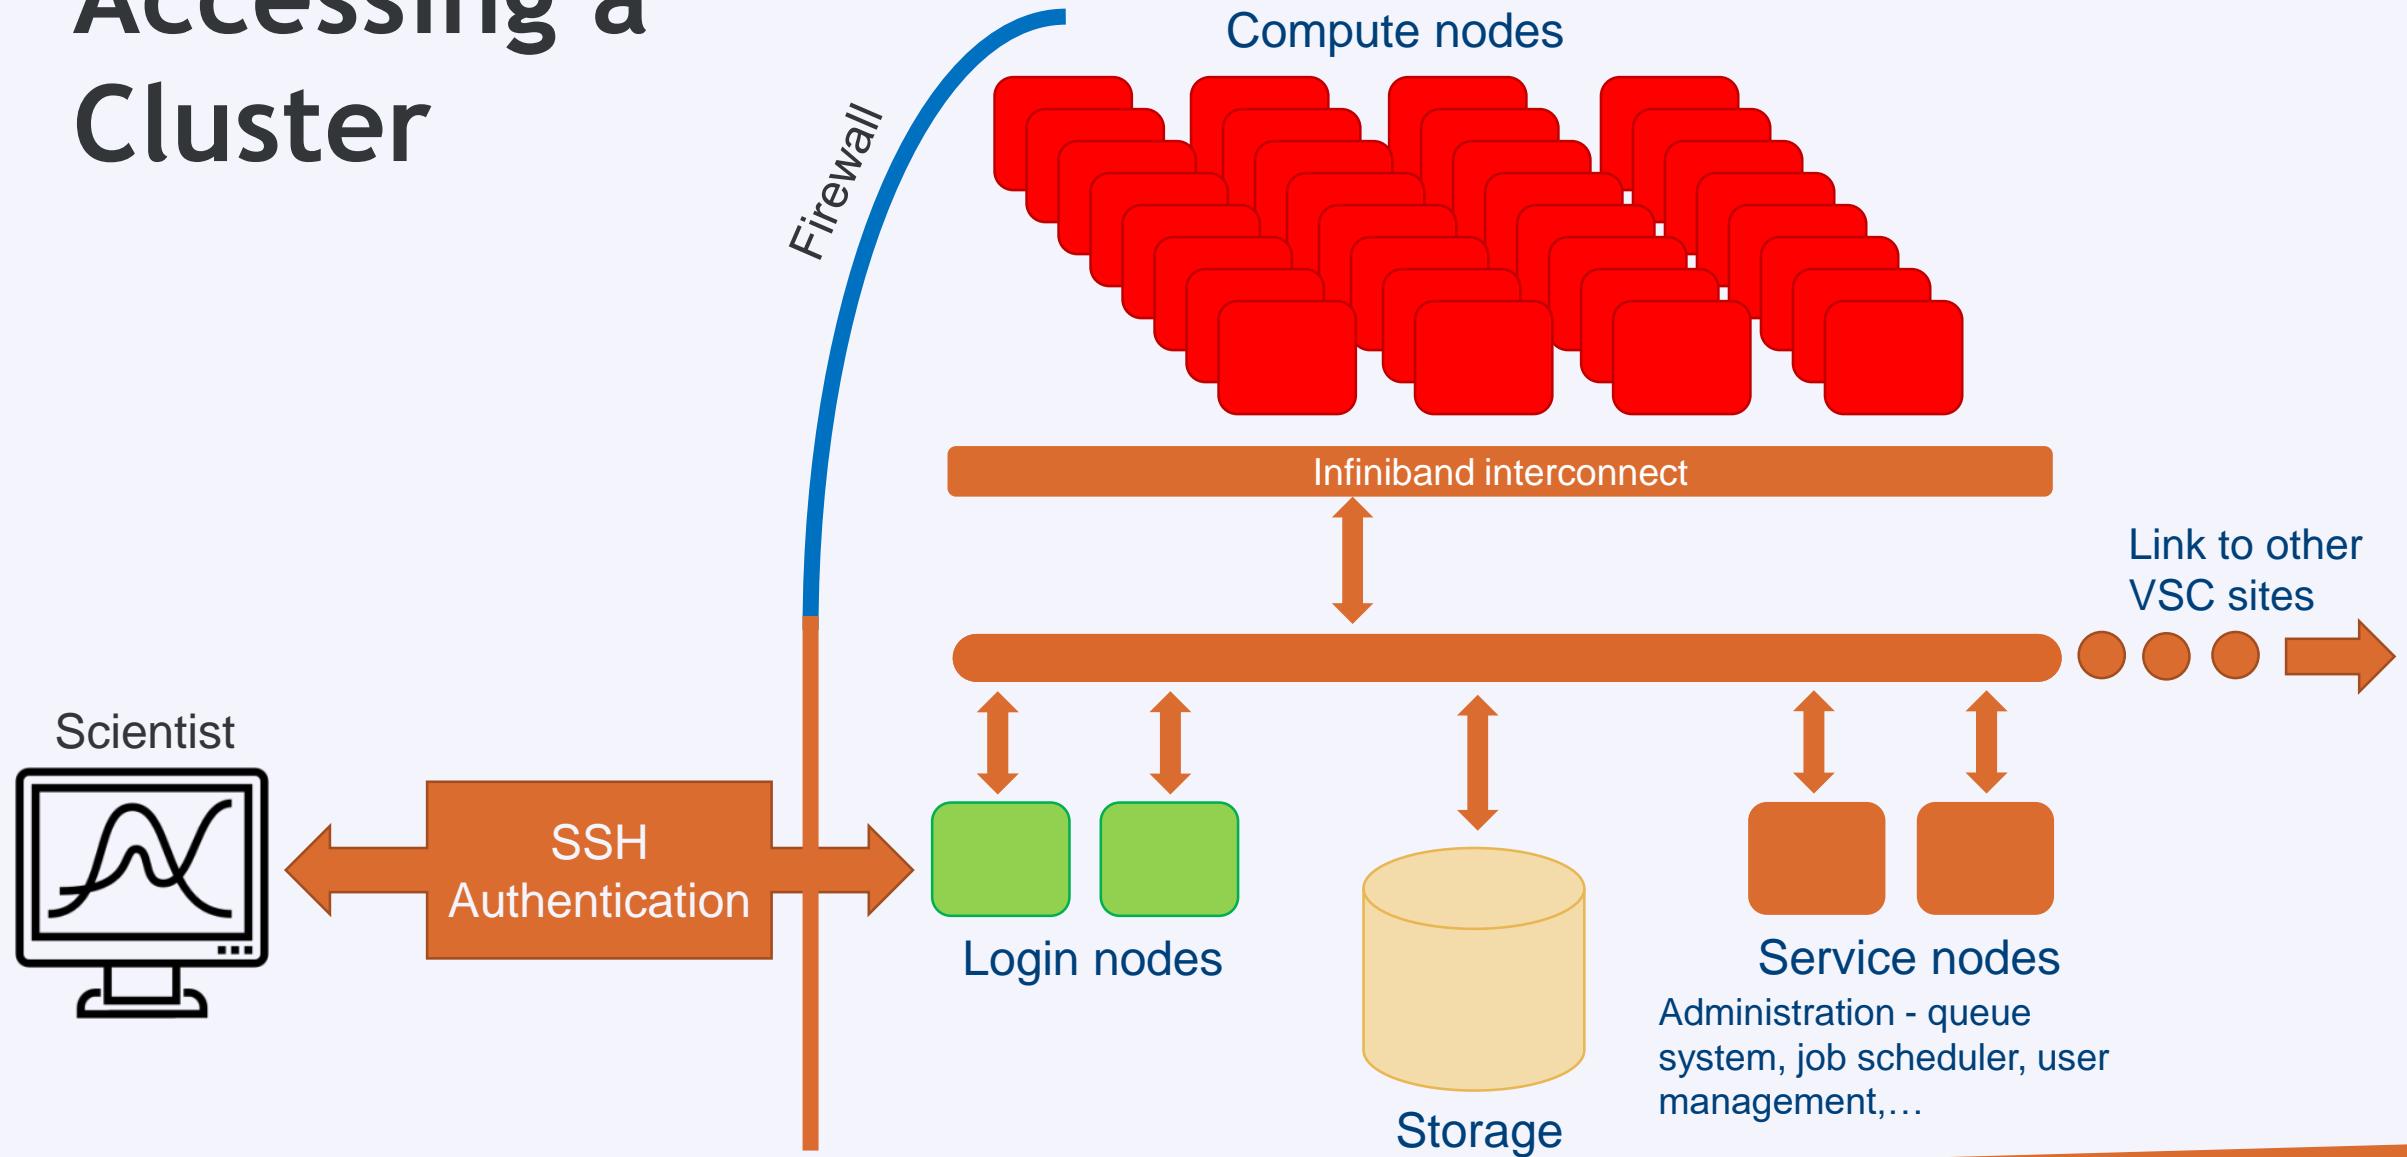
- [Request form for extra storage](#)
- [More information](#)
- Do not use /tmp**  
It is only 10 GB and is reserved for the OS  
and root processes  
Your application can crash if using /tmp
- You are automatically logged into your home  
folder upon login.  
Make sure you immediately go to your other  
storages, e.g.  
`$> cd $VSC_DATA`
- Always check your storage balance using  
`myquota` command

Example

```
$> myquota
file system $VSC_HOME
    Blocks: 1479M of 3072M
    Files: 12934 of 100k
file system $VSC_DATA
    Blocks: 102G of 225G
    Files: 1043k of 10000k
file system $VSC_SCRATCH
    Blocks: 15M of 1.5T
```

# Login Nodes

# Accessing a Cluster



# Using Login Nodes

- To develop and/or compile code and/or software
- To check your storage and credit balance
- To manage jobs (submit, check status, debug, resubmit, ...)
- To move data around
  - within VSC: use data, scratch, staging, archive
  - outside VSC: copy/sync from/to your local storage
- To pre-process or post-process your data/jobs
- To visualize your data
- To share files/folders

Tips

- Login nodes are shared resources
- Do not** execute heavy-lifting tasks (core, memory)
- Instead, submit jobs

Warning

# Login Hosts on Different Machines/partitions

- Windows: [PuTTY](#) or [MobaXterm](#) or NX  
Linux/Mac: [terminal](#) or NX
- To login, you need an active VSC number and a hostname  
\$> ssh -X vscXXXXX@<hostname>

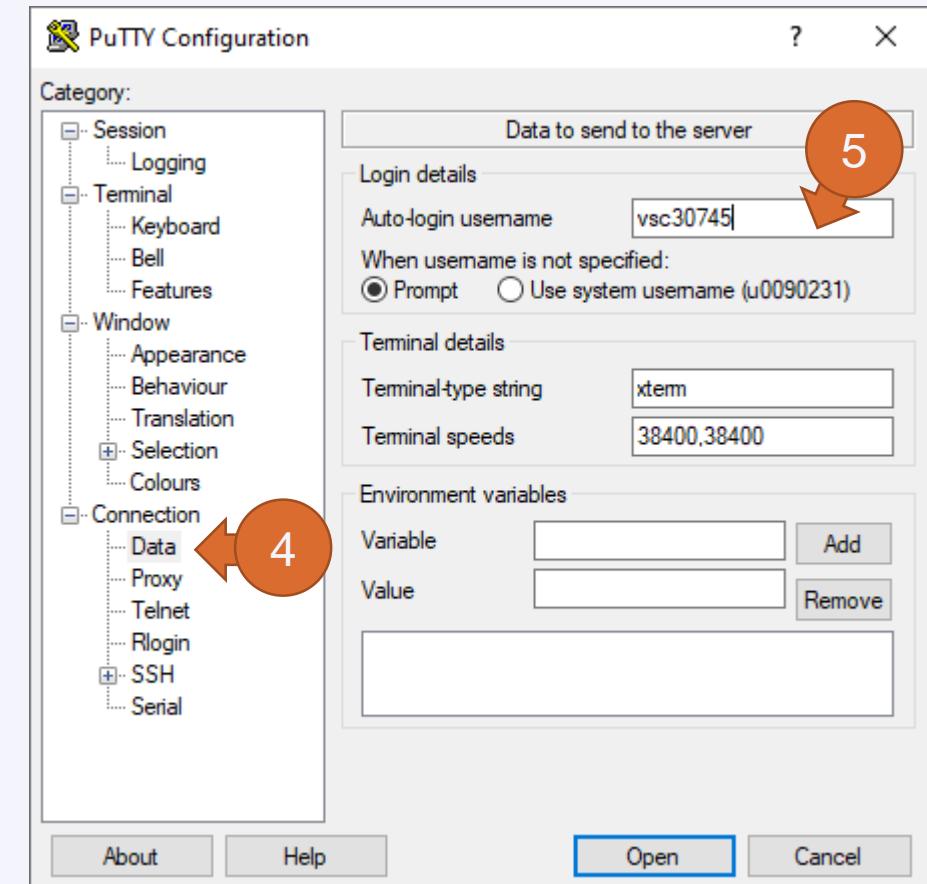
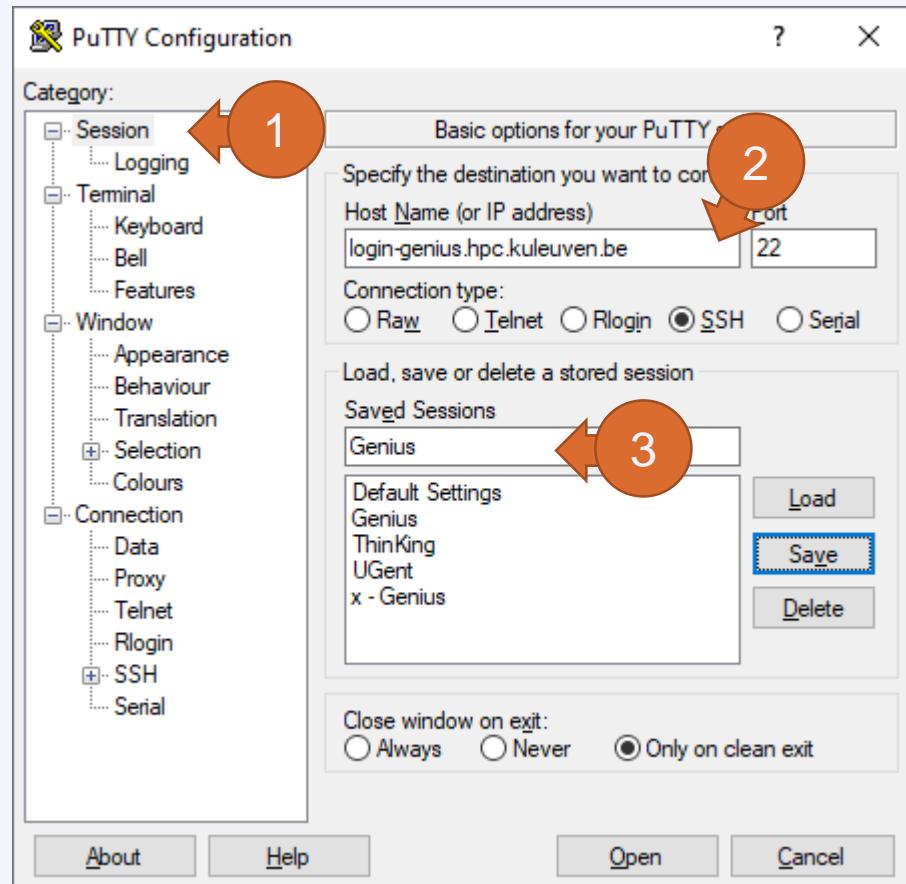
Cluster / Partition	<hostname>	Remark(s)
ThinKing	login-thinking.hpc.kuleuven.be	Recommended
	login7-tier2.hpc.kuleuven.be	Haswell partition
	login8-tier2.hpc.kuleuven.be	
Genius	<b>login.hpc.kuleuven.be</b>	Recommended
	login-genius.hpc.kuleuven.be	
	login{1,2}-tier2.hpc.kuleuven.be	No GPU
	login{3,4}-tier2.hpc.kuleuven.be	Nvidia Quadro P6000
Superdome	Any Genius login node (above)	module load superdome; partition=superdome
AMD	Any Genius login node (above)	Partition=amd

## Login Setup

- PuTTY
- Terminal

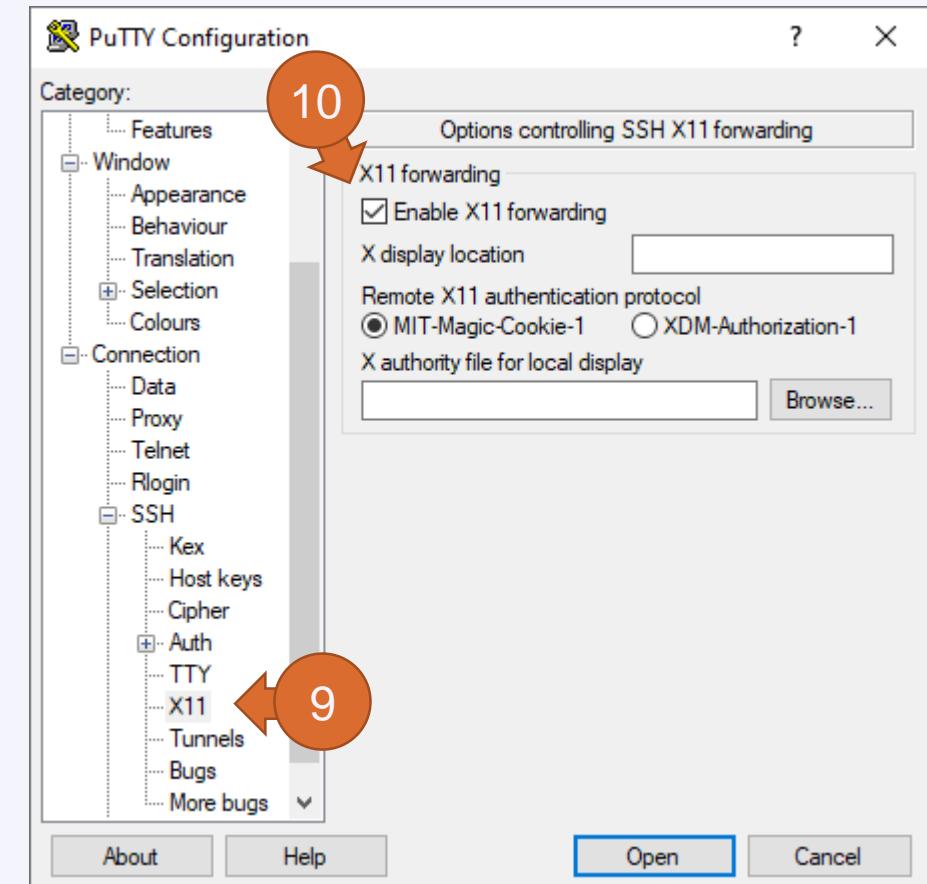
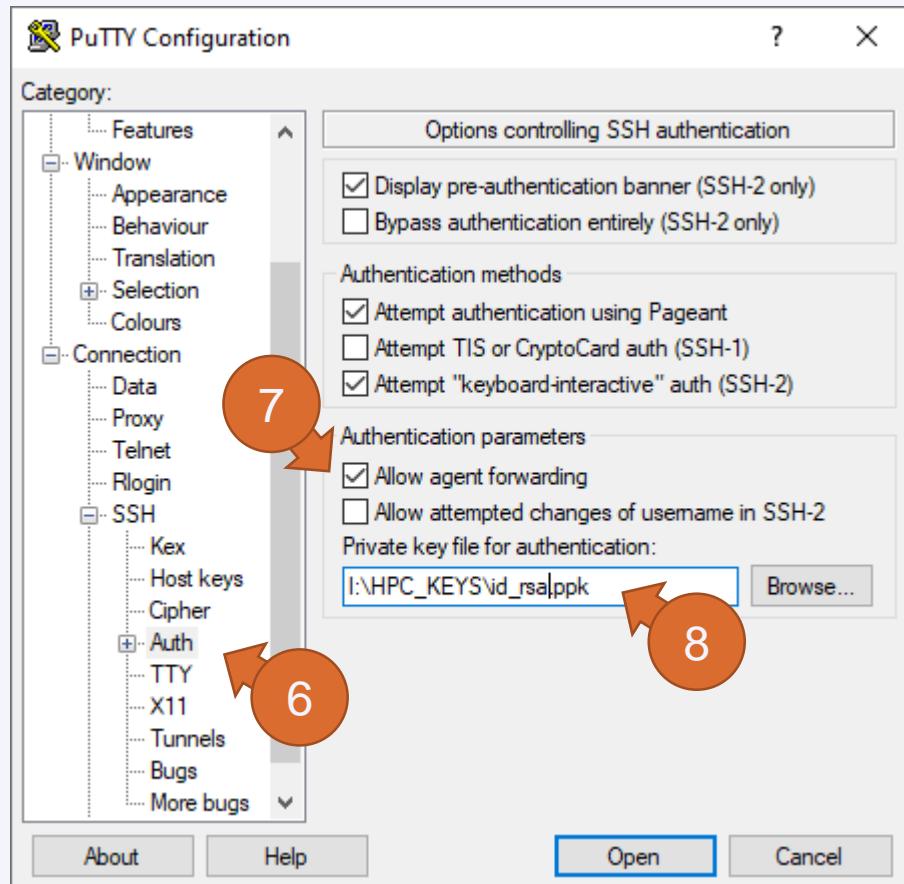
# Setup PuTTY in 12 Clicks

(Windows only)



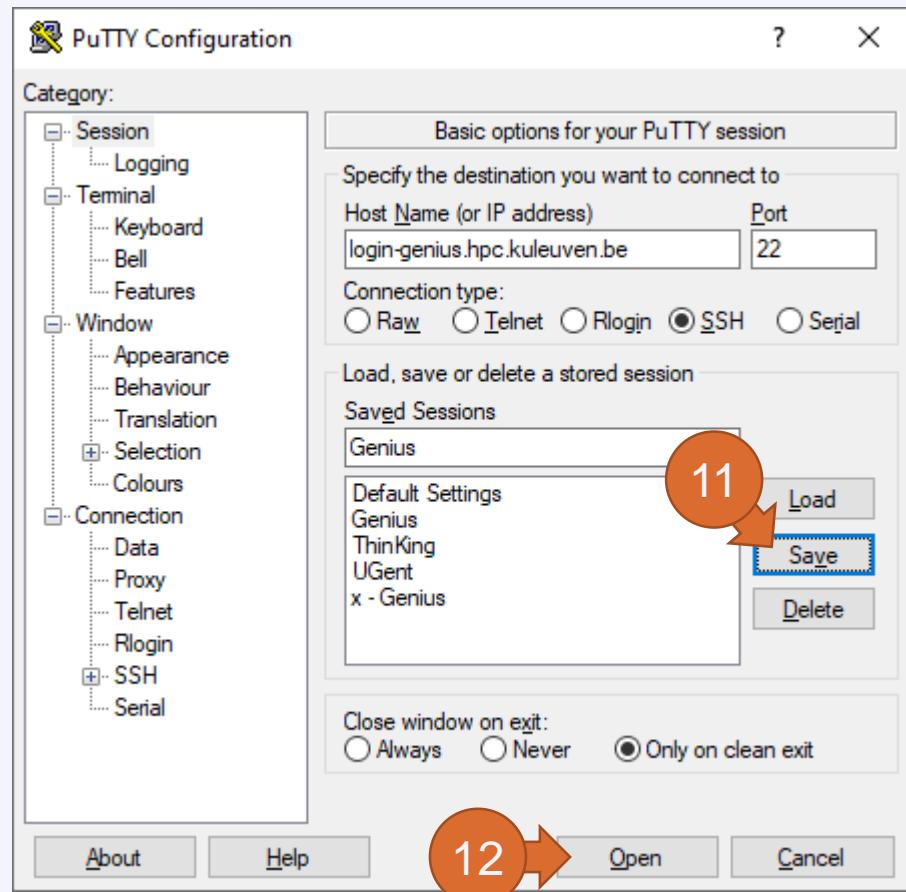
# Setup PuTTY in 12 Clicks

(Windows only)



# Setup PuTTY in 12 Clicks

(Windows only)



If PuTTY asks for **password**, exit immediately, and check the path to your private key. Else you will be blocked for 24h.

A screenshot of a PuTTY terminal window titled 'login4-tier2.hpc.kuleuven.be - PuTTY'. The window displays the following text:  
Using username "vsc30745".  
Authenticating with public key "rsa-key-20170705"  
Passphrase for key "rsa-key-20170705":

A screenshot of a PuTTY terminal window titled 'login4-tier2.hpc.kuleuven.be - PuTTY'. The window displays the following text:  
Using username "vsc30745".  
Authenticating with public key "rsa-key-20170705"  
Passphrase for key "rsa-key-20170705":  
Last login: Wed Jan 9 15:23:52 2019  
[User prompt for password]  
Deze server wordt (deels) met Puppet beheerd; alle wijzigingen die niet via Puppet worden aangebracht, riskeren (automatisch) ongedaan te worden gemaakt.  
This server is managed by Puppet; all changes not imp

# Connecting via Terminal

(Linux and Mac)

- Check your SSH Agent. Is your SSH key found?

```
$> ssh-add -l
```

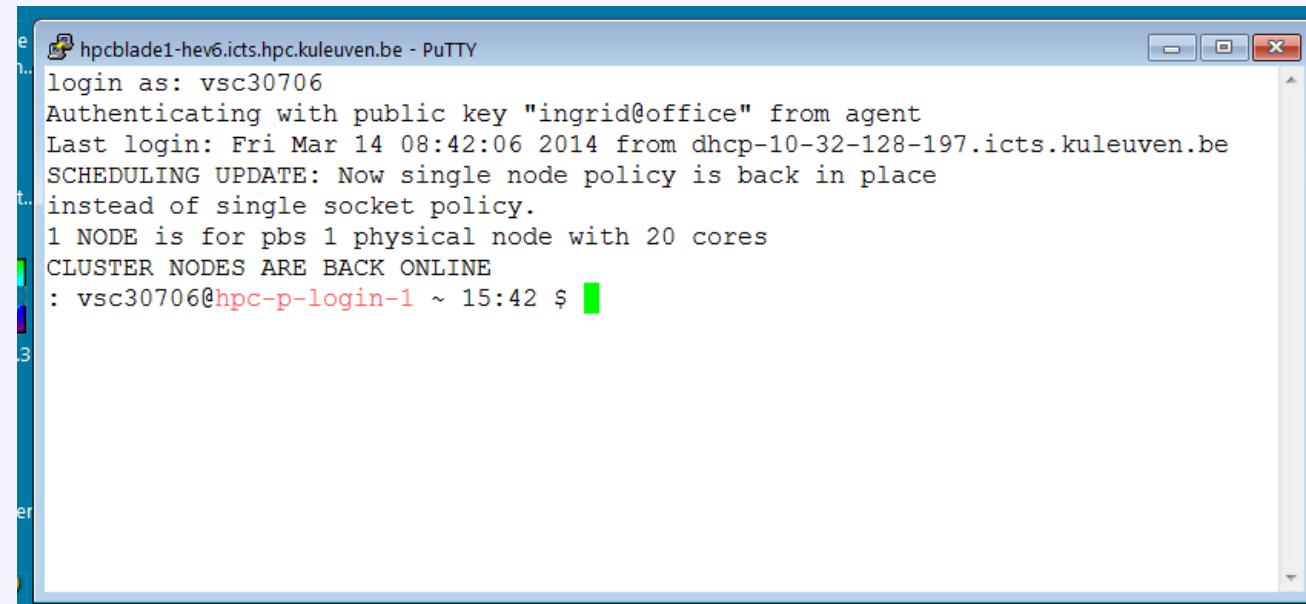
- If your key is not found, add it to the agent:

```
$> ssh-add </path/to/my/private/keyfile>
```

- Try to ssh now

```
$> ssh vscXXXXX@<hostname>
```

If asked for **password**, please stop connecting and contact support, otherwise after a few attempts you will be blocked for 24h.



The screenshot shows a PuTTY terminal window titled "hpcblade1-hev6.icts.hpc.kuleuven.be - PuTTY". The session log displays the following text:  
login as: vsc30706  
Authenticating with public key "ingrid@office" from agent  
Last login: Fri Mar 14 08:42:06 2014 from dhcp-10-32-128-197.icts.kuleuven.be  
SCHEDULING UPDATE: Now single node policy is back in place  
instead of single socket policy.  
1 NODE is for pbs 1 physical node with 20 cores  
CLUSTER NODES ARE BACK ONLINE  
: vsc30706@hpc-p-login-1 ~ 15:42 \$



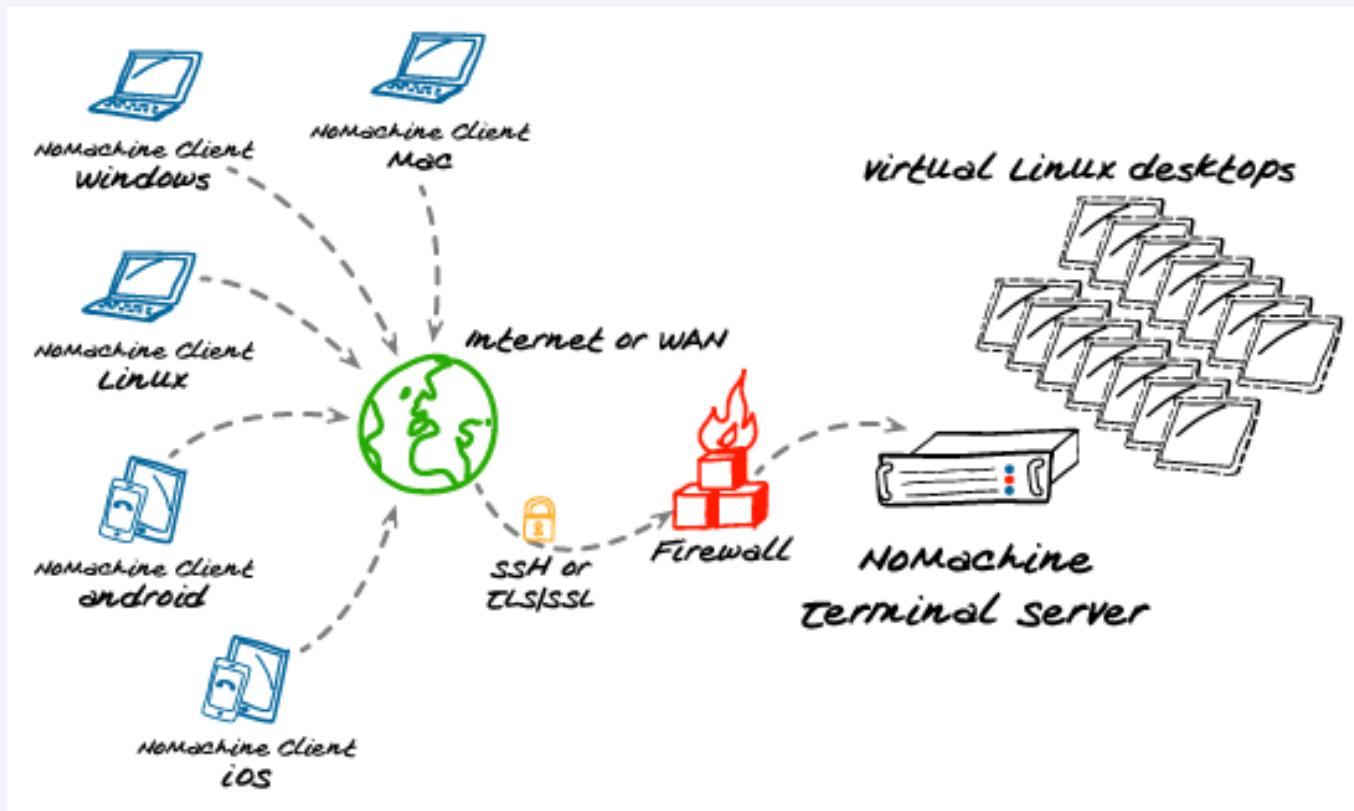
NX

- About
- Setup

# NX – The Graphical Login

NX is:

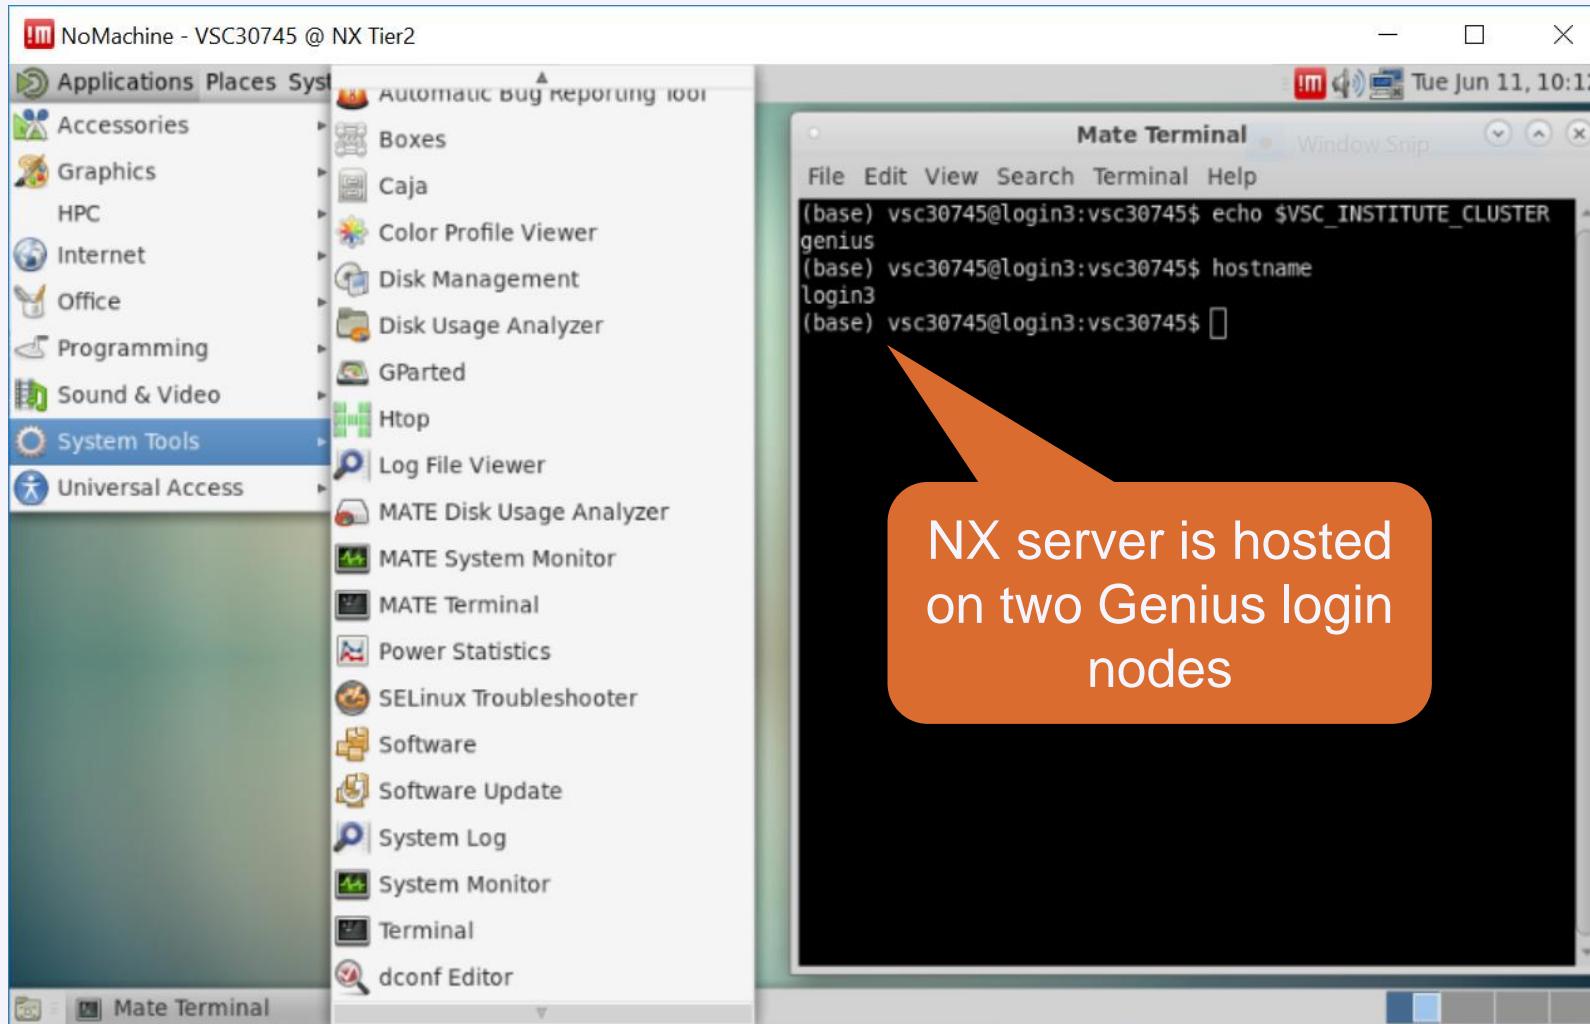
- a graphical remote desktop server
- login to VSC
- start a virtual Linux instance(s)
- Client: [NoMachine](#)
- [Configuration Guide](#)
- Using OpenSSH key  
[Convert your key](#)



# Advantages of NX

- ❑ Graphical login
- ❑ A session stays alive/active even if you close your client (e.g. laptop)
- ❑ Resume a session, and immediately keep working where you left off
- ❑ More interactive jobs
- ❑ Available apps:
  - Internet browser, terminal, ...
  - Text editor, PDF reader, Image viewer, ...
  - Matlab, RStudio, SAS
- ❑ **Remark:**
  - NX is shared among tens of users
  - Do not compute/test your code there**
  - Instead, submit jobs from NX to compute nodes

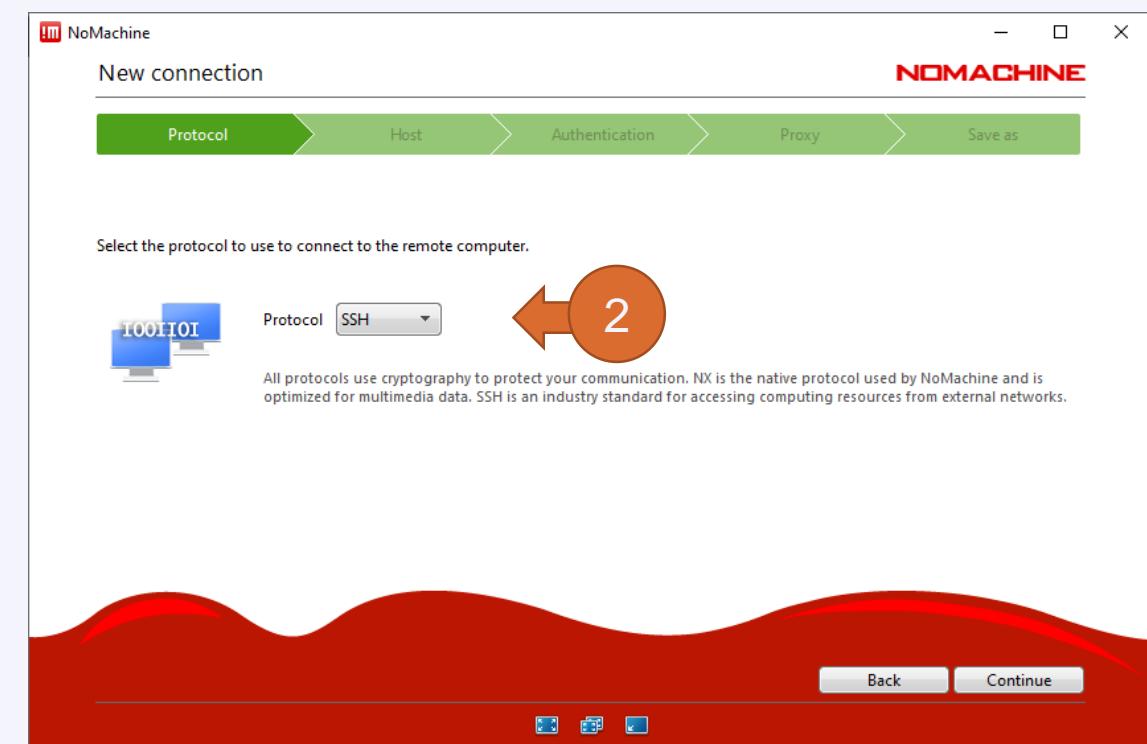
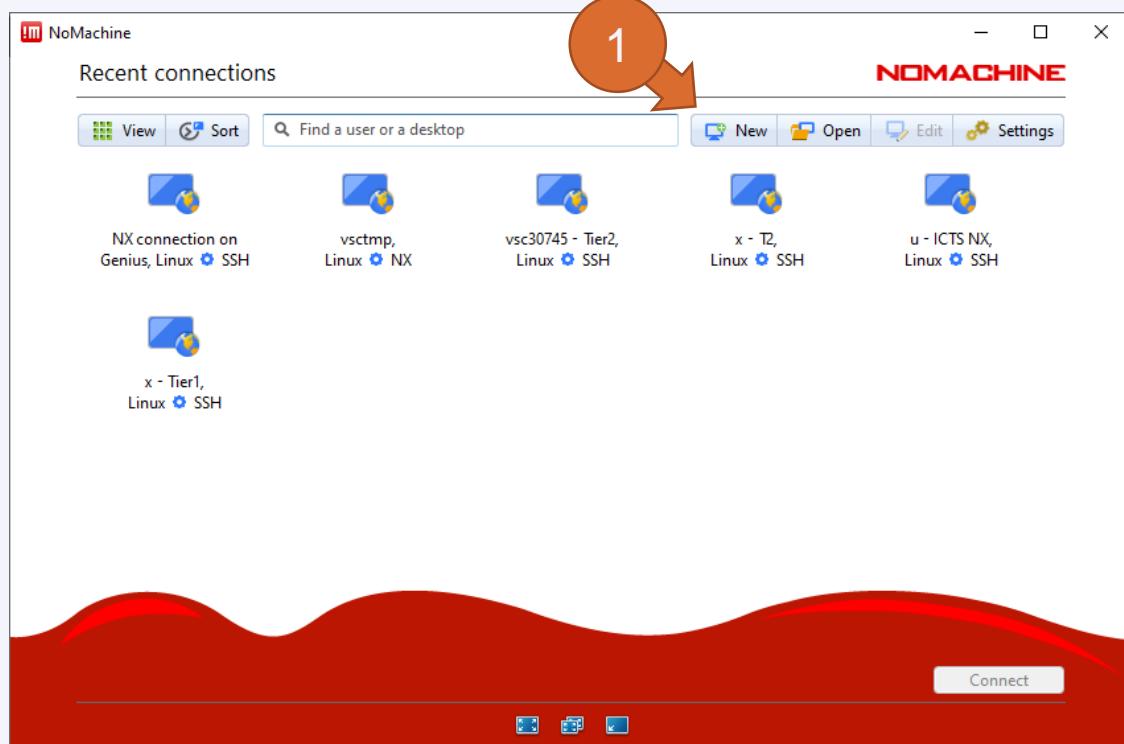
# NX virtual desktop



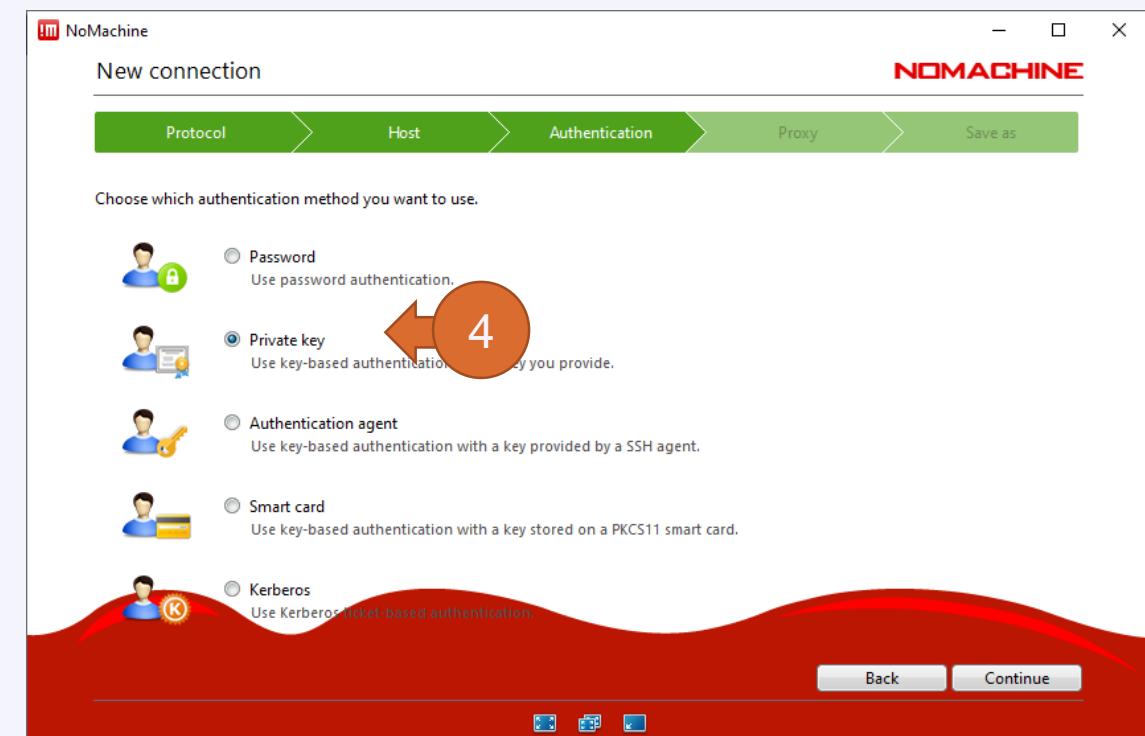
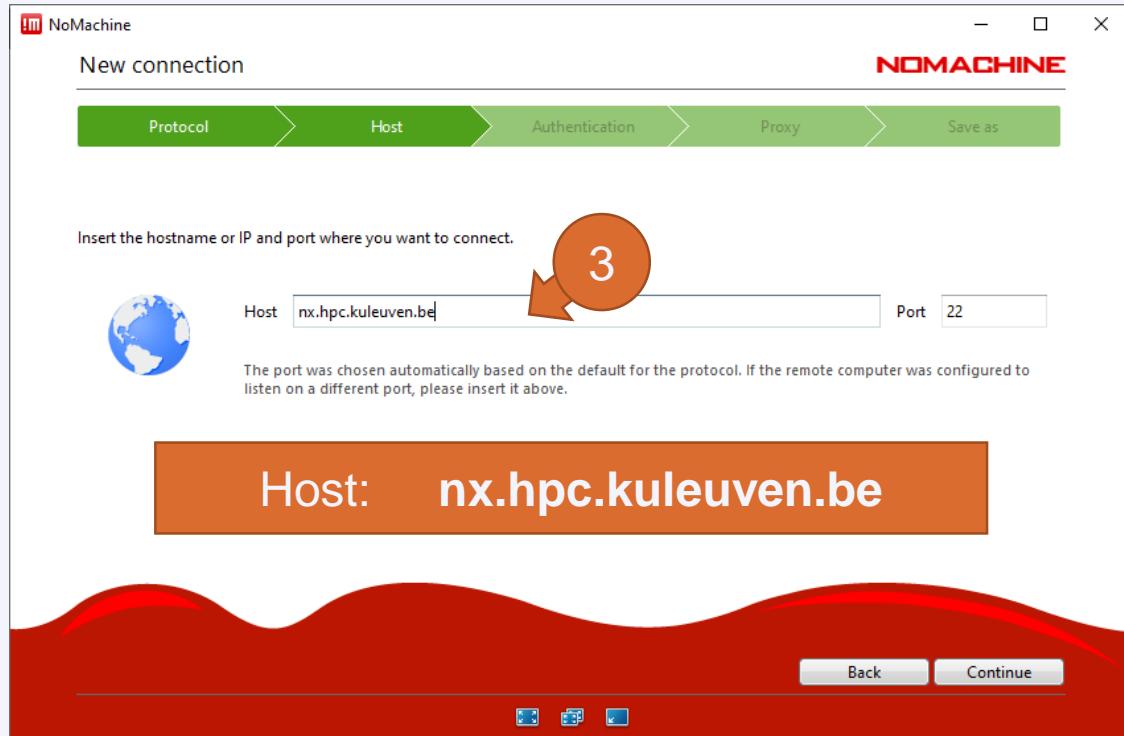
# NX: available software

- Accessories:** Gedit, Vi IMproved, Emacs (dummy version), Calculator
- Graphics:** gThumb (picture viewer), Xpdf Viewer
- Internet:** Firefox
- HPC: Computation:** Matlab (2018a), RStudio, SAS
- Visualisation:** Paraview, VisIt, VMD
- Programming:** Meld Diff Viewer (visual diff and merge tool)
- System tools:** File Browser, Terminal
- Additionally:** Gnuplot (graphing utility), Filezilla (file transfer tool), Evince (PDF, PostScript, TIFF, XPS, DVI Viewer)
- Software launched through modules from Terminal.

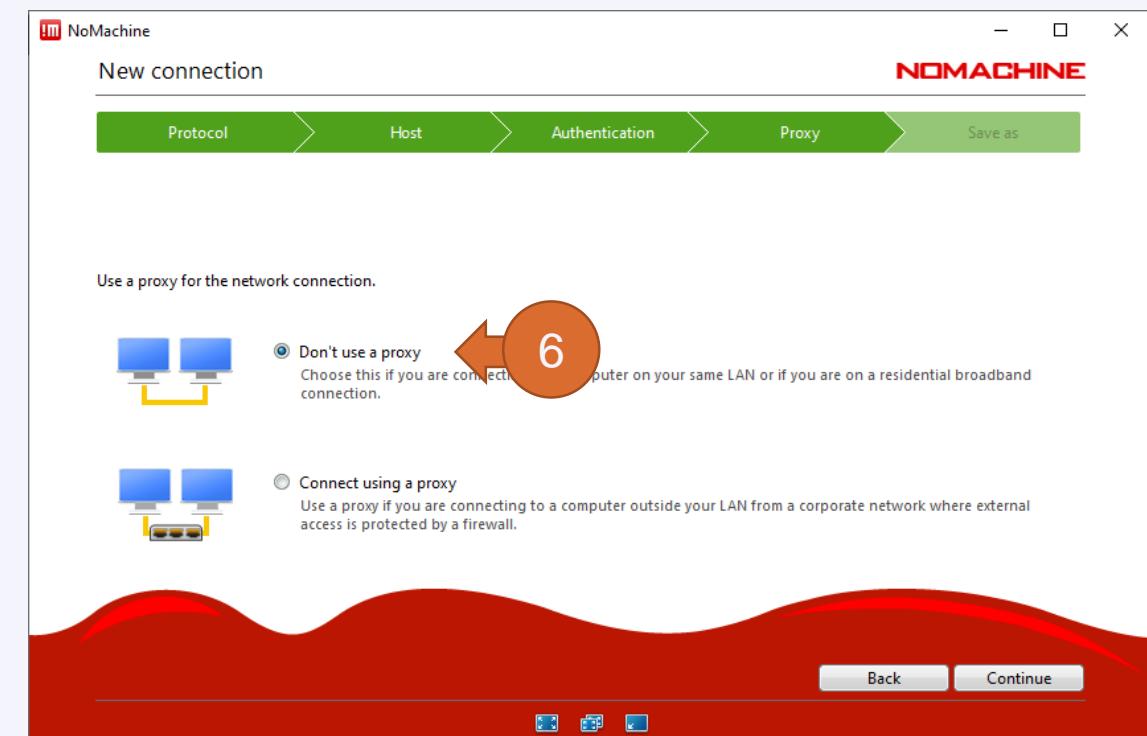
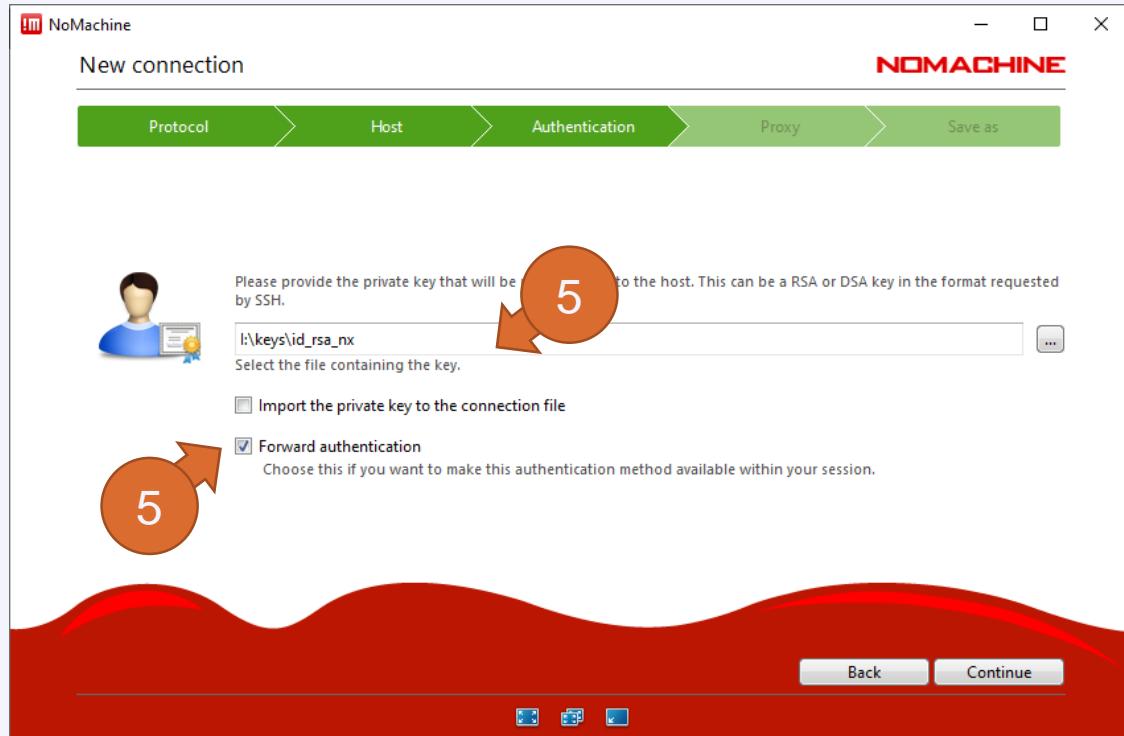
# Setup NX in 10 Steps



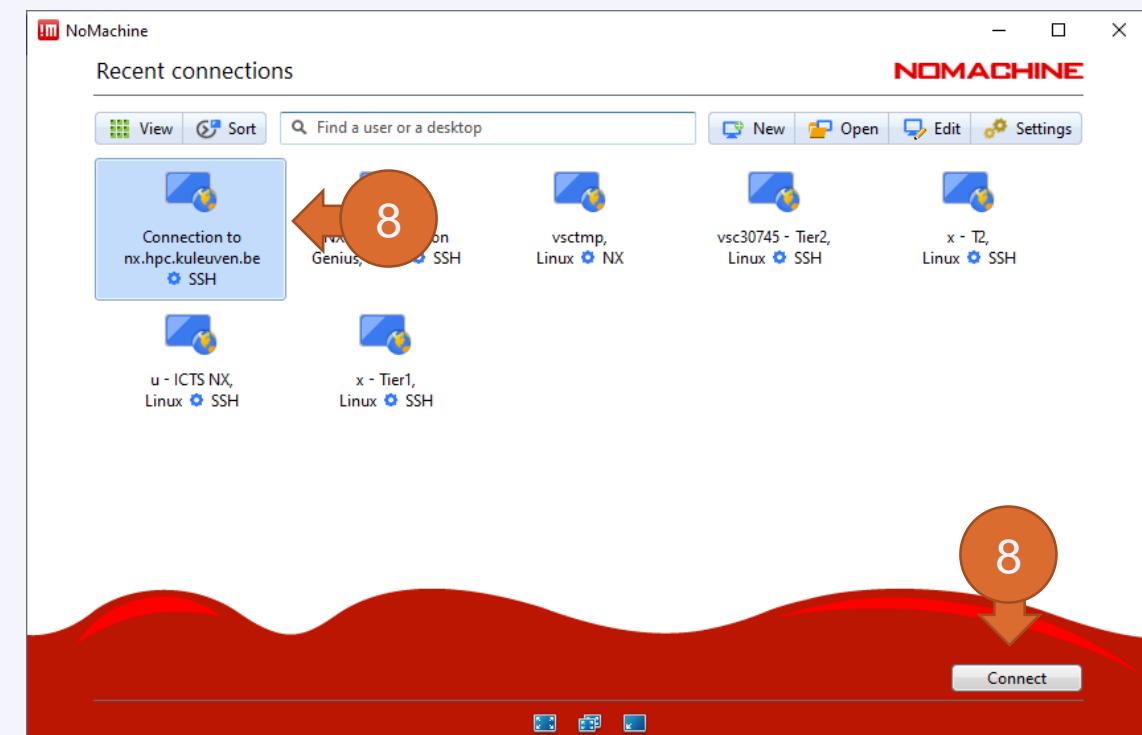
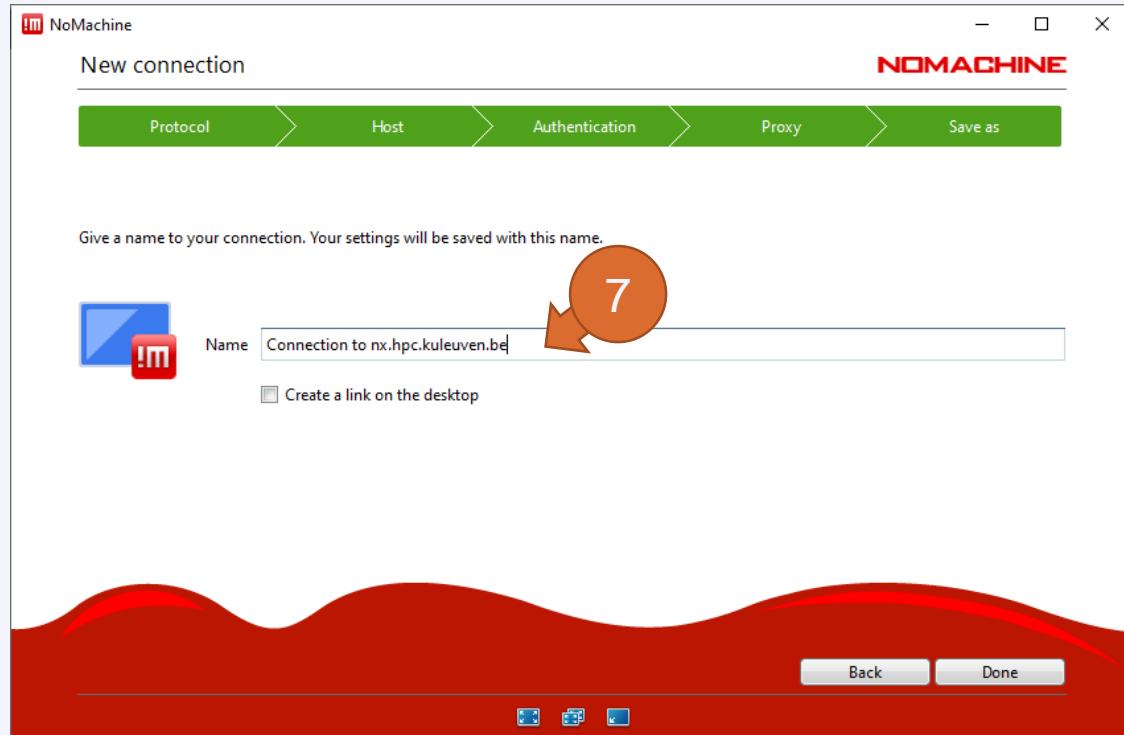
# Setup NX in 10 Steps



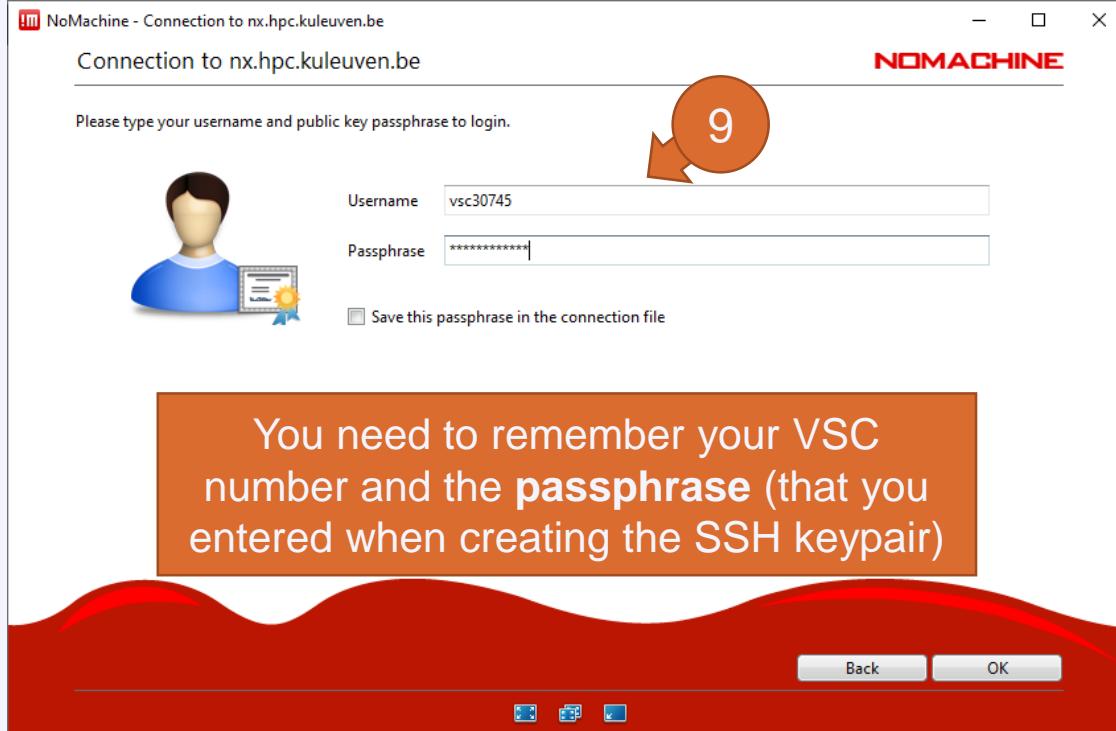
# Setup NX in 10 Steps

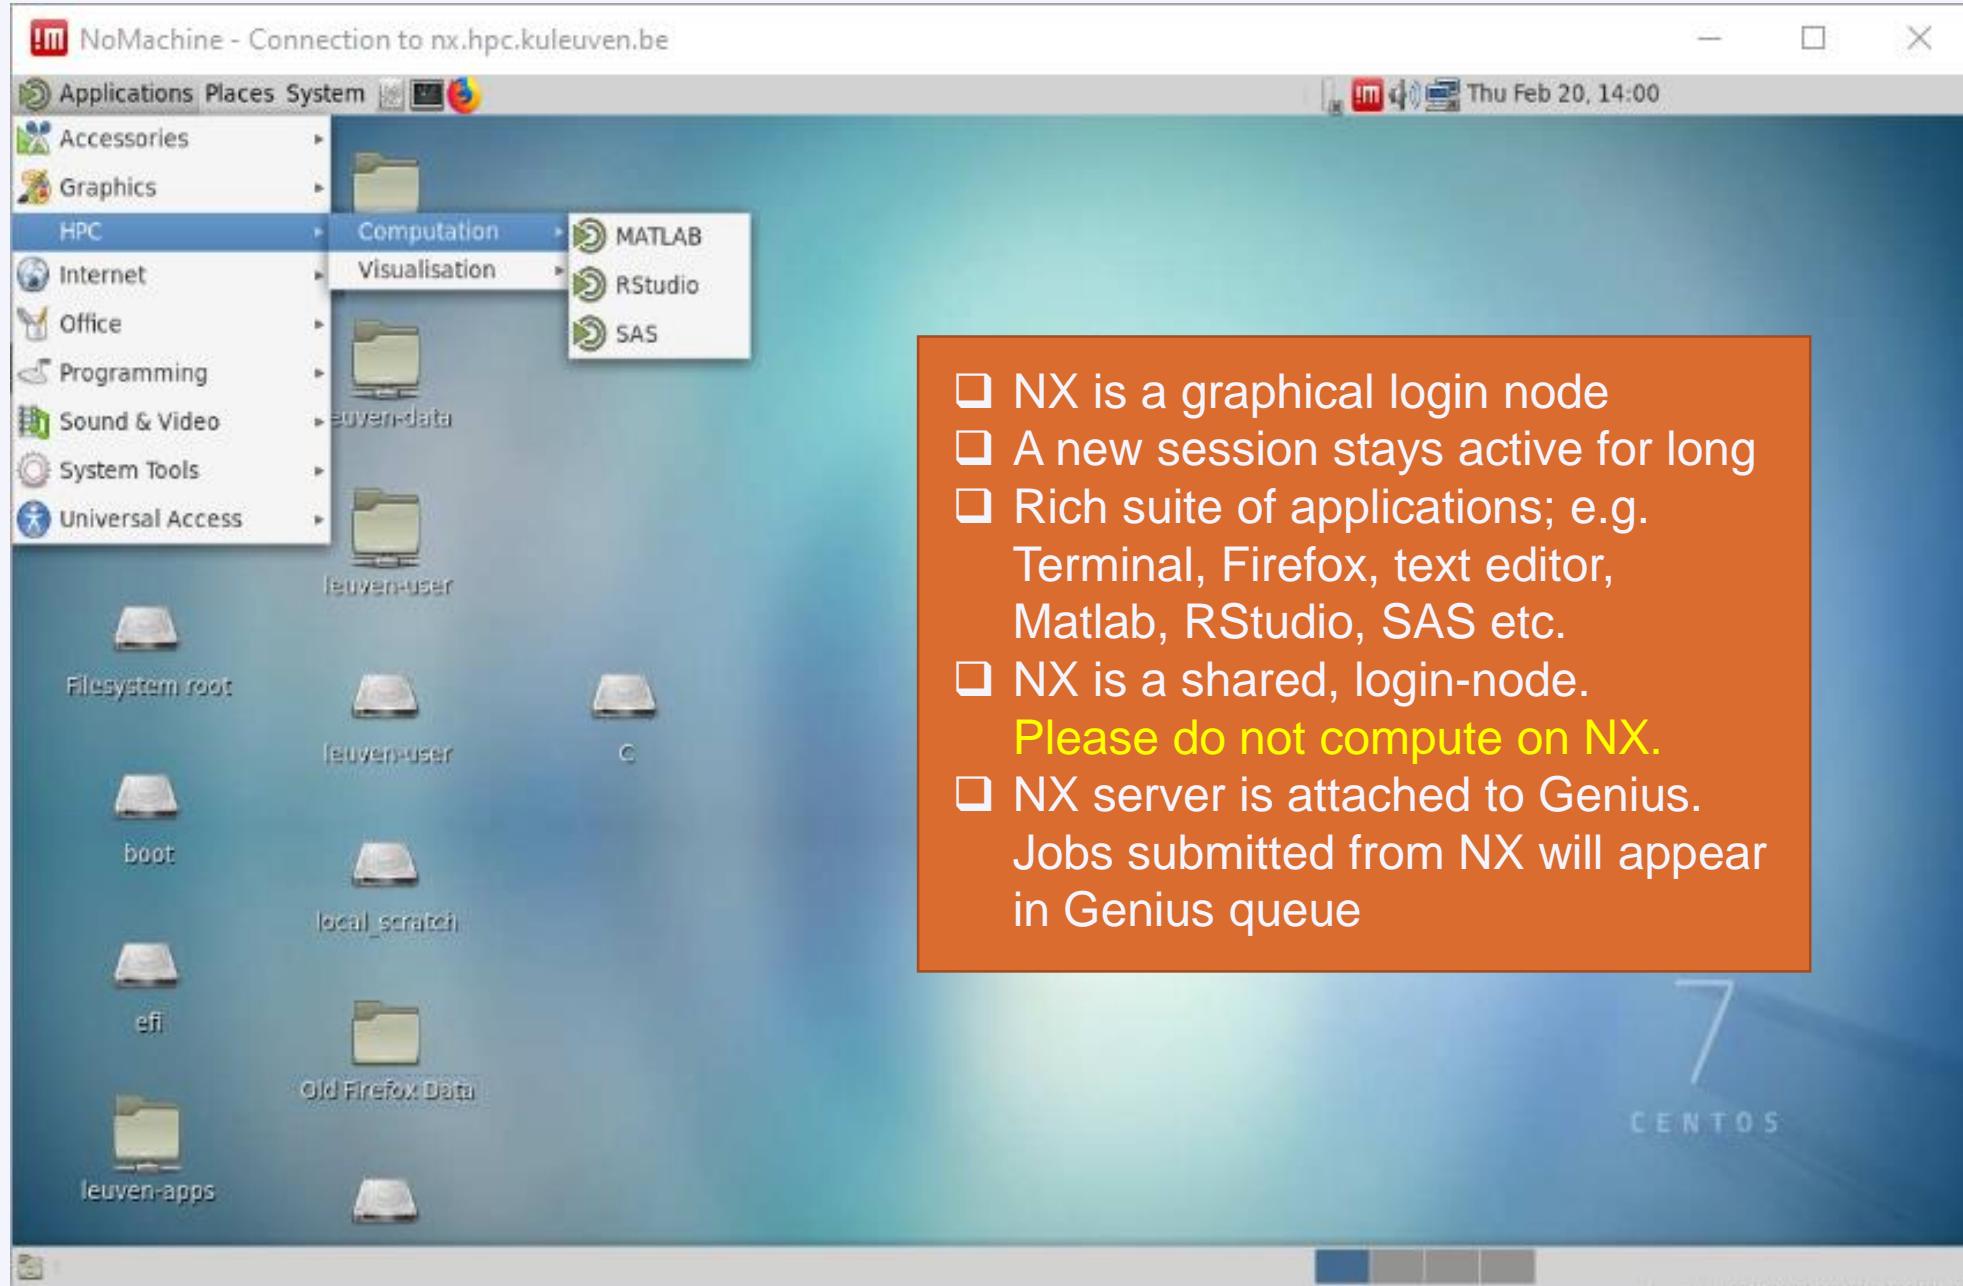


# Setup NX in 10 Steps



# Setup NX in 10 Steps

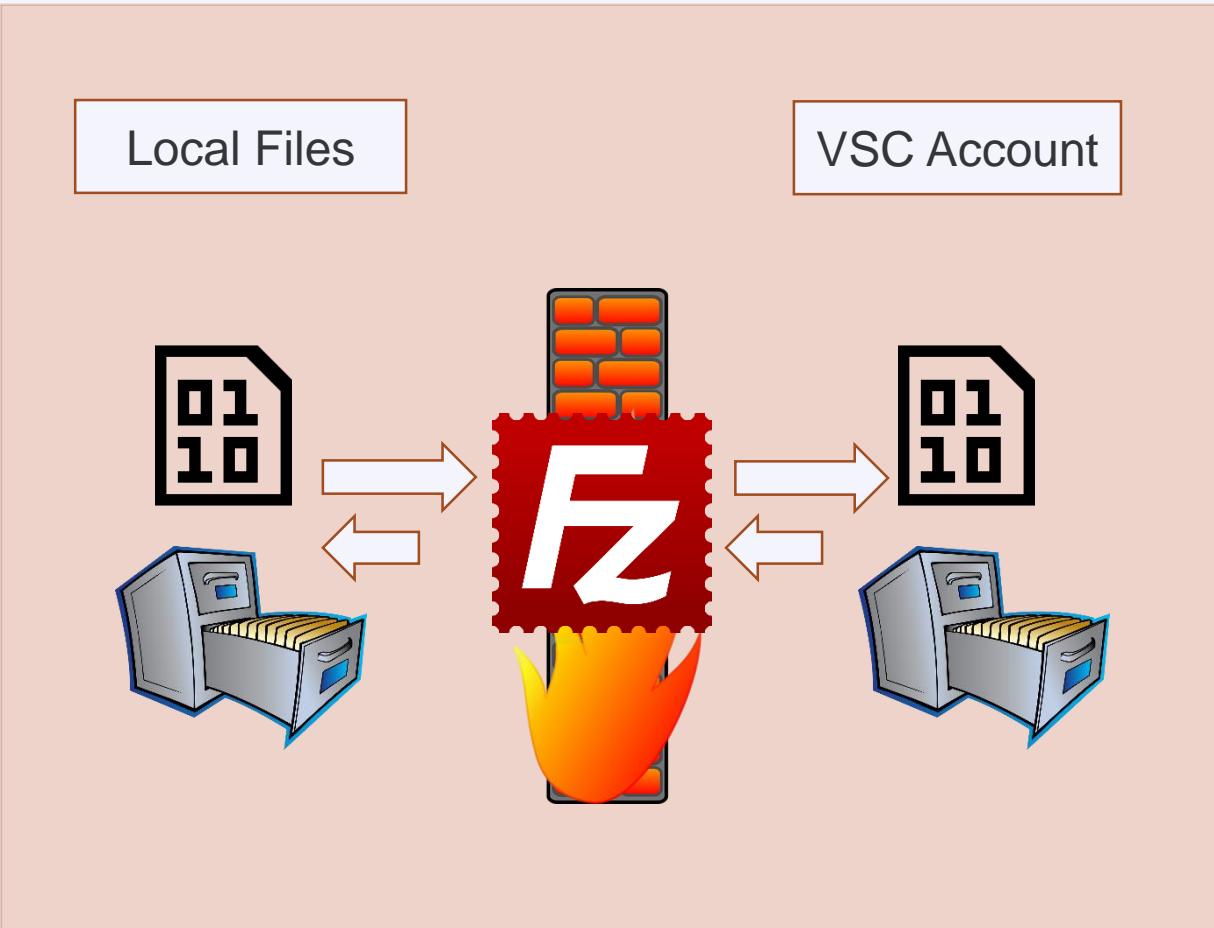




# File Transfer with FileZilla

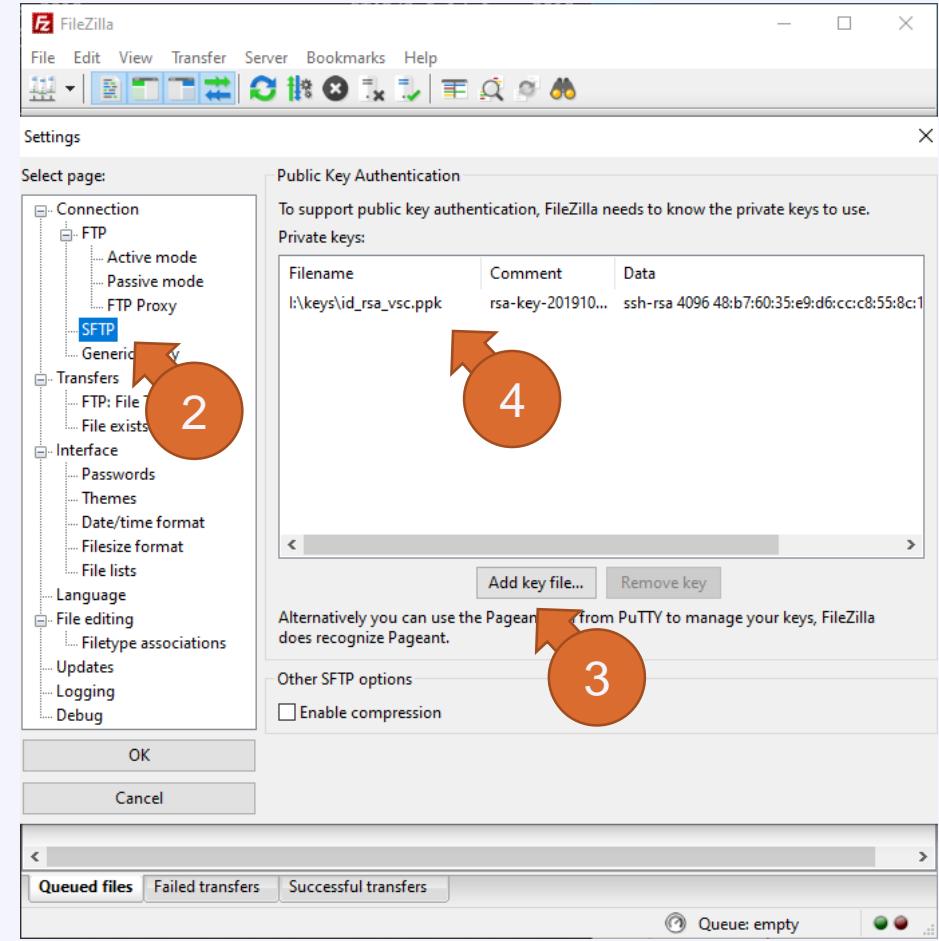
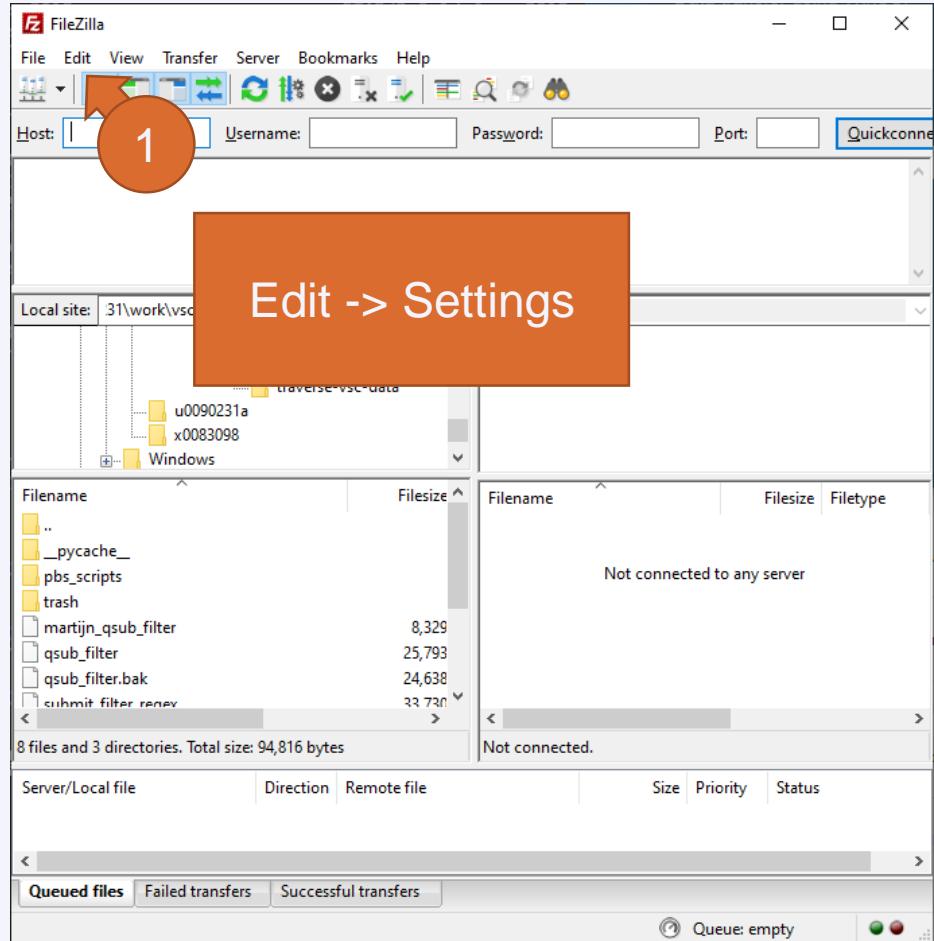


# File transfer

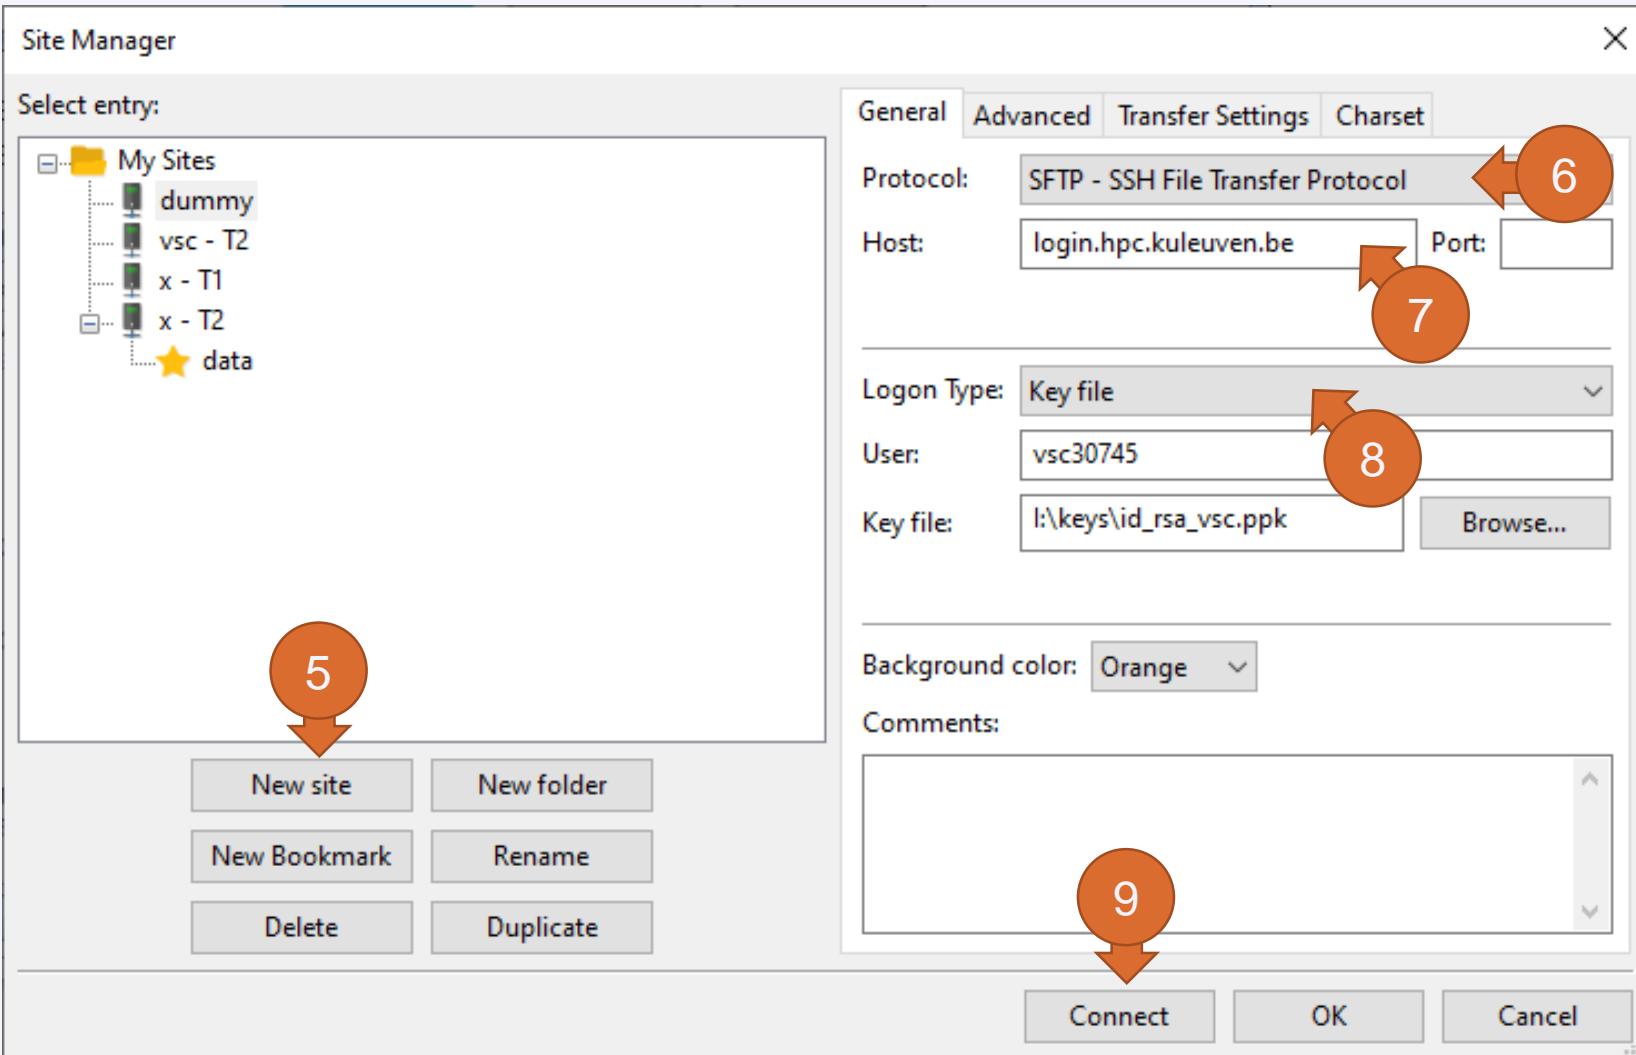


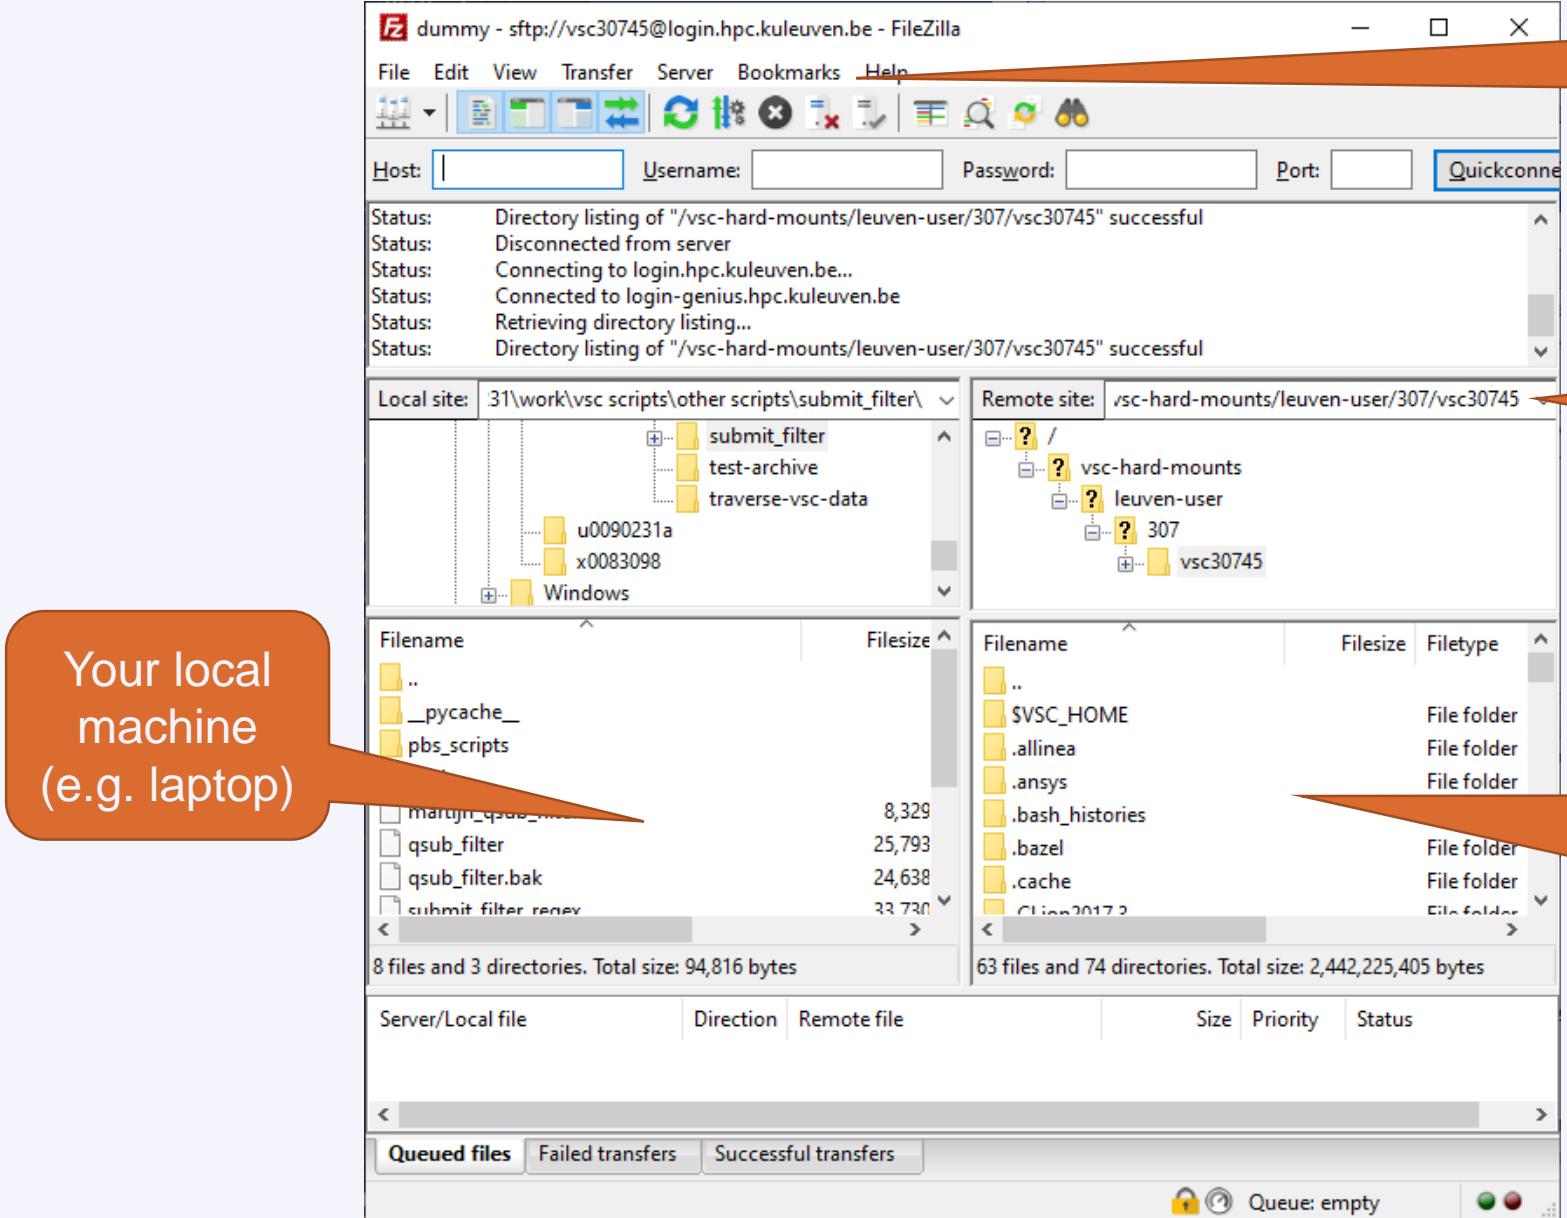
Application	OS
<a href="#">FileZilla</a>	Windows, Linux, Mac
<a href="#">WinSCP</a>	Windows
rsync, scp	Linux, Mac
<a href="#">Globus</a>	Window, Linux, Mac

# File transfer



# File transfer





Your local machine  
(e.g. laptop)

For convenience, you'd better bookmark your data and scratch folders!

Instead, go to your  
\$VSC\_DATA folder, e.g.  
/data/leuven/399/vsc39934

Your VSC storage  
(e.g. \$VSC\_DATA,  
\$VSC\_SCRATCH).

**Note:** by default you arrive at your own \$VSC\_HOME. Copy nothing there!

# Software Stack

# Software: Available Modules

- OS: Linux CentOS 7.6, Kernel 3.10.0-1062.1.2.el7.x86\_64
- Compilers:  
Intel 2018a (icc, icpc, ifort; Intel MPI; Intel MKL)  
FOSS 2018a (gcc, g++, gfortran; OpenMPI; ScaLAPACK, OpenBLAS, FFTW)
- Note: Never mix FOSS and Intel compilers (gives dependency conflict)

Command	Remark
module av	List all installed modules
module av Python	List all Python-related modules
module load Python/3.6.4-intel-2018a	Load a specific module
module list	List all loaded modules and their dependencies
module unload Python/3.6.4-intel-2018a	Unload a module (but dependencies still stay)
module purge	Remove all modules from your work session

# Software: Your Specific Needs

- You can always install your desired software in your \$VSC\_DATA  
Use Intel or FOSS toolchains
- If you cannot, ask us for help
- Python/R packages for AI and ML must be installed by the users themselves
- Install miniconda in your \$VSC\_DATA

```
$> wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh  
$> bash Miniconda3-latest-Linux-x86_64.sh -p $VSC_DATA/miniconda3
```

- Create conda environment per every project

```
conda create -n EnvTF tensorflow-gpu=2.0.0 pytorch
```

- Use your own conda/Python in your jobs/scripts

```
conda activate EnvTF
```



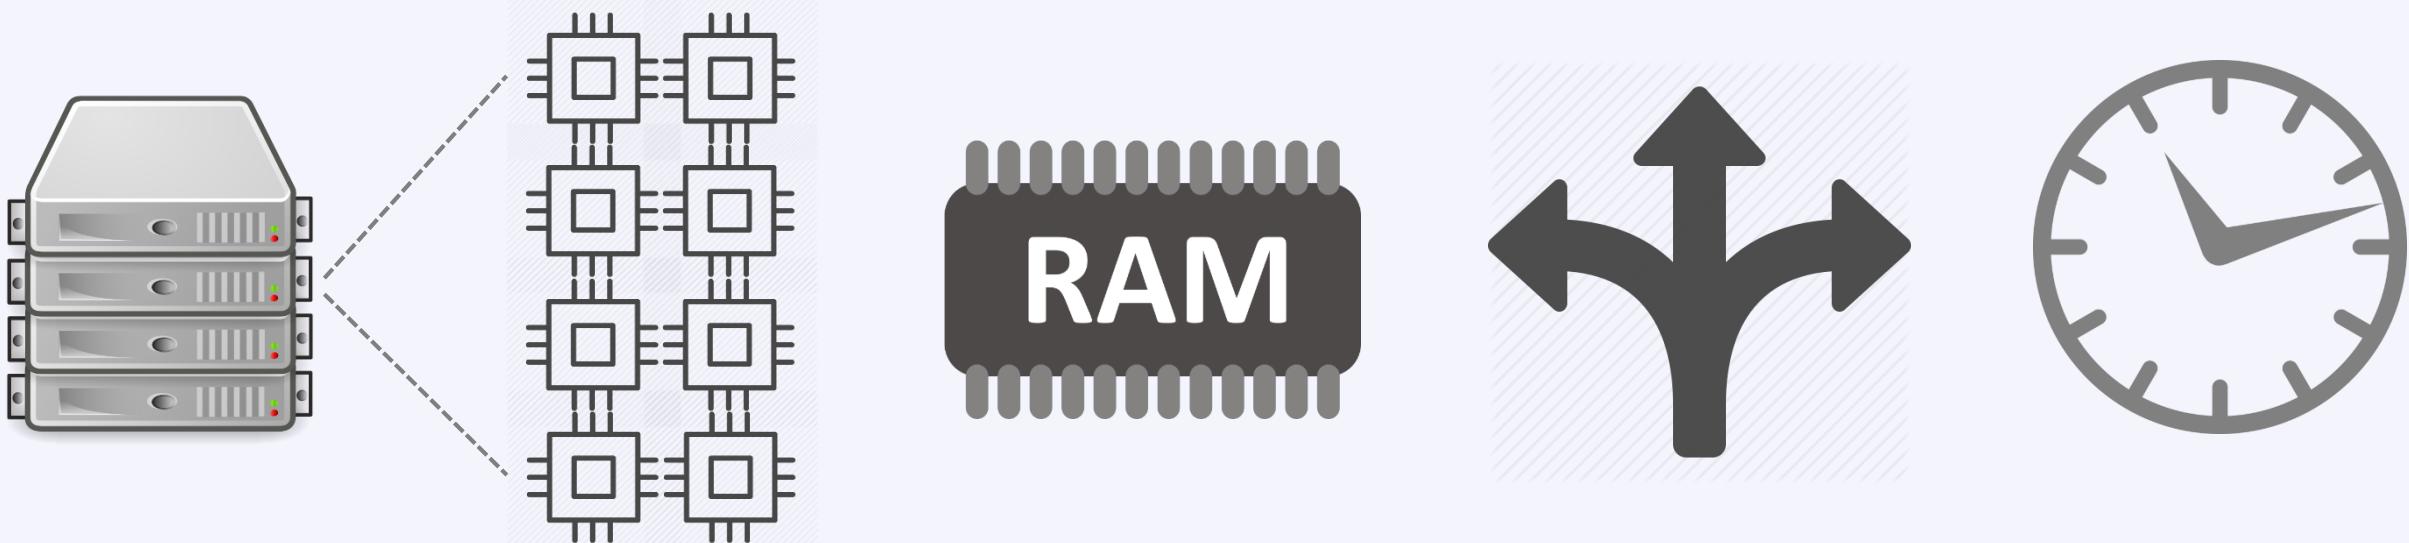
# Starting to Compute

# Extra services

- Request introduction credits  
<https://admin.kuleuven.be/icts/onderzoek/hpc/request-introduction-credits>
- Request project credits  
<https://admin.kuleuven.be/icts/onderzoek/hpc/request-project-credits>  
[https://icts.kuleuven.be/sc/forms/Aanvraagformulier\\_HPC\\_Credits](https://icts.kuleuven.be/sc/forms/Aanvraagformulier_HPC_Credits)
- Extra project credits (to add to existing project)  
<https://admin.kuleuven.be/icts/onderzoek/hpc/extra-project-credits>
- Request extra storage  
<https://admin.kuleuven.be/icts/onderzoek/hpc/hpc-storage>

# Resource Glossary

- Nodes**: how many compute servers to request?
- Cores**: how many cores per node to use?
- Memory requirement**: how much memory each core needs?
- Partition**: GPU, BigMem, Superdome, AMD
- Walltime**: how long to use resources?



# Backfill

Nodes

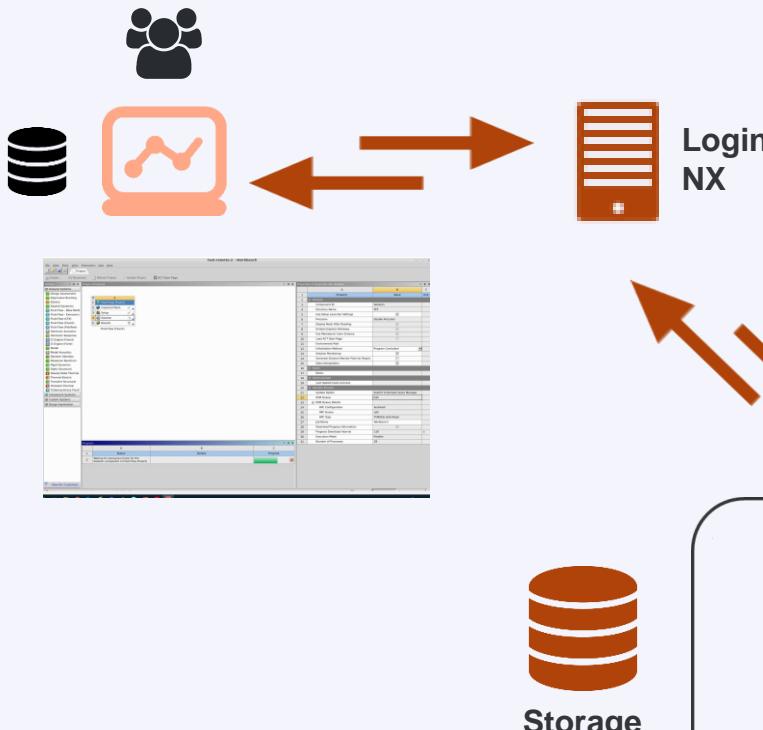


Job queue



# Ways to use the cluster

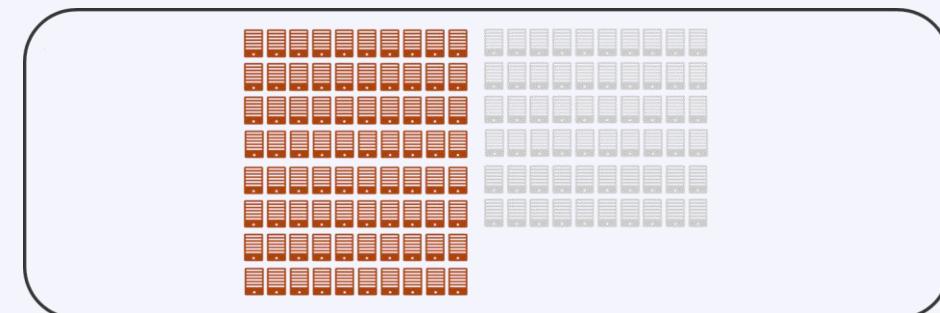
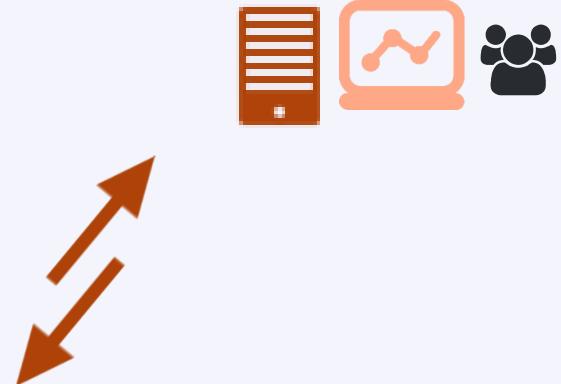
## INTERACTIVE GUI WORKLOAD



## BATCH WORKLOAD



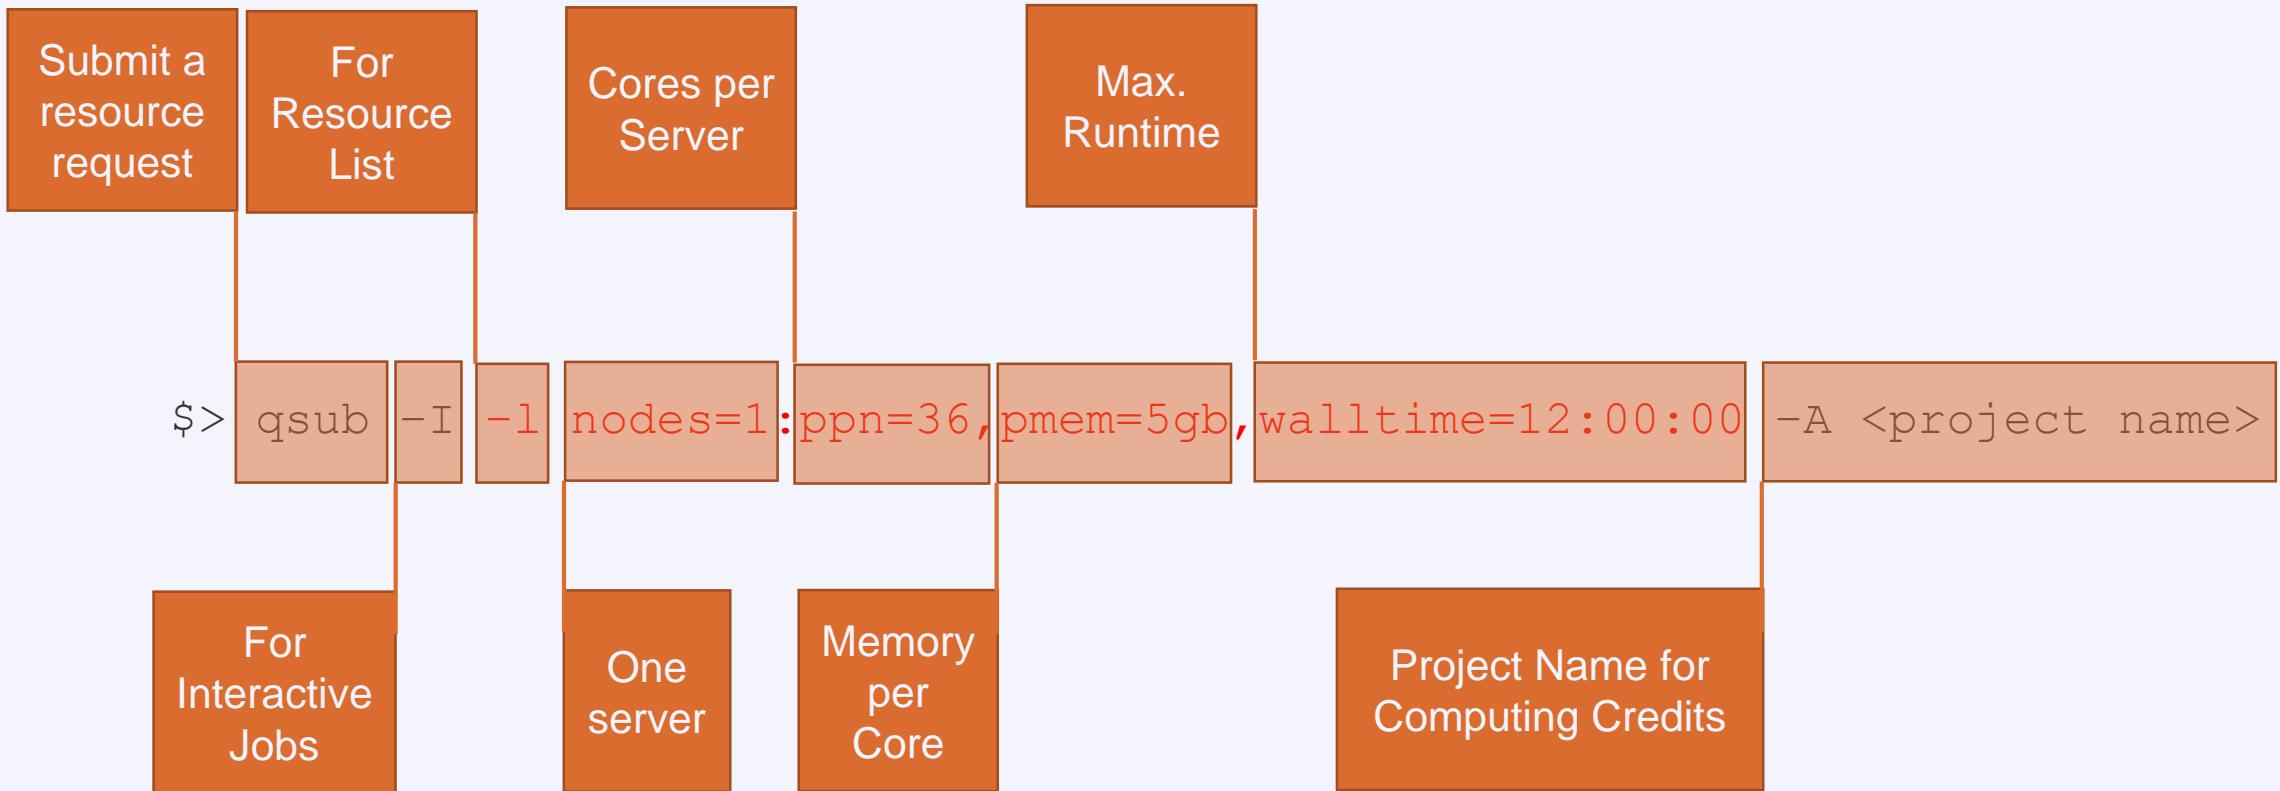
## INTERACTIVE WORKLOAD



# Types of Jobs

- 1) **Batch jobs** are by far the most common, and allow for the most efficient use of the infrastructure. Essentially, a batch job is a bash shell script that is executed on a compute node, and that can spawn a parallel computation on many nodes. These jobs are placed in the queue, and the user can forget about it until it is finished.
- 2) **Interactive jobs** are intended to work on one or more compute nodes interactively. This can be useful in the context of software development for debugging and profiling applications, or for interactive calculations or visualizations. Basically, one gets a shell on one of the compute nodes.
- 3) To get GPU(s), you can use **JupyterHub** with your internet browser

# Requesting Compute Resources Interactively



## Remarks

- Specifying `<project name>` for credits is mandatory, e.g.: `-A lp_hpcinfo_training`
- If you request 2000 introductory credits, you use them like: `-A default_project`
- Implicit defaults are `nodes=1:ppn=1,pmem=5gb,walltime=1:00:00`

# Interactive Jobs

- Interactive job: 1 core for 1 hour (default)

```
$> qsub -I -A lp_hpc
```

- Interactive job with X-forwarding

```
$> qsub -I -X -A lp_hpc
```

- Ask 2 SkyLake nodes, 36 cores each, 4 GB RAM per core for 12 hours

```
$> qsub -I -l nodes=5:ppn=36:skylake,pmem=4gb,walltime=12:00:00 -A lp_hpc
```

- Ask 2 CascadeLake nodes, 36 cores each, 4GB RAM per core for 30 min

```
$> qsub -I -l nodes=2:ppn=36:cascadelake,pmem=4gb,walltime=30:00 -A lp_hpc
```

- Request fraction of a node

```
$> qsub -I -l nodes=1:ppn=8:cascadelake -A lp_hpc
```

# Batch Workload

Job Script

- Create a new text (ASCII) file
- Give it a good name!  
E.g. simulation.pbs
- The first line is the shebang!  
#!/bin/bash
- Put all commands line by line, e.g.  
echo "Hello World!"
- Submit the job to the batch server
- Receive a unique JobID
- Error and output files

```
#!/bin/bash  
echo "Hello World!"
```

Command Line

```
$ qsub simulation.pbs -A default_project
```

JobID

**50041238.tier2-p-**  
moab2.tier2.hpc.kuleuven.be

stderr, stdout

```
$ ls simulation*  
simulation.pbs  
simulation.pbs.e50041238  
simulation.pbs.o50041238
```

# Standard Error File

- Always created
- <JobScript>.e<JobID>
- Contains all errors and warnings
- If empty: everything went well
- Always study it
- Address all warnings and errors (if you can)
- Typical error examples ...

stderr

```
$ ls *.e*
simulation.pbs.e50041238
```

Job Crash

```
forrtl: error (78): process killed (SIGTERM)
Stack trace terminated abnormally.
```

Short Walltime

```
PBS: job killed: walltime 432031 exceeded
limit 432000
```

Low Disk Space

```
IOError: [Errno 122] Disk quota exceeded
```

# Standard Output File

- Always created
- <JobScript>.o<JobID>
- Contains all standard output (instead of screen)
- Always study it

```
$ ls *.o*
simulation.pbs.o50041238
```

stdout

time: 3600  
nodes: 1  
procs: 1  
account string: lpt2\_sysadmin  
queue: q1h  
=====  
Hello World!  
=====

Date: Fri Mar 20 14:32:01 CET 2020  
Allocated nodes:  
r26i27n11  
Job ID: 50240161.tier2-p-moab-2.tier2.hpc. ...  
Resource List:  
pmem=5gb,walltime=01:00:00,neednodes=1:ppn=36  
Resources Used:  
cput=00:00:01,vmem=0kb,walltime=00:00:04,mem=0kb,energy\_used=0  
Queue Name: q1h  
-----

time: 4  
nodes: 1  
procs: 1  
account: lpt2\_sysadmin

Output File

Summary

stdout

Resources

summary

# Example: PBS Job Script

```
#!/bin/bash
#PBS -l walltime=00:10:00
#PBS -l nodes=1:ppn=36
#PBS -l pmem=5gb
#PBS -N testjob
#PBS -A lp_hpcinfo_training
#PBS -m abe
#PBS -M my.name@kuleuven.be
```

Shebang

Resource List

```
module load intel/2018a
which icc
```

Module load(s)

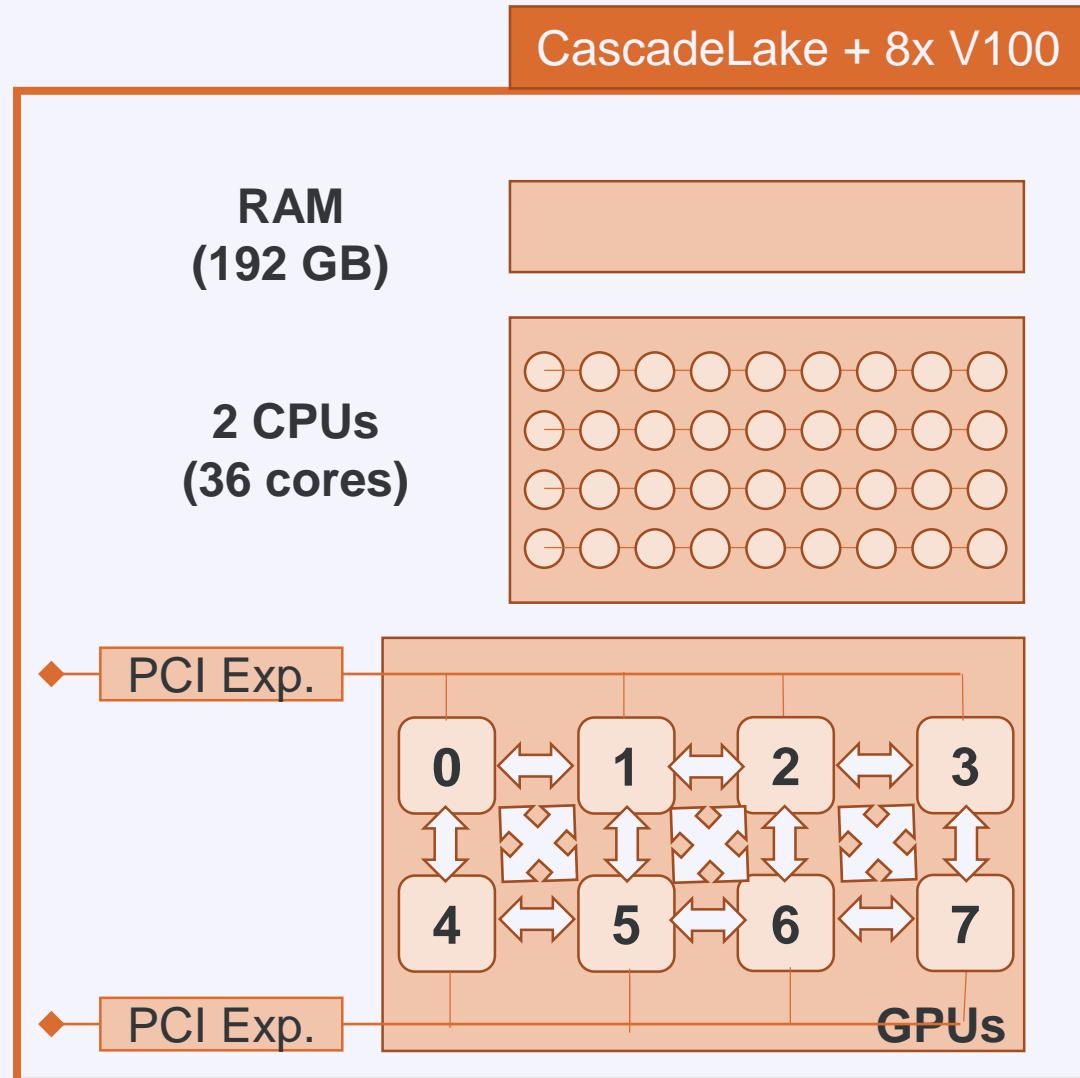
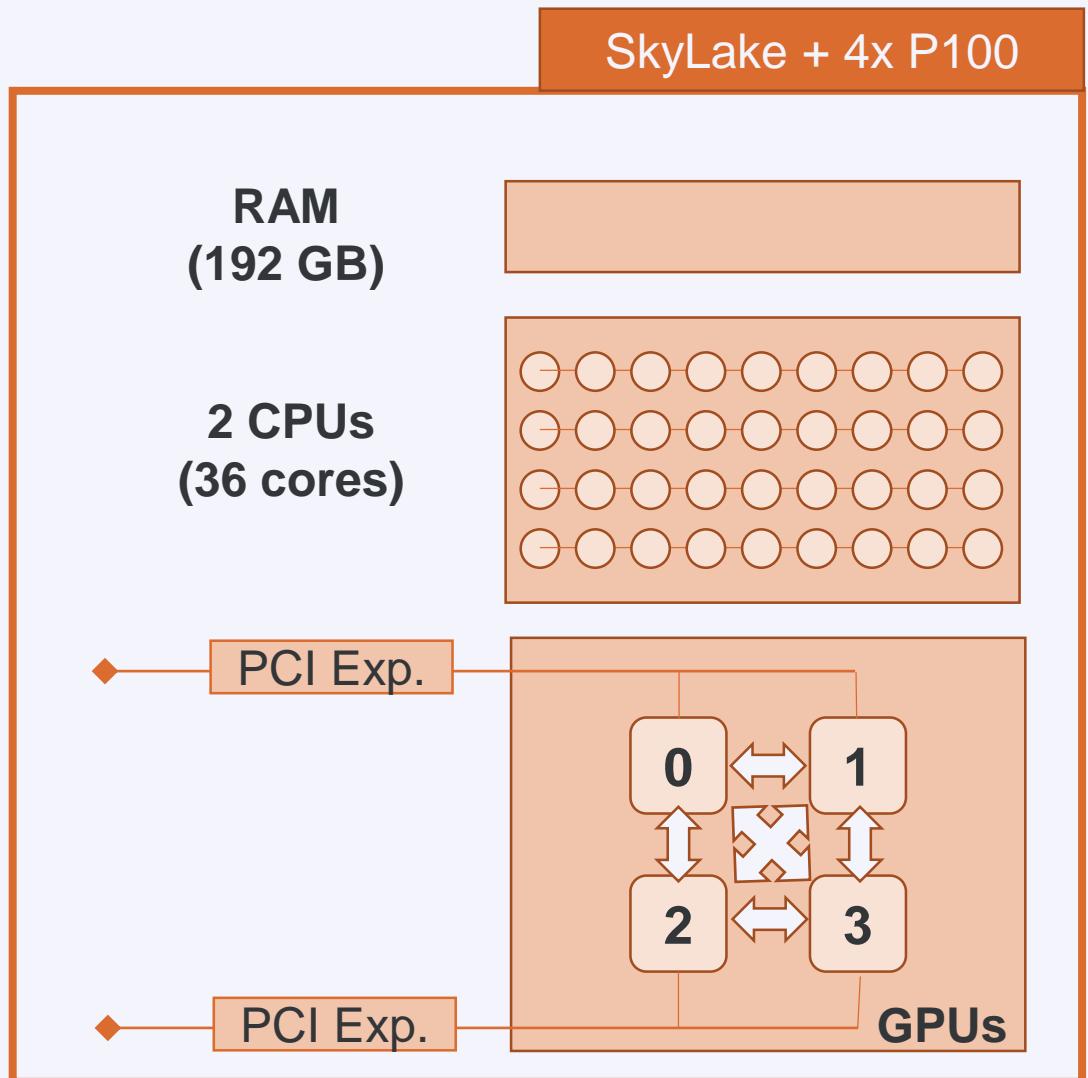
```
cd $PBS_O_WORKDIR
cp /apps/leuven/training/HPC-intro/cpujob.pbs $VSC_SCRATCH
touch output.log
echo I am done
```

Execute commands

Move data

# GPU Partition

# GPU Nodes



# Using P100 GPU(s)

Using 1 GPU

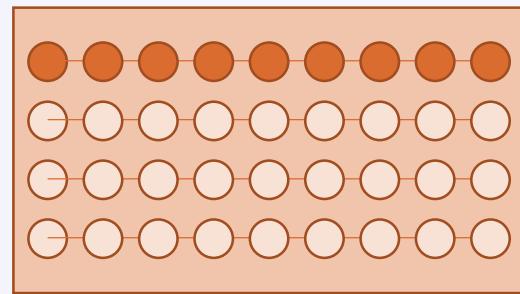
```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l nodes=1:ppn=9:gpus=1:skylake
#PBS -l pmem=5gb
#PBS -N testjob
#PBS -A default_project
```

You get ...

RAM  
(45 GB)

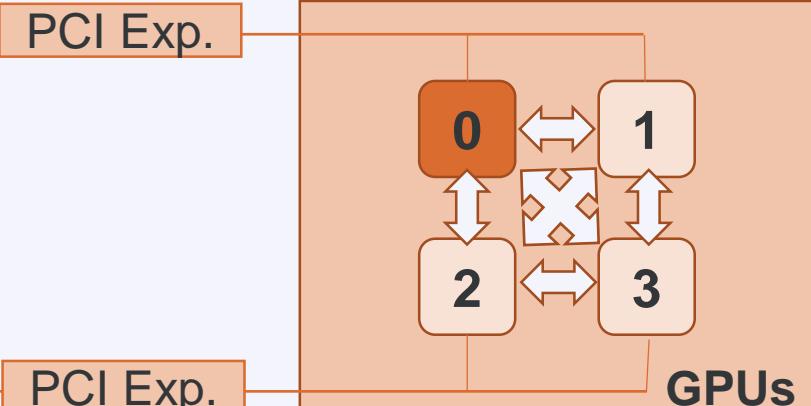


1 CPUs  
(9 cores)



PCI Exp.

PCI Exp.



# Using P100 GPU(s)

Using 2 GPUs

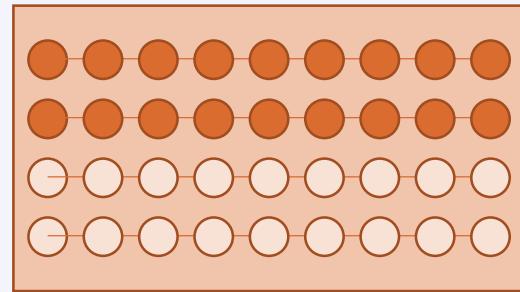
```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l nodes=1:ppn=18:gpus=2:skylake
#PBS -l pmem=5gb
#PBS -N testjob
#PBS -A default_project
```

You get ...

RAM  
(90 GB)

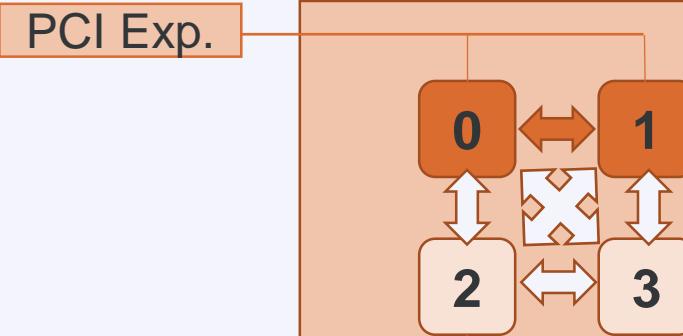


1 CPUs  
(18 cores)



PCI Exp.

PCI Exp.

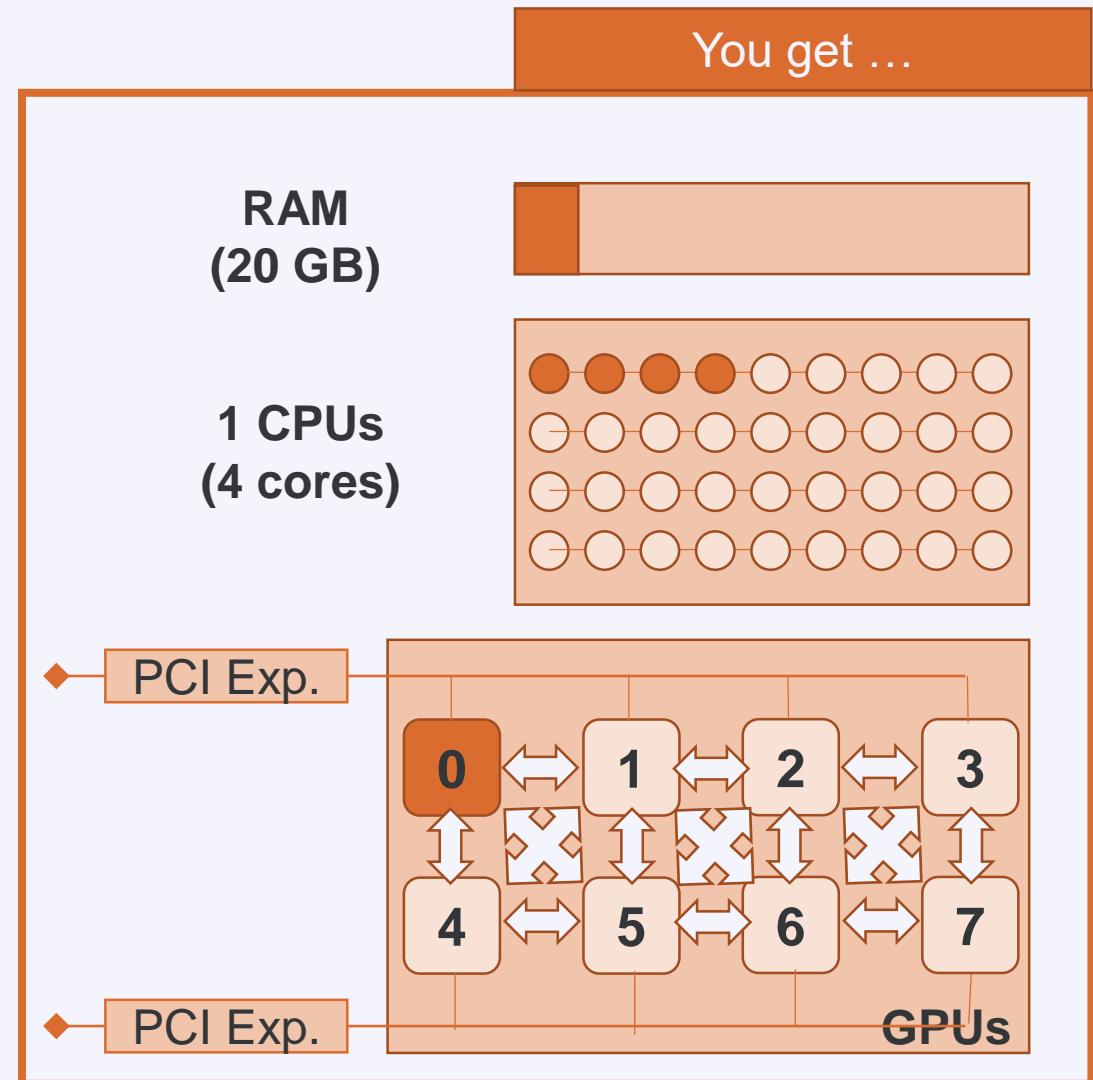


GPUs

# Using V100 GPU(s)

```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l
nodes=1:ppn=4:gpus=1:cascadelake
#PBS -l pmem=5gb
#PBS -N testjob
#PBS -A default_project
```

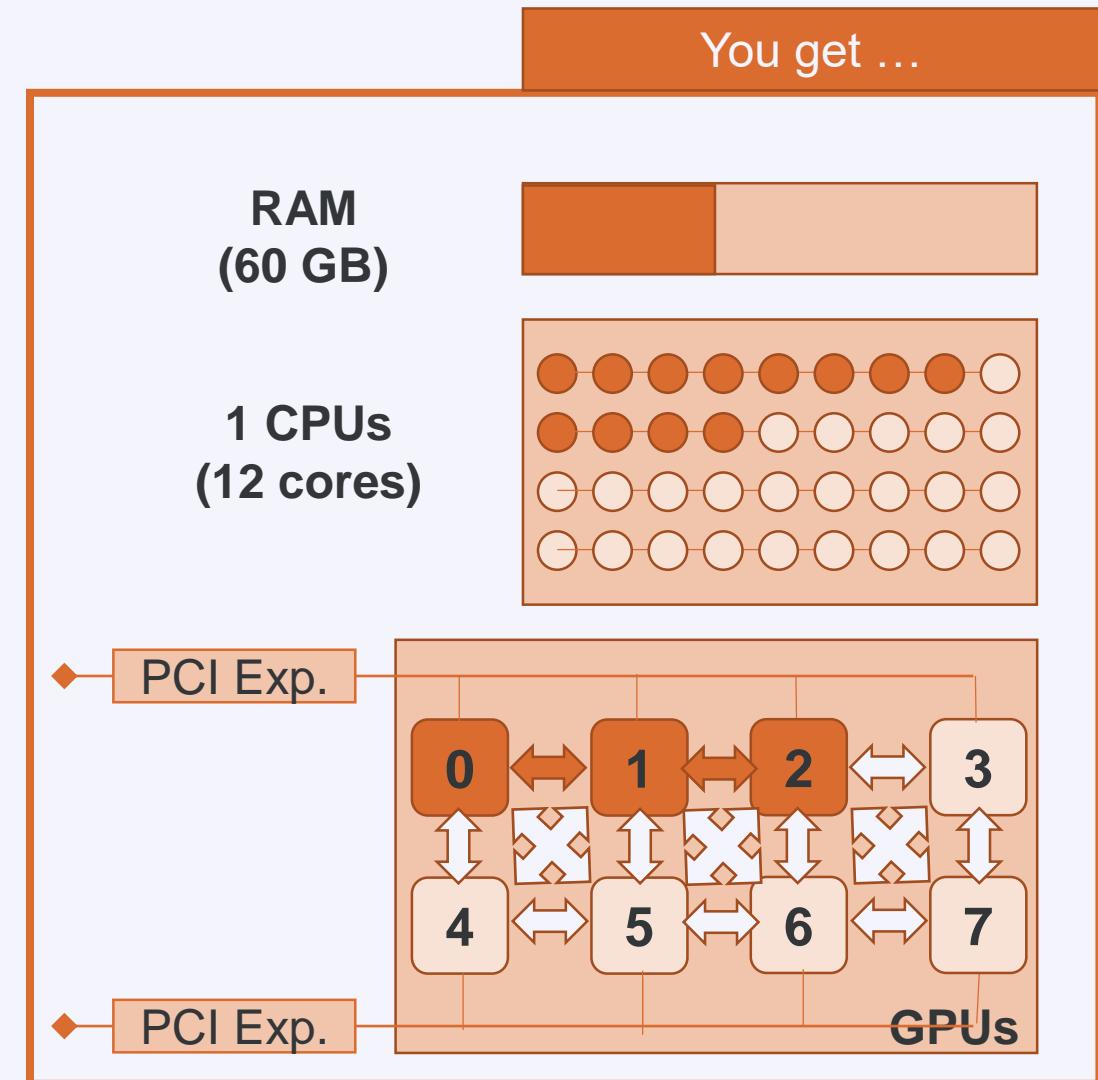
Using 1 GPU



# Using V100 GPU(s)

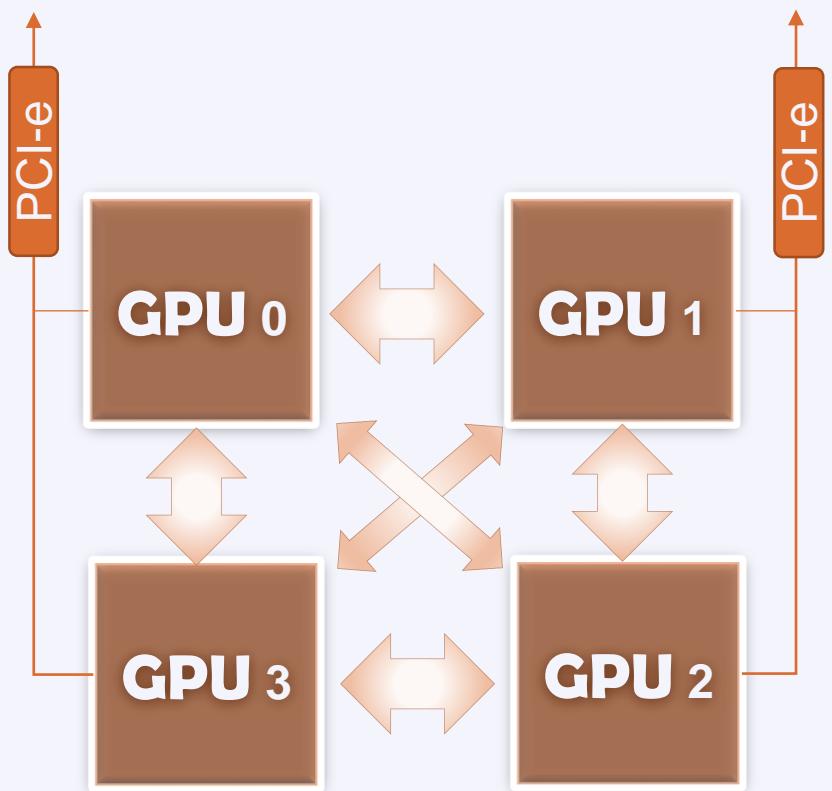
```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l partition=gpu
#PBS -l
nodes=1:ppn=12:gpus=3:cascadelake
#PBS -l pmem=5gb
#PBS -N testjob
#PBS -A default_project
```

Using 3 GPU



# Peer-to-Peer Bandwidth (GB per seconds)

High    Med    Low

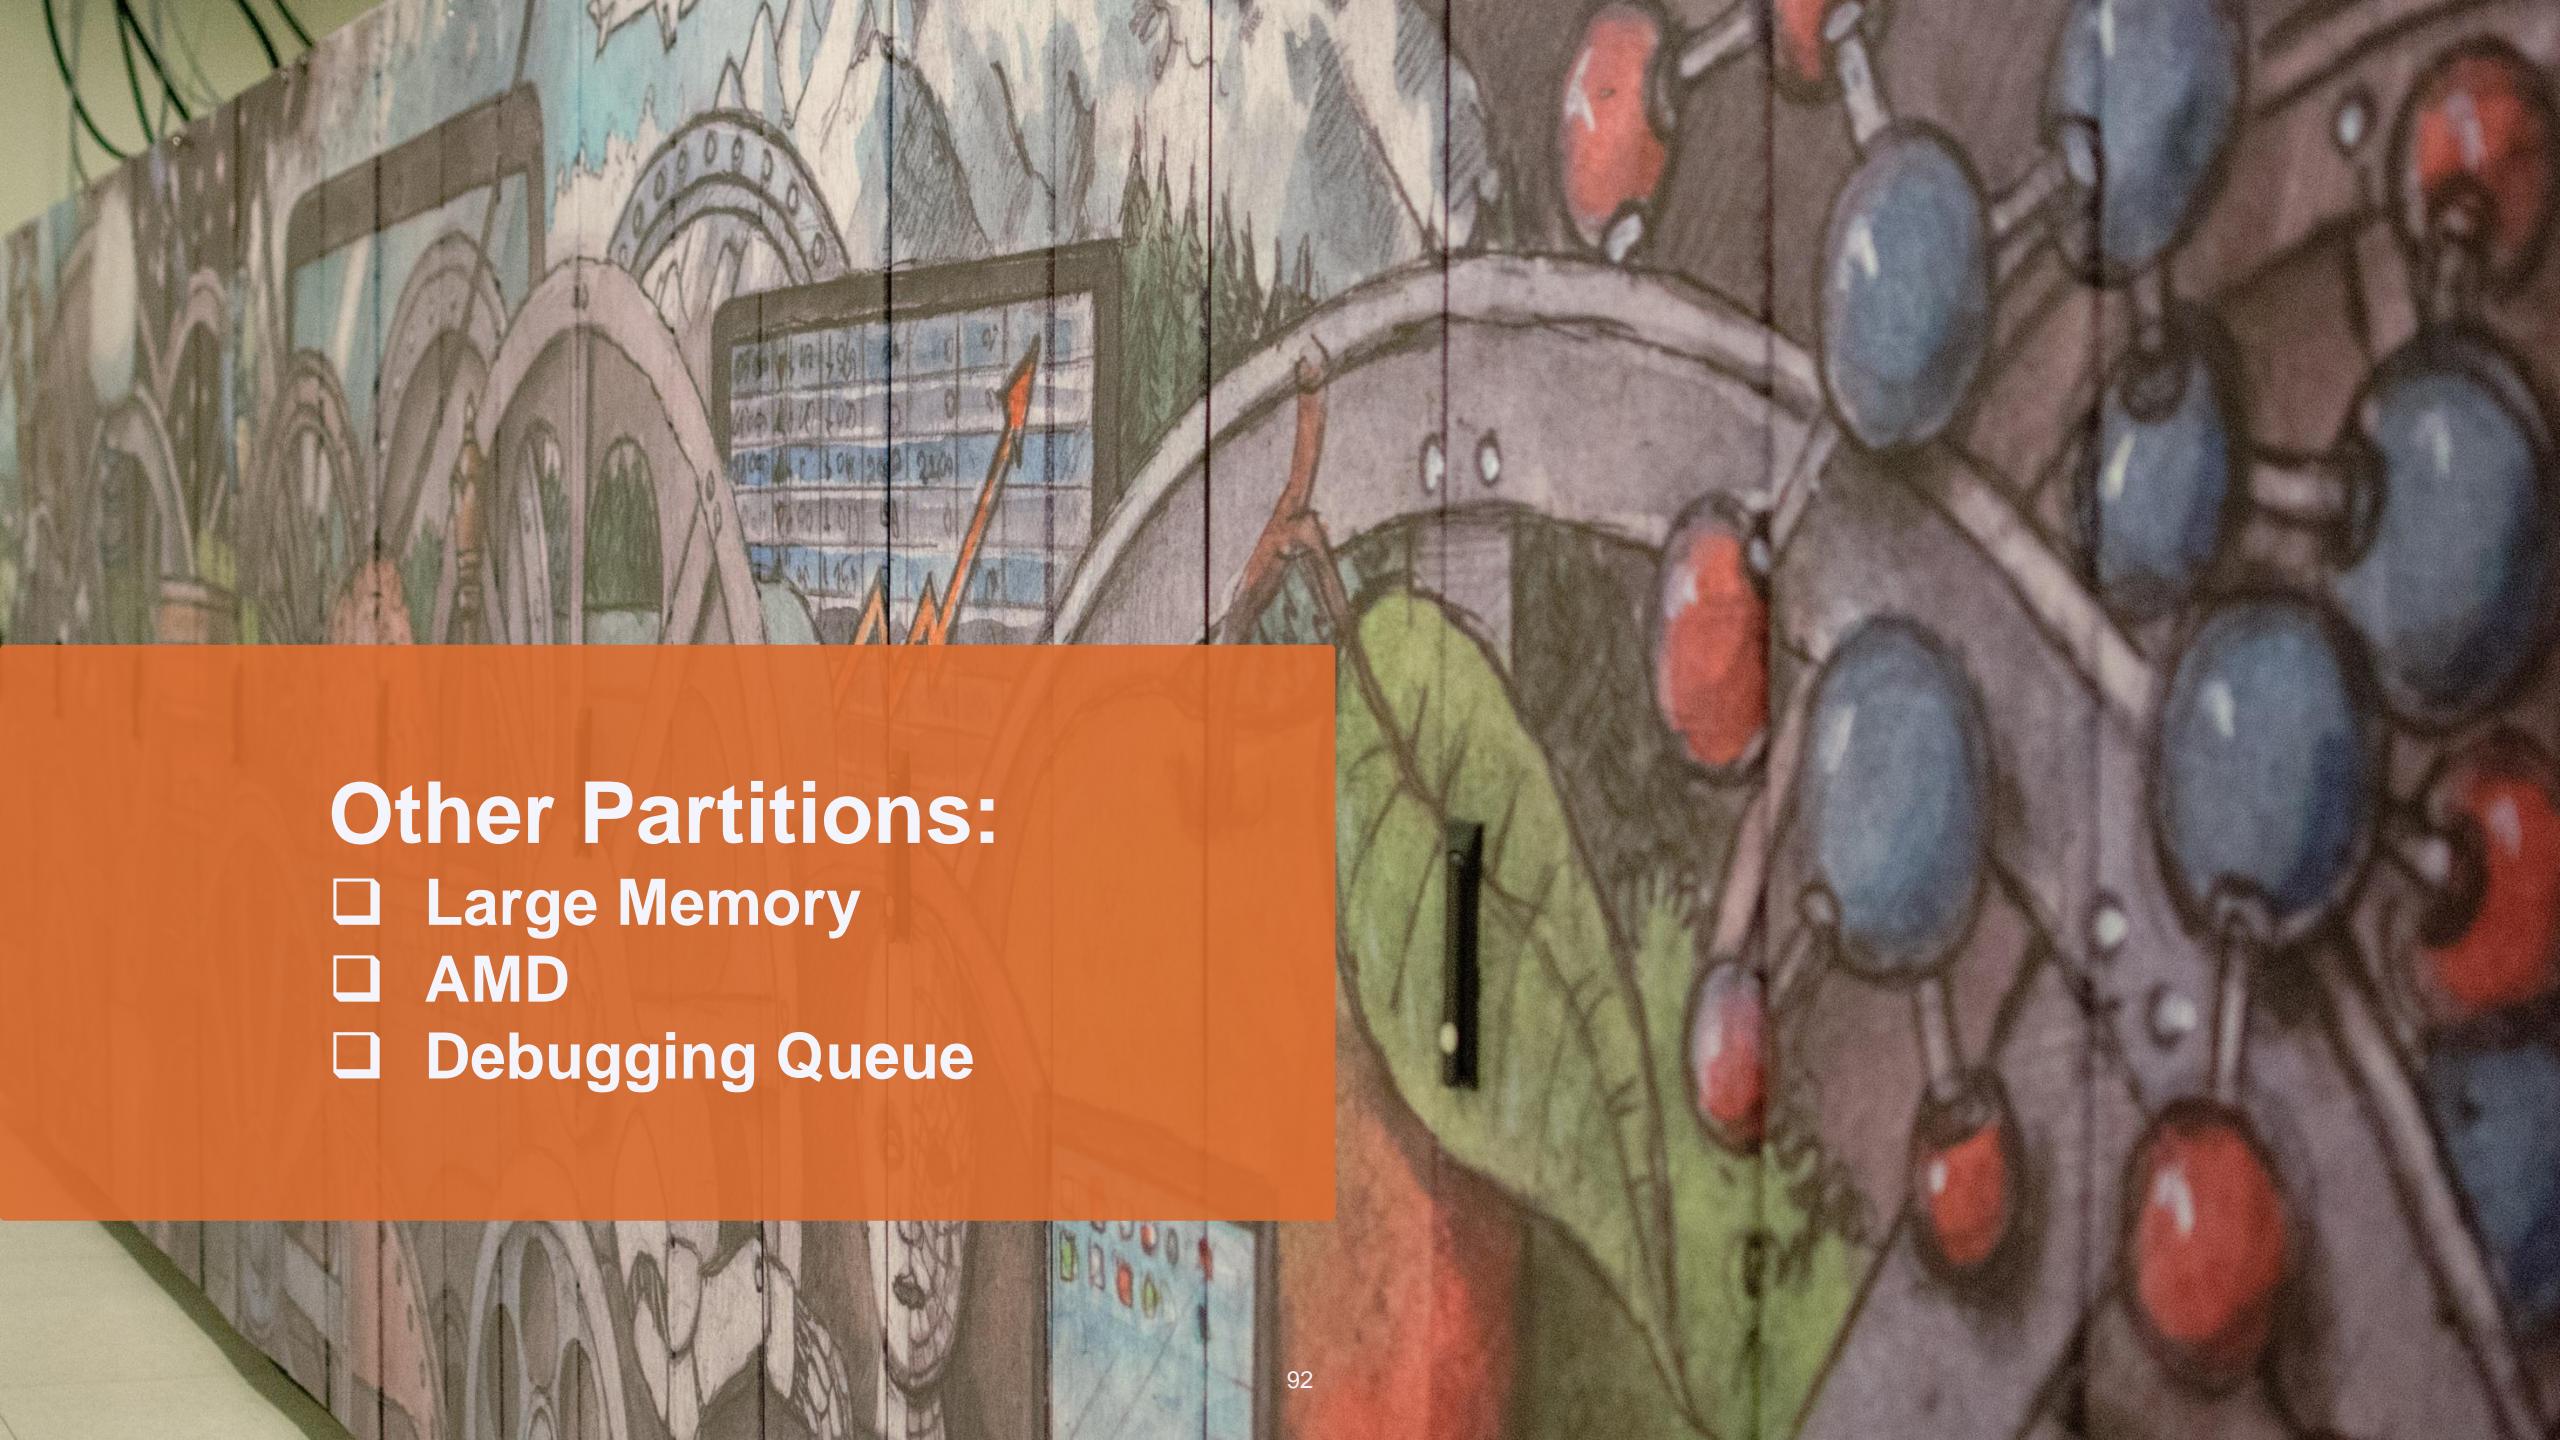


Dev. ID	0	1	2	3
0	509	10	19	18
1	10	508	18	18
2	19	18	508	10
3	18	18	10	507

Bi-directional P2P: **PCIe**

Dev. ID	0	1	2	3
0	508	37	37	61
1	37	507	61	37
2	36	61	508	37
3	62	37	37	506

Bi-directional P2P: **NVLink**  
(Nvidia P100)



## Other Partitions:

- Large Memory
- AMD
- Debugging Queue

# Requesting Large Memory Nodes

- All Thin Nodes have 192 GB memory each  
You need to leave 6 GB aside for the Operating System, and use the rest
- By default, every core can have  
 $\text{pmem} = (192\text{GB} - 6\text{GB}) / (36 \text{ cores}) \approx 5\text{GB per-core}$
- There are 3 methods you can request larger memory than normal:
  - Use less cores (ppn), but larger memory-per-core (pmem)
  - Use `bigmem` partition
  - Use Superdome machine

First Method

- You sacrifice ppn, in favor of pmem
- E.g. if you need 20GB memory per core:  
`$> qsub -I -l nodes=1:ppn=8,pmem=20gb ...`
- **Note:** the total memory requested must be less than  $192\text{GB} - 6\text{GB} = 186\text{GB}$ , i.e.  
 $\text{ppn} \times \text{pmem} \leq 186\text{GB}$

# Requesting Large Memory Nodes

Second Method

- We have 10 Large Memory nodes, with 768 GB memory each
- Specify `-l partition=bigmem` with `qsub` command
- Each core can have

$$\text{pmem} = (768 \text{ GB} - 6 \text{ GB}) / (36 \text{ cores}) \approx 21 \text{ GB per-core}$$

- E.g.  
`$> qsub -I -l partition=bigmem, nodes=1:ppn=36, pmem=21gb ...`
- You can also compromise ppn for pmem  
E.g.  
`$> qsub -I -l partition=bigmem, nodes=1:ppn=10, pmem=76gb ...`
- Make sure  $\text{ppn} \times \text{pmem} \leq 760\text{GB}$

# Requesting Large Memory Nodes

Third Method

- Superdome: 8 nodes, 14 cores-per-node, **6 TB** total RAM
- first do module load superdome
- Default pmem=50gb
- numanode is total CPUs to use (max. 8)
- lprocs is total number of cores to use
- E.g.

```
$> module load superdome
```

```
$> qsub -l partition=superdome -q qsuperdome \
      -L tasks=1:lprocs=14:place=numanode=1
```

Request **1/8th** of the machine,  
with 50 GB per core

- E.g.

```
$> qsub -l partition=superdome -q qsuperdome \
      -L tasks=1:lprocs=42:place=numanode=3
```

Request **3/8th** of the machine,  
with 50 GB per core

```
$> qsub -l partition=superdome -q qsuperdome \
      -L tasks=1:lprocs=7:place=numanode=1:memory=700gb
```

If you want to use > 50GB/core  
Eg 7 cores with 100GB/core

# (Experimental) AMD Nodes

- 4 nodes
- 2x EPYC 7501 CPUs (2.0 GHz, 32 cores)
- 8 memory channels / CPU
- Max Mem BW 158.95 GiB/s
- RAM: 256 GB
- Remark: pmem=3800mb  
If not specified, the job will not start

## Job Script

```
#!/bin/bash -l
#PBS -l walltime=01:30:00
#PBS -l partition=amd
#PBS -l nodes=1:ppn=64
#PBS -l pmem=3800mb
#PBS -N testjob
#PBS -A lp_my_project
```

# Debugging / Testing Jobs

- To quickly test/debug your (parallel) application
- 2 dedicated nodes on ThinKing and Genius  
One GPU node and one CPU node
- Such jobs do **not** go to the normal queue, so they start faster
- Max. walltime is **30 minutes**
- You must specify Quality of Service (**qos**)

Request Debugging Nodes

```
$> qsub -l nodes=1:ppn=10 -l qos=debugging -l \
    walltime=30:00 -A default_project
```

# Managing & Monitoring Jobs

Command	Purpose
\$> qsub ....	Submit a job (batch/interactive)
\$> qdel <JobID>	Delete a specific job
\$> checkjob -vvv <JobID>	Very detailed job info (very useful to diagnose issues)
\$> qstat -n	Status of all recent jobs
\$> qstat -Q -f	Info about available queues
\$> showstart <JobID>	Give a <i>rough</i> estimate of start time
\$> showq \$> showq -p gpu	Show minimal info about a queue or partition (-p)
\$> pbstop	Overview of the cluster
\$> module load accounting \$> mam-balance	Overview of different credit projects that you can use (qsub -A <Project>)
\$> mam-list-allocations	Detailed overview of your credit projects

JupyterHub

https://jupyterhub.hpc.kuleuven.be/hub/spawn?next=%2Fhub%2F

jupyter Home Token Logout

www.jupyterhub.hpc.kuleuven.be

Choose what you want

# Spawner Options

Select a job profile:

Genius - 9 cores, 45 GB, 5 minutes, 1 GPU

Select a project:

lp\_sys (20212 credits available)

Which credit project?

For further information on the service (status, bugs, announcements) please visit:  
<https://jupyterhub-doc.readthedocs.io>

Documentation

Current announcements: None

Last updated: 2018-11-12

Projects queried and page generated at 2019-06-12 09:37:06

Spawn

# Jupyter Notebook on NX

- Install Miniconda in your \$VSC\_DATA

```
$> wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh  
$> bash Miniconda3-latest-Linux-x86_64.sh -p $VSC_DATA/miniconda3
```

- Create a conda env including Jupyter

```
$> conda create -n JupNtbk jupyter scikit-learn
```

- Activate this env

```
$> conda activate JupNtbk
```

- Start an interactive GPU job (qsub -I ...)

- In another terminal, launch the notebook to get URL

```
$> jupyter notebook --ip=`hostname -i` --port=${USER}:3 --no-browser  
--allow-root
```

- Copy-paste the URL into NX Firefox; done

NoMachine - vsc30745 - Tier2

Applications Places System HPE\_AI\_LABS/ - Mozilla Firefox

HPE\_AI\_LABS/ - Mozilla Firefox

Jupyter

Files Running Clusters

Select items to perform actions on them.

Upload New

0 / HPE\_AI\_LABS

	Name	Last Modified	File size
..		seconds ago	
lab1		13 minutes ago	
lab2		a year ago	
lab3_1		a year ago	
lab3_2		a year ago	
lab4		a year ago	
hpe_ai_labs.yml		a year ago	2.71 kB
start.labs.sh		18 minutes ago	201 B

Mate Terminal

```
[I 16:13:45.368 NotebookApp] Adapting to protocol v5.1 for kernel 0c5741e4-af29-4470-b4d6-1094608d72c6
[I 16:13:45.369 NotebookApp] Restoring connection for 0c5741e4-af29-4470-b4d6-1094608d72c6:33c3dc16c95749548e4d34a1c49d2fd8
[I 16:13:45.369 NotebookApp] Replaying 6 buffered messages
[I 16:15:32.628 NotebookApp] Saving file at /HPE_AI_LABS/lab1/Building_a_first_NN_v10_EASY.ipynb
[I 16:26:27.136 NotebookApp] Starting buffering for 0c5741e4-af29-4470-b4d6-1094608d72c6:33c3dc16c95749548e4d34a1c49d2fd8

[I 16:27:12.084 NotebookApp] 302 GET /notebooks/HPE_AI_LABS/lab1/images/classification_kiank.png (10.118.230.5) 3.11ms
[I 16:27:12.561 NotebookApp] 302 GET /notebooks/HPE_AI_LABS/lab1/images/grad_summary.png (10.118.230.5) 2.54ms
[I 16:27:12.667 NotebookApp] 302 GET /notebooks/HPE_AI_LABS/lab1/images/sgd.gif (10.118.230.5) 2.79ms
[I 16:27:12.670 NotebookApp] 302 GET /notebooks/HPE_AI_LABS/lab1/images/sgd_bad.gif (10.118.230.5) 2.40ms
[I 16:27:13.307 NotebookApp] Adapting to protocol v5.1 for kernel 0c5741e4-af29-4470-b4d6-1094608d72c6
[W 16:27:15.398 NotebookApp] 404 GET /static/components/MathJax/fonts/HTML-CSS/TeX/otf/MathJax_AMS-Regular.otf (10.118.230.5) 6.36ms
referer=http://10.118.230.174:30745/notebooks/HPE_AI_LABS/lab1/Building_a_first_NN_v10_EASY.ipynb
[I 16:27:42.846 NotebookApp] Starting buffering for 0c5741e4-af29-4470-b4d6-1094608d72c6:6474c0264f7c496283a04a3b138e6c3e
[I 16:27:42.848 NotebookApp] 302 GET /notebooks/HPE_AI_LABS/ (10.118.230.5) 0.85ms
[I 16:27:42.855 NotebookApp] 302 GET /notebooks/HPE_AI_LABS/ (10.118.230.5) 1.10ms
```

# Singularity Containers

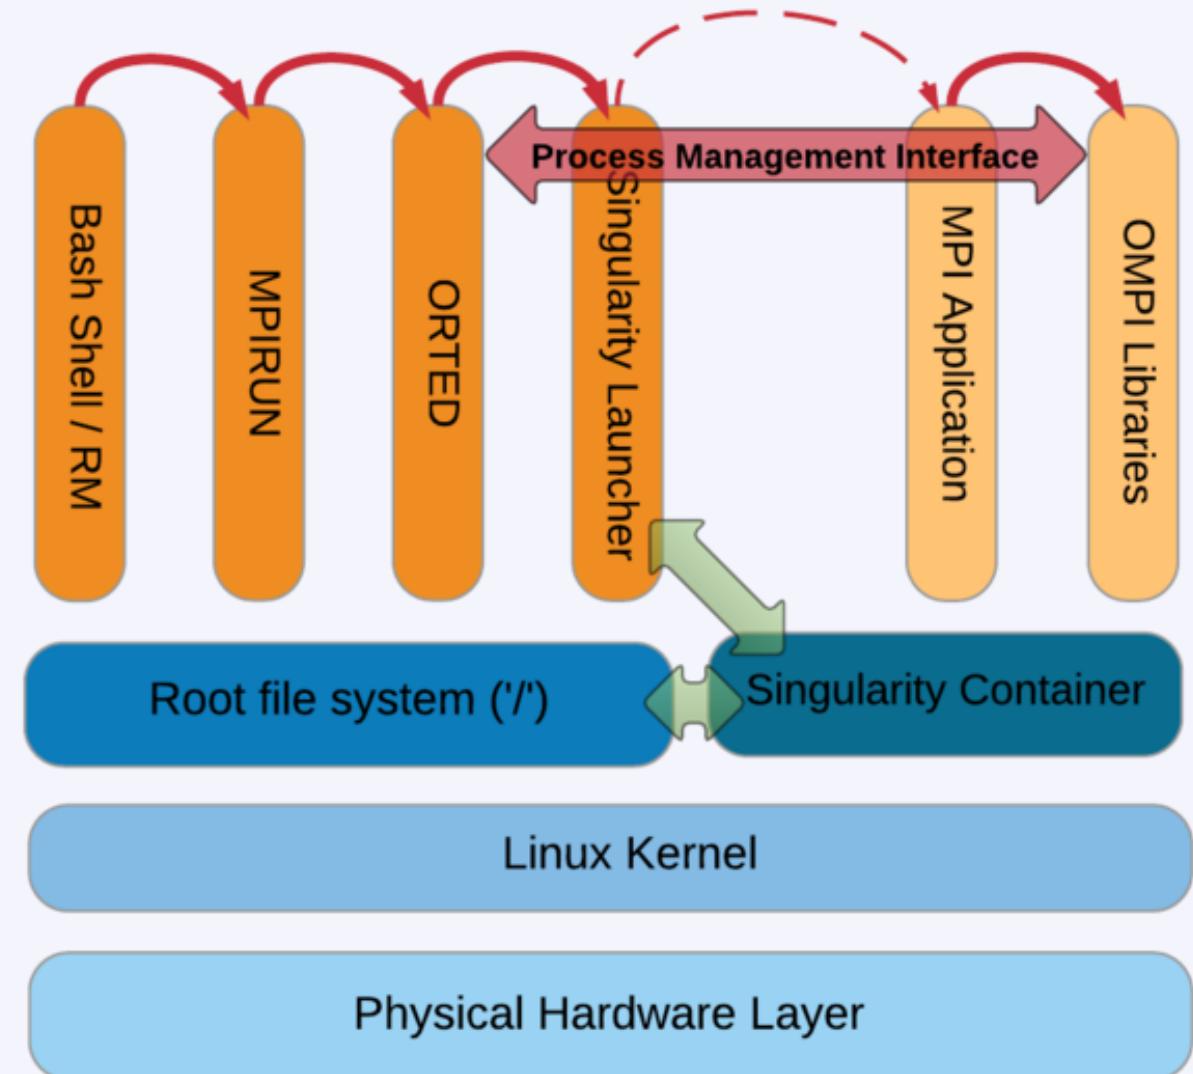
- What?
  - + Self-contained OS & software & data
- Why?
  - + fully resolved dependency chains
  - + portable workflow
- How?
  - + You create the image
  - + Run it on Genius
  - + MPI/OpenMP is supported

```
#!/bin/bash -l
#PBS -l walltime=00:30:00
#PBS -l nodes=4:ppn=36:pmem=5gb
#PBS -A lp_hpcinfo_training

cd $PBS_O_WORKDIR

singularity run Project.simg ./model.exe
```

PBS Script



The background of the slide is a hand-drawn illustration of an industrial setting. It features a network of brown pipes and valves in the foreground. Behind them is a large, circular metal drum with a green valve at the bottom. Further back, there's a wall with several blue circular valves and a control panel with a digital display showing various numbers and a graph. The style is sketchy and artistic.

Where to run your jobs

# Which system should I use?

Knowing yourself is  
the beginning  
of all wisdom



# What kind of work I do?

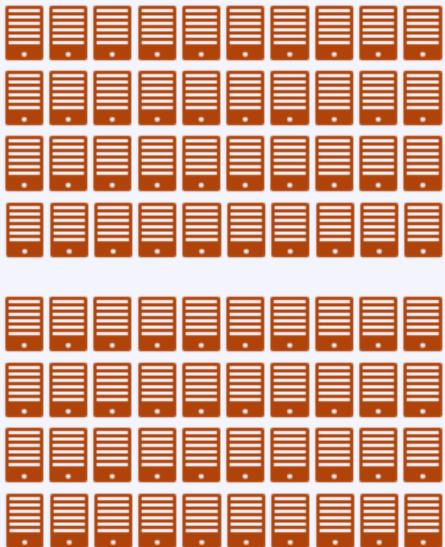
- CPU usage
  - Do you run single core jobs or parallel?
    - Single core
      - Is your code optimized to AVX ?
    - Parallel:
      - single node or multinode?
        - Single node: does my application scale? Up to how much cores ?
      - Multinode:
        - Does my application scale?
  - Memory usage
    - How much memory (per core) my jobs use?
    - Is your processing memory bound?

# What kind of work I do?

- How to find out
  - Know your algorithm
  - Benchmark your typical workload on the different architectures
  - Use tools to track behavior of your code
    - Use [Monitor](#) tool
    - Use profiling tools eg. [ARM MAP](#) or [Scalasca](#)
  - Consult training agenda
    - Upcoming Allinea ARM Forge training

# Which system for what?

Compute nodes (thin nodes)



## Thinking:

- Haswells 24 cores  
64/128 GB

Single core jobs  
Worker jobs  
Single node parallel jobs

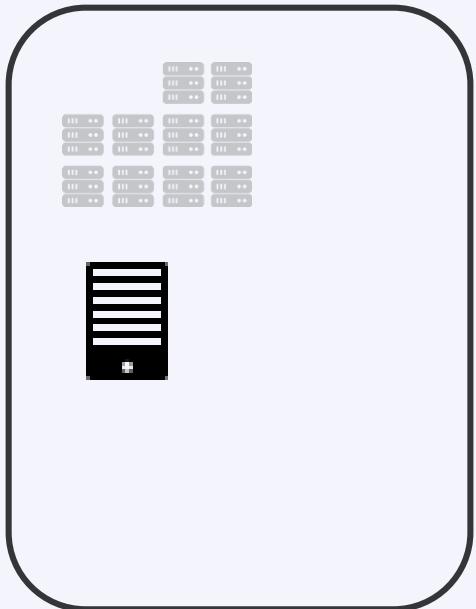
## Genius:

- Skylake: 36 cores
- Cascade Lake: 36 cores
- 192 GB
- IB EDR

Multinode parallel jobs  
AVX 512  
Higher core/memory ratio  
Higher memory bandwidth

# Which system for what?

Compute nodes (Large memory)



## Genius:

- Big memory 36 cores
- 768 GB

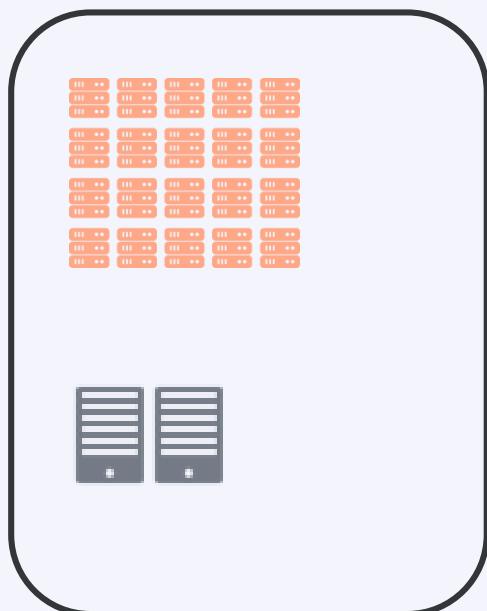
5GB<mem/core<20GB

- Superdome 112 cores
- 6TB

> 20GB mem/core  
> 700 GB of total RAM  
Single node scaling > 36 cores

# Which system for what?

Compute nodes (GPU nodes)



**Genius Skylake GPU: P100**

Single GPU workload  
Multi GPU workload  
with limited scaling

**Genius Cascadelake GPU: V100**

> 4 GPU workload  
Specific ML workloads  
(convolutional networks)

# Credits

Credits card concept:

- **Pre-authorization**: holding the balance as unavailable until the merchant clears the transaction
- Balance to be held as unavailable: based on requested resourced (walltime, nodes)
- **Actual charge** based on what was really used: used walltime (you pay only what you use, e.g. when job crashes)
- See output file

How to check available credits?

```
$ mam-balance
```

How to check the rates?

Check service catalogue

```
$ mam-list-chargerates
```

# Credit Pricing

- You pay as you go!
- For academic projects:  
1000 credits = 3.5 EUR
- Credits needed for a job:

#credits = Walltime (hr) × #nodes × Factor

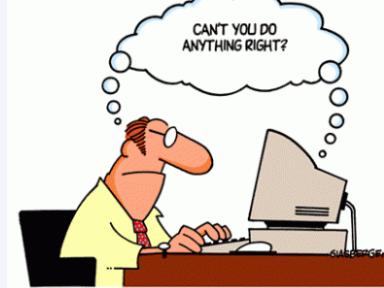
- For shared nodes, you pay a fraction of the costs, based on the ppn specified

Cluster / Partition	Credits/h
Genius Cascadelake	11.3
Genius Cascadelake 8 GPUs	40
Genius Skylake 4 GPUs	20
Genius Skylake BigMem	12
Genius Skylake	10
Genius / Superdome	10
Genius / AMD nodes	10
ThinKing / Haswell	6.68
ThinKing / IvyBridge	4.76

# How to Manage Credits?

Command	Purpose
mam-balance	List active projects and available credits
mam-list-allocations	List the validity dates of different projects/allocations
mam-list-chargerates	List current charge rates for different nodes
mam-list-accounts -a <account-name>	List the members in an account

# Hands-on 2



- Copy /apps/leuven/training/HPC\_bioinf to your \$VSC\_DATA directory and go to this directory

```
cp -r /apps/leuven/training/HPC_bioinf $VSC_DATA;  
cd $VSC_DATA/HPC_bioinf
```

- Submit jobscript cpujob.pbs to the cluster

```
qsub cpujob.pbs
```

- Check the status of your job(s) (qstat)

- Analyze outputs (i.e. display the output file – commands cat, more, less filename)

## Optional Slides:

- (Brief) Linux Intro
- Python & R with Conda
- Worker Framework

# (Very) Basic Linux Commands

# ls command

Lists the files in the current directory, in alphanumeric order, except files starting with the “.” character.

- \$ ls -a (all)  
Lists all the files (including .\* files)
- \$ ls -l (long)  
Long listing (type, date, size, owner, permissions)
- \$ ls -t (time)  
Lists the most recent files first
- \$ ls -S (size)  
Lists the biggest files first
- \$ ls -r (reverse)  
Reverses the sort order
- \$ ls -ltr (options can be combined)  
Long listing, most recent files at the end

# Moving around

- Print Working Directory
  - displays your current location within the file system.
  - `$ pwd`
- Change Directories
  - changes the position to the specific directory
  - `$ cd dir_name`
  - You can specify directory names in two ways:
    - Absolute pathname (starts from the root of the tree)  
`$ cd /u/home/hpc/test/bin`
    - Relative pathname (relative to your current directory)  
`$ cd`  
`$ cd .`  
`$ cd ..`  
`$ cd test/bin`
  - Use environment variables  
`$ cd $VSC_DATA`

# Special directories

## □ .

- The current directory.
- Useful for commands taking a directory argument.
- Useful to run commands in the current directory
- ./readme.txt and readme.txt are equivalent.

## □ ..

- The parent (enclosing) directory. Always belongs to the . Directory
- Typical usage:  
`cd ..`
- ~
  - Shells just substitute it by the home directory of the current user.
- -
  - `cd -` – jump back to the previous directory

# Directories

- Create directories
  - The `mkdir` command is used to create directories
  - `$ mkdir dir1 dir2 dir3`
- Remove directories
  - The `rmdir` command removes directories
  - `$ rmdir dir1`
  - `rmdir` will only remove empty directories.  
To remove a non-empty directory, use  
`$ rm -r [DIRECTORY]` instead.

# Copy a file

- The `cp` command copies files and directories
- The default behavior will overwrite any existing file(s). The `-i` option overrides this behavior and prompts the user before overwriting the destination file.
- **syntax: `cp [OPTIONS] [SOURCE] [DESTINATION]`**
  - `$ cp <source_file> <target_file>`  
Copies the source file to the target.
  - `$ cp file1 file2 file3 ... dir`  
Copies the files to the target directory (last argument).
  - `$ cp -i (interactive)`  
Asks for user confirmation if the target file already exists
  - `$ cp -r <source_dir> <target_dir> (recursive)`  
Copies the whole directory.
  - `$ cp -v (verbose)`  
Displays what has been copied

# Move or rename files

- Move or rename files and directories: `mv`
- The default behavior will overwrite any existing file(s).
- **syntax:** `mv [OPTIONS] [SOURCE] [DESTINATION]`
  - `$ mv old_name new_name`  
Renames the given file or directory.
  - `$ mv -i (interactive)`  
If the new file already exists, asks for user confirm
- The `mv` command can also be used to move or rename directories
  - `$ mv NewFiles/ OldFiles/`
  - -r option is not necessary



# Remove files

- The rm command removes files.
- `$ rm file1 file2 file3 ...`  
Removes the given files.
- `$ rm -i (interactive)`  
Always ask for user confirmation.
- `$ rm -r dir1 dir2 dir3 (recursive)`  
Removes the given directories with all their contents.
- Tip:  
Whenever you use wildcards with `rm` (besides carefully checking your typing!), test the wildcard first with `ls`. This will let you see the files that will be deleted. Then press the up arrow key to recall the command and replace the `ls` with `rm`.

# Auto-Completion

- Have the shell automatically complete commands or file paths.
- Activated using the **<TAB>** key on most systems
- examples
  - \$ whe<TAB>
  - \$ whereis
  - \$ ls -l /etc/en<TAB>
  - \$ ls -l /etc/environment
- When more than one match is found, the shell will display all matching results (use **<TAB>** twice)
  - \$ ls -l /etc/host<TAB>

# Command history: Arrow Up

- Previously executed commands can be recalled by using the **Up Arrow** key on the keyboard.
- Most Linux distributions remember the last five hundred commands by default.
- Display commands that have recently been executed
  - The `history` command displays a user's command line history.
  - You can execute a previous command using `! [NUM]` where NUM is the line number in history you want to recall.

# Displaying file contents

Several ways of displaying the contents of files.

- `$ cat file1`  
displays the contents of the given file.
- `$ cat file1 file2 file3 ...` (concatenate)  
Concatenates and outputs the contents of the given files.
- `$ more file1`  
Display the output of a command or text file one page at a time.
  - Can also jump to the first occurrence of a keyword (/ command).
- `$ less file1`  
Does more than more.  
Doesn't read the whole file before starting.  
Press q to exit

# The head and tail commands

- `$ head [-<n>] <file>`

Displays the first `<n>` lines (or 10 by default) of the given file.  
Doesn't have to open the whole file to do this!

- `$ tail [-<n>] <file>`

Displays the last `<n>` lines (or 10 by default) of the given file.  
No need to load the whole file in RAM! Very useful for huge files.

- `$ tail -f <file>` (follow)

Displays the last 10 lines of the given file and continues to display new lines when they are appended to the file.

Very useful to follow the changes in a log file, for example.

# The grep command

- `$ grep <pattern> <files>`  
Scans the given files and displays the lines which match the given pattern.
- `$ grep error *.log`  
Displays all the lines containing error in the \*.log files
- `$ grep -i error *.log`  
Same, but case insensitive
- `$ grep -ri error .`  
Same, but recursively in all the files in the current directory and its subdirectories
- `$ grep -v info *.log`  
Outputs all the lines in the files except those containing info.
- <http://www.thegeekstuff.com/2009/03/15-practical-unix-grep-command-examples/>

# wget

- Instead of downloading files from your browser,  
just copy and paste their URL and download them with wget
- main features
  - http and ftp support
  - Can resume interrupted downloads
  - Can download entire sites or at least check for bad links
  - Very useful in scripts or when no graphics are available  
(system administration, embedded systems)

•\$ wget -c http://microsoft.com/customers/dogs/winxp4dogs.zip  
Continues an interrupted download.

# Conda Package Management: Python and R

# Installing your own packages using conda

## Installing Miniconda

- Download the Bash script that will install it from conda.io using, e.g., wget:  
`$ wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh`
- Once downloaded, run the installation script:  
`$ bash Miniconda3-latest-Linux-x86_64.sh -p  
$VSC_DATA/miniconda3`
- Optionally, you can add the path to the Miniconda installation to the PATH environment variable in your .bashrc file. This is convenient, but may lead to conflicts when working with the module system or OS, so make sure that you know what you are doing in either case.  
The line to add to your .bashrc file would be:  
`export PATH=$VSC_DATA/miniconda3/bin:$PATH`

# Installing Python packages using conda

## Creating an environment

- First, ensure that the Miniconda installation is in your PATH environment variable. The following command should return the full path to the conda command:

```
$ which conda
```

- If the result is blank, or reports that conda can not be found, modify the `PATH` environment variable appropriately by adding miniconda's bin directory to PATH.
- Creating a new conda environment is straightforward:

```
$ conda create -n bioinf numpy biopython
```

- This command creates a new conda environment called science, and installs a number of Python packages.
- This will default to the latest Python 3 version, if you need a specific version, e.g., Python 2.7.x, this can be specified as follows:

```
$ conda create -n bioinf python=2.7 numpy biopython
```

# Installing Python packages

## Creating an environment

- First, ensure that the Miniconda is in your system path. Check this by running the command:

```
$ which conda
```

- If the result is blank, or reports that conda is not found, set the environment variable appropriately.
- Creating a new conda environment:

```
$ conda create -n bioinf
```

- This command creates a new conda environment with a default number of Python packages.
- This will default to the latest Python version. If you prefer e.g., Python 2.7.x, this can be specified:

```
$ conda create -n bioinf
```

The screenshot shows a PuTTY terminal window titled "login1-tier2.hpc.kuleuven.be - PuTTY". The command entered was "conda create -n bioinf numpy biopython". The output shows the package plan, download details, and the list of new packages to be installed. The terminal prompt at the bottom asks "Proceed ([y]/n)?".

```
: vsc30468@login1 ~ 16:21 $ conda create -n bioinf numpy biopython
Solving environment: done

## Package Plan ##

environment location: /data/leuven/304/vsc30468/bioinf/envs/bioinf

added / updated specs:
- biopython
- numpy

The following packages will be downloaded:

          package                               build
          -----|-----
numpy-1.15.3                            py37h1d66e8a_0   36 KB
setuptools-40.4.3                         py37_0          556 KB
blas-1.0                                 mkl             6 KB
python-3.7.1                             h0371630_3      36.4 MB
mkl_random-1.0.1                          py37h4414c95_1  372 KB
intel-openmp-2019.0                        118            721 KB
wheel-0.32.2                             py37_0          35 KB
certifi-2018.10.15                         py37_0          138 KB
sqlite-3.25.2                            h7b64a7c_0      1.9 MB
numpy-base-1.15.3                          py37h81de0dd_0  4.2 MB
biopython-1.72                            py37h04863e7_0  2.5 MB
openssl-1.1.1                            h7b64a7c_0      5.0 MB
mkl_fft-1.0.6                            py37h7dd41cf_0  150 KB
mkl-2019.0                                118            204.4 MB
libgfortran-ng-7.3.0                       hdf63c60_0      1.3 MB

Total: 257.6 MB

The following NEW packages will be INSTALLED:

biopython:      1.72-py37h04863e7_0
blas:          1.0-mkl
ca-certificates: 2018.03.07-0
certifi:        2018.10.15-py37_0
intel-openmp:   2019.0-118
libedit:        3.1.20170329-h6b74fdf_2
libffi:         3.2.1-hd88cf55_4
libgcc-ng:      8.2.0-hdf63c60_1
libgfortran-ng: 7.3.0-hdf63c60_0
libstdcxx-ng:   8.2.0-hdf63c60_1
mkl:           2019.0-118
mkl_fft:        1.0.6-py37h7dd41cf_0
mkl_random:    1.0.1-py37h4414c95_1
ncurses:        6.1-hf484d3e_0
numpy:          1.15.3-py37h1d66e8a_0
numpy-base:     1.15.3-py37h81de0dd_0
openssl:        1.1.1-h7b64a7c_0
pip:            10.0.1-py37_0
python:         3.7.1-h0371630_3
readline:       7.0-h7b64a7c_5
setuptools:     40.4.3-py37_0
sqlite:         3.25.2-h7b64a7c_0
tk:             8.6.8-hbc83047_0
wheel:          0.32.2-py37_0
xz:             5.2.4-h14c3975_4
zlib:           1.2.11-ha838bed_2

Proceed ([y]/n)? y
```

# Installing Python packages using conda

## Working with the environment

- To work with an environment, you have to activate it. This is done with, e.g.,  
  \$ source activate bioinf  
Here, science is the name of the environment you want to work in.

## Install an additional package

- To install an additional package, e.g., `pandas`, first ensure that the environment you want to work in is activated.  
  \$ source activate bioinf
- Next, install the package:  
  \$ conda install pandas  
Note that conda will take care of all dependencies, including non-Python libraries. This ensures that you work in a consistent environment.

# Installing Python packages using conda

## Updating/removing

- Using conda, it is easy to keep your packages up-to-date. Updating a single package (and its dependencies) can be done using:

```
$ conda update biopython
```

- Updating all packages in the environment is trivial:

```
$ conda update --all
```

- Removing an installed package:

```
$ conda remove pandas
```

## Deactivating an environment

- To deactivate a conda environment, i.e., return the shell to its original state, use the following command

```
$ source deactivate
```

# Installing Python packages - alternatives

## Checking for installed packages

- Pip utility will list all packages that are installed for the Python distribution you are using, including those installed by you, i.e., those in your `PYTHONPATH` environment variable.
- Load the module for the Python version you wish to use, e.g.:  
`$ module load Python/2.7.14-foss-2018a`
- Run pip:  
`$ pip freeze`
- Note that some packages, e.g., `mpi4py`, `h5py`, `pytables`, ..., are available through the module system, and have to be loaded separately. These packages will not be listed by pip unless you loaded the corresponding module.
- If you have any packages installed in `.local` directory, it will always take priority on whatever the Python version used (conda, module, system). That can lead to strange problems, so please avoid using that location.

# Installing Python packages - pip

1. Load the appropriate Python module, i.e., the one you want the python package to be available for:

```
$ module load Python/2.7.14-foss-2018a
```

2. Create a directory to hold the packages you install, the last three directory names are mandatory:

```
$ mkdir -p "${VSC_DATA}/python_lib/lib/python2.7/site-packages/"
```

3. Add that directory to the PYTHONPATH environment variable for the current shell to do the installation:

```
$ export PYTHONPATH="${VSC_DATA}/python_lib/lib/python2.7/site-packages/:$PYTHONPATH"
```

4. Add the following to your .bashrc so that Python knows where to look next time you use it:

```
export PYTHONPATH="${VSC_DATA}/python_lib/lib/python2.7/site-packages/:$PYTHONPATH"
```

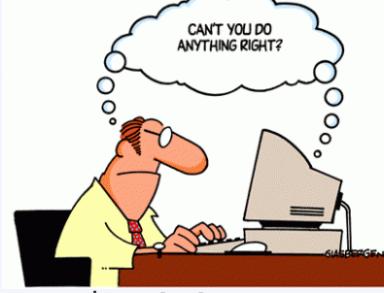
5. Install the package, using the prefix install option to specify the install path (this would install the biopython package):

```
$ pip install --user biopython
```

or

```
$ pip install --install-option="--prefix=${VSC_DATA}/python_lib" biopython
```

# Hands-on 3



1. Install conda environment with the Biopython package in your \$VSC\_DATA directory
2. Submit jobscript pythonjob1.pbs to the cluster. Python script checks how many times string ATA occurs in a given dna1.txt file. Check the answer in output file.
3. Copy the file pythonjob1.pbs to pythonjob2000.pbs and modify it so that:
  - You will be notified by e-mail when job finishes or aborts
  - Name of the job is count-string-2000
  - Instead of Python from 2018a toolchain conda Python is used
  - First action that the job does is running multiply.sh script
  - Instead of count\_substr1.py count\_substr2000.py script is used
  - The script looks for ATG string instead of ATASubmit the job and check how many times string ATG occurs in dna2000.txt file.

# Installing R packages using conda

## Creating an environment

- First, ensure that the Miniconda installation is in your PATH environment variable. The following command should return the full path to the conda command:

```
$ which conda
```

- If the result is blank, or reports that conda can not be found, modify the `PATH` environment variable appropriately by adding miniconda's bin directory to PATH.

- Creating a new conda environment is straightforward:

```
$ conda create -n science -c r r-essentials r-rodbc
```

This command creates a new conda environment called science, and installs essentials and required packages.

# Installing R packages using conda

## Working with the environment

- To work with an environment, you have to activate it. This is done with, e.g.,  
  \$ source activate science  
Here, science is the name of the environment you want to work in.

## Install an additional package

- To install an additional package, e.g., `r-ggplot2`, first ensure that the environment you want to work in is activated.  
  \$ source activate science
- Next, install the package:  
  \$ conda install -c r r-ggplot2  
Note that conda will take care of all independencies. This ensures that you work in a consistent environment.

# Installing R packages using conda

## Updating/removing

- Using conda, it is easy to keep your packages up-to-date. Updating a single package (and its dependencies) can be done using:

```
$ conda update r-rodbc
```

- Updating all packages in the environment is trivial:

```
$ conda update --all
```

- Removing an installed package:

```
$ conda remove r-mass
```

## Deactivating an environment

- To deactivate a conda environment, i.e., return the shell to its original state, use the following command

```
$ source deactivate
```

# Installing other R packages using conda

- Installing CRAN package:

```
$ conda skeleton cran readr  
$ conda build r-readr  
$ conda install --use-local r-readr
```

**Doing that for the first time you need to install conda-build before:**

```
$ conda install conda-build
```

- Some packages not available in r-essentials are still available on conda channels, in that case, it's simple:

```
$ conda config --add channels r; conda install r-readxl
```

# Installing R packages – alternatives

1. Load the appropriate R module, i.e., the one you want the package to be available for:

```
$ module load R/3.5.0-foss-2018a-bare
```

2. start R and install the package from there:

```
> install.packages ("DEoptim")
```

3. Alternatively you can download the desired package:

```
$ wget cran.r-project.org/src/contrib/Archive/DEoptim/DEoptim_2.0-0.tar.gz
```

And Install the package from the command line:

```
$ R CMD INSTALL DEoptim_2.2-3.tar.gz -l /$VSC_HOME/R/
```

4. These packages might depend on the specific R version, so you may need to reinstall them for the other version.

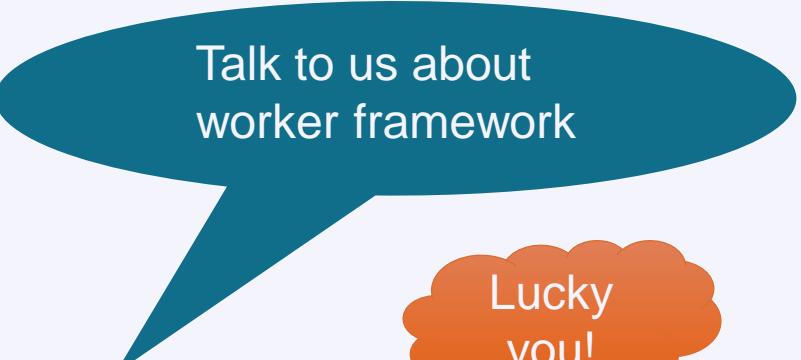
# Worker Framework

# Glossary

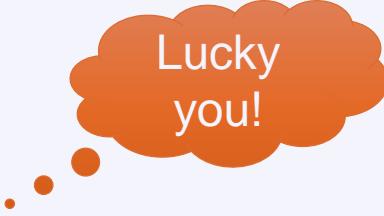
- **Serial application:** a program that runs a single process, with a single thread. All computations are done sequentially, i.e., one after the other, no explicit parallelism is used.
- **Multi-core application:** a multi-core application uses more than one core of processor during its execution by running multiple threads, also called a shared memory application.
- **Distributed application:** an application that uses multiple compute nodes for its computations, concurrent computations are executed as processes. These processes communicate by exchanging messages, typically implemented by calls to an MPI library. Messages can be used to exchange data and coordinate the execution.

# Parallel Computing

- Serial:
  - one program, on one core
- 'Embarrassingly parallel' problems:
  - lots of runs of one program, with different parameters
- Problems that require 'real' parallel algorithms
  - OpenMP
  - MPI : Message Passing Interface



Talk to us about  
worker framework



Lucky  
you!

# Use case: parameter exploration

temperature	pressure	humidity
293.0	1.0e05	87
...	...	...
313.0	1.3e05	75

```
#!/bin/bash -l
#PBS -l nodes=1:ppn=1 -l walltime=00:10:00
    #!/bin/bash -l
cd   #PBS -l nodes=1:ppn=1 -l walltim
wea .00
    cd   #!/bin/bash -l
wea #PBS -l nodes=1:ppn=1 -l walltime=00:10:00
    cd $PBS_O_WORKDIR
weather -p 1.3e05 -t 313.0 -h 75
```

Many single core computations

# Solution: worker with -data

temperature	pressure	humidity
293.0	1.0e05	87
...	...	...
313.0	1.3e05	75

data.csv

```
#!/bin/bash -l  
#PBS -l nodes=5:ppn=20 -l walltime=01:20:00  
  
cd $PBS_O_WORKDIR  
weather -p $pressure -t $temperature -h $humidity
```

job.pbs

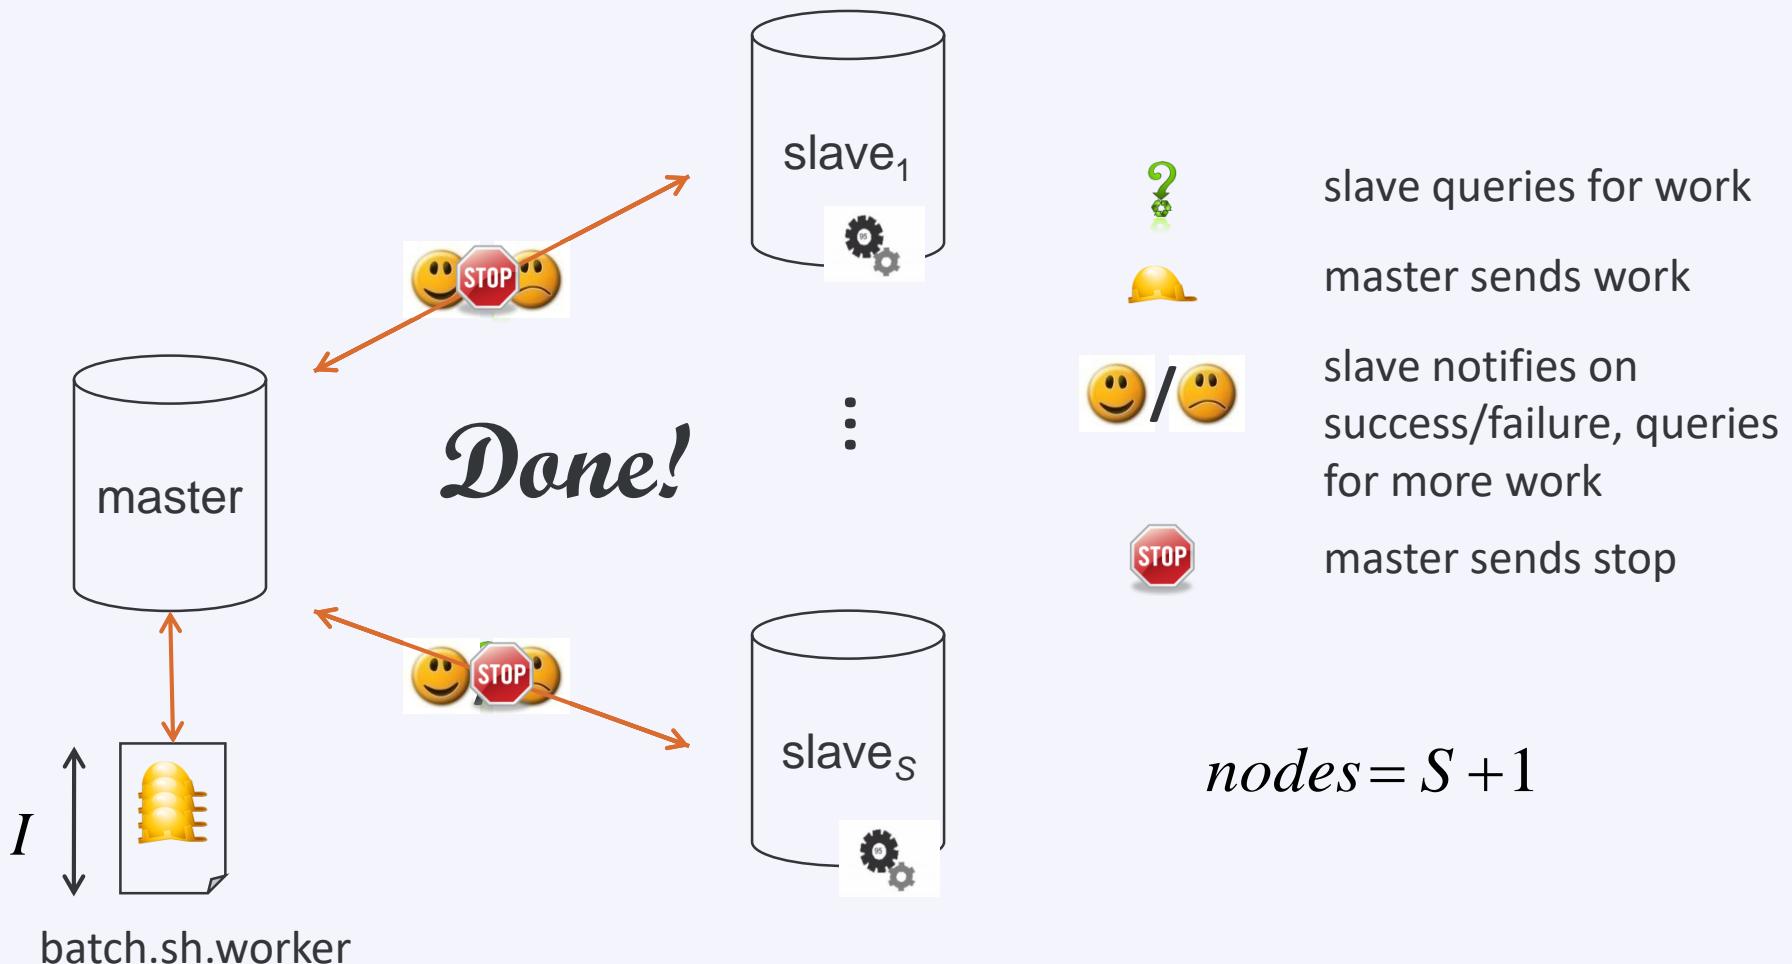
```
$ module load worker  
$ wsub -data data.csv -batch job.pbs
```

# Data exploration: steps

- Write PBS script with parameters
- Create Excel sheet with data
  - Convert to CSV format
- Submit with `wsub`
  - walltime is time to complete all work items

$$\text{walltime}_{\text{job}} \geq \frac{N \cdot \text{walltime}_{\text{work item}}}{\text{nodes} \cdot \text{ppn}}$$

# worker processing: informally



# Resuming jobs: wresume

- Resuming a job
  - \$ **wresume -l walltime=1:30:00 -jobid 445948**

- Redoing failed work items
  - \$ **wresume -jobid 445948 -retry**

# How to use worker well?

- Many work items, i.e., `#work items/#proc >> 1`
- `time(work item) > 1 minute`
- Work item is not multithreaded
- Work item is multithreaded
  - will work, but user must be careful to request the right resources
  - Use `-threaded <n>` flag with `wsub`
- There be dragons: licensing!

# worker & conflicts

- worker module only required for
  - job submission, i.e., wsub, wresume
  - data aggregation, ..., e.g., wcat, wreduce, ...
- No need to load in PBS script
  - use module purge
  - minimizes conflicts
  - work items run in own Bash shell
- However, MPI may be problematic
  - e.g., mpi4py

# worker & multithreading

- Some software uses multithreading automatically, e.g.,
  - R
  - Matlab
- Will use as many threads as there are cores, regardless of system load
  - 36 cores/node
  - 36 work items/node

} **36 × 36 threads/node**

Oversubscription: ***very inefficient!!!***

# Controlling number of threads

- R, most of the time: OMP\_NUM\_THREADS=1

```
#!/bin/bash -l
#PBS -N my-pe
#PBS -l
walltime=1:00:00, nodes=5:ppn=20

module load R
cd $PBS_O_WORKDIR

export OMP_NUM_THREADS=1
Rscript program $a $b
```

- Matlab
  - Use maxNumCompThreads(1) function call
  - Use compiler flag: mcc -singleCompThread ...

# Help on worker

- See documentation  
<http://worker.readthedocs.io/>
- Each command has help, use -h



# Epilogue

# Questions

- Helpdesk:  
[hpcinfo@kuleuven.be](mailto:hpcinfo@kuleuven.be) or  
[https://admin.kuleuven.be/icts/HPCinfo\\_form/HPC-info-formulier](https://admin.kuleuven.be/icts/HPCinfo_form/HPC-info-formulier)
- If you need support:  
Explain your problem in detail; provide jobID; path to your job scripts
- VSC web site: <http://www.vscentrum.be/>
- VSC documentation: <https://vlaams-supercomputing-centrum-vscdocumentation.readthedocs-hosted.com/en/latest/>  
VSC agenda: training sessions, events
- Systems status page:  
<http://status.kuleuven.be/hpc>  
<https://wwsw.vscentrum.be/en/user-portal/system-status>