



Vlaanderen  
is supercomputing

# VSC HPC Introduction

ICTS KU Leuven

VSC staff:

Ehsan Moravveji

Mag Selwa

Geert Jan Bex (UHasselt)

Jan Ooghe

Wouter van Assche

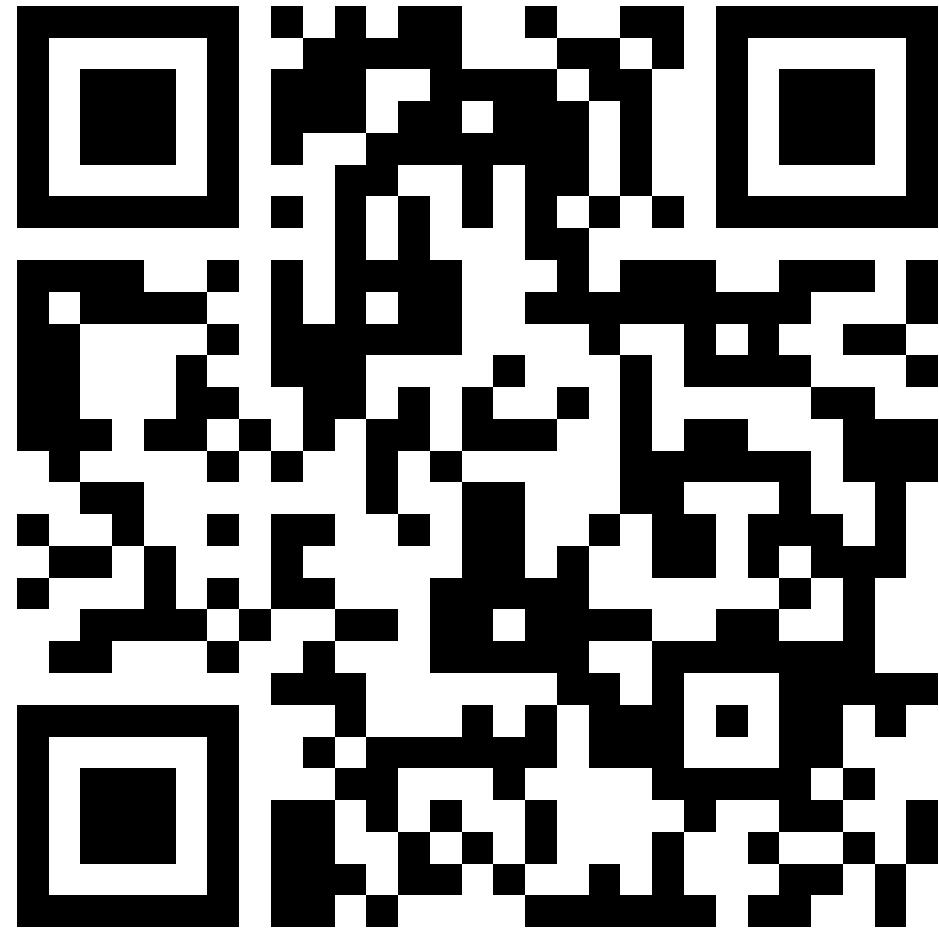
# Questions and problems troubleshooting

- Blackboard Collaborate session:  
<https://eu.bbcollab.com/guest/ce4229a2f9eb4cc5b90f490d660eb1c2>



# Material

- Everything is on Github:  
<https://hpcleuven.github.io/HPC-intro/>
- Video Recordings
  - Scan the QR code
  - Recommended videos: ~ 2 hrs
  - Optional videos: ~1 hr



# What is High Performance Computing?

- using supercomputers to solve advanced computation problems
- Reduce the computation time from days, years, decades, or centuries to minutes, hours, days, or weeks
- The key is parallelism



In practice, it is more like ...



The concept is simple: **Parallelism** = employing multiple processors for a single problem

# Outline

- What is the VSC?
- What is a cluster?
- Genius Cluster
- Storage
- Login nodes & MFA
- Connection Setup
- Software environment
- How to submit jobs?
- Dedicated hardware
- How to choose resources?
- Optional material
  - Linux in brief
  - Conda for Python and R
  - Worker Framework



The background of the slide features a large, colorful mural painting. The mural depicts a futuristic cityscape with various floating spheres of different sizes and colors (red, blue, green). In the foreground, there is a large, stylized green leaf-like shape. A central element is a tablet or screen displaying a grid of numbers and data points, suggesting a theme of data processing or supercomputing.

# VSC

(Vlaams Supercomputer Centrum)

# VSC PARTNERSHIP



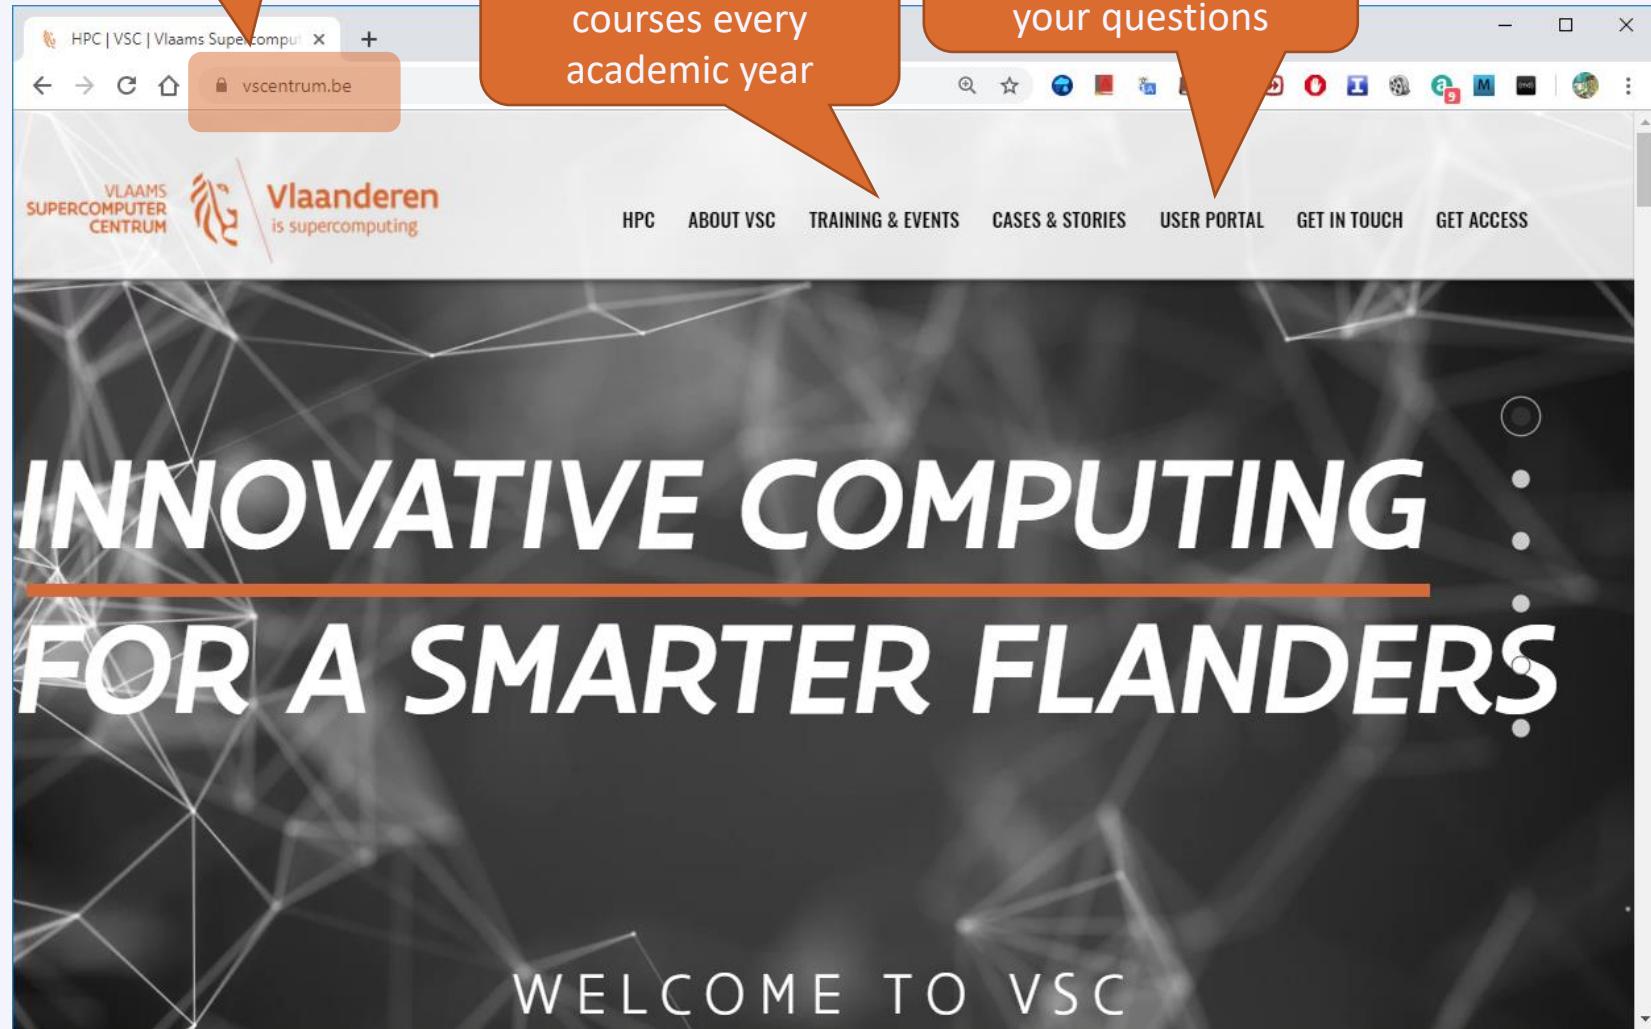
Supported by



# VSC HPC Environments



[www.vscentrum.be](http://www.vscentrum.be)



Search your keywords here;  
e.g. qsub

Welcome to VSC documentation

VSC documentation 1.0 documentation »

next | index

Next topic  
Getting access

This Page  
Show Source

Quick search

Go

- Getting access
  - VSC accounts
  - How to request an account?
  - Next steps
  - Additional information
- Access and data transfer
  - Logging in to a cluster
  - Data storage
  - Transferring data
  - GUI applications on the clusters
  - VPN
- Software stack
  - Using the module system
  - Specialized software stacks
- Running jobs
  - Job script
  - Submitting and monitoring a job
  - Job output
  - Troubleshooting
  - Advanced topics
- Software development

v: latest

# Support and Services

## Basic support

- Helpdesk ([hpcinfo@kuleuven.be](mailto:hpcinfo@kuleuven.be))
- Monitoring and reporting

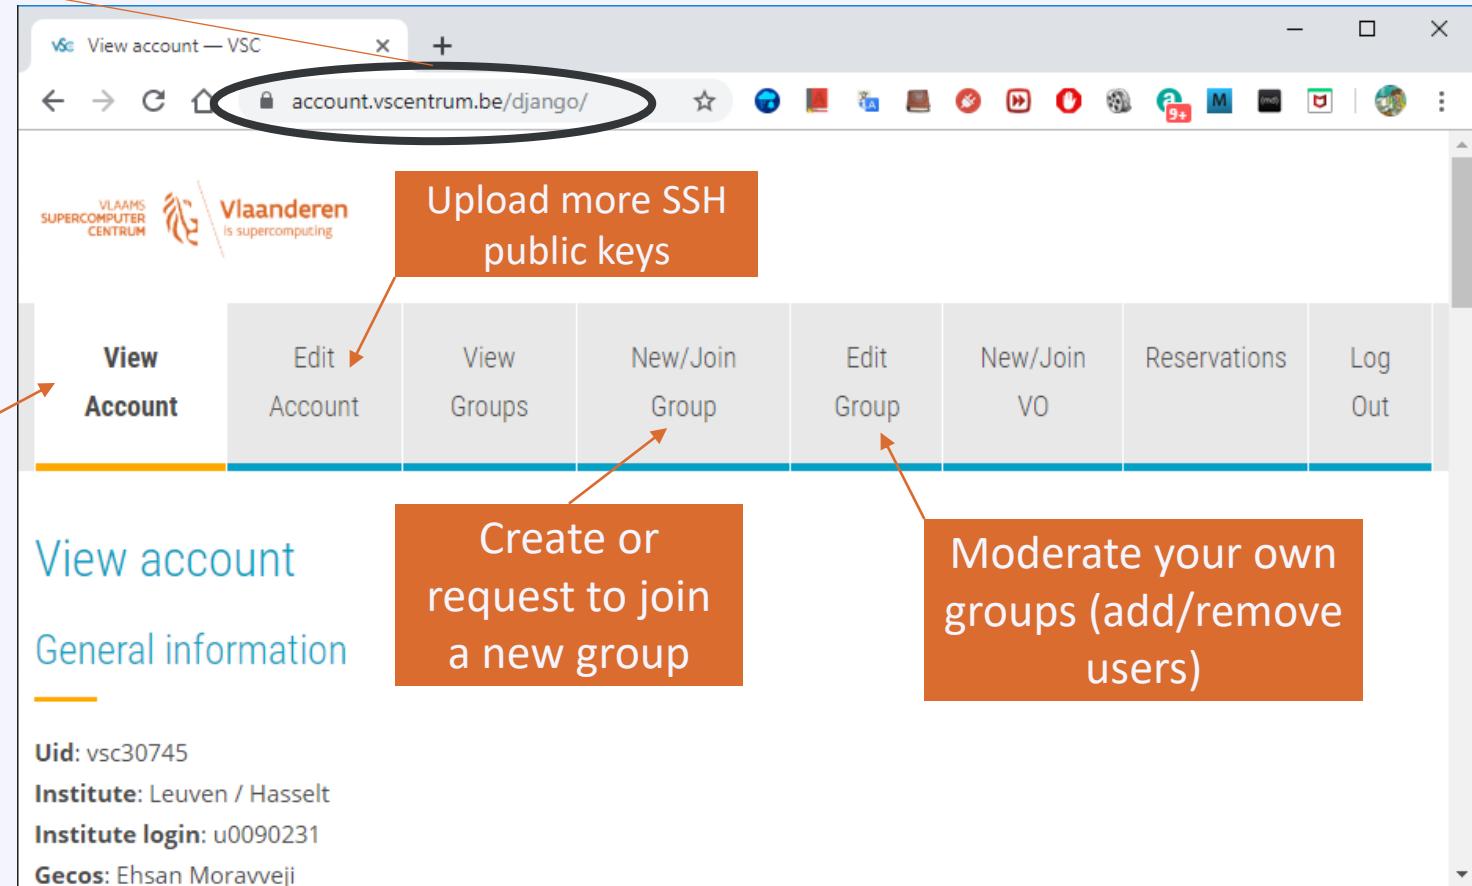
## Application support

- Installation and porting
- Optimisation and debugging
- Benchmarking
- Workflows and best practices

## Training

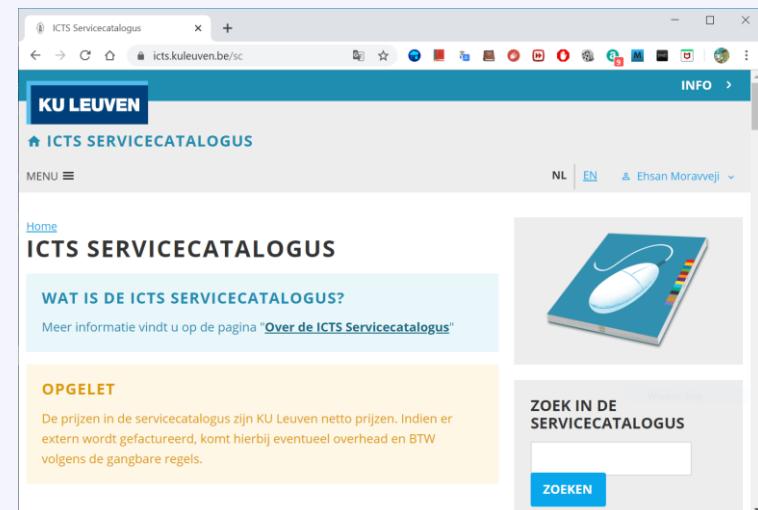
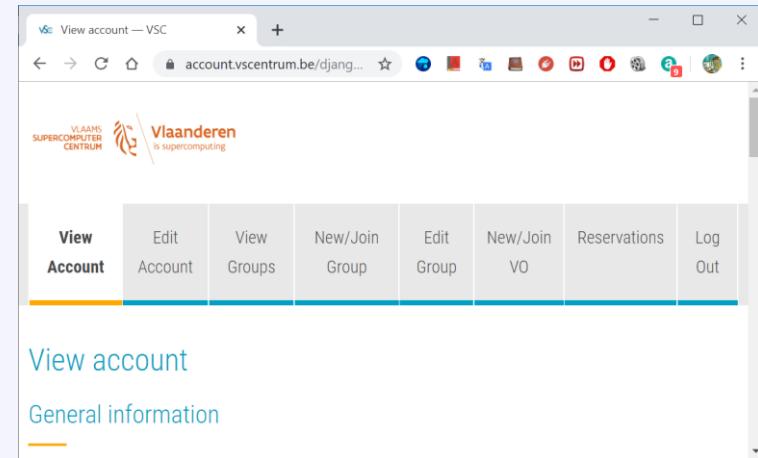
- Documentation and tutorials
- Scheduled trainings / workshops
- On request workshops
- One-to-one sessions

To manage your  
VSC account:  
[account.vscentrum.be](http://account.vscentrum.be)



# Become a VSC user

- Create a secure (4096 bit) [SSH key pairs](#)  
Upload it on the account page: [www.account.vscentrum.be](http://www.account.vscentrum.be)
- You need to [request a VSC account](#)  
Normally processed swiftly
- Request [introductory credits](#) (2M free credits for 6 months)
- Request [project credits](#) (for supervisors and project leaders)
  - You need to create a VSC group
  - Add users to the group to give them access to use credits
  - Fill out the request form
- Extra storage requests
  - Scratch extension: free of charge
  - Staging fileset: 20 € per TB per year
- All service costs (compute and storage) are all explained  
Go to ICTS service catalogus: <https://icts.kuleuven.be/sc>  
Click on [High Performance Computing \(NL/EN\)](#)



# VSC training 2020/2021

- Introductory

Matlab (Supcalculator)  
Matlab Programming

Linux

HPC intro

Linux for HPC

Make intro

Infosessions:

- Containers
- Notebooks

Linux scripting

Linux tools

Version control systems

- Intermediate

C++ for scientific computing

Fortran for programmers

C

- Python as a second language
- Python: System programming
- Scientific Python
- Python for Software engineering
- Python for data science
- Python for machine learning

worker/atools

- Advanced

High Performance Python

?

MPI

OpenMP

Debugging techniques

Code optimization

- Specialized track

# To Acknowledge VSC in publications

## Why?

- a contractual obligation for the VSC
- helps VSC secure funding
- you will benefit from it in the long run

## At KU Leuven

- add the relevant papers to the virtual collection "High Performance Computing" in Lirias

## In het nederlands

*De rekeninfrastructuur en dienstverlening gebruikt in dit werk, werd voorzien door het VSC (Vlaams Supercomputer Centrum), gefinancierd door het FWO en de Vlaamse regering – departement EWI.*

## In English

*The computational resources and services used in this work were provided by the VSC (Flemish Supercomputer Center), funded by the Research Foundation - Flanders (FWO) and the Flemish Government – department EWI.*



## Tier-2 Clusters

# Tier-2 Clusters @ KU Leuven

**Genius** (since 2018)  
250 nodes; 8,936 cores



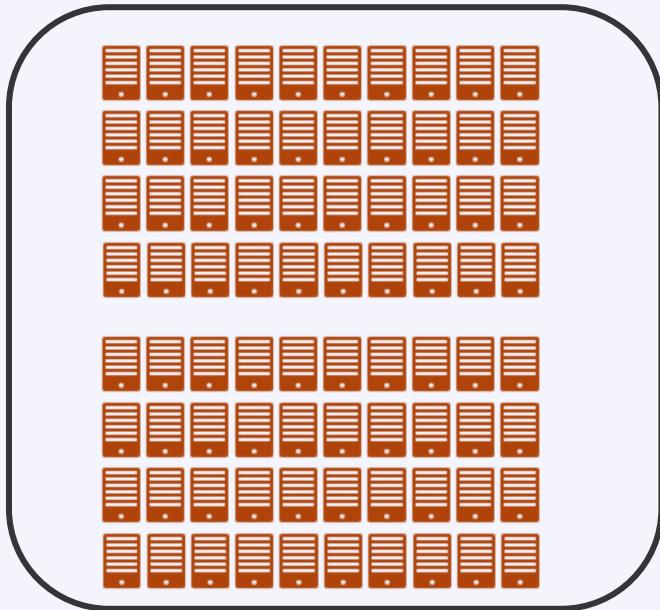
**wICE** (since 9/2022)  
186 nodes; 13,392 cores



# Genius Overview



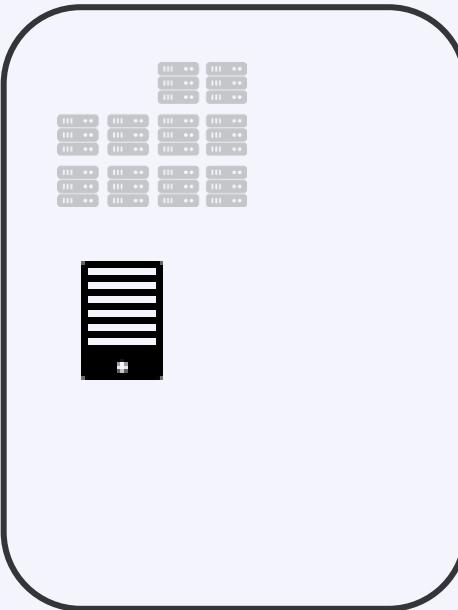
Compute nodes



## Thin nodes

- + 96x Skylake 36c 192 GB
- + 144x CascadeLake 36c 192 GB

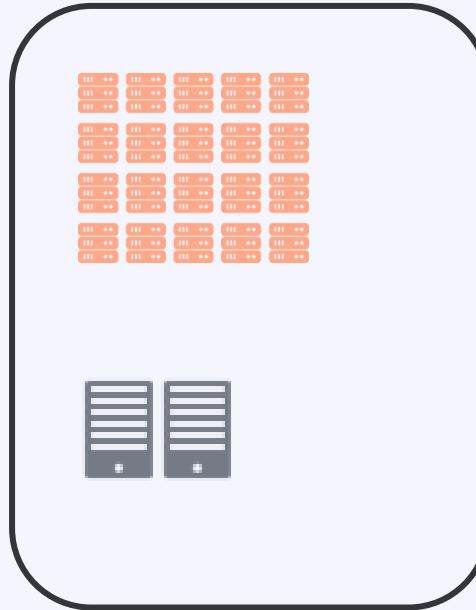
Large memory nodes



## BigMem

- + 10x Skylake 36c 768GB
- + 1x Superdome 112c 6 TB

GPU nodes



## GPUs

- + 20x Skylake 36c 192 GB
- 4x P100 16GB
- + 2x CascadeLake 36c 768GB
- 8x V100 32GB

Exp nodes



## AMD

- 4x AMD Naples 64c 256 GB

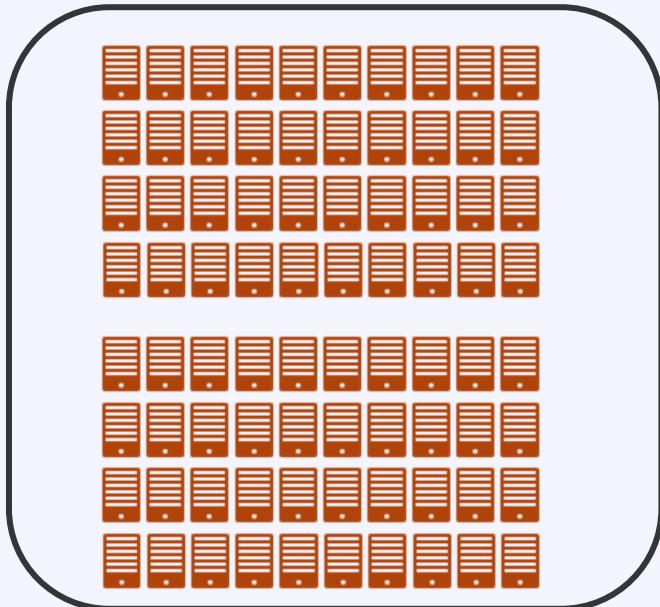
# Tier-2 Cluster: Genius

Type of node	CPU type	Inter-connect	# cores	installed mem	local discs	# nodes
SkyLake	Xeon 6140	IB-EDR	36	192 GB	800 GB	86
SkyLake large mem	Xeon 6140	IB-EDR	36	768 GB	800 GB	10
SkyLake GPU	Xeon 6140 <b>4xP100 SXM2</b>	IB-EDR	36	192 GB	800 GB	20
CascadeLake	Gold 6240	IB-EDR	36	192 GB	800 GB	144
CascadeLake GPU	Gold 6240 <b>8xV100 SMX2</b>	IB-EDR	36	768 GB	800 GB	2
SkyLake Superdome	Gold 6132	Flex Grid	14	6 TB	6 TB	8

# wICE Overview



Compute nodes



Thin nodes

+ 172x Icelake 72c 256 GB

Large memory nodes



BigMem

+ 5x Icelake 72c 2 TB

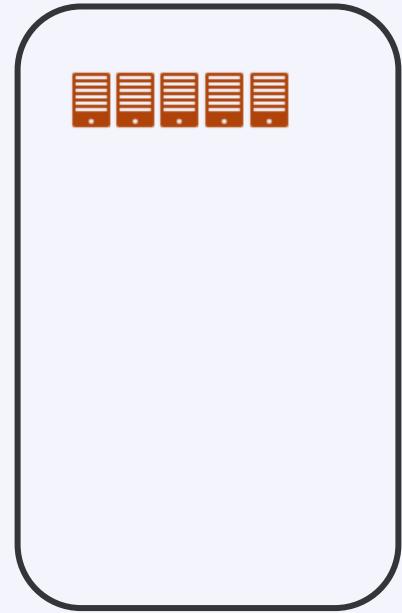
GPU nodes



GPUs

+ 4x Icelake 72c 512 GB  
4x A100 80 GB

Interactive nodes



+ 5x Icelake 64c 512 GB  
1x A100 80 GB

# Tier-2 Cluster: wICE

Type of node	CPU type	Inter-connect	# cores	installed mem	local discs	# nodes
Icelake	Xeon 8358	IB HDR-100	72	<b>256 GB</b>	960 GB	<b>172</b>
Icelake large mem	Xeon 8358	IB HDR-100	72	<b>2048 GB</b>	960 GB	<b>5</b>
Icelake GPU	Xeon 8358 <b>4xA100 SXM2</b> 80GB	IB HDR-100	72	<b>512 GB</b>	960 GB	<b>4</b>
Icelake Interactive	Xeon 8358 <b>1xA100 SXM2</b> 80GB	IB HDR-100	64	<b>512 GB</b>	960 GB	<b>5</b>

# Technical Hardware Specifications

	Tier 2					
Cluster name	Genius			wICE		
Processor type	SkyLake		Cascade Lake	Ice Lake		
Cores per node	36		36	72		
Base Clock Speed	2.3 GHz		2.6 GHz	2.6 GHz		
Total nodes	86	10	144	172	5	5 + 4
Node memory (GB)	192	768	192	256	2048	512
Memory per core (GB)	5.2	21.2	5.3	3.4	25	6.8
Total cores	3,456		4,320	13,392		
Peak performance (Flops/cycle)	16 DP FLOPs/cycle: 8-wide FMA (fused multiply-add) instructions AVX-512		16 DP FLOPs/cycle: 8-wide FMA (fused multiply-add) instructions AVX-512			
Network	Infiniband EDR		Infiniband EDR	Infiniband HDR-100		
Cache (L1 KB/L2 KB/L3 MB)	18x(32i+32d) / 18x1024 / 25 MB		18x(32i+32d) / 18x1024 / 25 MB	72x(48i+32d) / 72x 1280K / 48 MB		



# Storage

# Overview of the storage infrastructure

- Your files are owned only by you.  
Other VSC users have no permission to read/write/execute your files (POSIX)
- A VSC account has 3 default storages (free of charge)
  - \$VSC\_HOME
  - \$VSC\_DATA
  - \$VSC\_SCRATCH
- You can additionally request staging storage
- Different storage volumes have different:
  - mount point
  - size and performance
  - use case
  - backup and maintenance policy
- More info on [ICTS Service Catalog](#) (EN/NL)

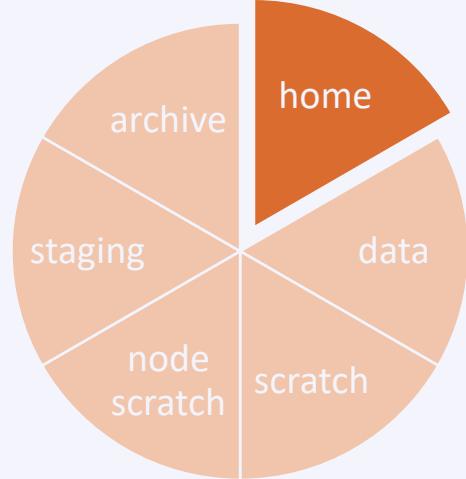
# Storage

- [Request form for extra storage](#)
- [More information](#)
- Do not use /tmp**  
It is only 10 GB and is reserved for the OS  
and root processes  
Your application can crash if using /tmp
- You are automatically logged into your home  
folder upon login.  
Make sure you immediately go to your other  
storages, e.g.  
`$ cd $VSC_DATA`
- Always check your storage balance using  
myquota command

Example

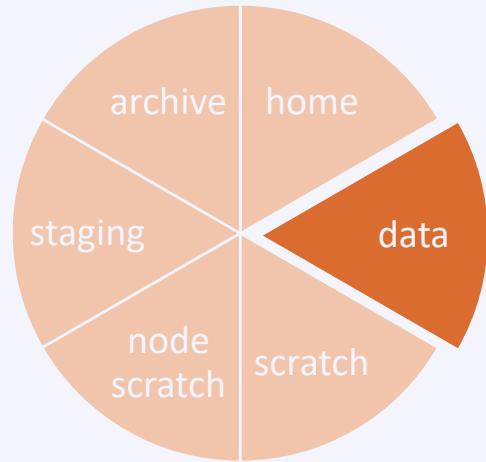
```
$ myquota
file system $VSC_HOME
    Blocks: 1479M of 3072M
    Files: 12934 of 100k
file system $VSC_DATA
    Blocks: 12G of 75G
    Files: 1043k of 10000k
file system $VSC_SCRATCH
    Blocks: 15M of 1.5T
```

# Storage



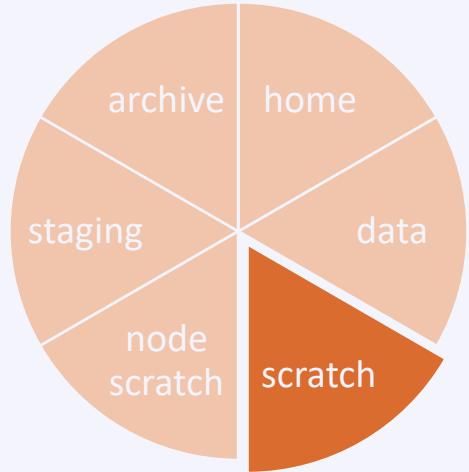
Storage	home folder
Env. Variable	\$VSC_HOME
Filesystem Type	NFS
Access	Global
Backup	Hourly, daily, weekly (until last month) inside the .snapshot folder.
Default Quota	3 GB
Extension	Not possible
Usage	Only storing SSH keys, config files
Remarks	<ul style="list-style-type: none"><li>- Stay away from using it</li><li>- Can easily overflow:<ul style="list-style-type: none"><li>+ Your jobs may crash</li><li>+ Login issues</li></ul></li></ul>

# Storage



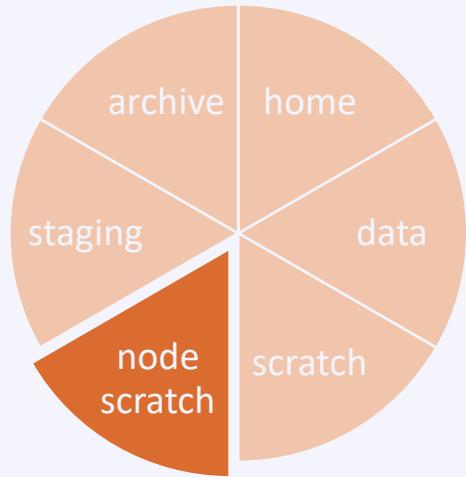
Storage	data folder
Env. Variable	\$VSC_DATA
Filesystem Type	NFS
Access	Global
Backup	Hourly, daily, weekly (until last month) inside the .snapshot folder.
Default Quota	75 GB
Extension	On purchase
Usage	Your data, code, software, results
Remarks	<ul style="list-style-type: none"><li>- Permanent storage for initial/final results</li><li>- Not optimal for intensive or parallel I/O</li></ul>

# Storage



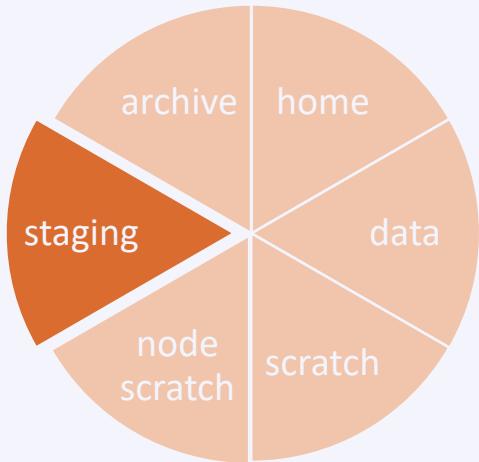
Storage	scratch folder
Env. Variable	\$VSC_SCRATCH
Filesystem Type	Lustre
Access	Global
Backup	delete after 30 days from last access
Default Quota	500 GB
Extension	Free
Usage	Intensive, parallel I/O, temporary files
Remarks	<ul style="list-style-type: none"><li>- Recommended storage for all jobs</li><li>- Copy scratch files to VSC_DATA or local storage after jobs are done</li><li>- Deleted files cannot be recovered</li></ul>

# Storage



Storage	Node scratch folder
Env. Variable	\$VSC_SCRATCH_NODE
Filesystem Type	Lustre
Access	On compute node, only at runtime
Backup	None
Default Quota	591 GB
Extension	<a href="#">Read about beeOND</a>
Usage	Temporary storage at runtime
Remarks	<ul style="list-style-type: none"><li>- Fastest I/O, attached to the node</li><li>- Is cleaned after job terminates</li><li>- Copy the data to your home, scratch, or staging before job ends</li></ul>

# Storage

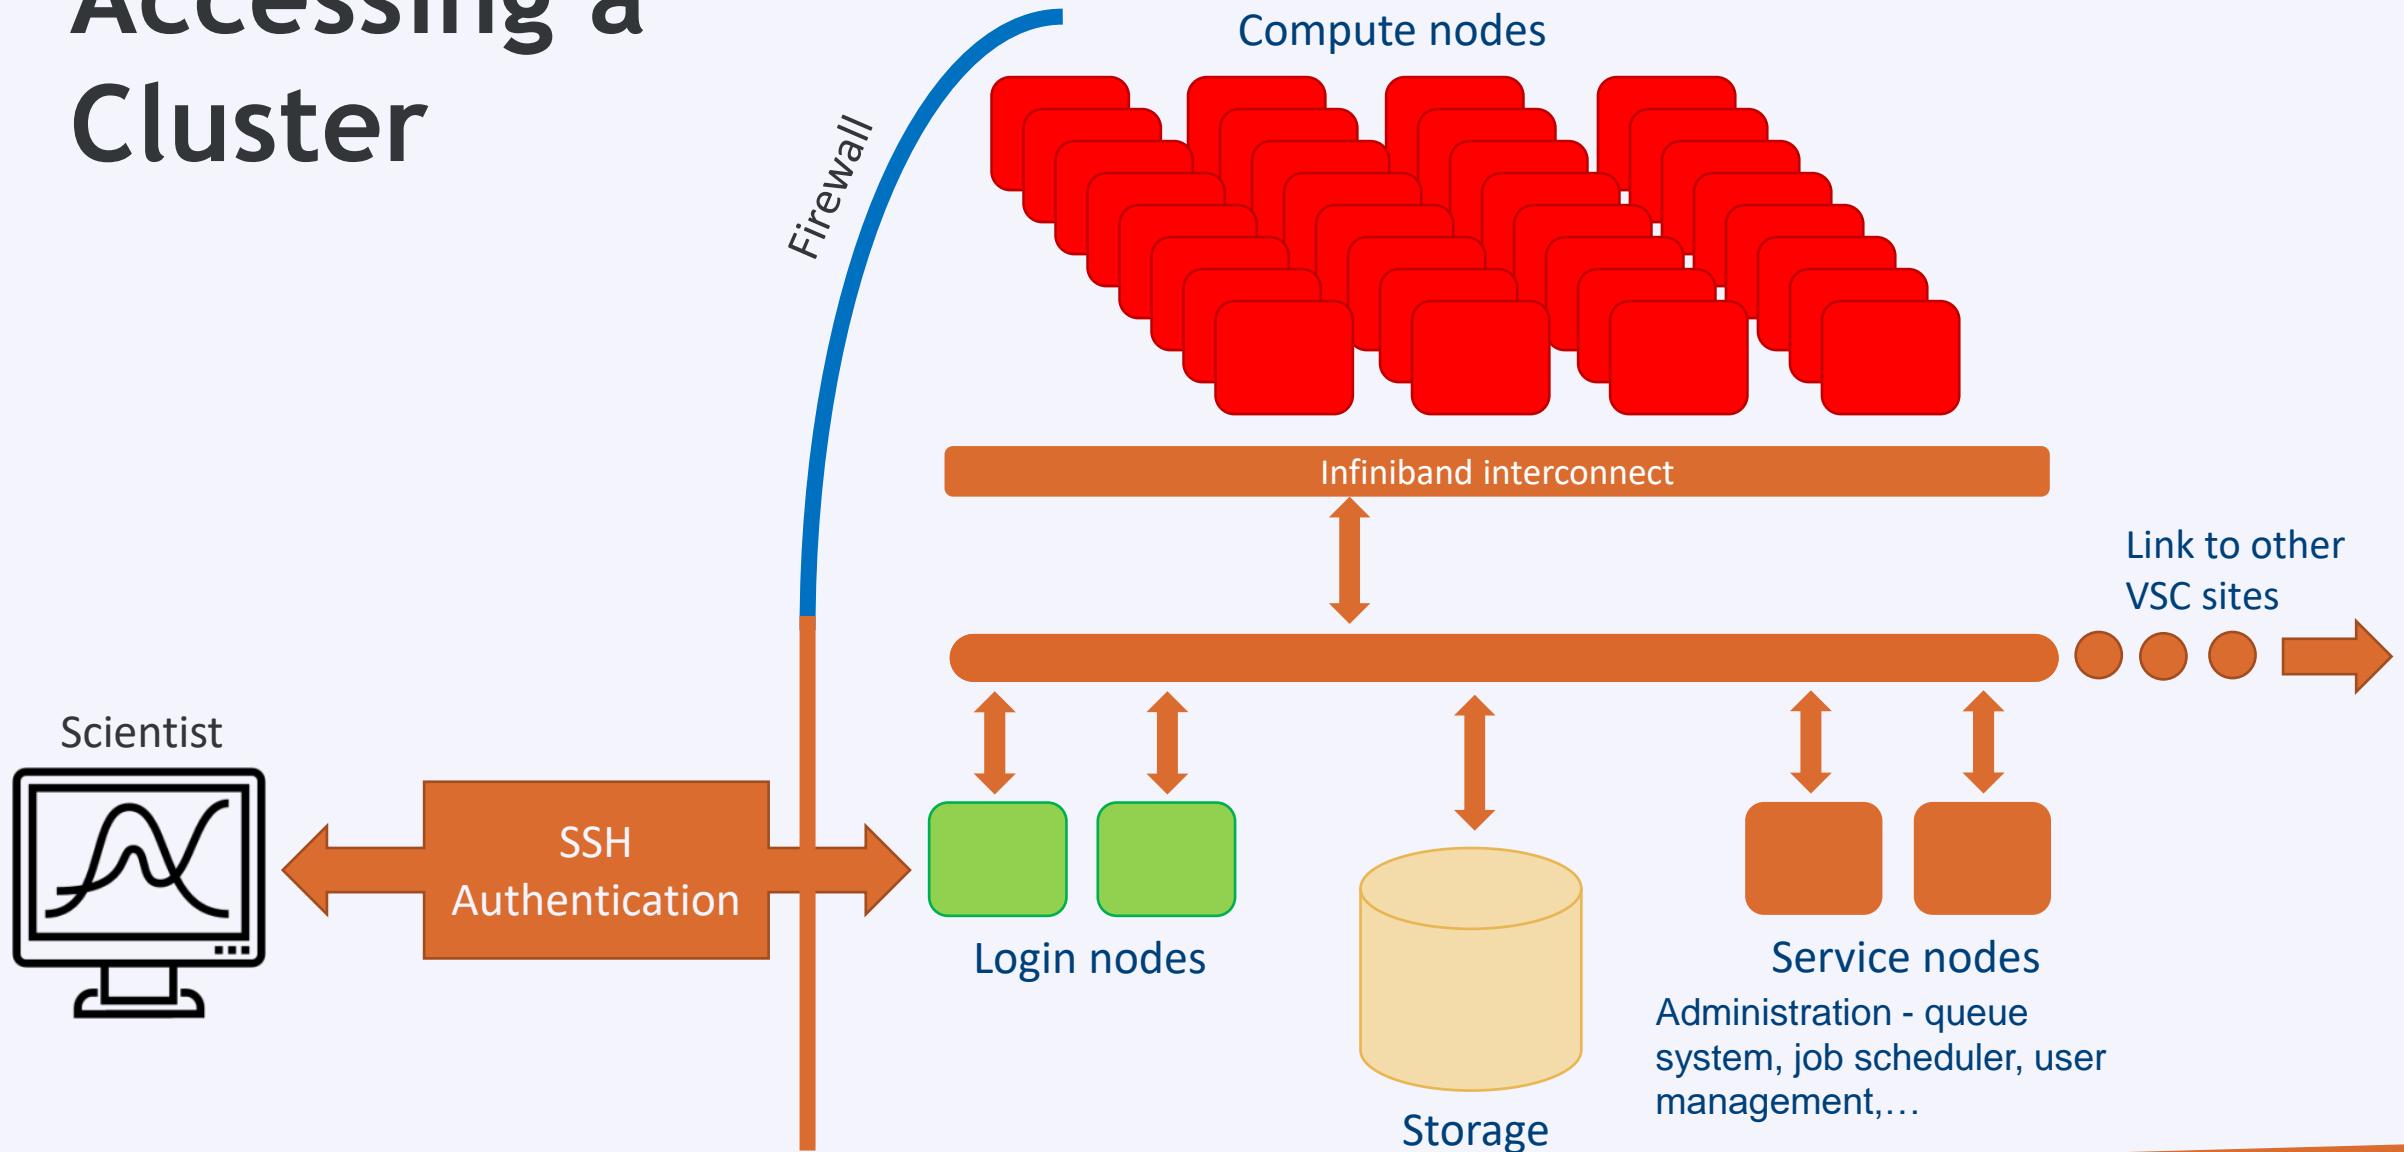


Storage	Staging folder
<b>Path</b>	/staging/leuven/stg_000XX
<b>Filesystem Type</b>	Lustre
<b>Access</b>	On demand, only Tier-2@KUL
<b>Backup</b>	None
<b>Default Quota</b>	None
<b>Extension</b>	On purchase, from 1 TB
<b>Usage</b>	Permanent; share with a group
<b>Remarks</b>	<ul style="list-style-type: none"><li>- Accessible from login/compute nodes</li><li>- Fast, parallel I/O</li></ul>



# Login Nodes & MFA

# Accessing a Cluster



# Using Login Nodes

- To develop and/or compile code and/or software
- To check your storage and credit balance
- To manage jobs (submit, check status, debug, resubmit, ...)
- To move data around
  - within VSC: use data, scratch, staging
  - outside VSC: copy/sync from/to your local storage
- To pre-process or post-process your data/jobs
- To visualize your data
- To share files/folders

Tips

- Login nodes are shared resources
- Do not** execute heavy-lifting tasks (core, memory)
- Instead, submit jobs

Warning

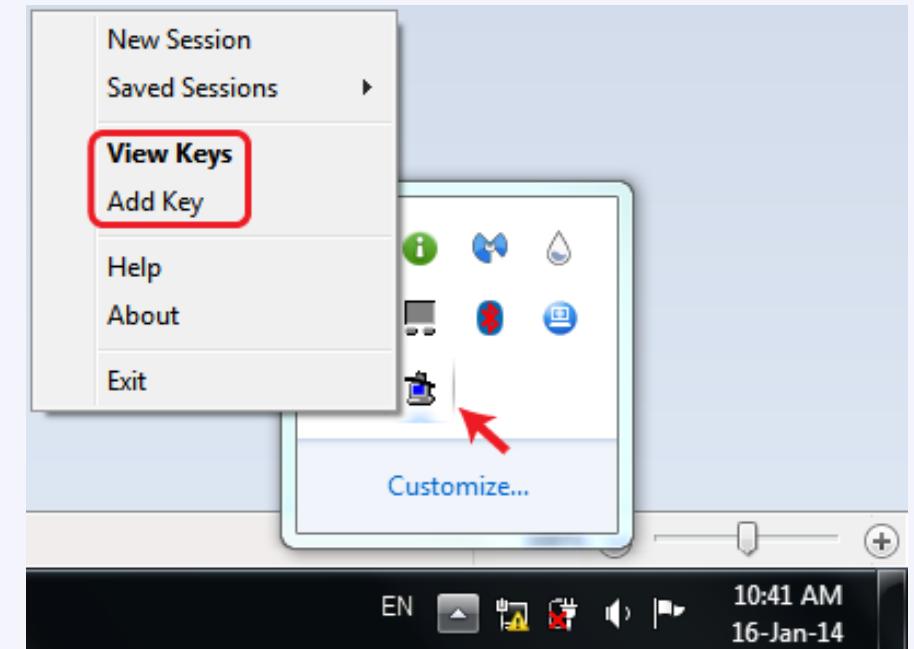
## Login Setup

- PuTTY
- Terminal

# Activate Your SSH Agent

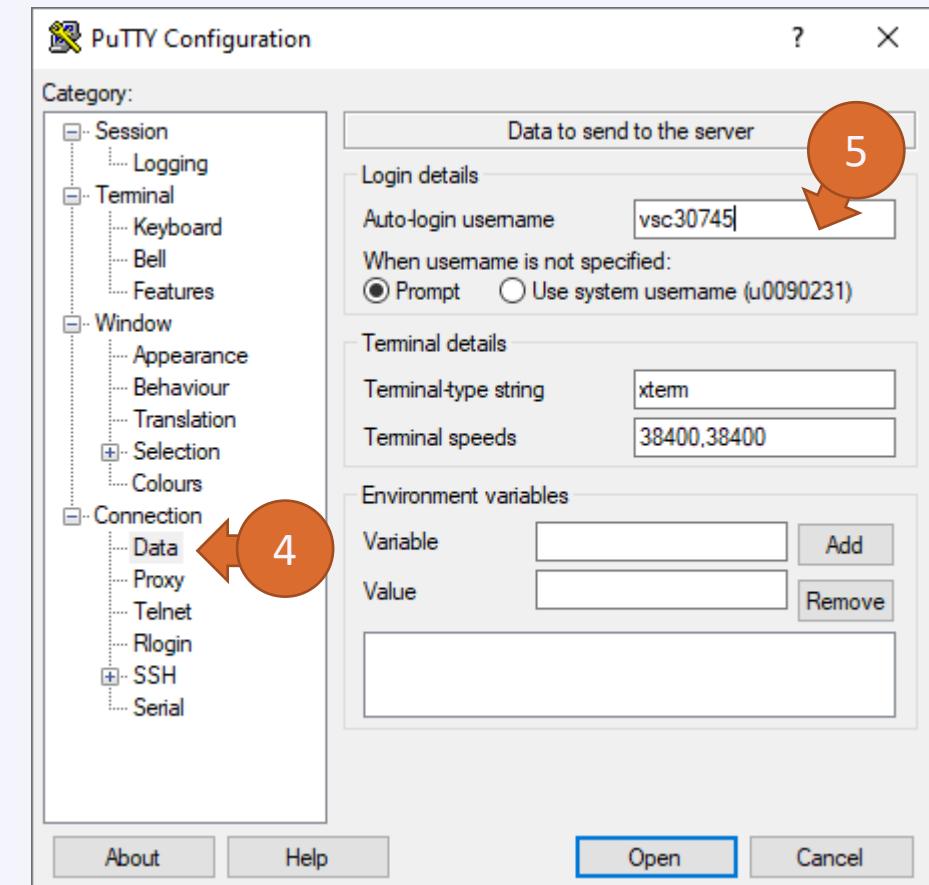
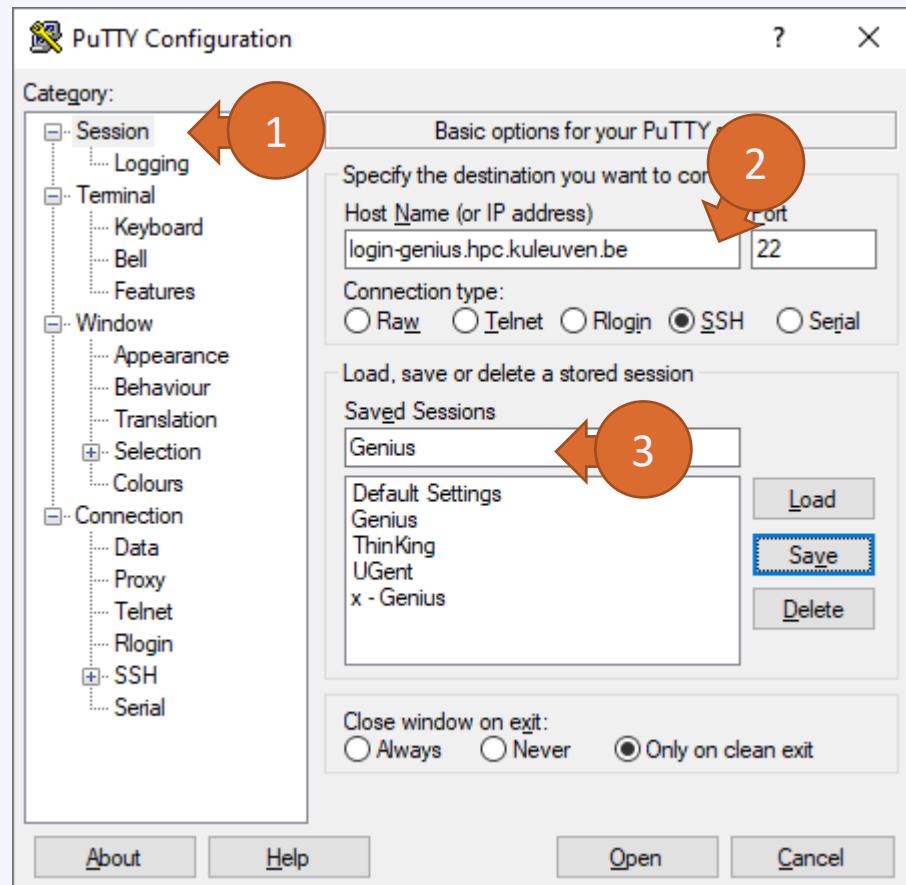
(Windows only)

- When you install PuTTY, the agent called **Pageant** is automatically installed
- Open Pageant; it hides inside your bottom-right tray
- Click on “Add Key” and brows to your private key(s) folder
- After choosing a key, you are asked for a passphrase
- If successful, agent always remembers your SSH keys and Certificates (Multi-Factor Authentication)



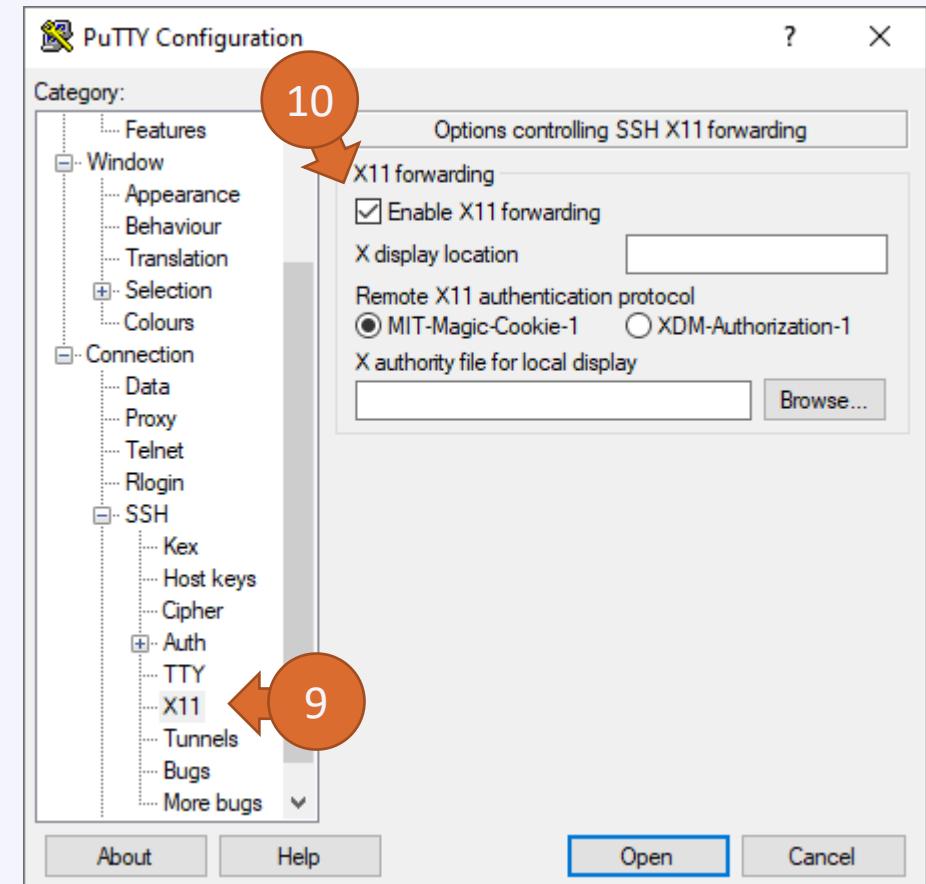
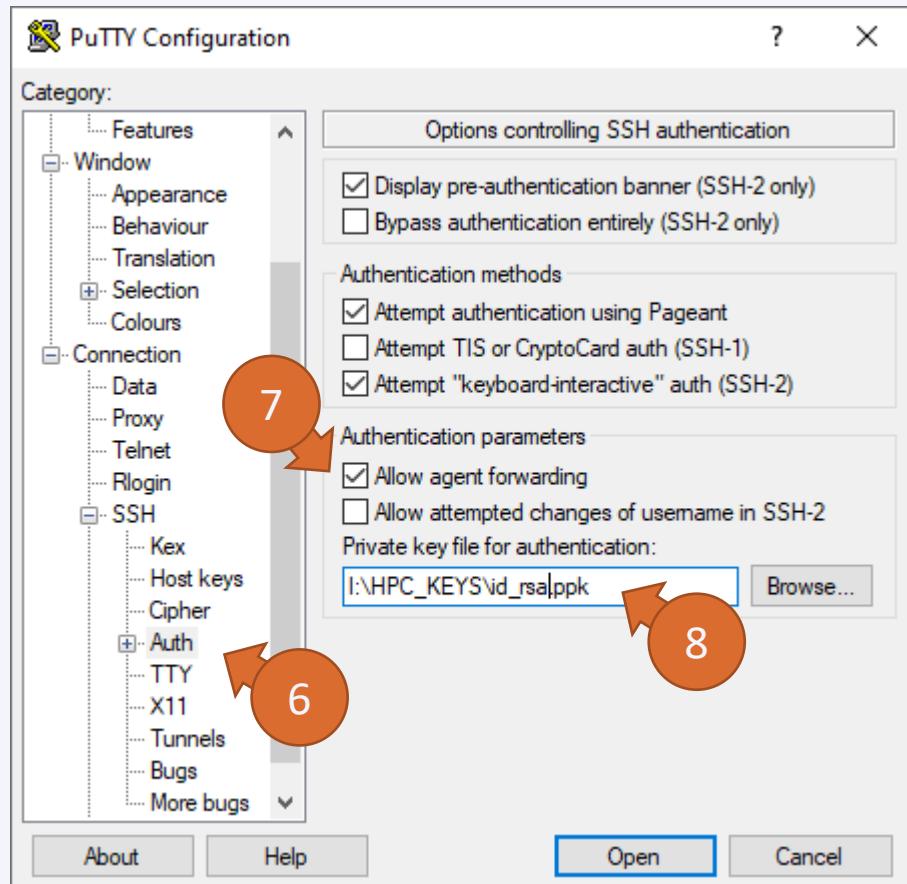
# Setup PuTTY in 12 Clicks

(Windows only)



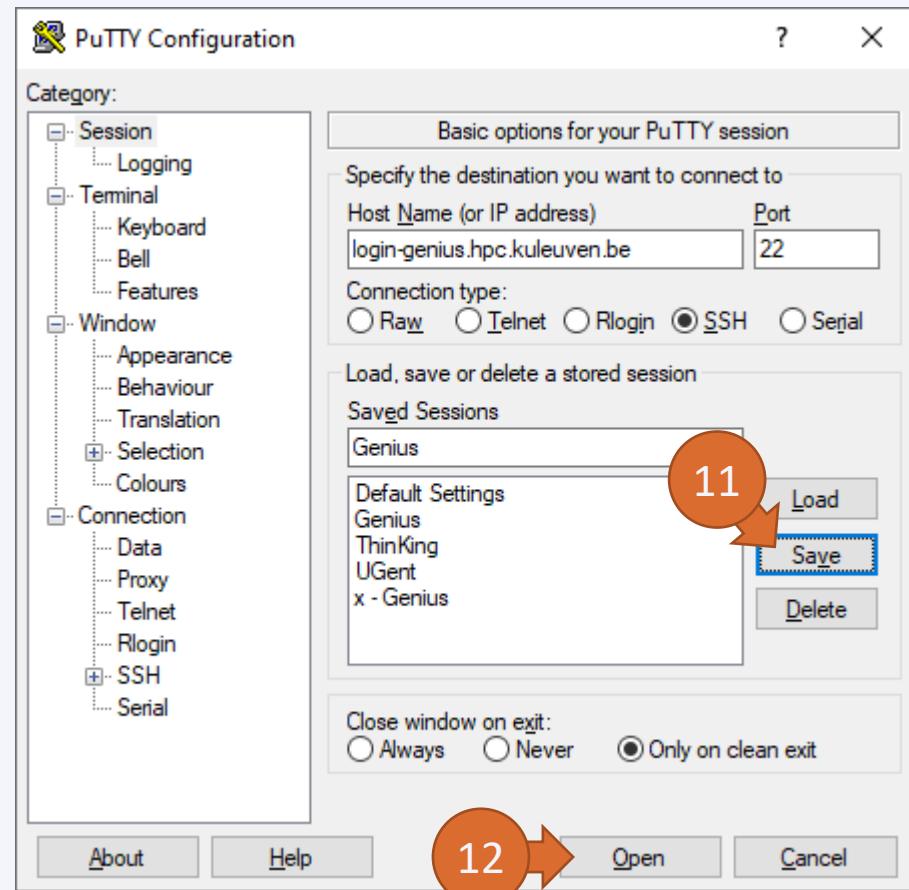
# Setup PuTTY in 12 Clicks

(Windows only)



# Setup PuTTY in 12 Clicks

(Windows only)



If PuTTY asks for **password**, exit immediately, and check the path to your private key. Else you will be blocked for 24h.

A PuTTY terminal window titled 'login4-tier2.hpc.kuleuven.be - PuTTY' is shown. The text in the window reads:  
Using username "vsc30745".  
Authenticating with public key "rsa-key-20170705"  
Passphrase for key "rsa-key-20170705":

A second PuTTY terminal window titled 'login4-tier2.hpc.kuleuven.be - PuTTY' is shown. The text in the window reads:  
Using username "vsc30745".  
Authenticating with public key "rsa-key-20170705"  
Passphrase for key "rsa-key-20170705":  
Last login: Wed Jan 9 15:23:52 2019  
[REDACTED]  
Deze server wordt (deels) met Puppet beheerd; alle wijzigingen die niet via Puppet worden aangebracht, riskeren (automatisch) ongedaan te worden gemaakt.  
This server is managed by Puppet; all changes not imp



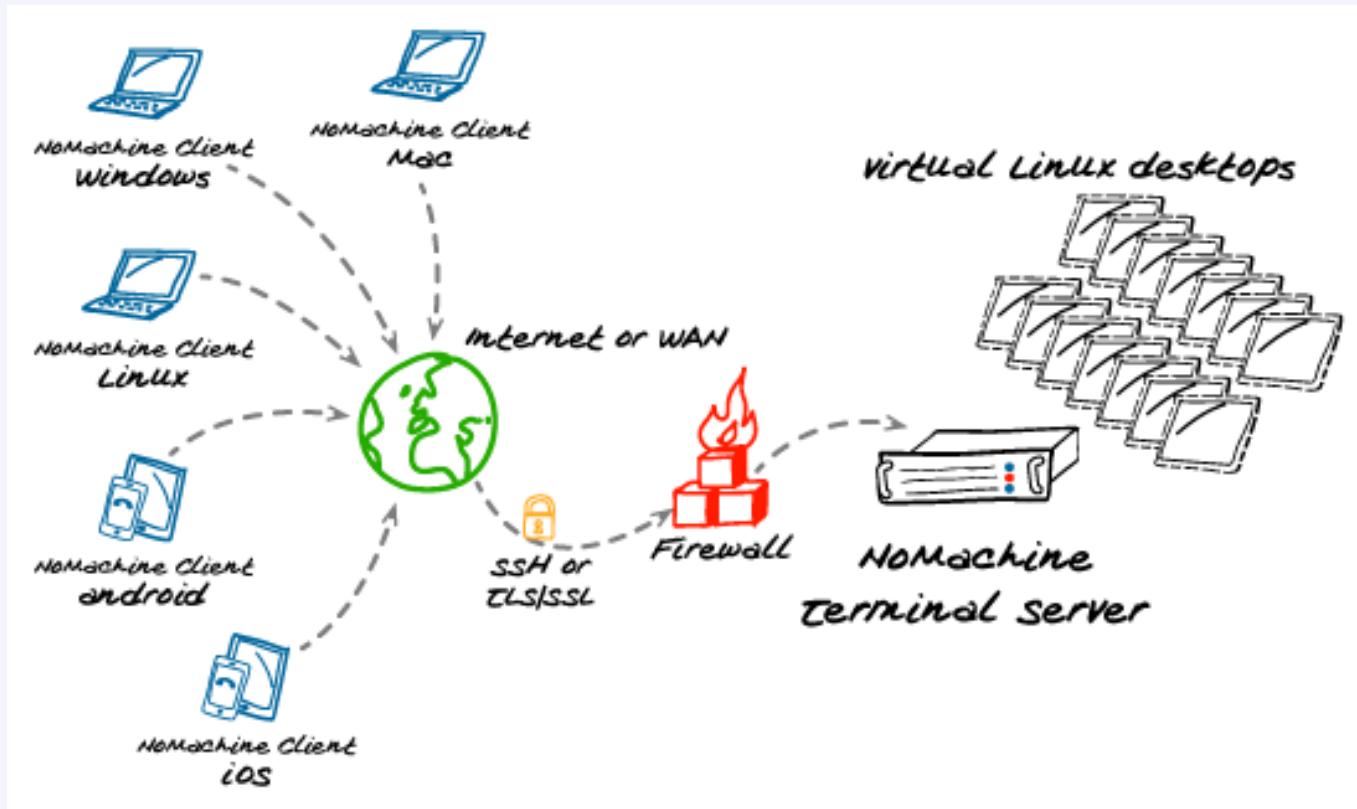
NX

- About
- Setup

# NX – The Graphical Login

NX is:

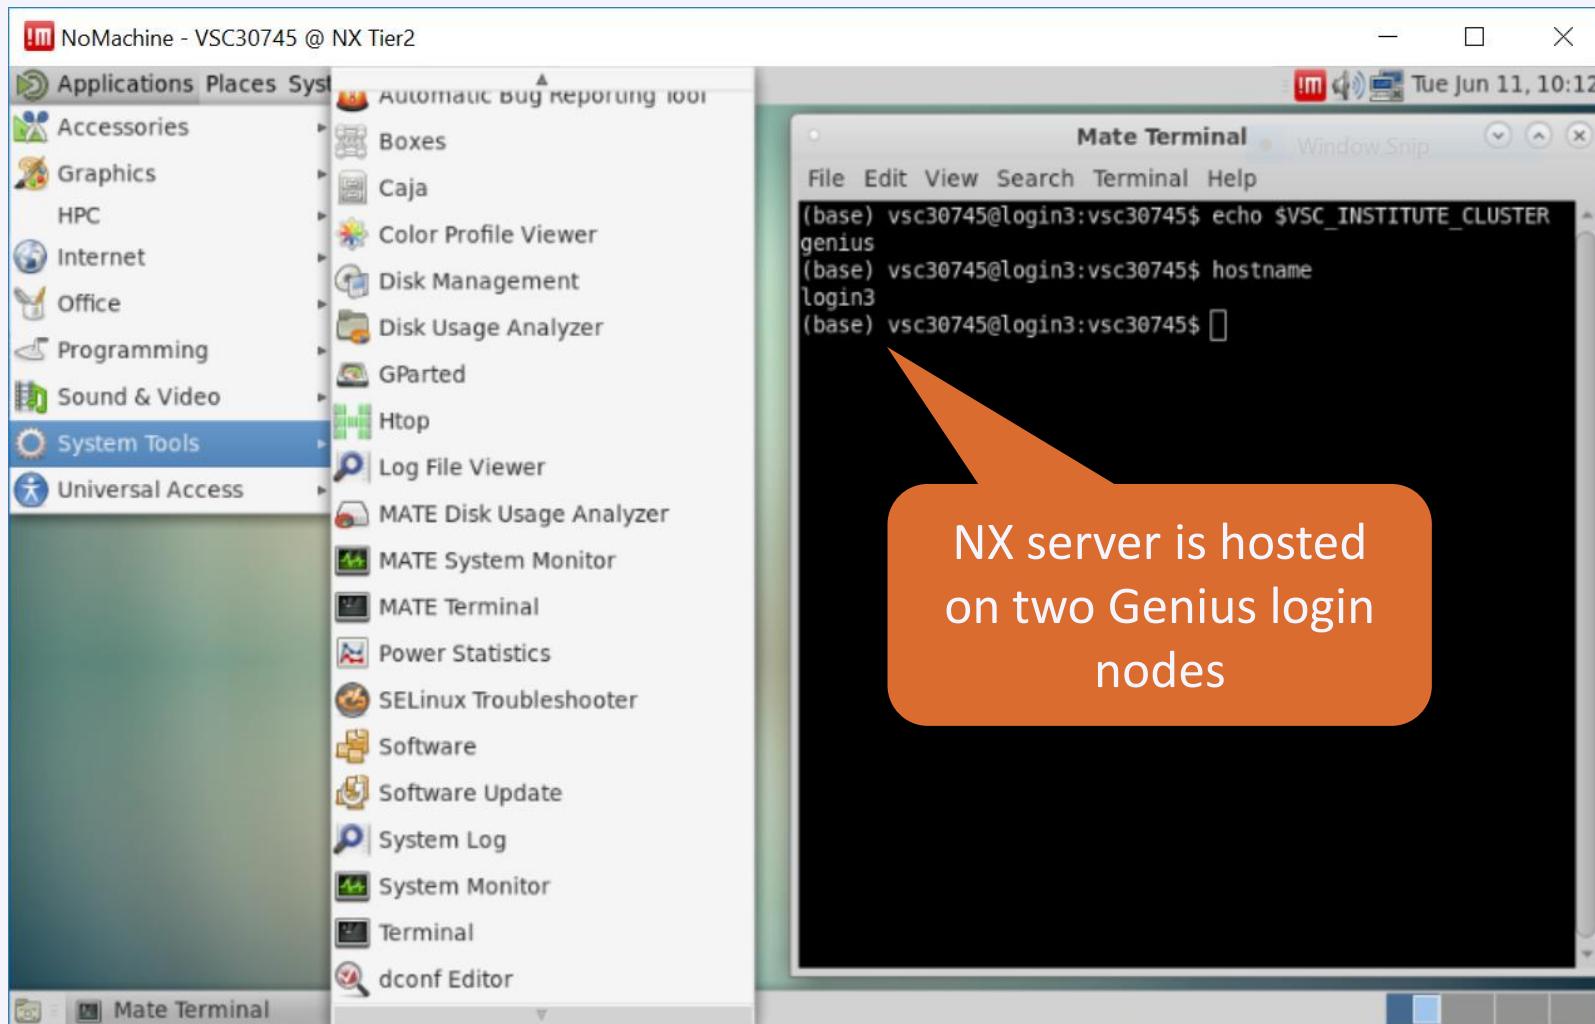
- a graphical remote desktop server
- login to VSC
- start a virtual Linux instance(s)
- Client: [NoMachine](#)
- [Configuration Guide](#)
- Using OpenSSH key  
[Convert your key](#)



# Advantages of NX

- ❑ Graphical login
- ❑ A session stays alive/active even if you close your client (e.g. laptop)
- ❑ Resume a session, and immediately keep working where you left off
- ❑ More interactive jobs
- ❑ Available apps:
  - Internet browser, terminal, ...
  - Text editor, PDF reader, Image viewer, ...
  - Matlab, RStudio, SAS
- ❑ **Remark:**
  - NX is shared among tens of users
  - Do not compute/test your code there**
  - Instead, submit jobs from NX to compute nodes

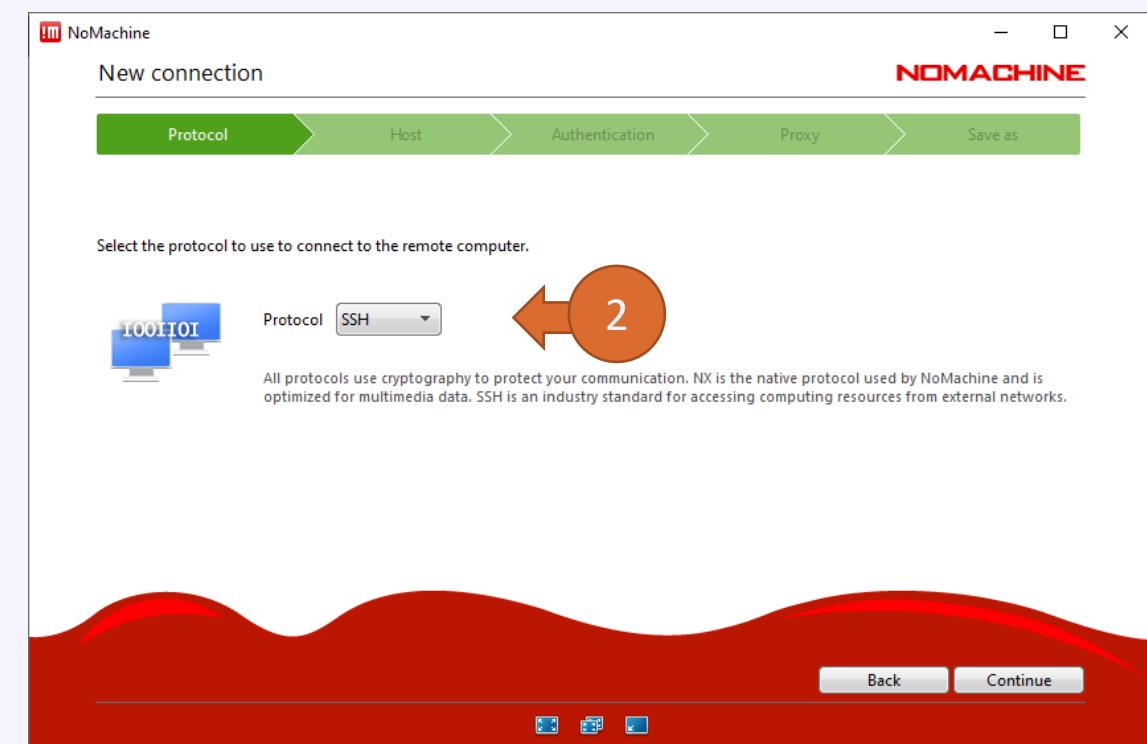
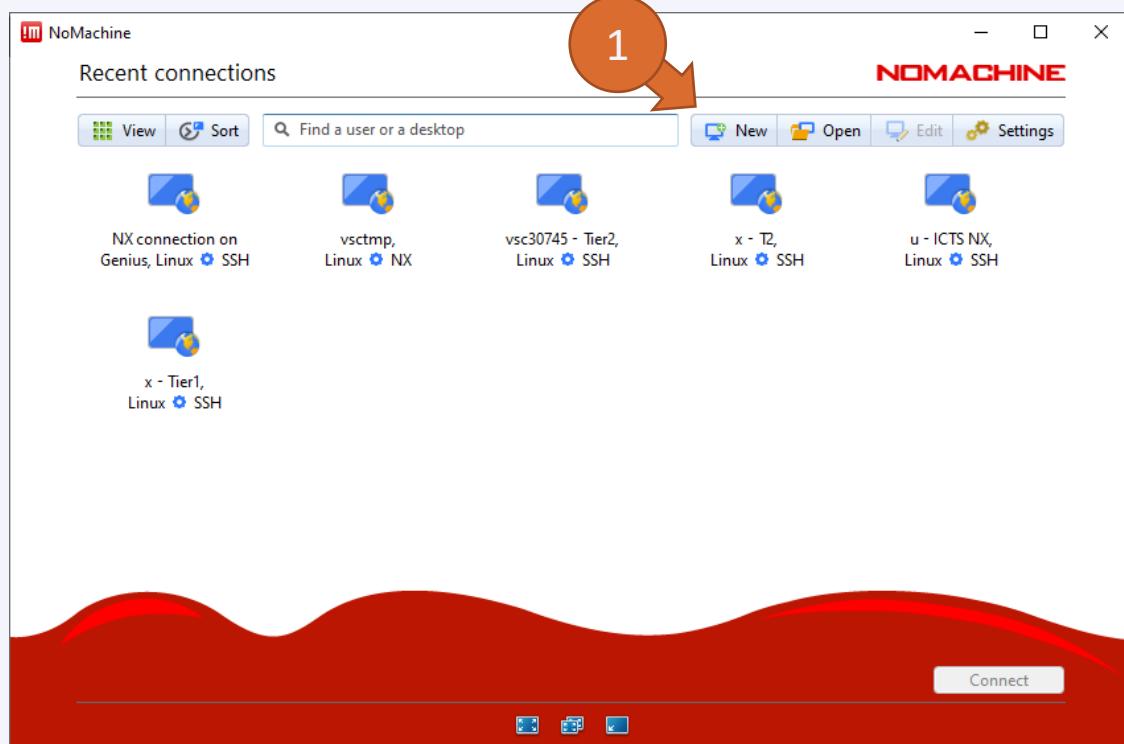
# NX virtual desktop



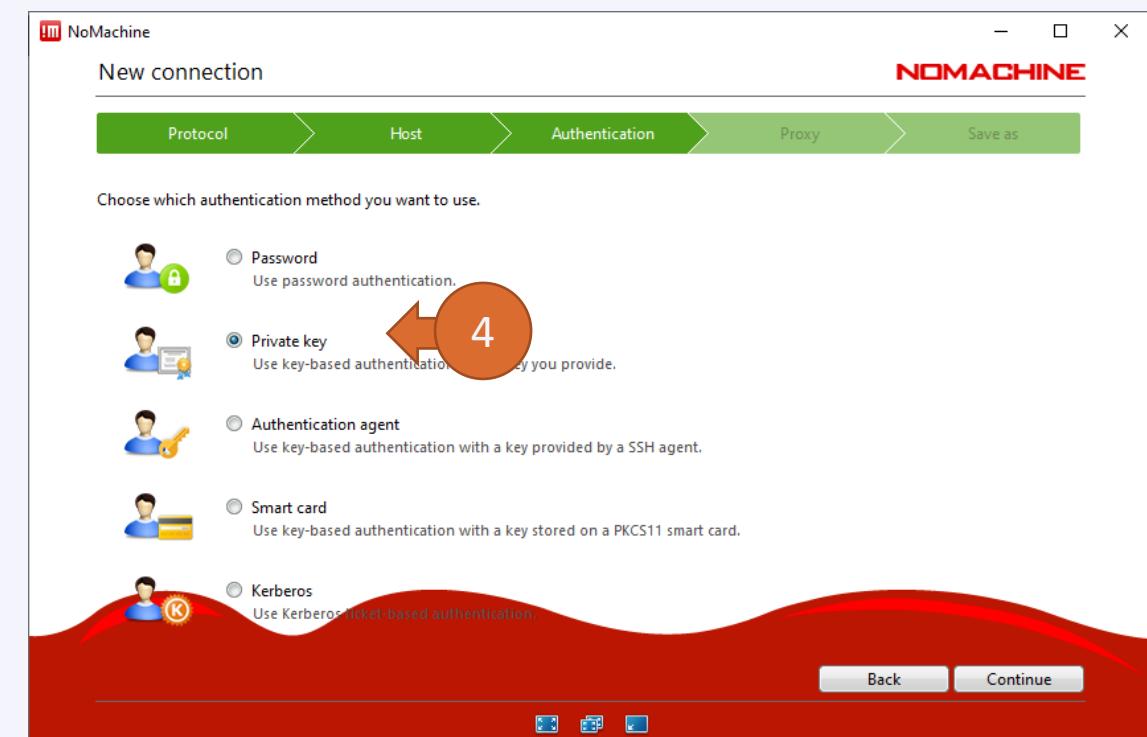
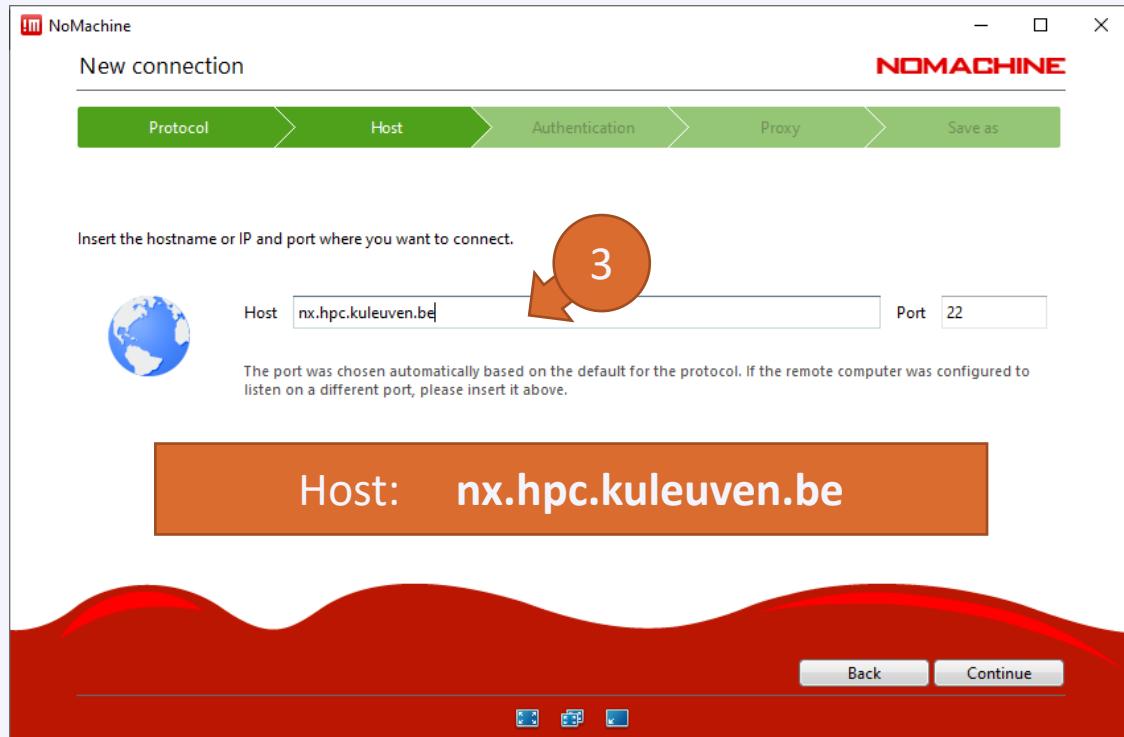
# NX: available software

- Accessories:** Gedit, Vi IMproved, Emacs (dummy version), Calculator
- Graphics:** gThumb (picture viewer), Xpdf Viewer
- Internet:** Firefox
- HPC: Computation:** Matlab (2018a), RStudio, SAS
- Visualisation:** Paraview, VisIt, VMD
- Programming:** Meld Diff Viewer (visual diff and merge tool)
- System tools:** File Browser, Terminal
- Additionally:** Gnuplot (graphing utility), Filezilla (file transfer tool), Evince (PDF, PostScript, TIFF, XPS, DVI Viewer)
- Software launched through modules from Terminal.

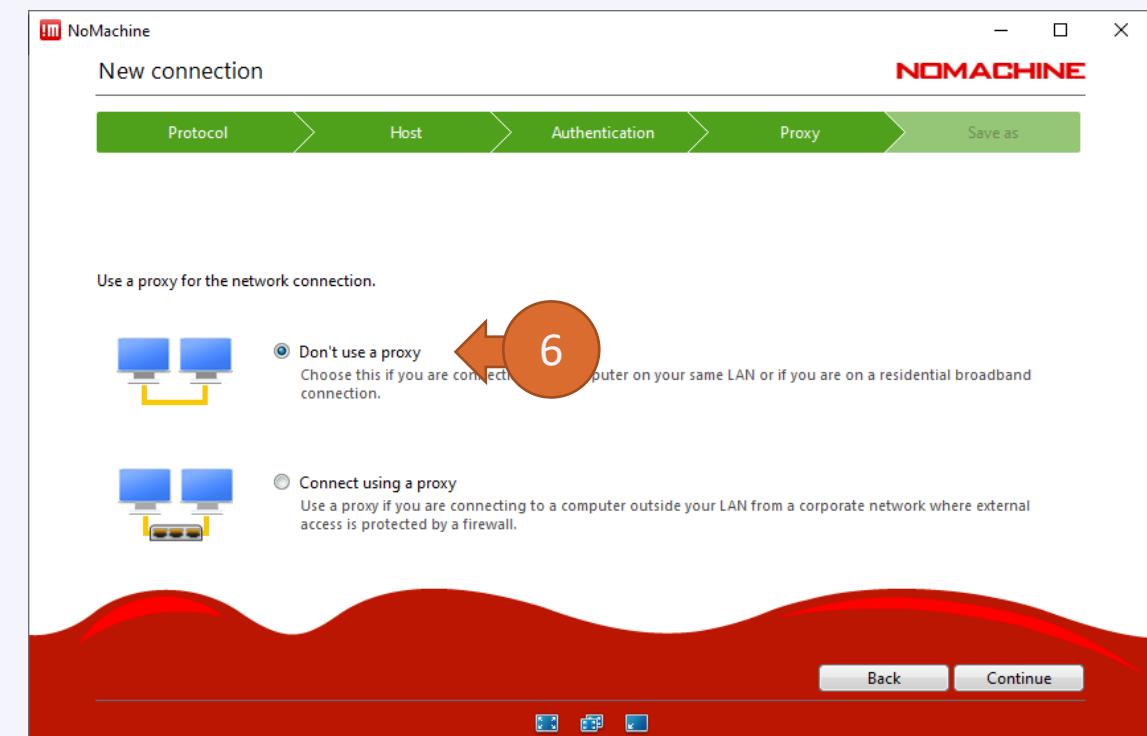
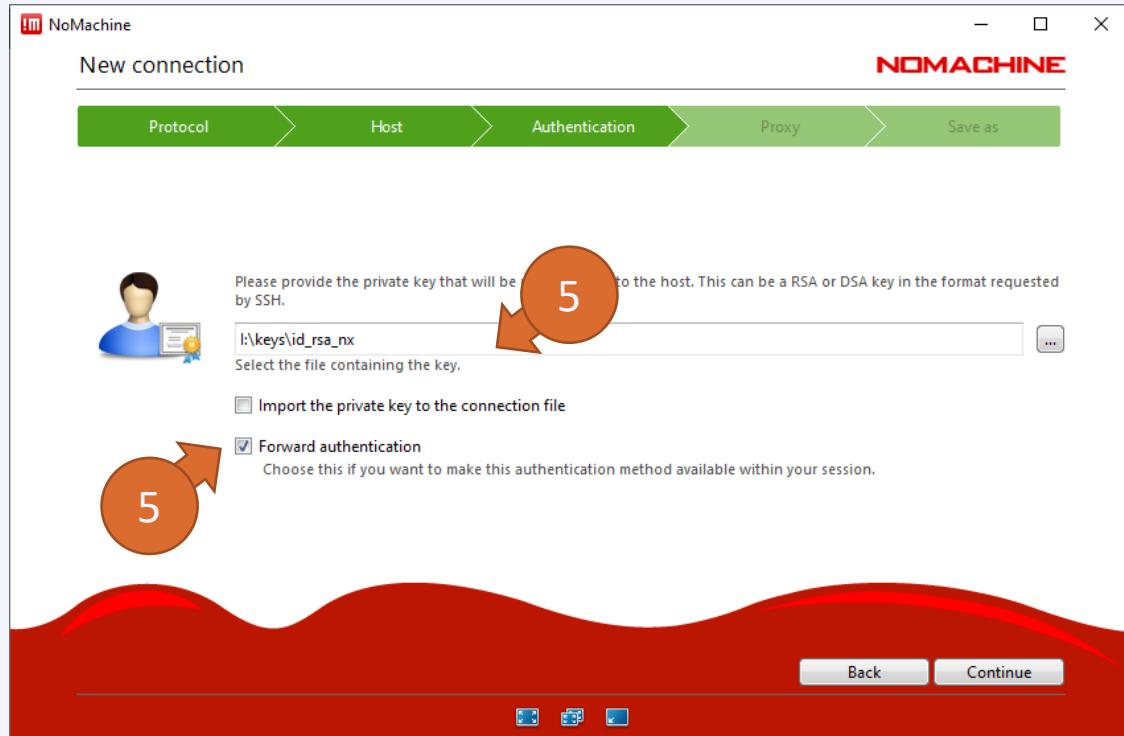
# Setup NX in 10 Steps



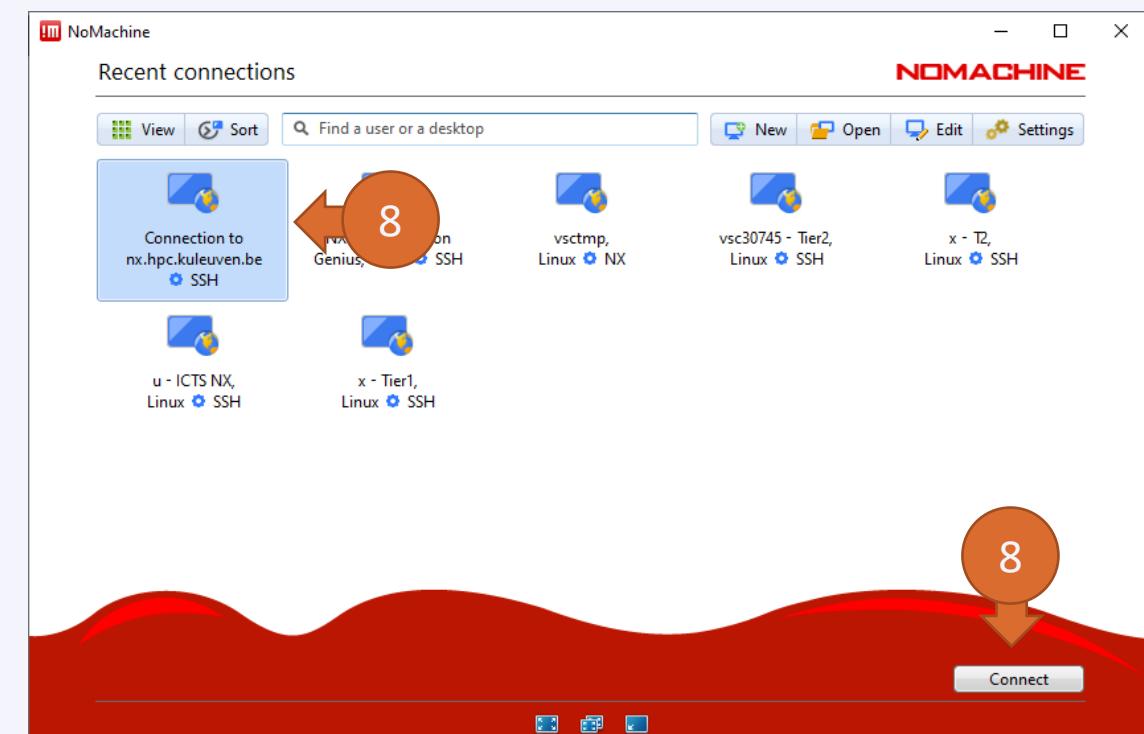
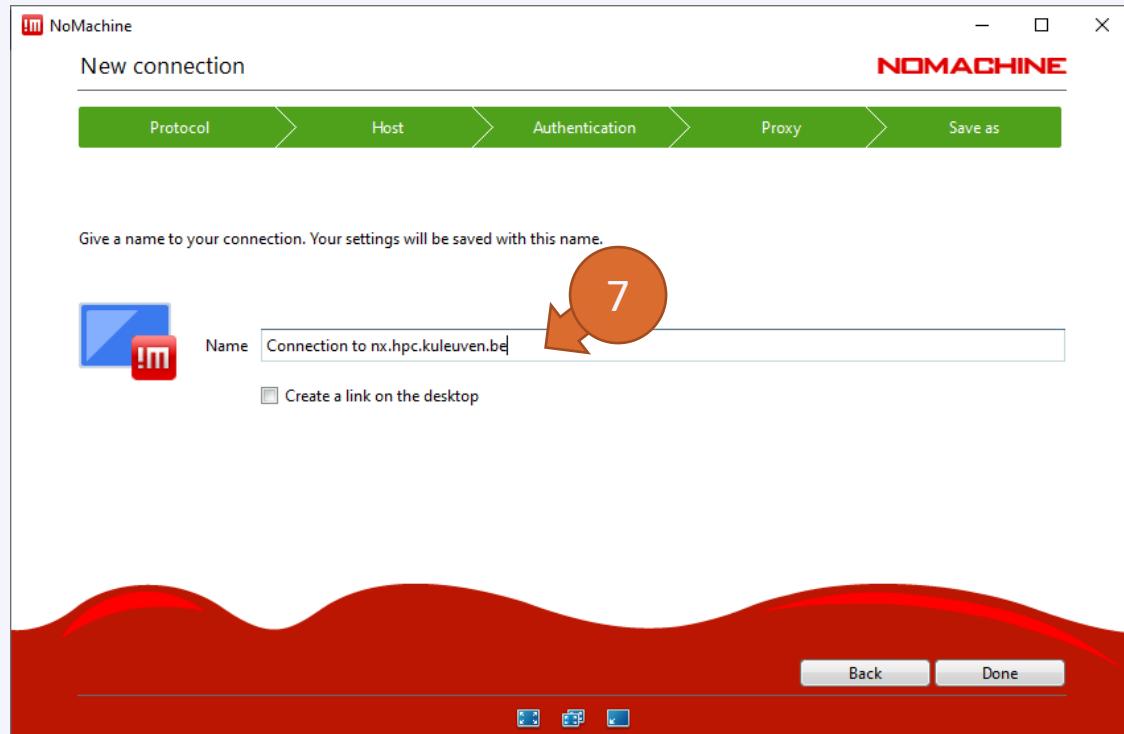
# Setup NX in 10 Steps



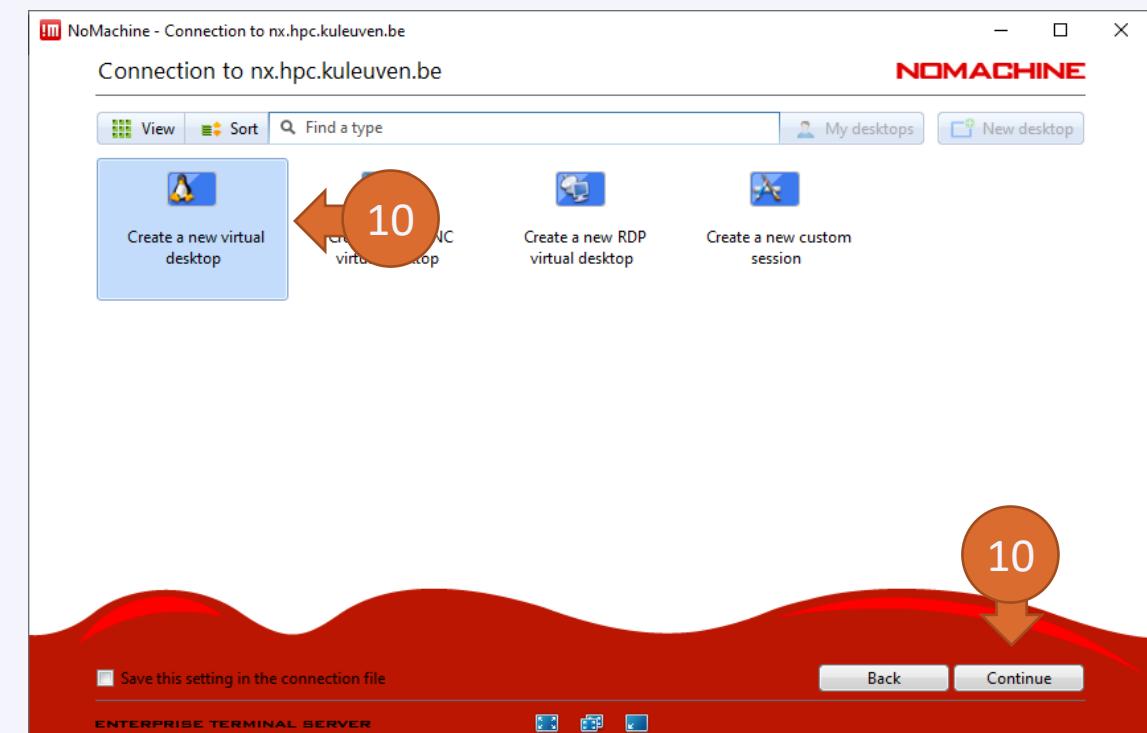
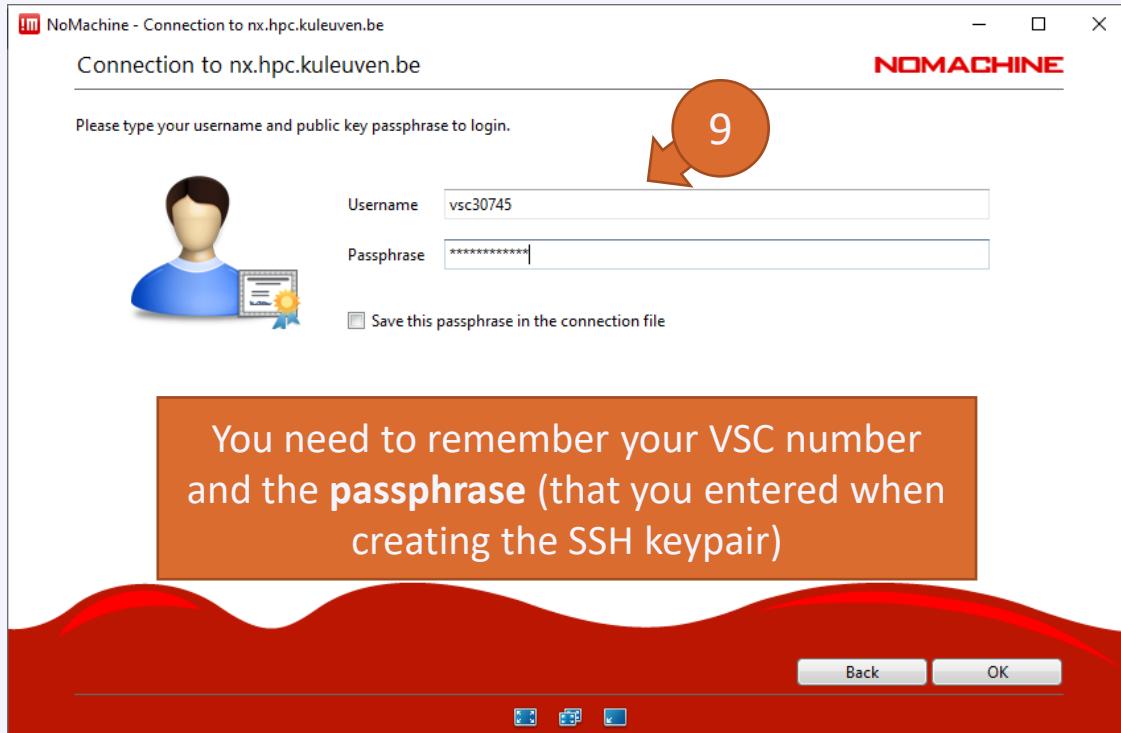
# Setup NX in 10 Steps

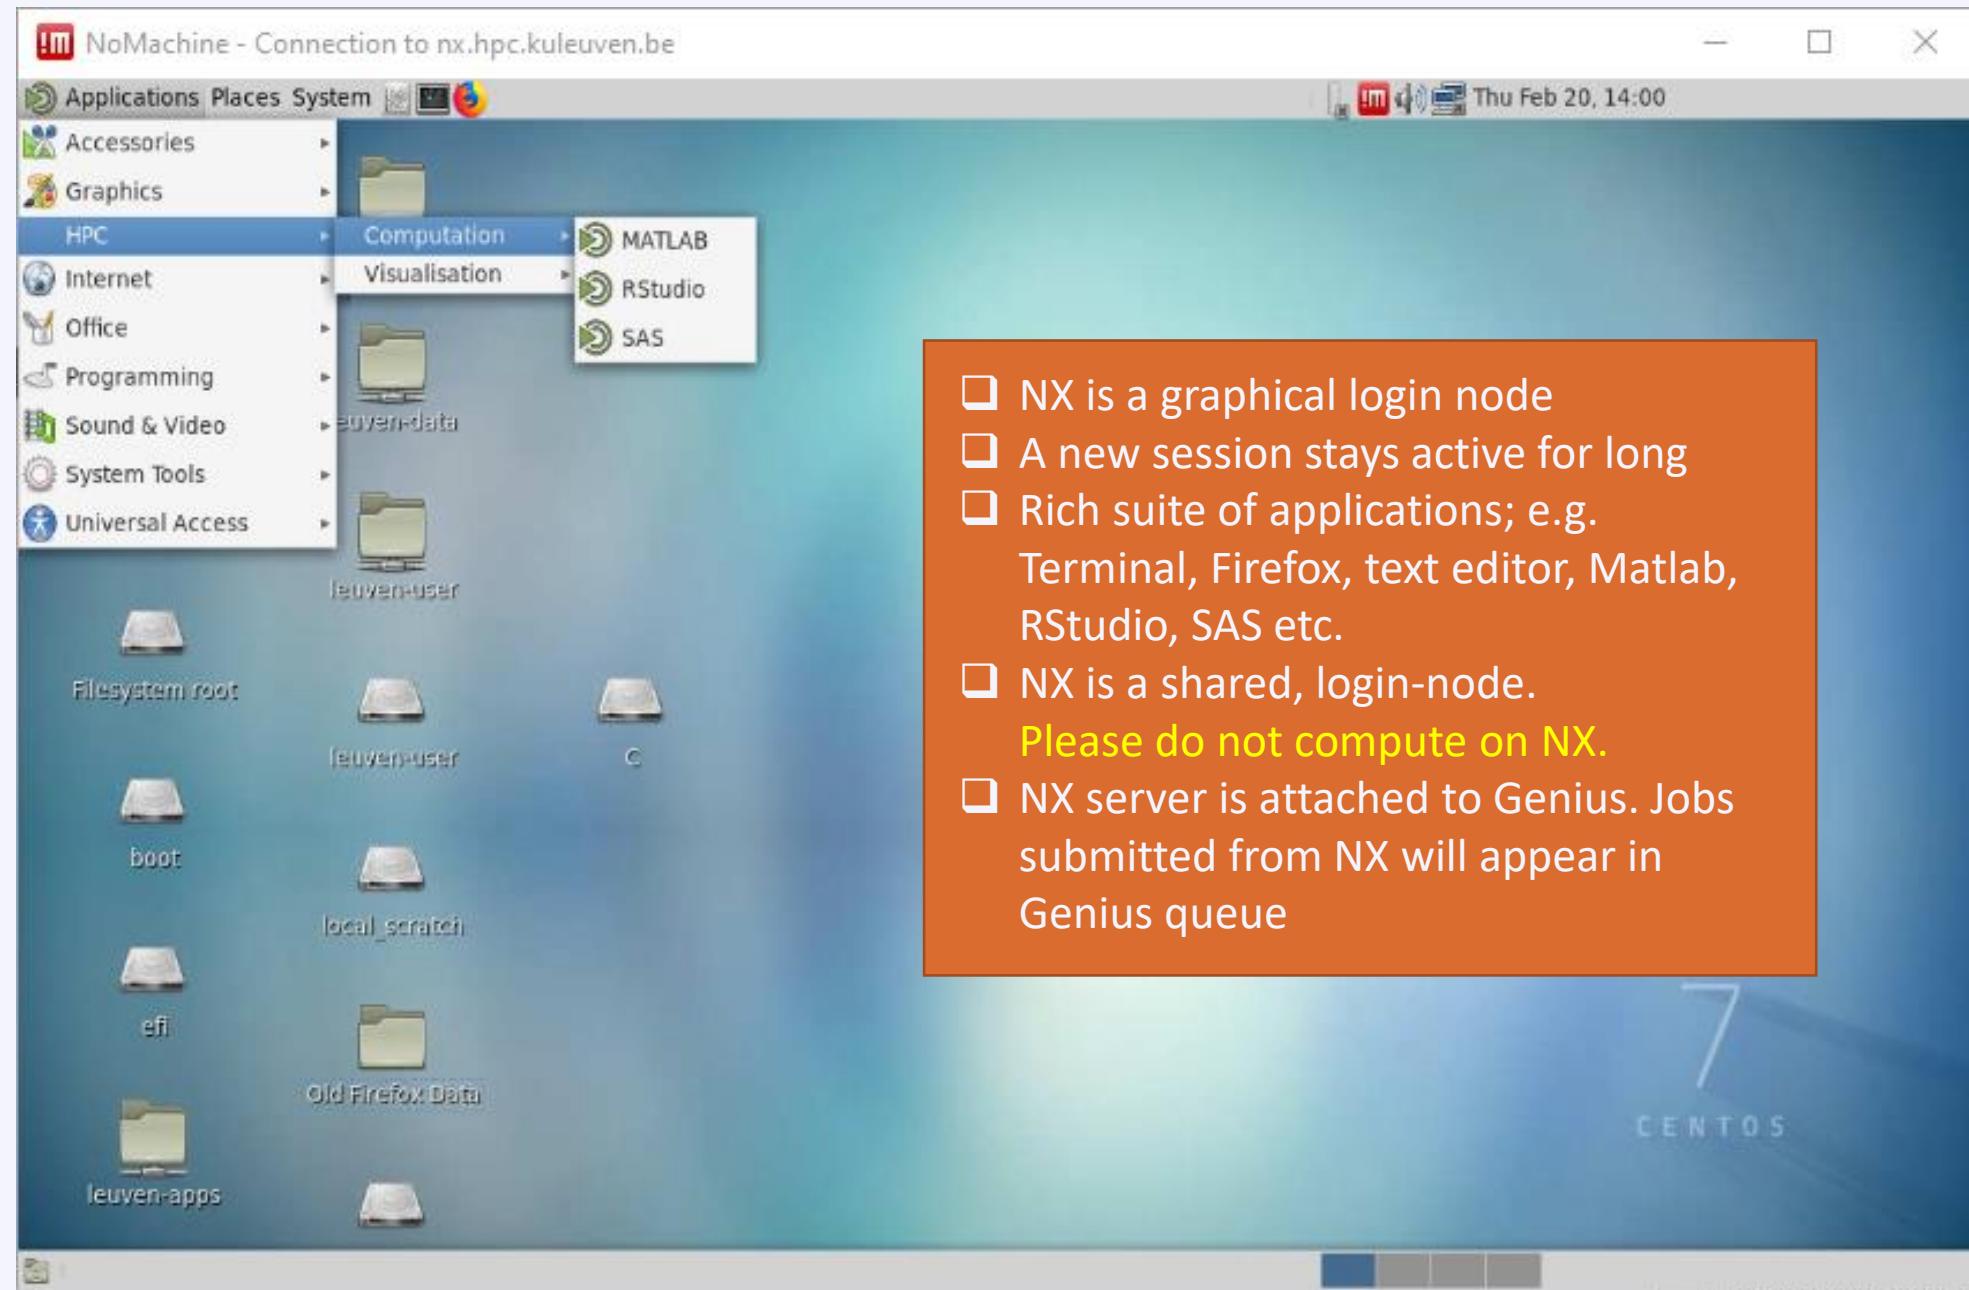


# Setup NX in 10 Steps



# Setup NX in 10 Steps







# Open OnDemand

- Access wICE via web browser
- <https://ondemand.hpc.kuleuven.be>
- Login via KULeuven MFA
- File browser
- Interactive apps + integrated shell
- Create, start and monitor jobs
- Develop, compile and test your code
- VSCode editor
- Jupyter Lab, RStudio and Tensorboard

KU LEUVEN

OnDemand provides an integrated, single access point for all of your HPC resources.

Pinned Apps A featured subset of all available apps

 Active Jobs System Installed App	 Home Directory System Installed App	 Job Composer System Installed App	 Login Server Shell Access System Installed App
 code-server System Installed App	 Interactive Shell System Installed App	 Jupyter Lab System Installed App	 RStudio Server System Installed App
 Tensorboard System Installed App			



# Open OnDemand

- E.g. to start a Jupyter Notebook
- Pick a valid Slurm credit account
- Default partition: batch  
interactive: Max 8 cores, 1 GPU, 16 hr
- Default resources:  
1 core, 1 hour, 3400MB RAM, no GPU
- (At the moment) you cannot use scratch and staging

**Interactive Apps**

Servers
● Interactive Shell
<b>Jupyter Lab</b>
● RStudio Server
● Tensorboard
● code-server
Work in progress
● ParaViewWeb - Work in progress
● cryosparc

## Jupyter Lab

This app will launch a Jupyter Lab server on one or more nodes.

### Account

lpt2\_sysadmin

### Partition

interactive

batch(\_long) or bigmem or interactive or gpu or dedicated...

### Number of hours

3

### Number of cores

1

### Required memory per core in megabytes

3400

### Number of nodes

1

### Number of gpu's

0

[type:] Specify the total number of GPUs slices for the job. An optional GPU type specification can be supplied. For example "A100:3" or "3".

### Reservation (optional)

Name of an existing reservation in which the job should run

I would like to receive an email when the session starts

VLAAMS  
SUPERCOMPUTER

Launch

# Software Stack

# Software: Available Modules

- OS: Linux CentOS 7.x (Genius) and Rocky Linux 8.x (wlCE)
- Toolchains (Genius):  
Intel 2018a (icc, icpc, ifort; Intel MPI; Intel MKL)  
FOSS 2018a (gcc, g++, gfortran; OpenMPI; ScaLAPACK, OpenBLAS, FFTW)
- wlCE: default is 2021a
- Note: **Never mix FOSS and Intel compilers (gives dependency conflict)**

Command	Remark
module av	List all installed modules
module av Python	List all Python-related modules
module use /apps/leuven/skylake/2021a/modules/all	Expose other modules from non-default toolchains, by including them in the search
module spider Python	Get more info
module load Python/3.6.4-intel-2018a	Load a specific module
module list	List all loaded modules and their dependencies
module unload Python/3.6.4-intel-2018a	Unload a module (but dependencies still stay)
module purge	Remove all modules from your work session

# Software: Your Specific Needs

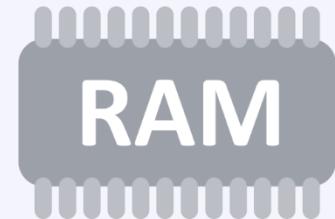
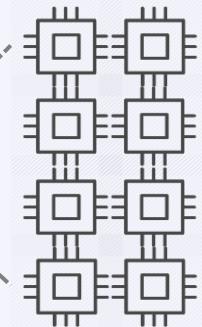
- You can always install your desired software in your \$VSC\_DATA  
Use Intel or FOSS toolchains
- Compile your code on a compute node (with interactive job)
- If you cannot, ask us for help
- Python/R packages for AI and ML must be installed by the users themselves
- Read more about [Python Package Management](#)
- Read more about [R Package Management](#)



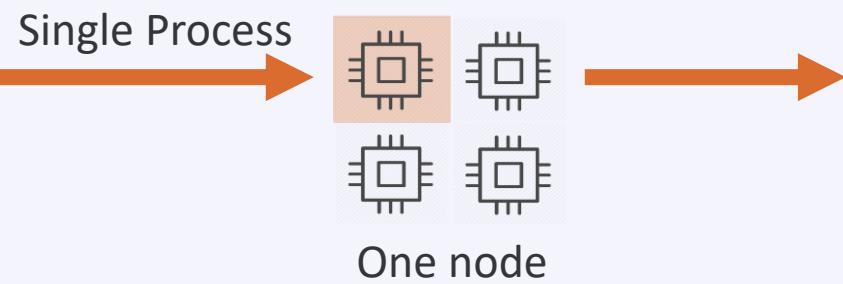
## Starting to Compute

# Resource Glossary

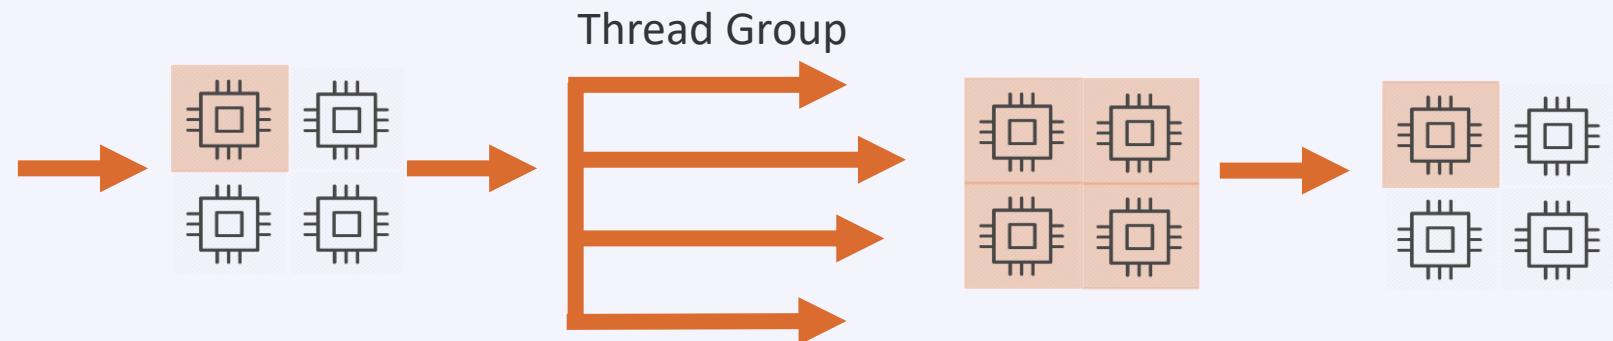
- Nodes:** how many compute servers to request?
- Cores:** how many cores per node to use?
- Memory requirement:** how much memory each core needs?
- Partition:** gpu, bigmem, superdome, amd
- Walltime:** how long to use resources?
- Storage:** how much storage (data, scratch, etc) the job needs?
- Credits:** how many compute credits will be consumed?



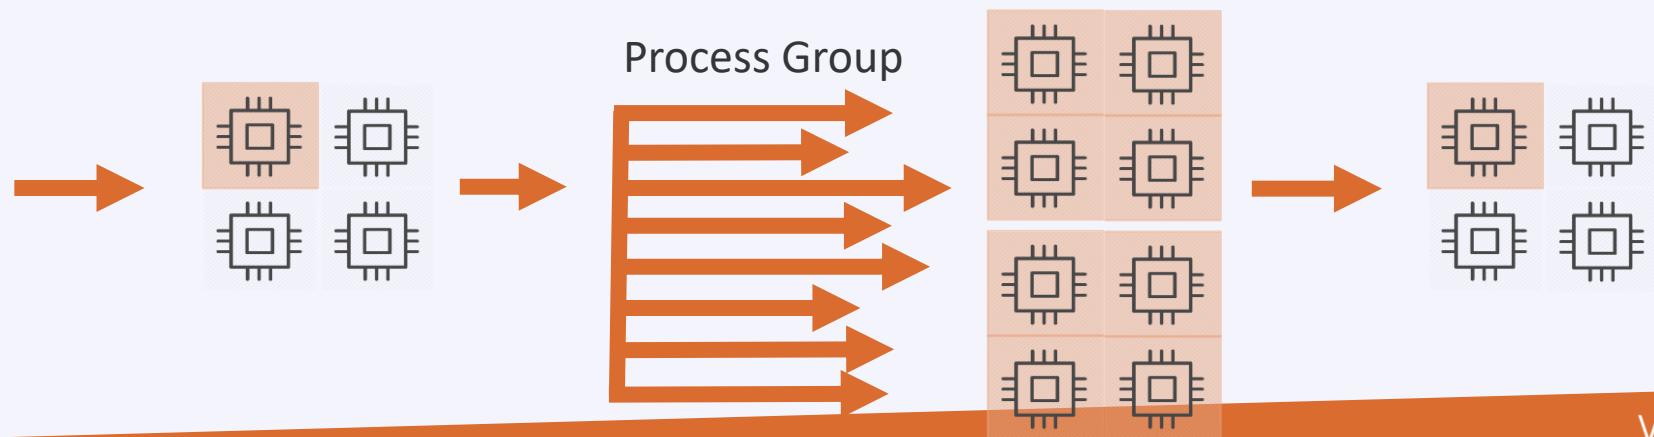
## Serial Application (1 process on 1 core)



## Multi-Core Appl. (N threads on N cores from 1 node)



## Distributed Appl. (many processes on multiple cores/nodes)

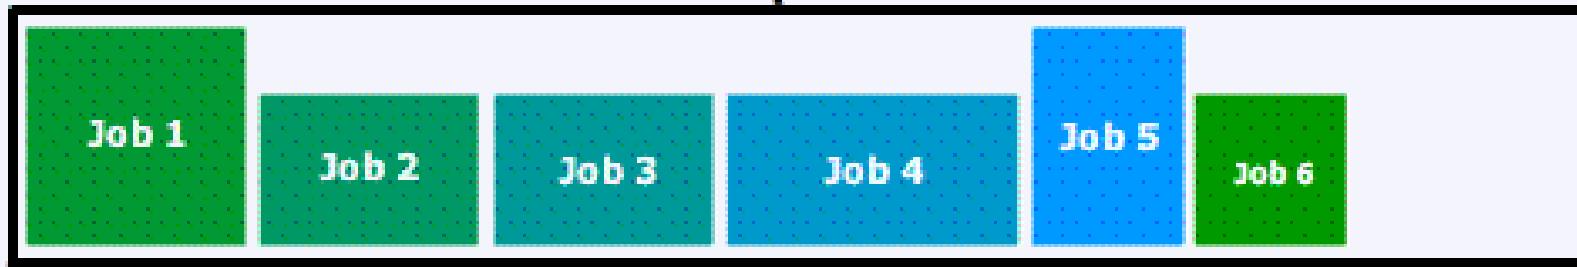


# Backfill

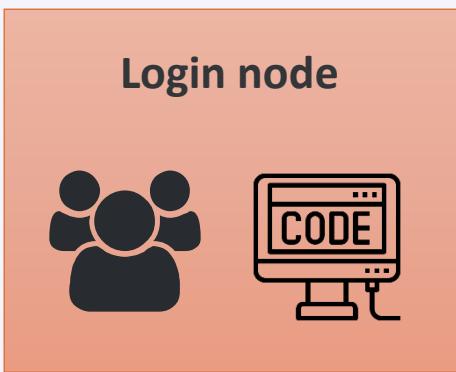
Nodes



Job queue



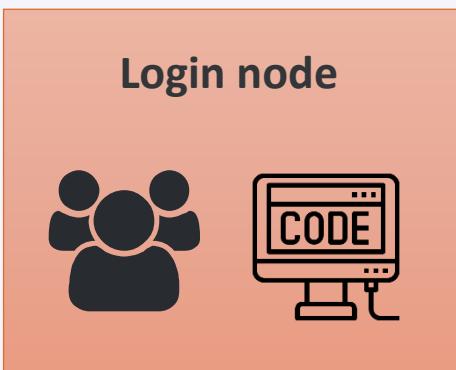
## Interactive Job



`srun ...`

Compute nodes

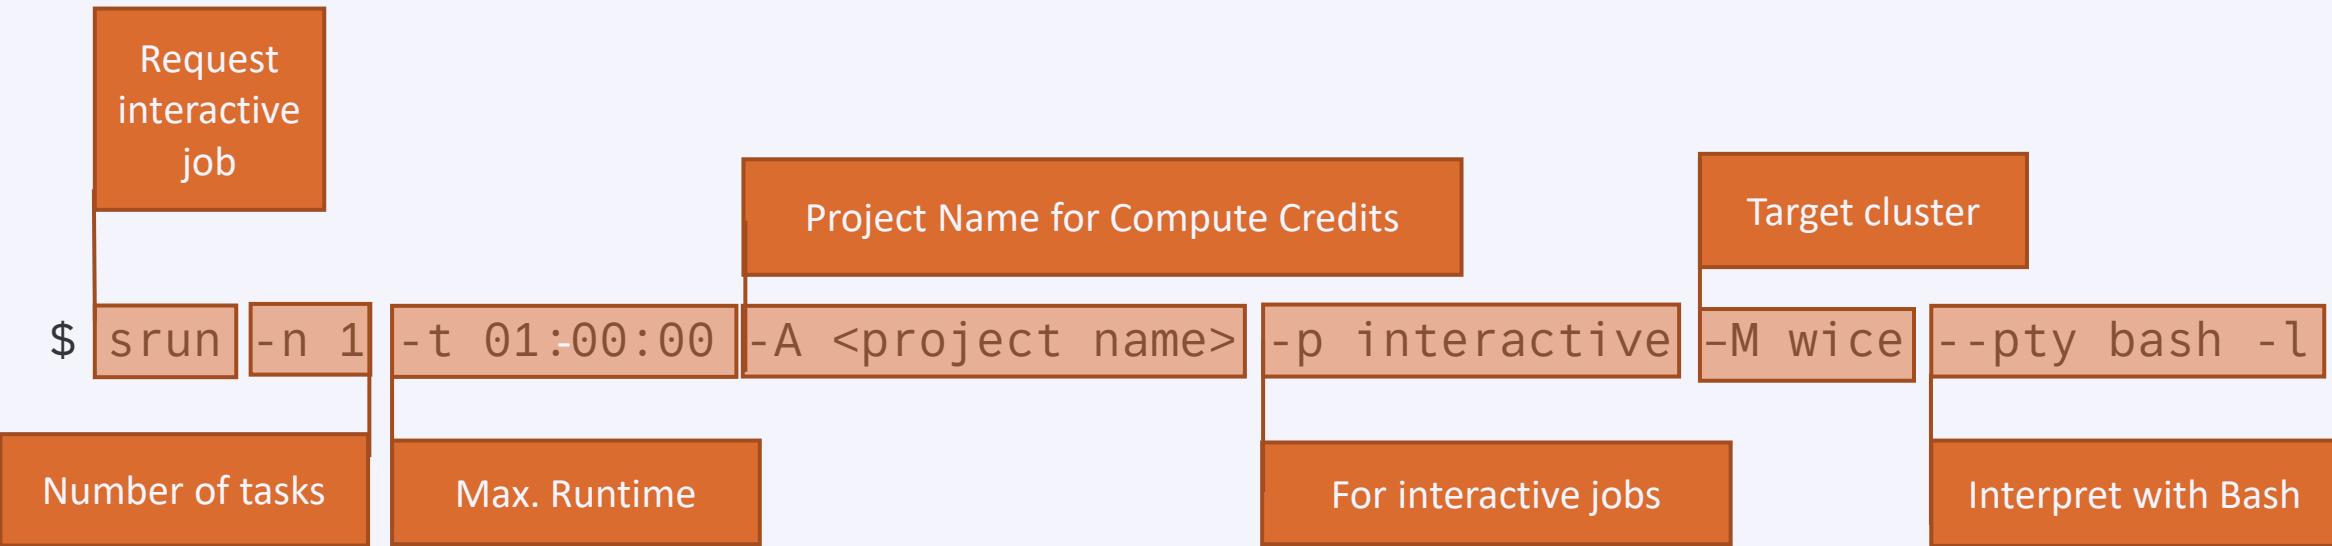
## Batch Job



`sbatch script.pbs`

Compute nodes

# Interactive Job on wICE



## Remark

- Specifying <project name> for credits is mandatory, e.g.: **-A lp\_hpcinfo**
- Implicit defaults for an interactive job on wICE are: **-n 1**, **-t 01:00:00**, **--mem-per-cpu=2000M**

# Interactive Jobs

- Interactive job: 1 core for 1 hour (default)

```
$ srun -M wice -A lp_hpc --pty /bin/bash -l
```

- Interactive job with X-forwarding

```
$ srun -M wice -A lp_hpc -x11 --pty /bin/bash -l
```

- Request fraction of a node from interactive partition

```
$ srun -M wice -A lp_hpc -N 1 -n 4 -p interactive --pty /bin/bash -l
```

- Request a GPU accelerator

```
$ srun -M wice -A lp_hpc -N1 -n 18 -p gpu -G 1 --pty /bin/bash -l
```

## Example: Slurm Job Script

```
#!/bin/bash  
  
#SBATCH --cluster wice  
#SBATCH --account lp_hpcinfo_training  
#SBATCH --nodes 1  
#SBATCH --ntasks 4  
#SBATCH --mem-per-cpu 4G  
#SBATCH --job-name hpc_workflow  
#SBATCH --mail-type BEGIN,END,FAIL  
#SBATCH --mail-user my.name@kuleuven.be  
  
Module --force purge  
module load intel/2021a  
module load Python/3.9.5-GCCcore-10.3.0  
  
which python3  
  
cd $VSC_SCRATCH/projects/simulations  
cp -r $VSC_DATA/input_data .  
  
python modelling.py  
  
cp -r output_data $VSC_DATA  
rm -rf ./input_data ./output_data
```

Shebang

Resource List

Module load(s)

Move data

Execute commands

Move data

# Batch Workload

Job Script

```
#!/bin/bash
```

```
...
```

- Submit the job to the batch server
- Receive a unique JobID
- Error and output files

Command Line

```
$ sbatch simulation.slurm
```

JobID

**Submitted batch job 60042478 on cluster  
wice**

stderr, stdout

```
$ ls *.out  
slurm-60042478.out
```

# Output File

- STDERR and STDOUT are by default redirected to a single file:
- slurm-<Job ID>.out
- Contains job info, all errors and warnings, and printouts
- Always study it
- Address all warnings and errors (if you can)
- Typical error examples ...

STDOUT + STDERR

```
$ ls slurm-* .out  
slurm-60042478.out
```

Out of Memory

```
slurmstepd: error: Detected 1 oom-kill event(s).  
Some of your processes may have been killed by the  
cgroup out-of-memory handler.
```

Short Walltime

```
slurmstepd: error: *** JOB 60042478 ON s28c11n2  
CANCELLED AT 2023-02-08T10:03:43 DUE TO TIME LIMIT  
***
```

Low Disk Space

```
IOError: [Errno 122] Disk quota exceeded
```

# Output File

STDOUT + STDERR

```
$ cat slurm-60042490.out
SLURM_JOB_ID: 60042490
SLURM_JOB_USER: vsc30745
SLURM_JOB_ACCOUNT: lp_hpcinfo
SLURM_JOB_NAME: hpc_workflow
SLURM_CLUSTER_NAME: wice
SLURM_JOB_PARTITION: batch
SLURM_NNODES: 1
SLURM_NODELIST: s28c11n2
SLURM_JOB_CPUS_PER_NODE: 8
Date: Wed Feb  8 10:41:21 CET 2023
Walltime: 00-01:00:00
=====
/apps/leuven/skylake/2018a/software/Python/3.6.4-intel-2018a/bin/python3
cp: cannot stat '/data/leuven/307/vsc30745/input_data': No such file or
directory
python: can't open file 'modelling.py': [Errno 2] No such file or directory
cp: cannot stat 'output_data': No such file or directory
```

# Inspecting Jobs

Example

```
$ scontrol show job -M wice 60049330
JobId=60049330 JobName=f70.slurm
  UserId=vsc3....(253...) GroupId=vsc3....(253...)
Account=lp_metacommunity_models QOS=lp_metacommunity_models
JobState=PENDING Reason=QOSGrpBillingMinutes
  SubmitTime=2023-02-16T00:24:58 EligibleTime=2023-02-16T00:24:58
  Partition=batch NodeList= NumNodes=4-4 NumCPUs=288 NumTasks=288 CPUs/Task=1
  TRES(cpu=288,mem=979200M,node=4,billing=733
  MinCPUsNode=72 MinMemoryCPU=3400M
  WorkDir=/vsc-hard-mounts/leuven-data/3.../vsc3.../Odonate_SOM
  StdErr=/vsc-hard-mounts/leuven-data/3.../vsc3.../Odonate_SOM/slurm-60049330.out
  StdIn=/dev/null
  StdOut=/vsc-hard-mounts/leuven-data/3.../vsc3.../Odonate_SOM/slurm-60049330.out
```

Diagnosis

Why is my job in pending/hold state?

Check out the “Reason” on Slurm docs: [https://slurm.schedmd.com/resource\\_limits.html](https://slurm.schedmd.com/resource_limits.html)

# Other Partitions on wICE

GPU

```
#SBATCH --partition=gpu
#SBATCH --nodes=1
#SBATCH --ntasks=18
#SBATCH --gpus-per-node=1
```

Big Memory

```
#SBATCH --partition=bigmem
#SBATCH --nodes=1
#SBATCH --tasks-per-node=72
#SBATCH --mem-per-cpu=28000M
```

# Managing & Monitoring Jobs

Command	Purpose
sbatch ...	Submit a batch job
srun ...	Submit an interactive job
scancel --cluster=wice <JobID>	Cancel a specific pending or running job
scontrol show job --cluster=wice <JobID> slurm_jobinfo <JobID>	Detailed job info (very useful to diagnose issues)
squeue --cluster=wice -long	Status of all recent jobs
squeue --cluster=wice -start	Give a <i>rough</i> estimate of start time
sinfo --cluster=wice	Info about the state of available partitions and nodes
sacct --cluster=wice --batch --job <JobID>	Show minimal info about a queue or partition
slurmtop --cluster=wice	Overview of the cluster

# Interactive Partition

- Accessible via command line and Open OnDemand
- To quickly compile, test, debug your (parallel) application
- To pre-/post-process your data and make visualizations
- Short queue time
- 5 dedicated nodes on wICE
- 64 cores per node, 1 GPU (=7 GPU instances), 512 GB memory
- Max resource per user: 8 cores, 1 GPU instance, 16 hour walltime

Interactive Job

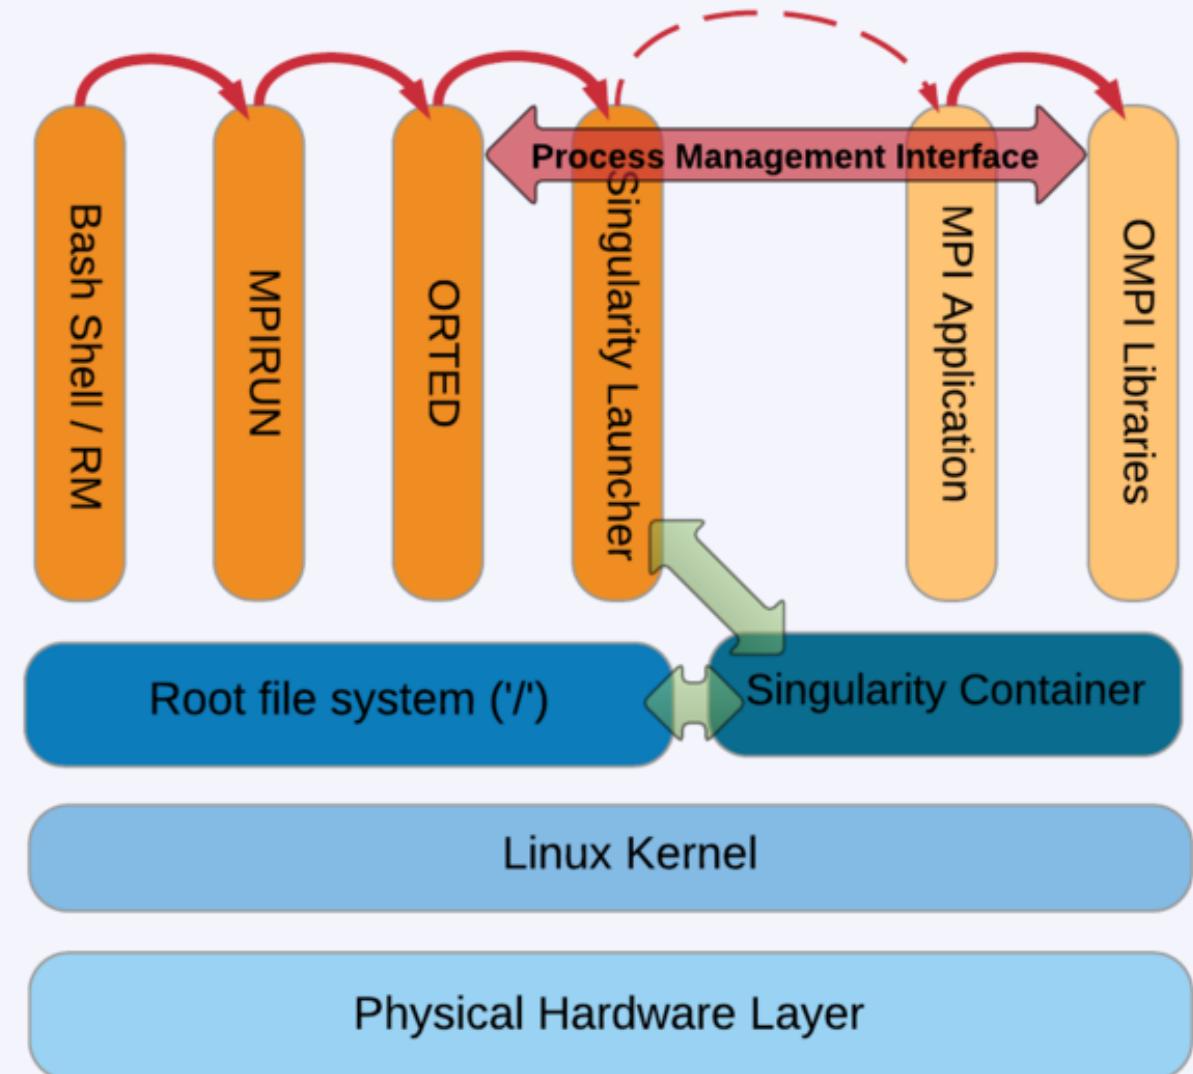
```
$ srun -n 1 -p interactive -t 16:00:00 -G 1 --pty bash -l
```

# Singularity Containers

- What?
  - + Self-contained OS & software & data
- Why?
  - + fully resolved dependency chains
  - + portable workflow
- How?
  - + You create the image
  - + Run it on Genius
  - + MPI/OpenMP is supported

```
#!/bin/bash -l
#SBATCH -t 30:00
#SBATCH -N 2
#SBATCH -n 72
#SBATCH -M wice
#SBATCH -A lp_hpcinfo_training
singularity run Project.simg ./model.exe
```

Slurm Script



# Credit Pricing

- You pay as you go!
- For academic projects:  
1000 credits = 3.5 EUR
- Credits needed for a job:  
$$\text{#credits} = \text{Walltime (hr)} \times \text{\#nodes} \times \text{Factor}$$
- For shared nodes, you pay a fraction of the costs,  
based on the ppn specified

Cluster / Partition	Credits/hr
Genius Cascadelake	11.3
Genius Cascadelake 8 GPUs	40
Genius Skylake 4 GPUs	20
Genius Skylake BigMem	12
Genius Skylake	10
Genius / Superdome	10
Genius / AMD nodes	10
ThinKing / Haswell	6.68
ThinKing / IvyBridge	4.76

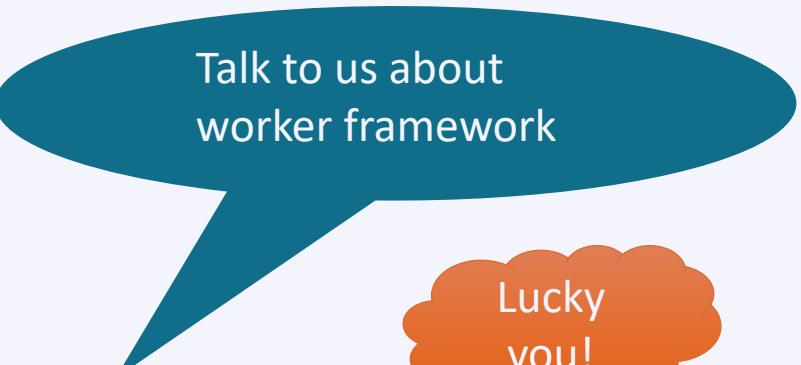
# How to Manage Credits?

Command	Purpose
sam-balance	List active projects and available credits
sam-list-allocations	List the validity dates of different projects/allocations
sam-list-usagerecords	List current charge rates for different nodes
sam-statement -A <account-name>	Job statements from an account

# Worker Framework

# Parallel Computing

- Serial:
  - one program, on one core
- 'Embarrassingly parallel' problems:
  - lots of runs of one program, with different parameters
- Problems that require 'real' parallel algorithms
  - OpenMP
  - MPI : Message Passing Interface



Talk to us about  
worker framework



Lucky  
you!

# Use case: parameter exploration

temperature	pressure	humidity
293.0	1.0e05	87
...	...	...
313.0	1.3e05	75

```
#!/bin/bash -l
#PBS -l nodes=1:ppn=1 -l walltime=00:10:00
    #!/bin/bash -l
cd
wea
:
cd
wea
:
cd $PBS_O_WORKDIR
weather -p 1.3e05 -t 313.0 -h 75
```

job\_001  
job\_030.pbs  
.00  
job\_600.pbs

Many single core computations

# Solution: worker with -data

temperature	pressure	humidity
293.0	1.0e05	87
...	...	...
313.0	1.3e05	75

data.csv

```
#!/bin/bash -l  
#PBS -l nodes=5:ppn=20 -l walltime=01:20:00  
  
cd $PBS_O_WORKDIR  
weather -p $pressure -t $temperature -h $humidity
```

job.pbs

```
$ module load worker  
$ wsub -data data.csv -batch job.pbs
```

# Data exploration: steps

- Write PBS script with parameters
- Create Excel sheet with data
  - Convert to CSV format
- Submit with `wsub`
  - walltime is time to complete all work items

$$\text{walltime}_{\text{job}} \geq \frac{N \cdot \text{walltime}_{\text{work item}}}{\text{nodes} \cdot \text{ppn}}$$

The background of the slide features a vibrant, hand-drawn style illustration of a train station. On the left, a train with several green and brown cylindrical cargo cars is visible. In the center, there's a large digital screen displaying a grid of numbers and some orange arrows pointing upwards. The overall scene is colorful and artistic.

**Demo: Test for yourself**

# demo/test yourself

- ✓ Request membership to lp\_hpcintro\_training group (account.vscentrum.be)
  
- ✓ Login with putty
- ✓ Filetransfer with Filezilla
- ✓ Login with NX
- ✓ Check disk quota
- ✓ Check the credits
- ✓ Check/load/list/unload/purge module

# demo/test yourself

- ✓ Copy intro training files (`/apps/leuven/training/HPC_intro/`) to your `$VSC_HOME`
- ✓ Submit cpujob to the cluster
- ✓ List all your jobs (`qstat`)
- ✓ Check the information about the cpujob (`checkjob`)
- ✓ Modify the mat.pbs script to request 1 node, 36 cores for 30 minutes and get the notification about job start/end by e-mail
- ✓ Check the status of all the jobs

# demo - monitoring

- ✓ Submit an interactive job

Run your program on a compute node

Open a new terminal and ssh to a compute node

Check the resources usage (`top`, `htop`)

# demo – conda usage

- ✓ Create a conda environment including Jupyter

```
$ conda create -n science jupyter numpy scipy
```

Activate this environment

- ✓ \$ source activate science

Add matplotlib package to this environment

- ✓ \$ conda install matplotlib

Return to original environment

```
$ conda deactivate
```

# demo – notebooks

- ✓ Start an interactive GPU job

```
$ qsub -I -l walltime=30:00 -l nodes=1:ppn=9:gpus=1 -l partition=gpu -A default_project
```

- ✓ Activate conda environment

```
$ source activate science
```

- ✓ Go to your working directory (you can use \$PBS\_O\_WORKDIR if you qsub from there)

Start notebook

```
$ jupyter notebook --port ${USER:3} --ip $(hostname)
```

```
$ jupyter notebook --port 30468 --ip $(hostname)
```

- ✓ Open the link in the browser in NX and test your notebook

# demo – worker

- ✓ Copy intro training files (/apps/leuven/training/worker/) to your \$VSC\_HOME
- ✓ Go to exercise1 directory
- ✓ Submit worker job
- ✓ Check the output file

# Questions

Helpdesk:

[hpcinfo@kuleuven.be](mailto:hpcinfo@kuleuven.be) or [https://admin.kuleuven.be/icts/HPCinfo\\_form/HPC-info-formulier](https://admin.kuleuven.be/icts/HPCinfo_form/HPC-info-formulier)

VSC web site:

<http://www.vscentrum.be/>

VSC documentation: <https://docs.vscentrum.be/en/latest/>

VSC agenda: training sessions, events

Systems status page:

<http://status.kuleuven.be/hpc>

*Stay Connected  
to vsc*

